

# Introduction to Kubernetes

Instructor:

**John J Rofrano**

Senior Technical Staff Member, DevOps Champion

IBM T.J. Watson Research Center

[rofrano@us.ibm.com](mailto:rofrano@us.ibm.com)

@JohnRofrano 



# What will you learn?

- What is Kubernetes
- Why would you use a Kubernetes for deployment
- Kubernetes Architecture Overview
- How to deploy your own containers in Kubernetes



# Kubernetes



# Kubernetes is the de facto Orchestrator for Containers

The screenshot shows the Amazon EKS homepage. At the top, there's a navigation bar with links for 'Contact Sales', 'Products', 'Solutions', 'Pricing', 'More', 'English', 'My Account', and 'Sign In to the Console'. Below the navigation is a banner with the text 'Amazon EKS' and 'Highly available and scalable Kubernetes service'. A large blue button says 'Sign up for the Preview'. The main content area contains a detailed description of Amazon EKS, mentioning its managed nature, Kubernetes integration, and support for containerized applications across multiple Availability Zones. It also highlights automated upgrades and patching.

**KUBERNETES ENGINE**  
Deploy, manage, and scale containerized applications on Kubernetes, powered by Google Cloud

**Containerized Application Management at Scale**  
Google Kubernetes Engine is a managed environment for deploying containerized applications. It brings our latest innovations in developer productivity, resource efficiency, automated operations, and open source flexibility to accelerate your time to market.

Every major cloud provider supports Kubernetes

The screenshot shows the IBM Cloud Kubernetes Service homepage. At the top, there's a navigation bar with links for 'Cloud', 'Why IBM', 'Products', 'Solutions', 'Garage', 'Pricing', 'Blog', 'Docs', and 'Support'. Below the navigation is a banner with the text 'IBM Cloud Kubernetes Service' and 'Explore now with a starter account on IBM Cloud to start creating your clusters'. A 'Sign up' button and a 'Get started with clusters' button are visible. The main content area features a video player showing a man speaking about the service. To the right, there's a section titled 'What is IBM Cloud Kubernetes Service?' which provides a brief overview of the service's features, including intelligent scheduling, self-healing, horizontal scaling, service discovery, and load balancing. It also mentions automated rollouts and rollbacks, and secret and configuration management. A 'Get the FAQs' button and a 'Let's talk' button are located at the bottom right.

**Microsoft Azure**

Azure Container Service (AKS)  
Simplify the deployment, management, and operations of Kubernetes  
Use a fully managed Kubernetes container orchestration service or choose other orchestrators.

**Announcing the public preview of Managed Kubernetes for Azure Container Service (AKS) >**

# Why choose Kubernetes?

- An open-source system for automating deployment, scaling, and management of containerized applications.
- No Vendor Lock-In
- Large Community of support
- Robust platform for container orchestration
- Based on 15 years of experience of running production workloads at Google, combined with best-of-breed ideas and practices from the community



# Kubernetes is an Orchestration Platform

- **Scheduling**  
decide where containers run
  - **Lifecycle and Health**  
keep containers running and restart them if they fail
  - **Scaling**  
grow and shrink deployments as needed
  - **Naming and Discovery**  
help containers find each other
  - **Load Balancing**  
distribute traffic across containers
- ...and a whole lot more



# Why Do You Need Container Orchestration?

- Deploy applications to servers without worrying about specific servers
- Scale the application horizontally up and down
- Restore the application if the server on which it worked fails
  - This is called container auto-healing or rescheduling



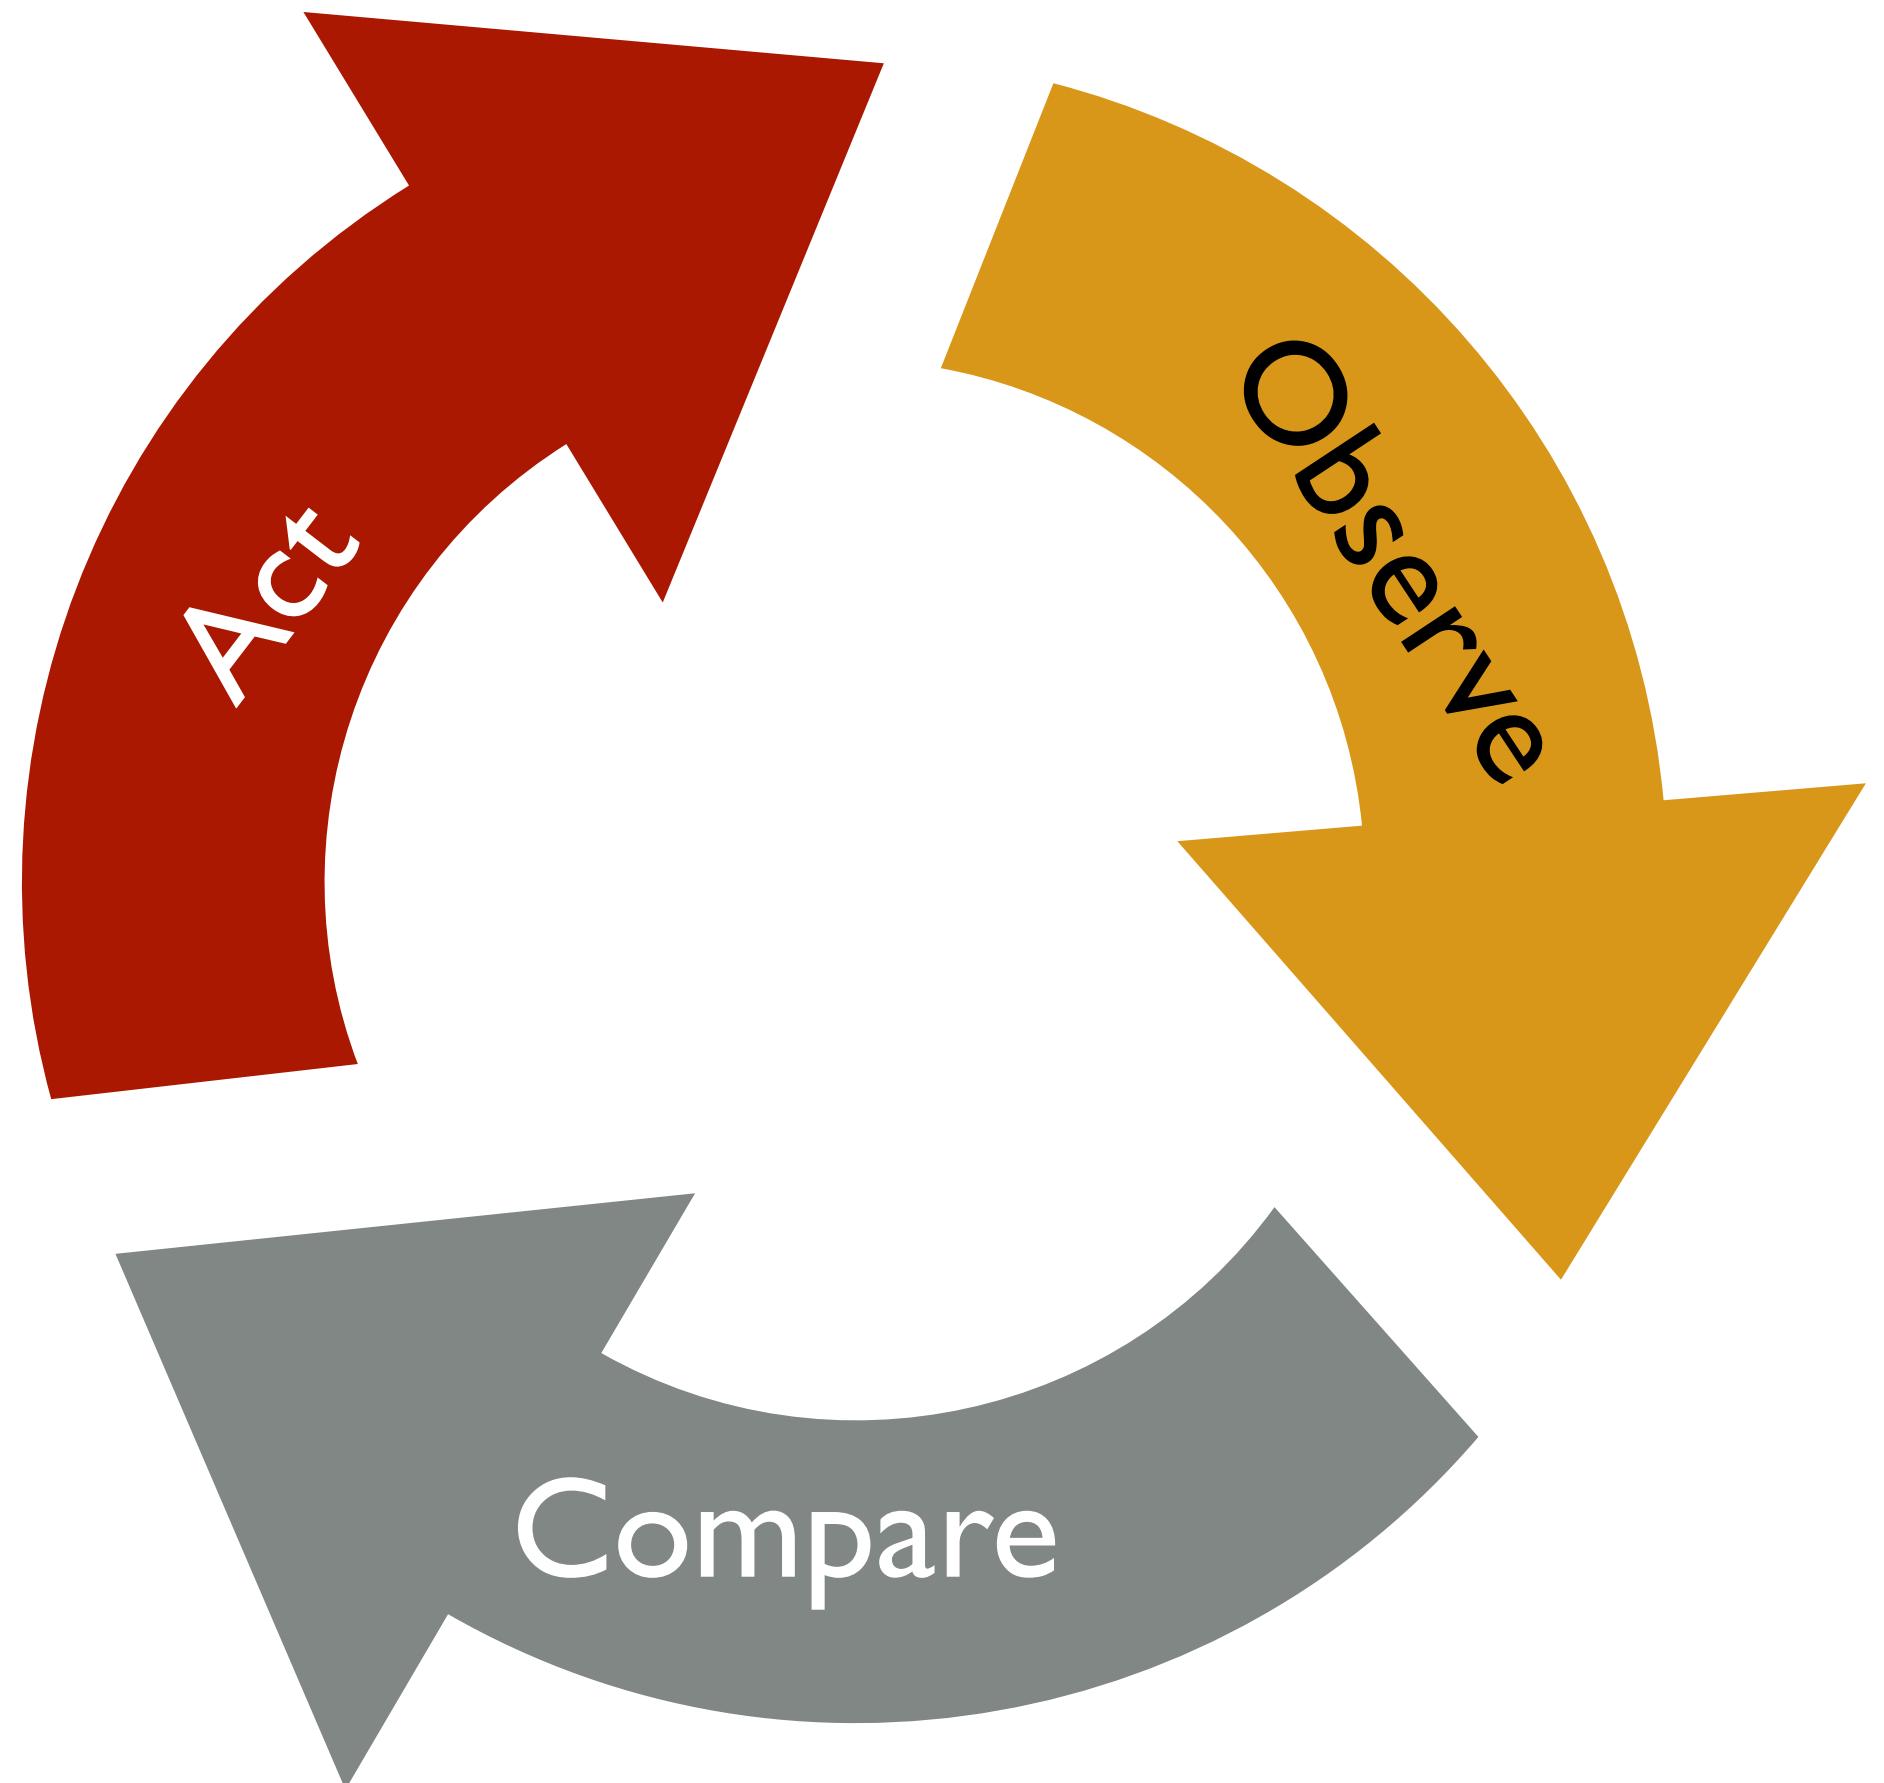
# Kubernetes has a Declarative API

Kubernetes is a Declarative Model

You express the desired state

Kubernetes maintains it

What could be simpler?

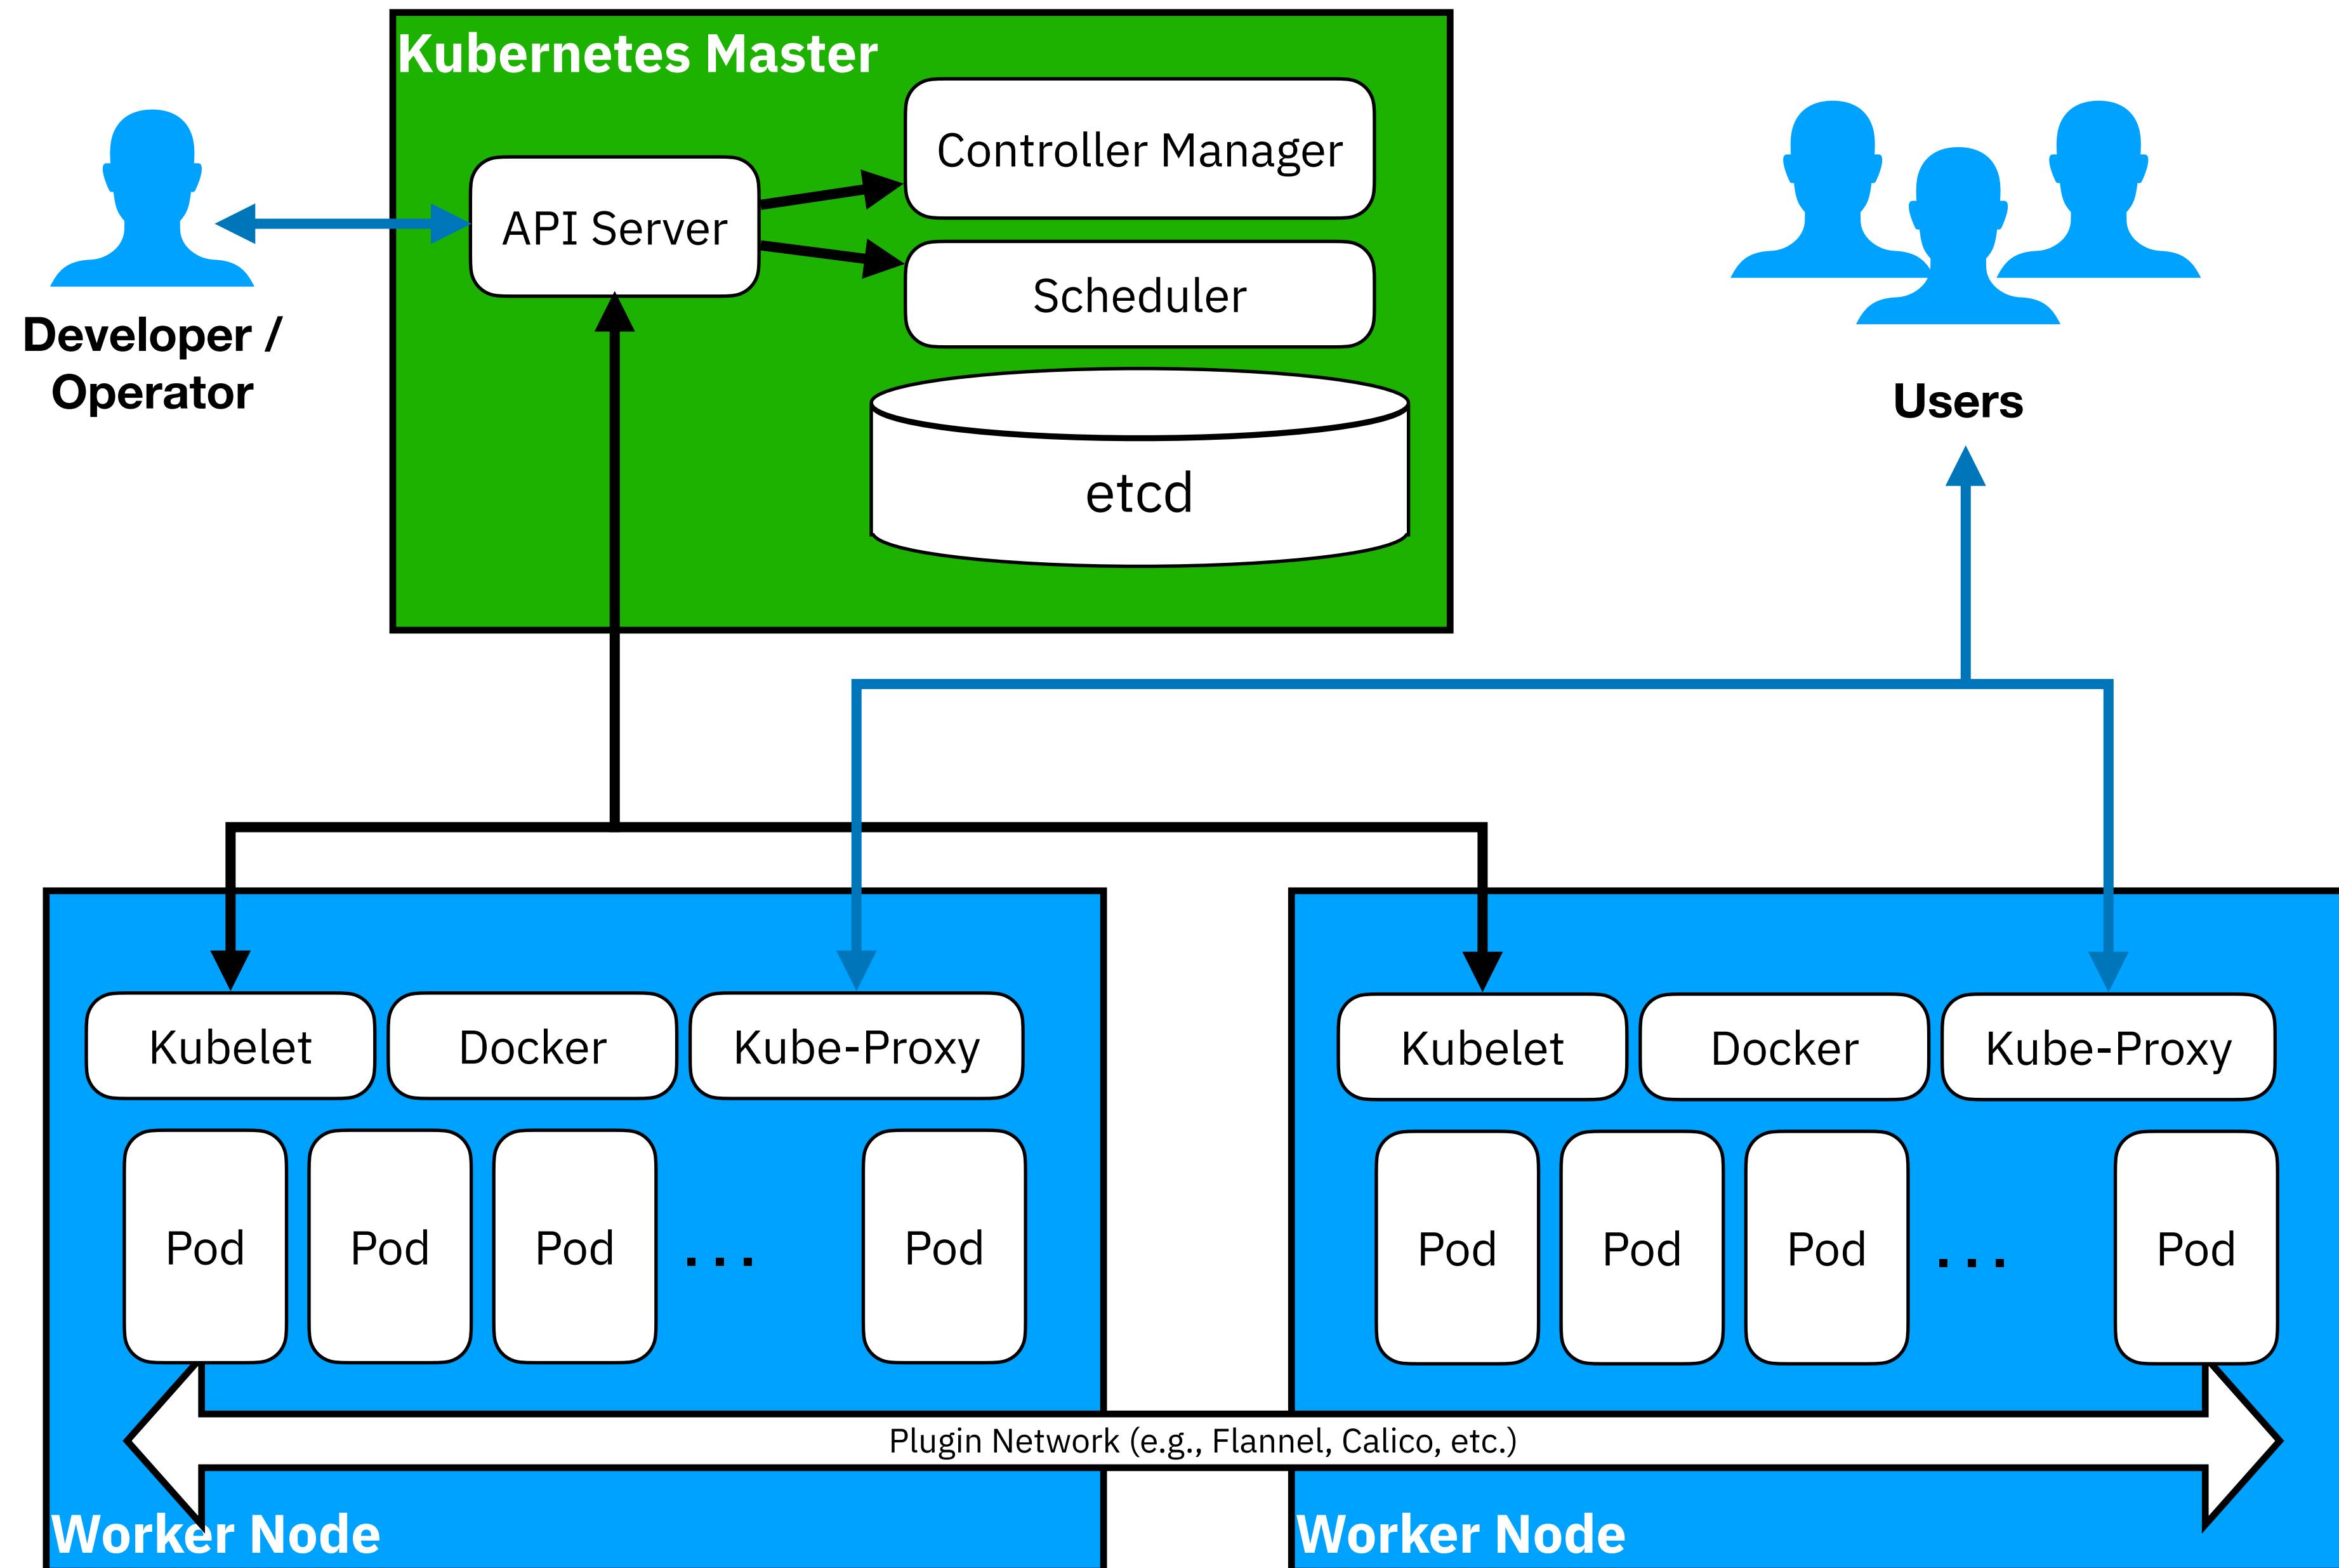


# Kubernetes Architecture



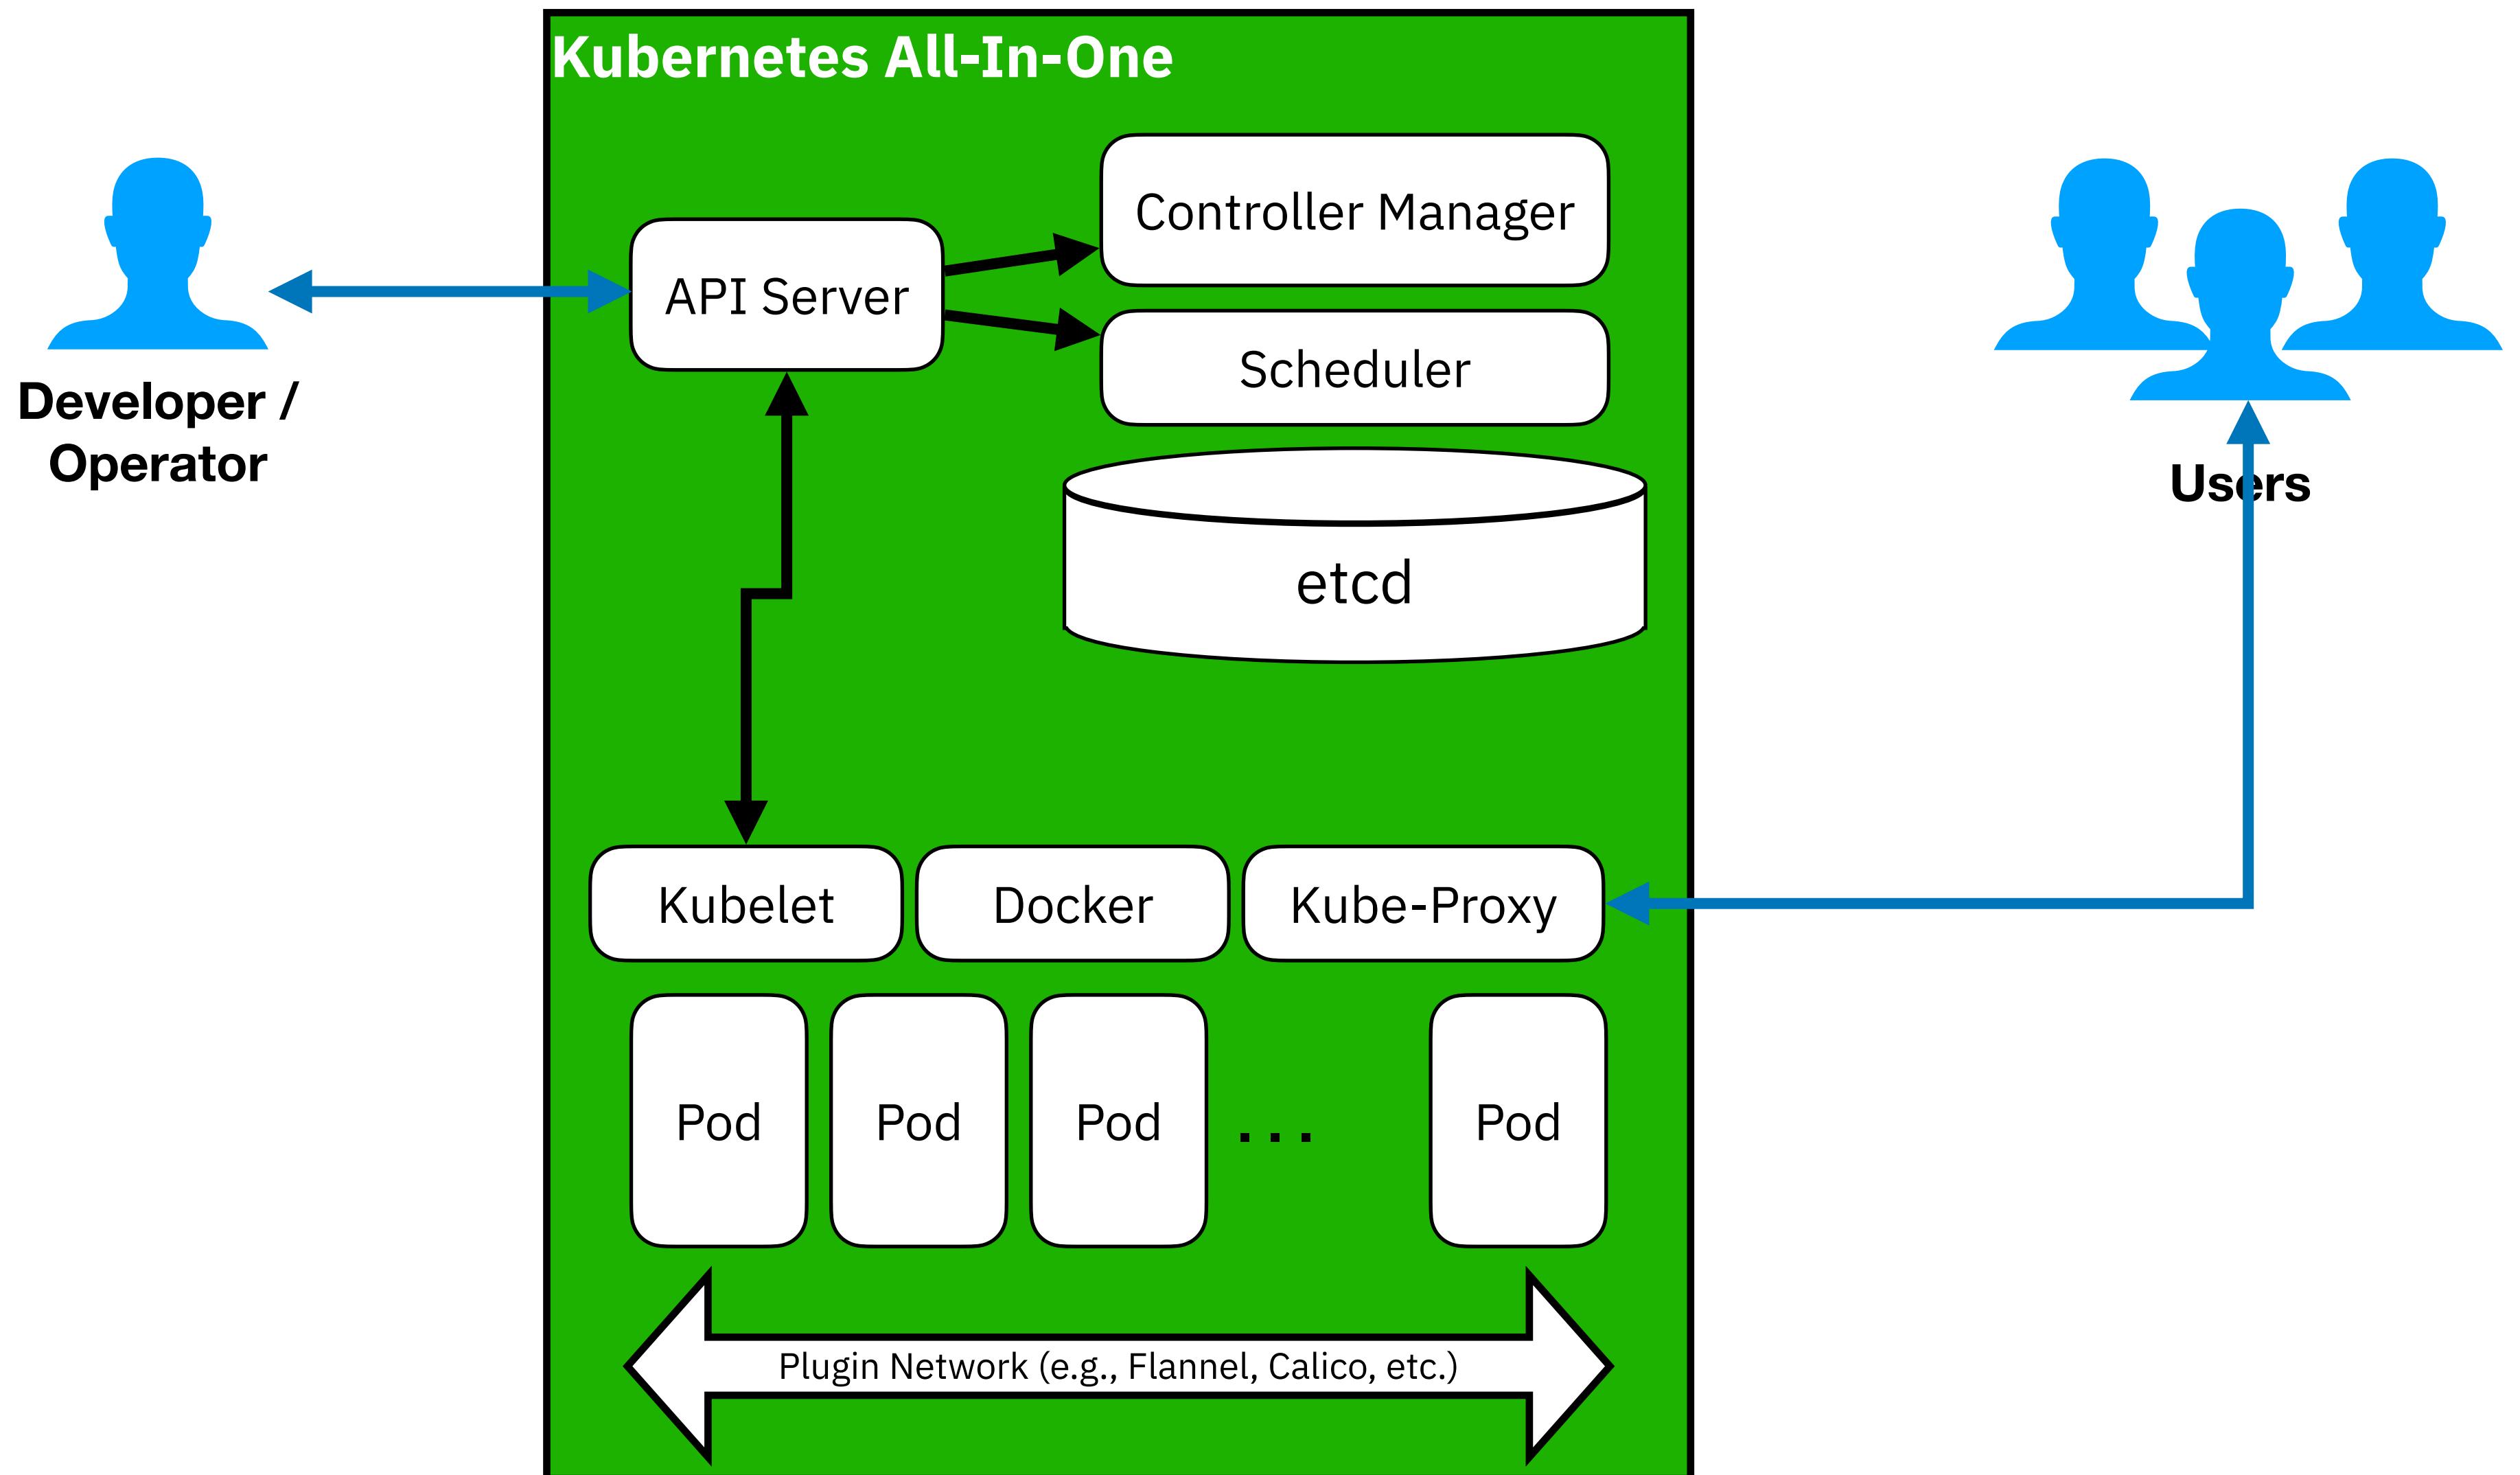
# Kubernetes Architecture

- There can be one or more Master Nodes (HA)
- There can be zero or more Worker Nodes
- Everyone communicates via the API Server
- Kubelet on each Worker Node acting as an agent
- Everything can be on one node for development use



# Kubernetes All-In-One

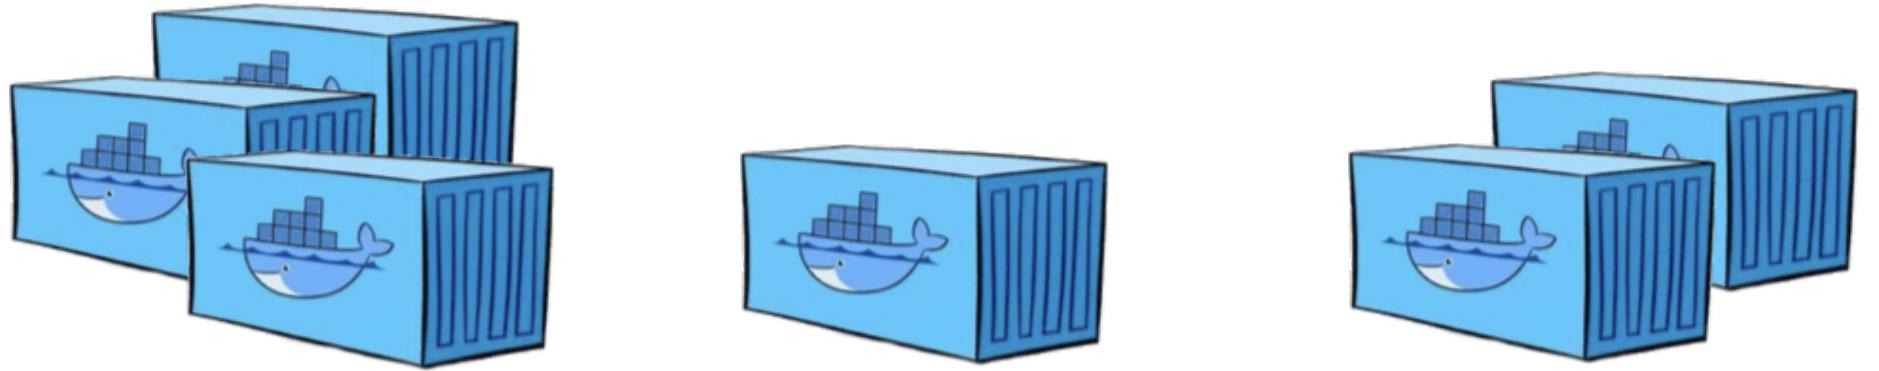
- You can run Kubernetes all in one VM for development work (not for production!)
- **Minikube** is great for setting this up
- **Minishift** will deploy a development version of RedHat OpenShift which uses Kubernetes



# Kubernetes is the new "Cloud OS"

## Managing Containers with VMs

- SysAdmins must decide where to place containers
- Workload balancing is manual



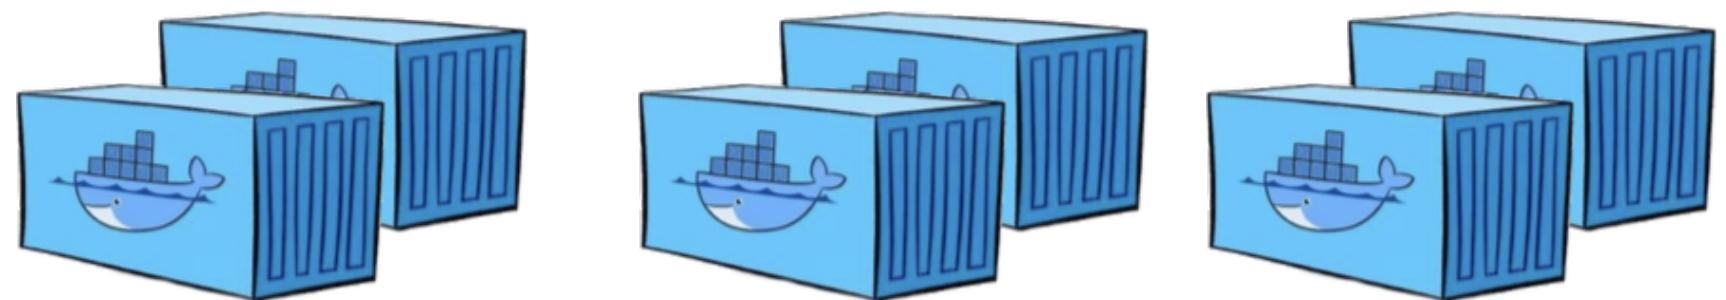
Virtual  
Machine

Virtual  
Machine

Virtual  
Machine

## Managing Containers with Kubernetes

- Kubernetes schedules and optimizes workloads automatically



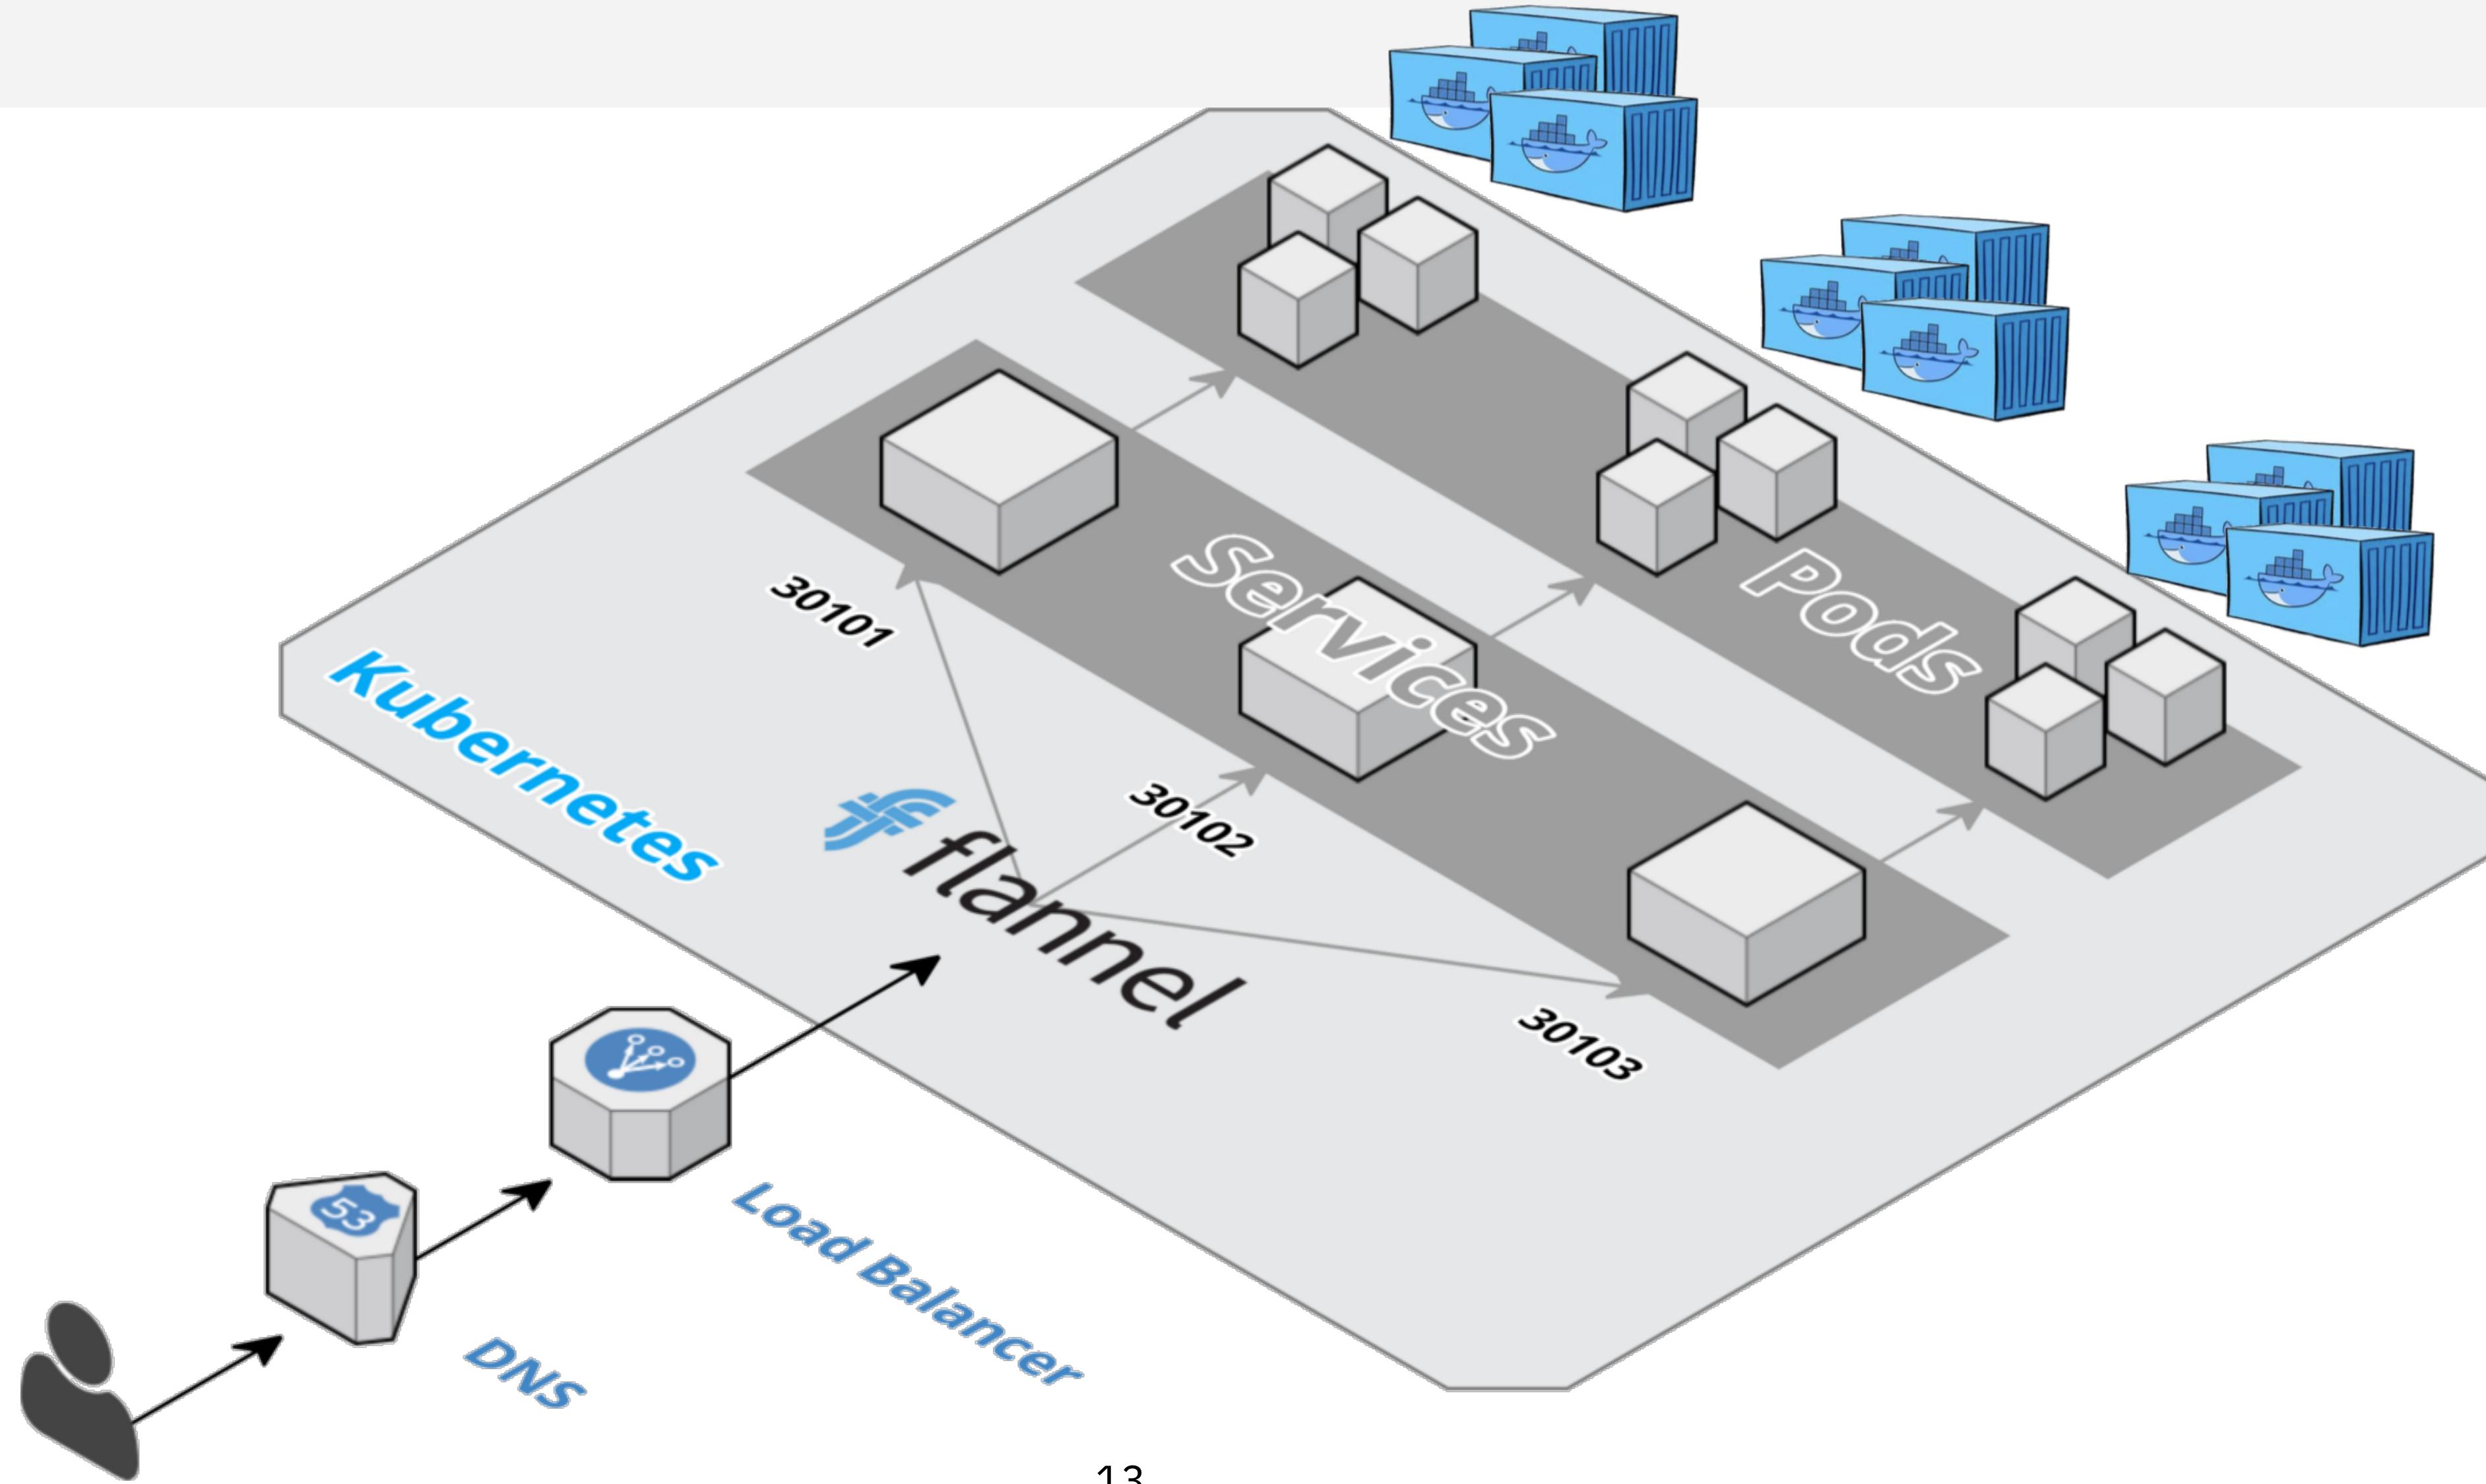
Kubernetes

Master VM

Node VM

Node VM

# Kubernetes @ 20,000 ft.



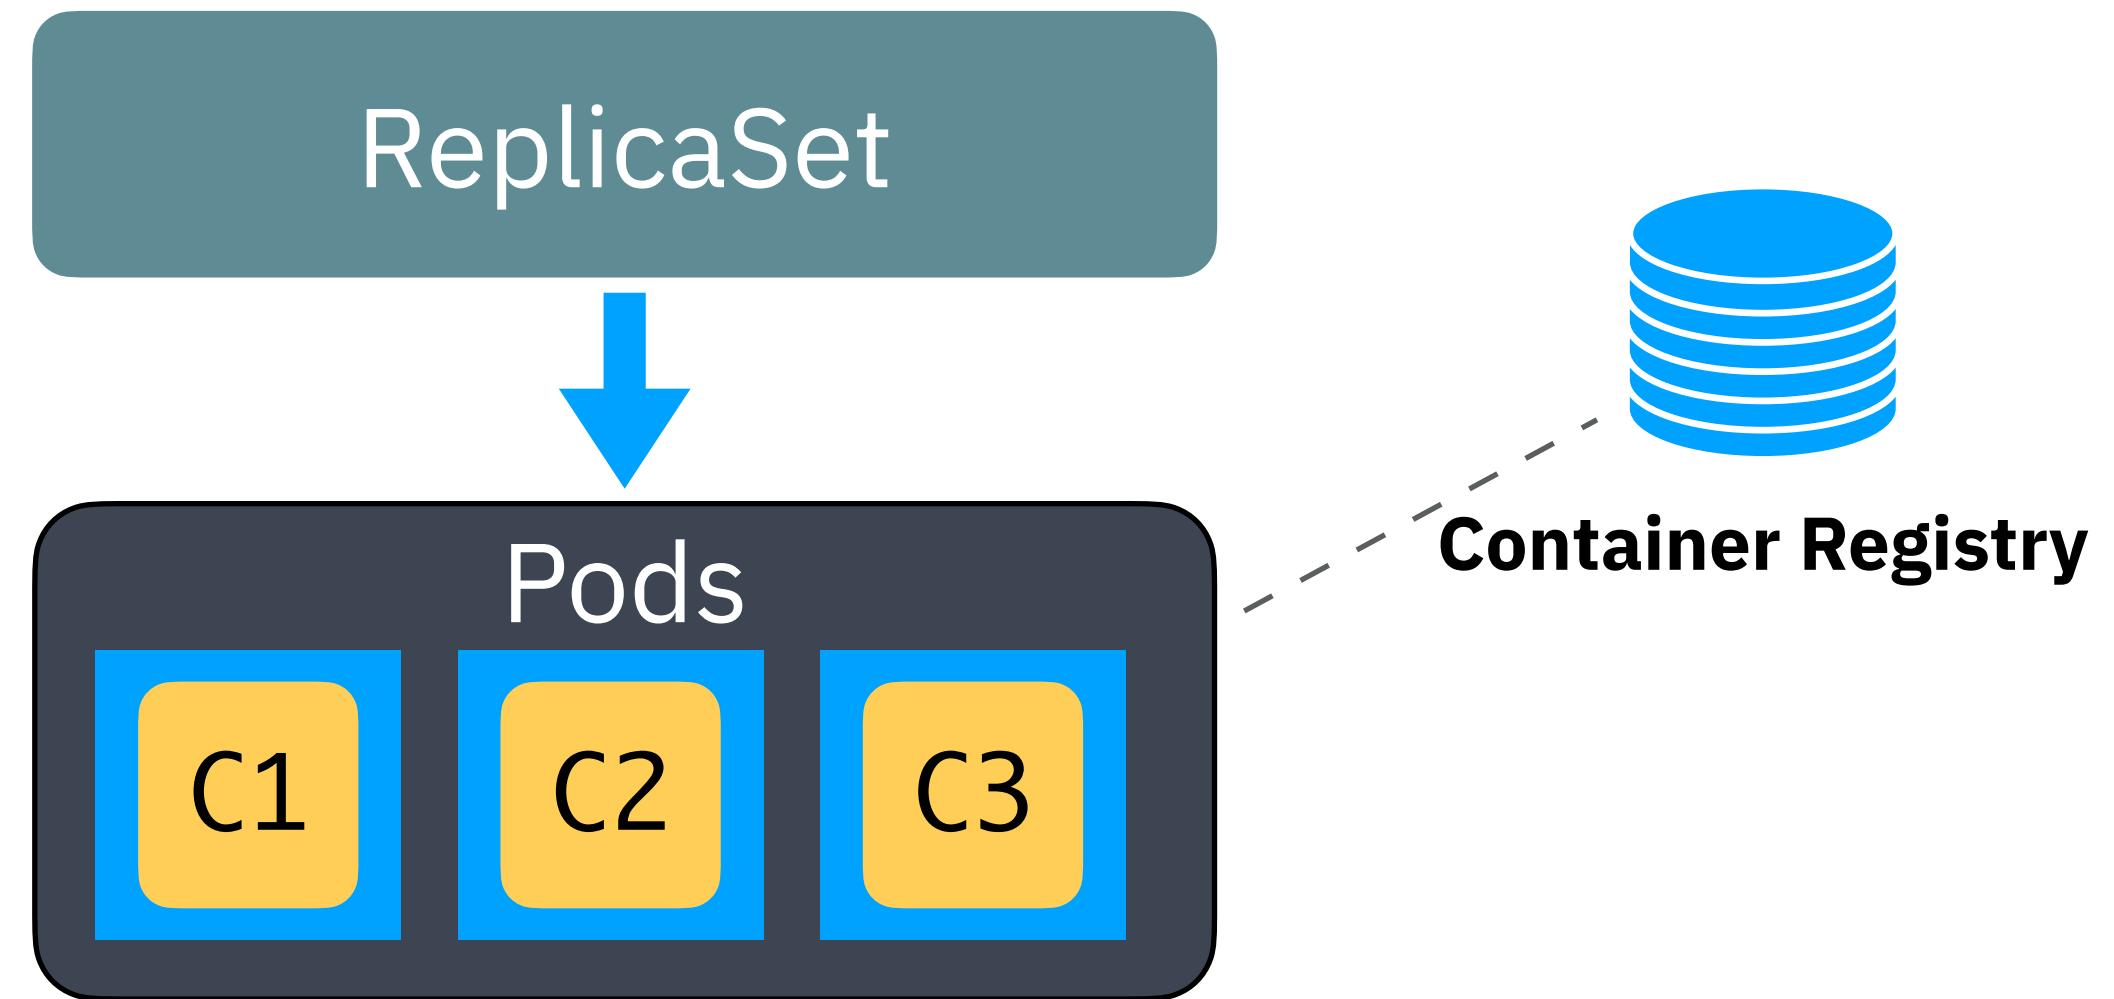
# Kubernetes Pods

- Pods
  - Containers run in Pods
  - The Pod is the smallest deployment possible
  - Containers are deployed from a container registry (public or private)



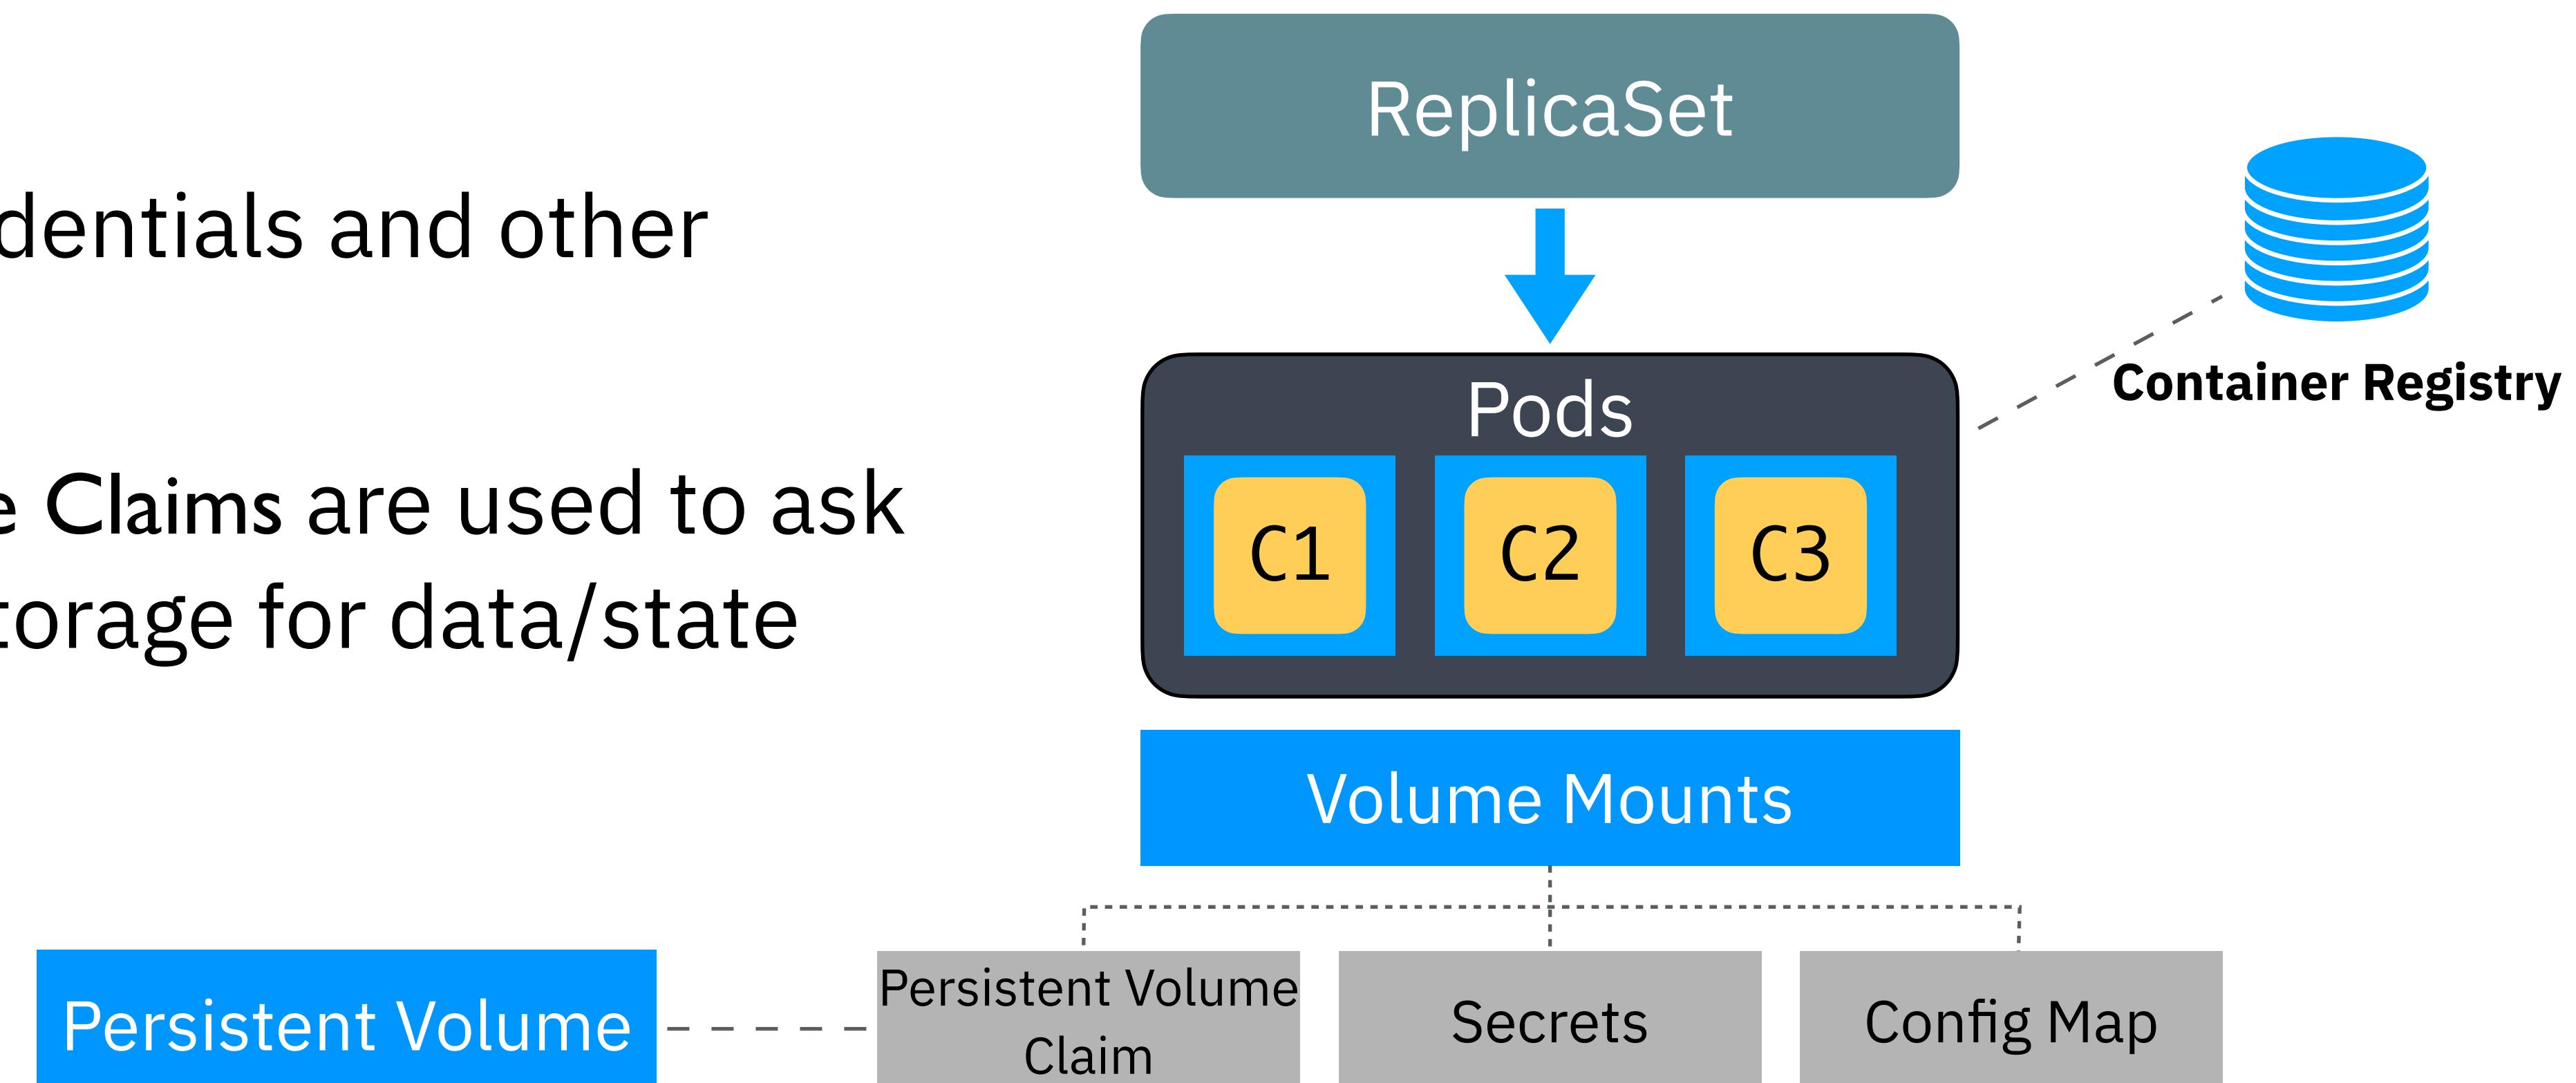
# Kubernetes ReplicaSets

- ReplicaSet
  - ReplicaSets allow multiple copies of containers to be deployed
  - Each one in its own Pod
  - If a container dies, the ReplicaSet will spawn a new one



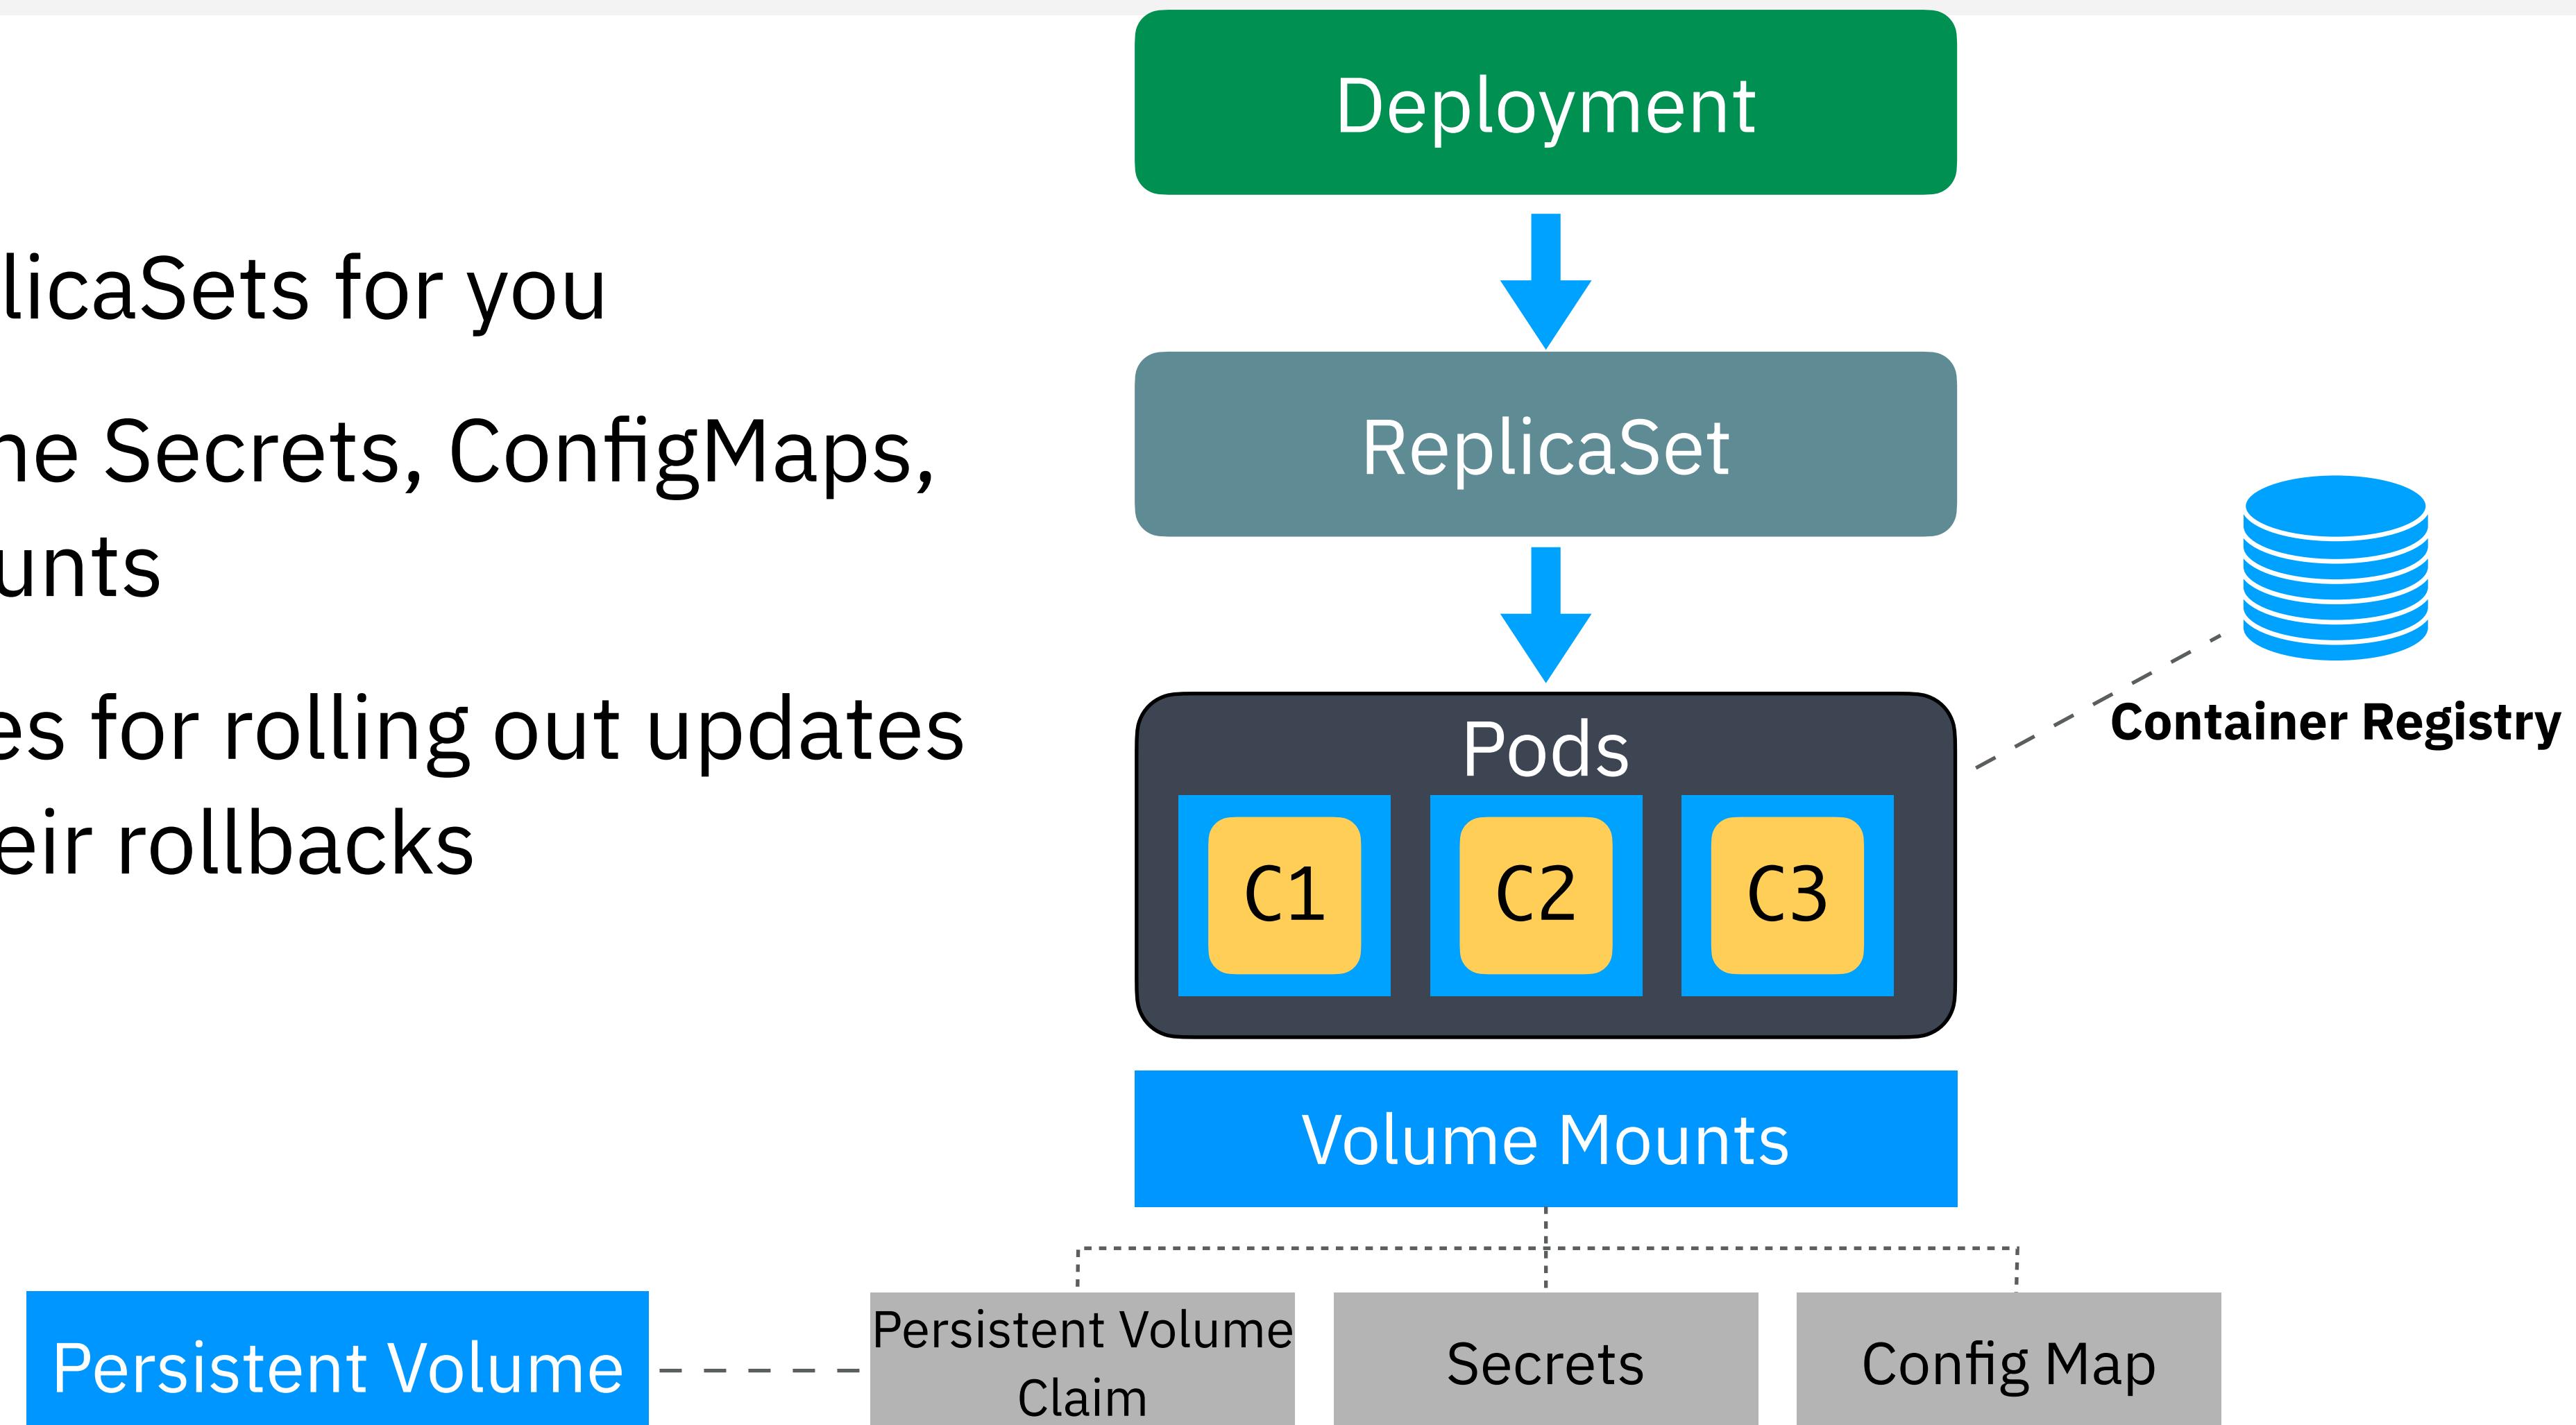
# Kubernetes Volume Mounts

- Volumes
  - ConfigMaps hold configuration parameters
  - Secrets hold credentials and other secrets
  - Persistent Volume Claims are used to ask for persistent storage for data/state



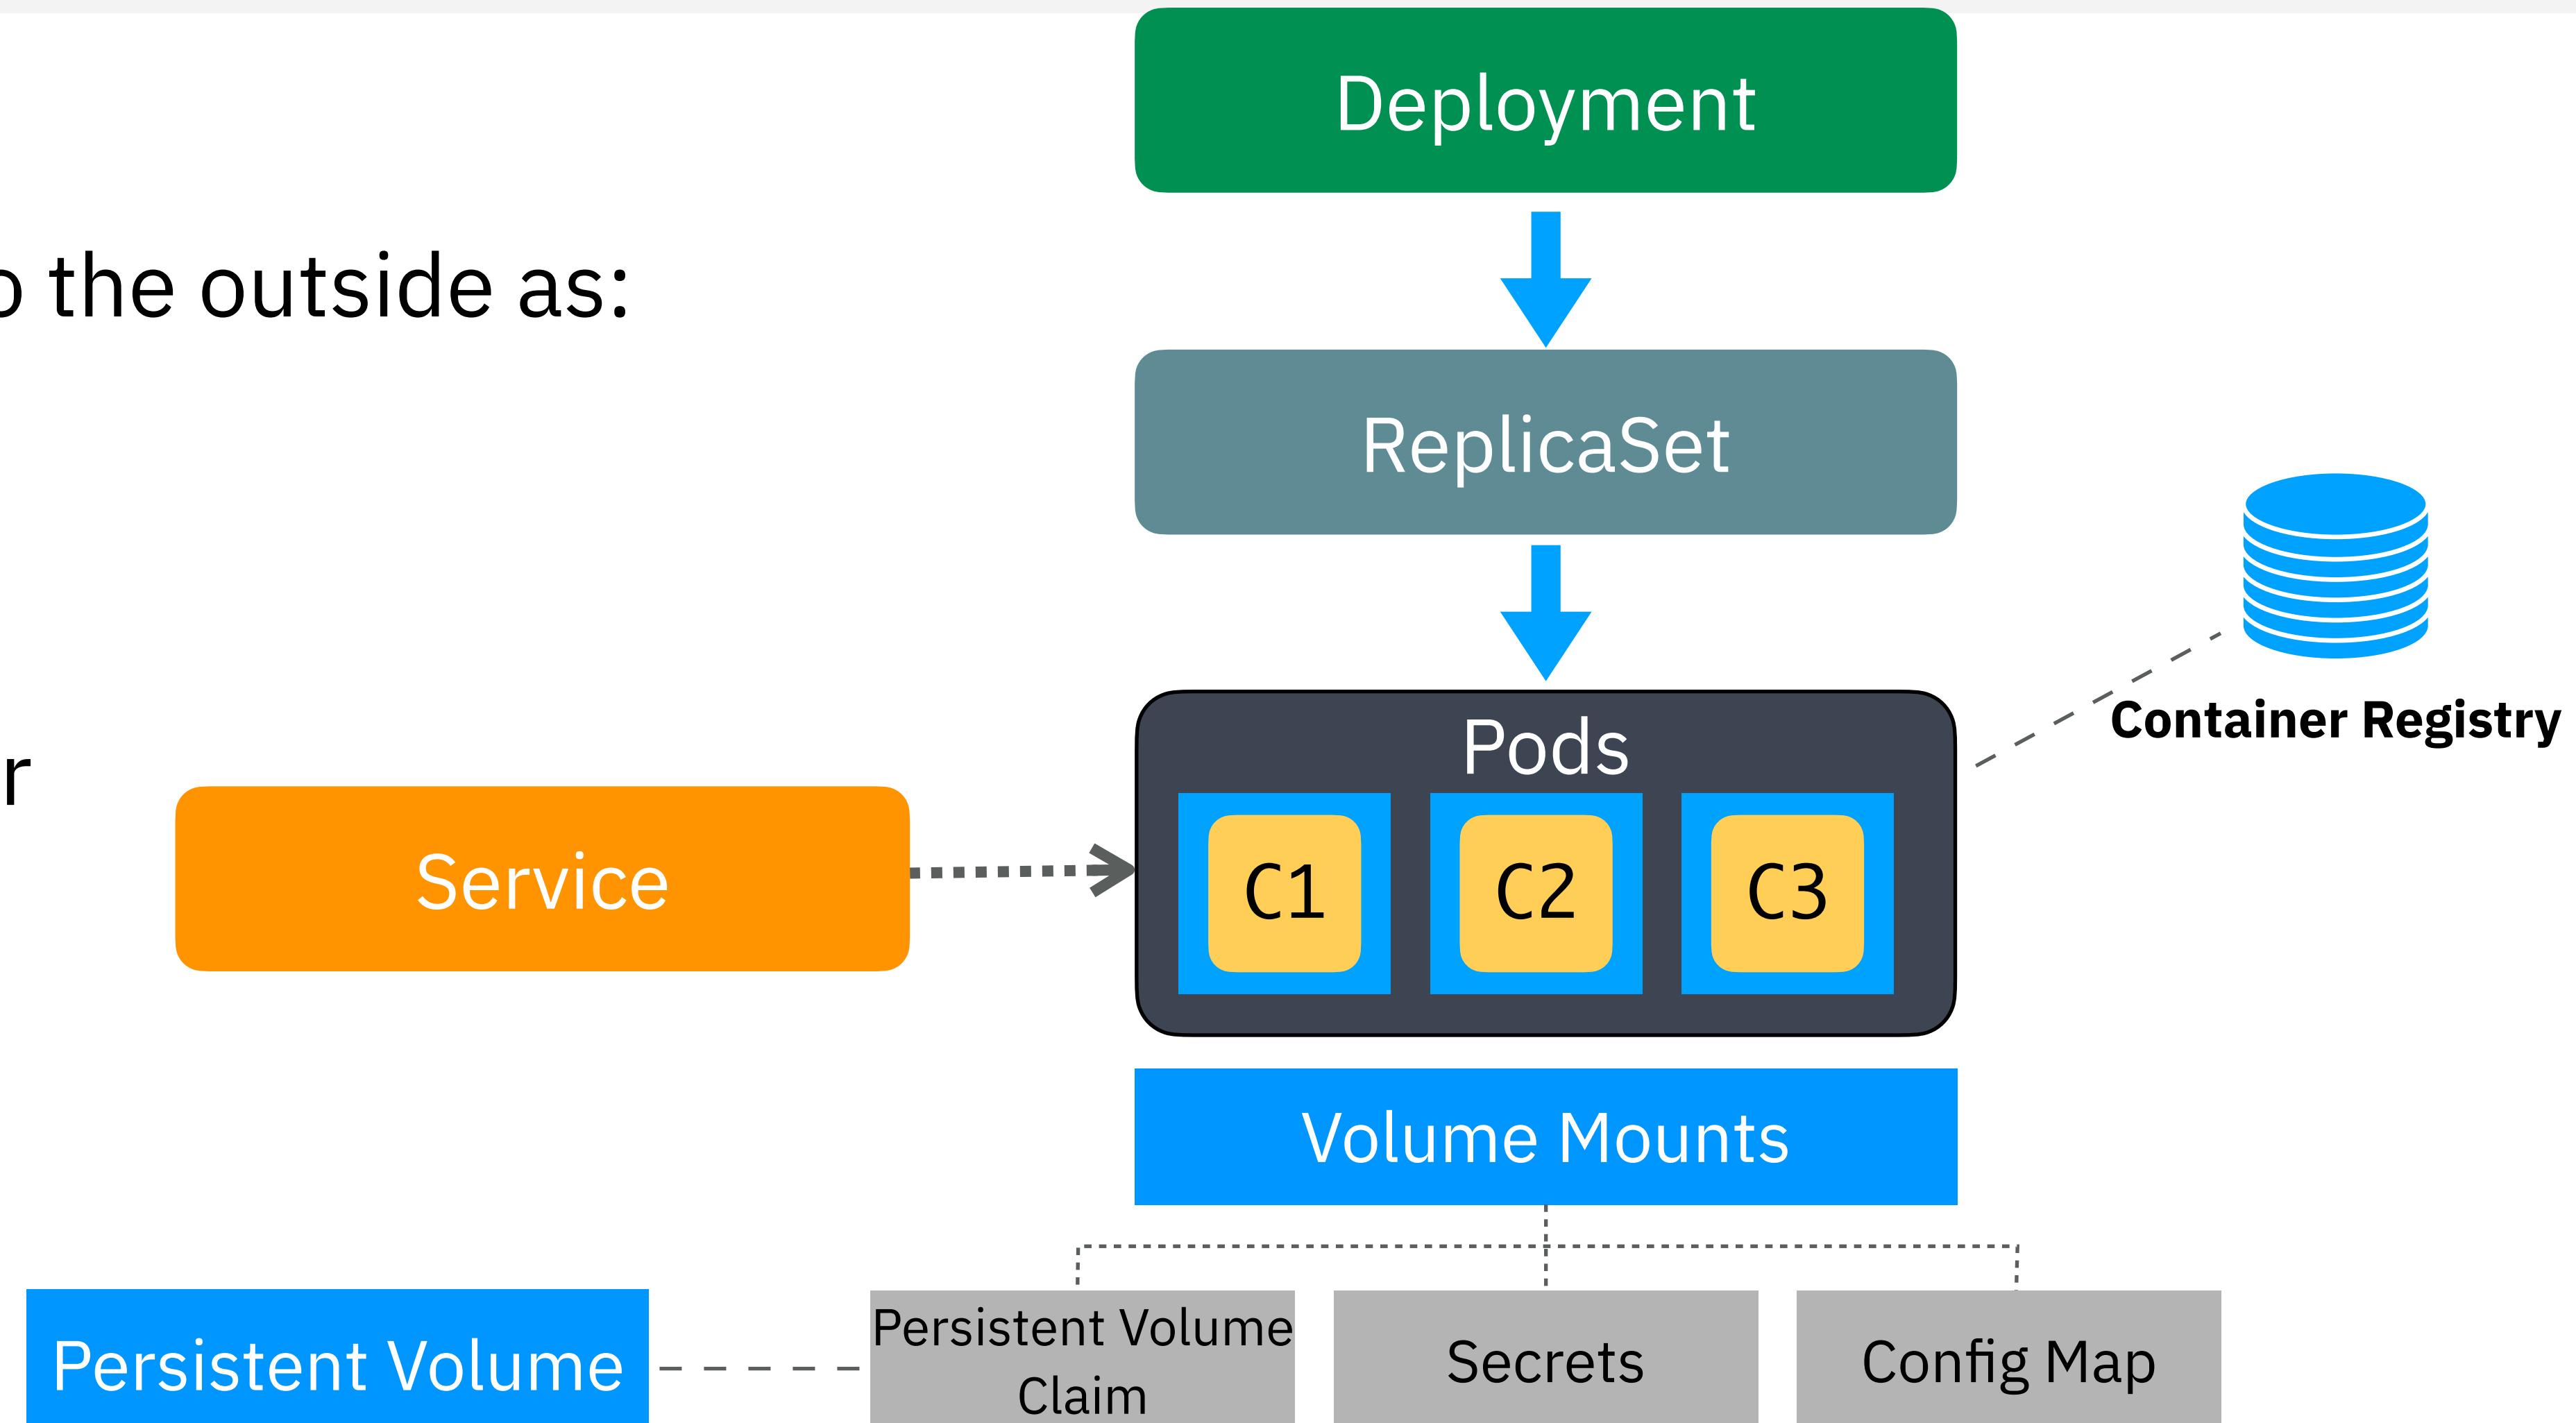
# Kubernetes Deployments

- Deployment
  - Sets up the ReplicaSets for you
  - Also specified the Secrets, ConfigMaps, and Volume Mounts
  - Provides features for rolling out updates and handling their rollbacks



# Kubernetes Service

- Service
  - Exposes Pods to the outside as:
    - ClusterIP
    - NodePort
    - Load Balancer

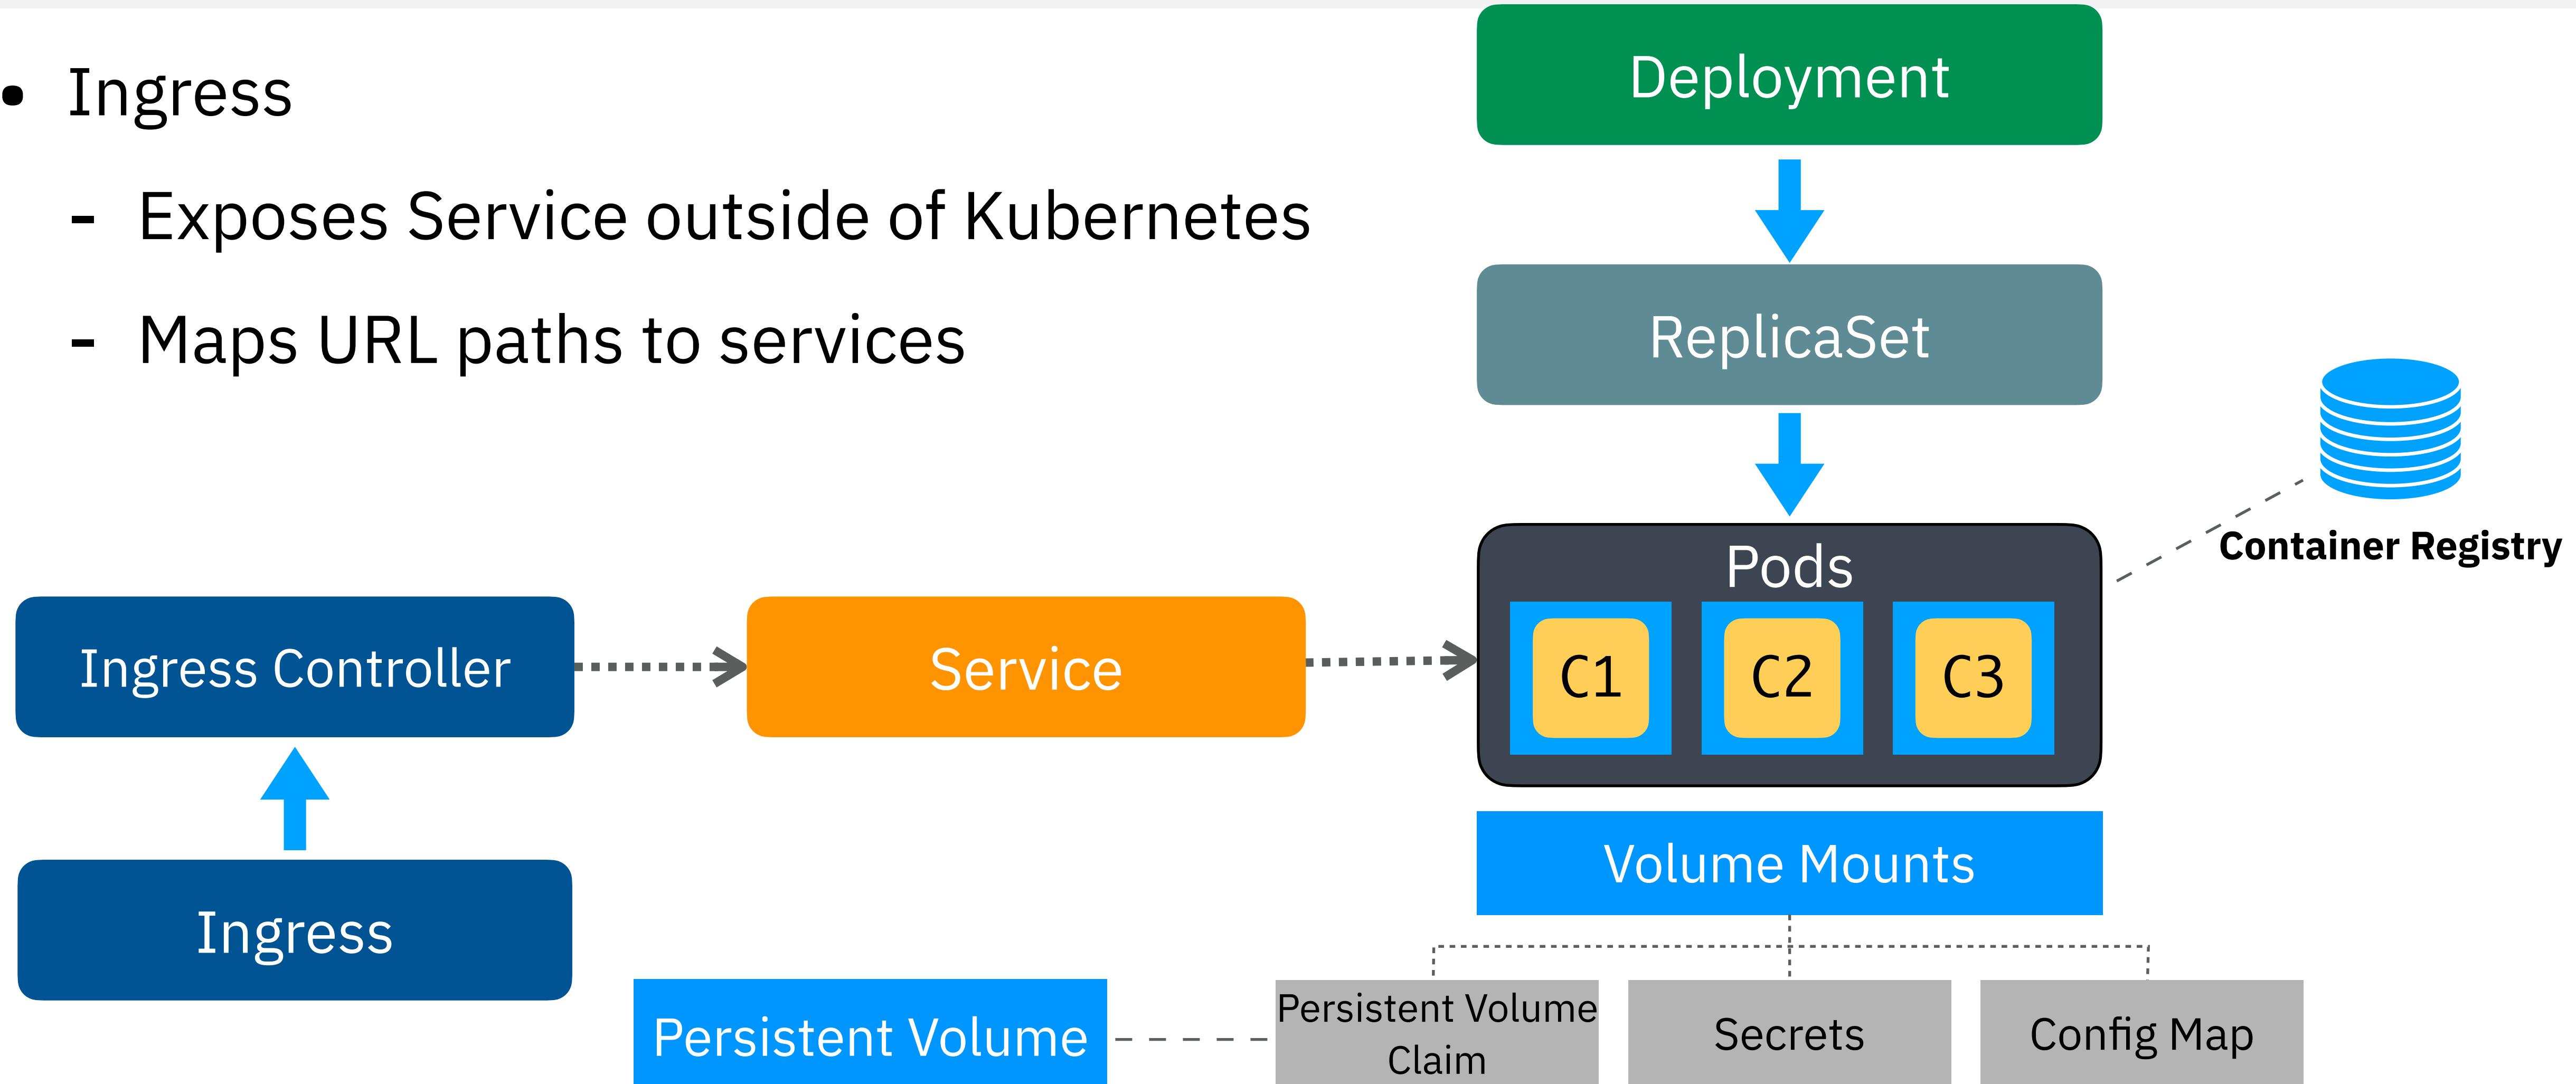


# Types of Services

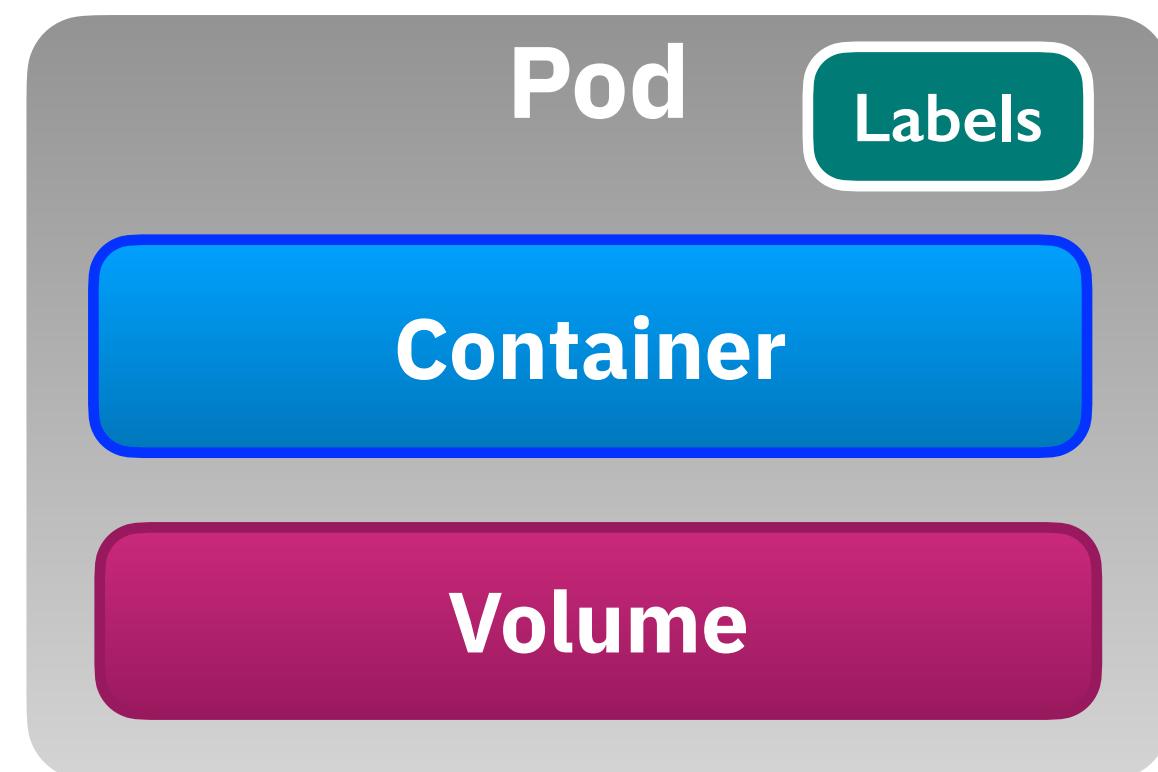
- **LoadBalancer** – Only available via cloud providers. Front end for service that balances the load across multiple backends from a single IP address
- **NodePort** – Exposes the service as an arbitrary port on every worker node in the cluster
- **ClusterIP** – Service is only accessible from other services within the cluster (no external exposure)

# Kubernetes Ingress Controller

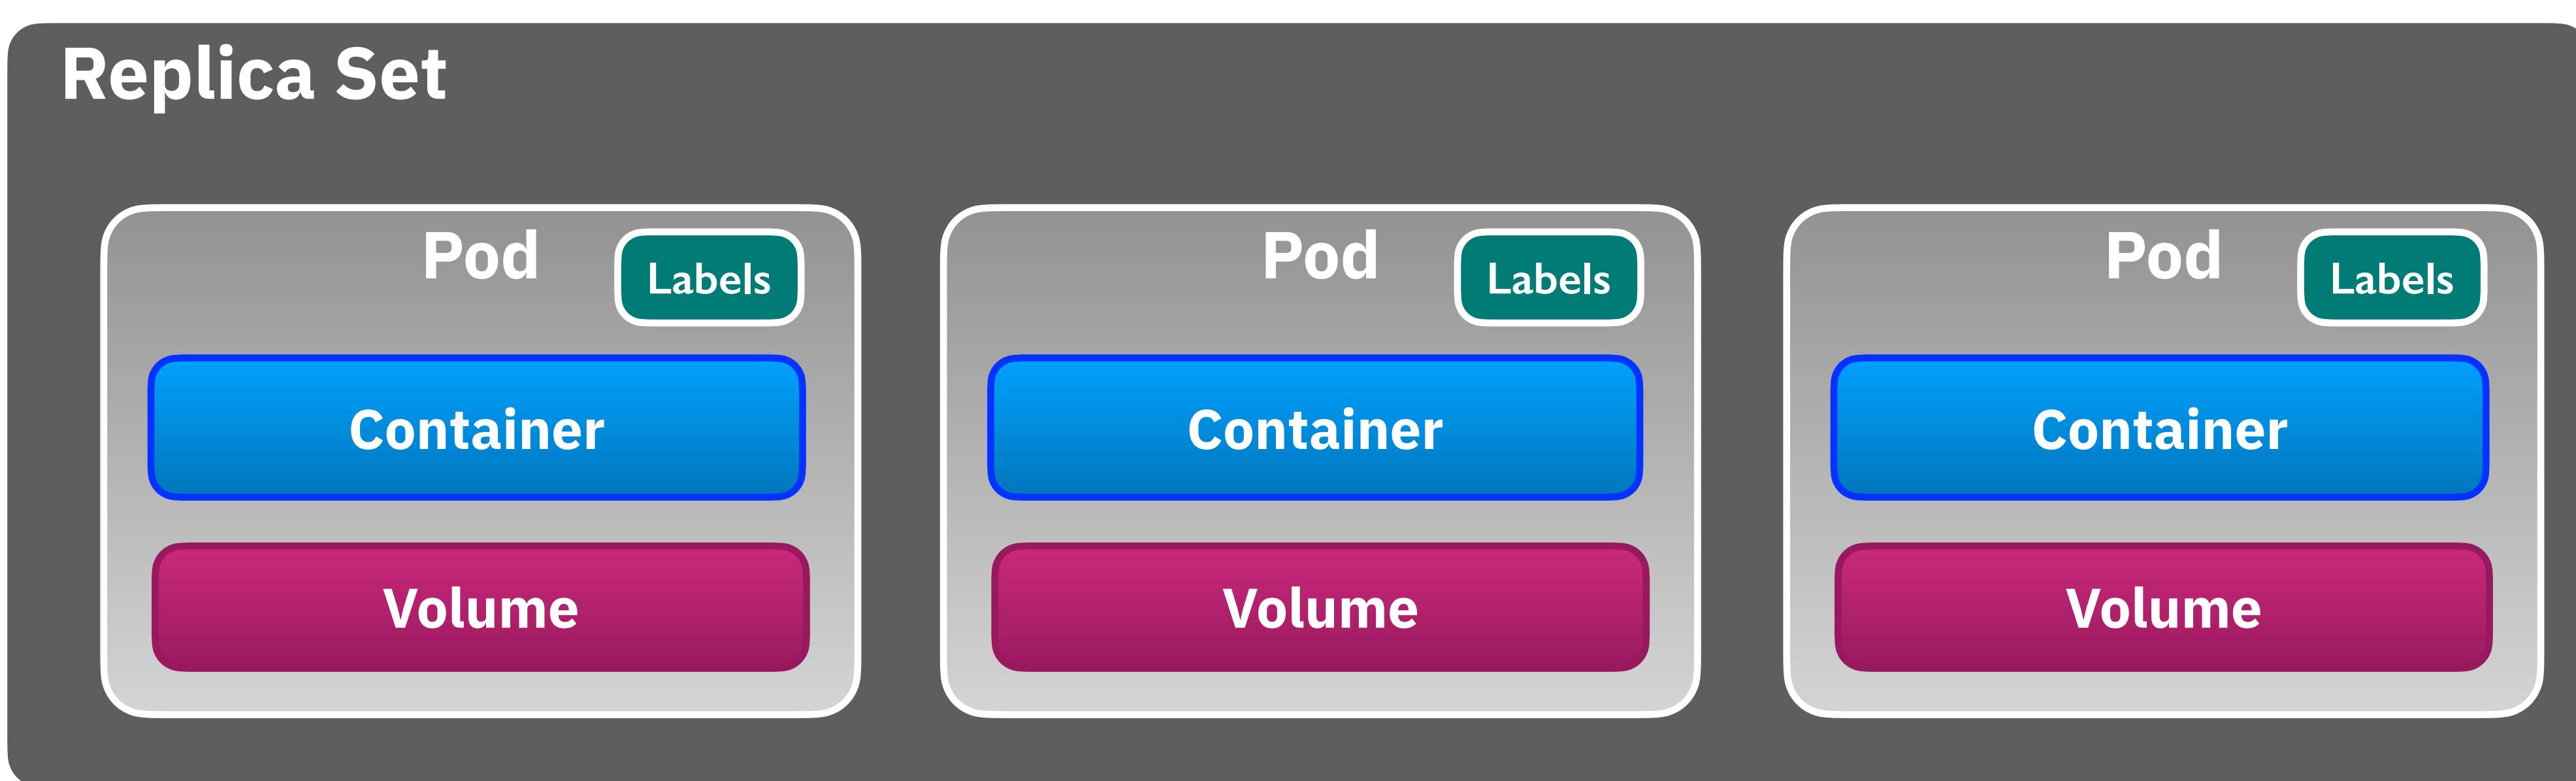
- Ingress
  - Exposes Service outside of Kubernetes
  - Maps URL paths to services



# Services Map to Pods via Labels



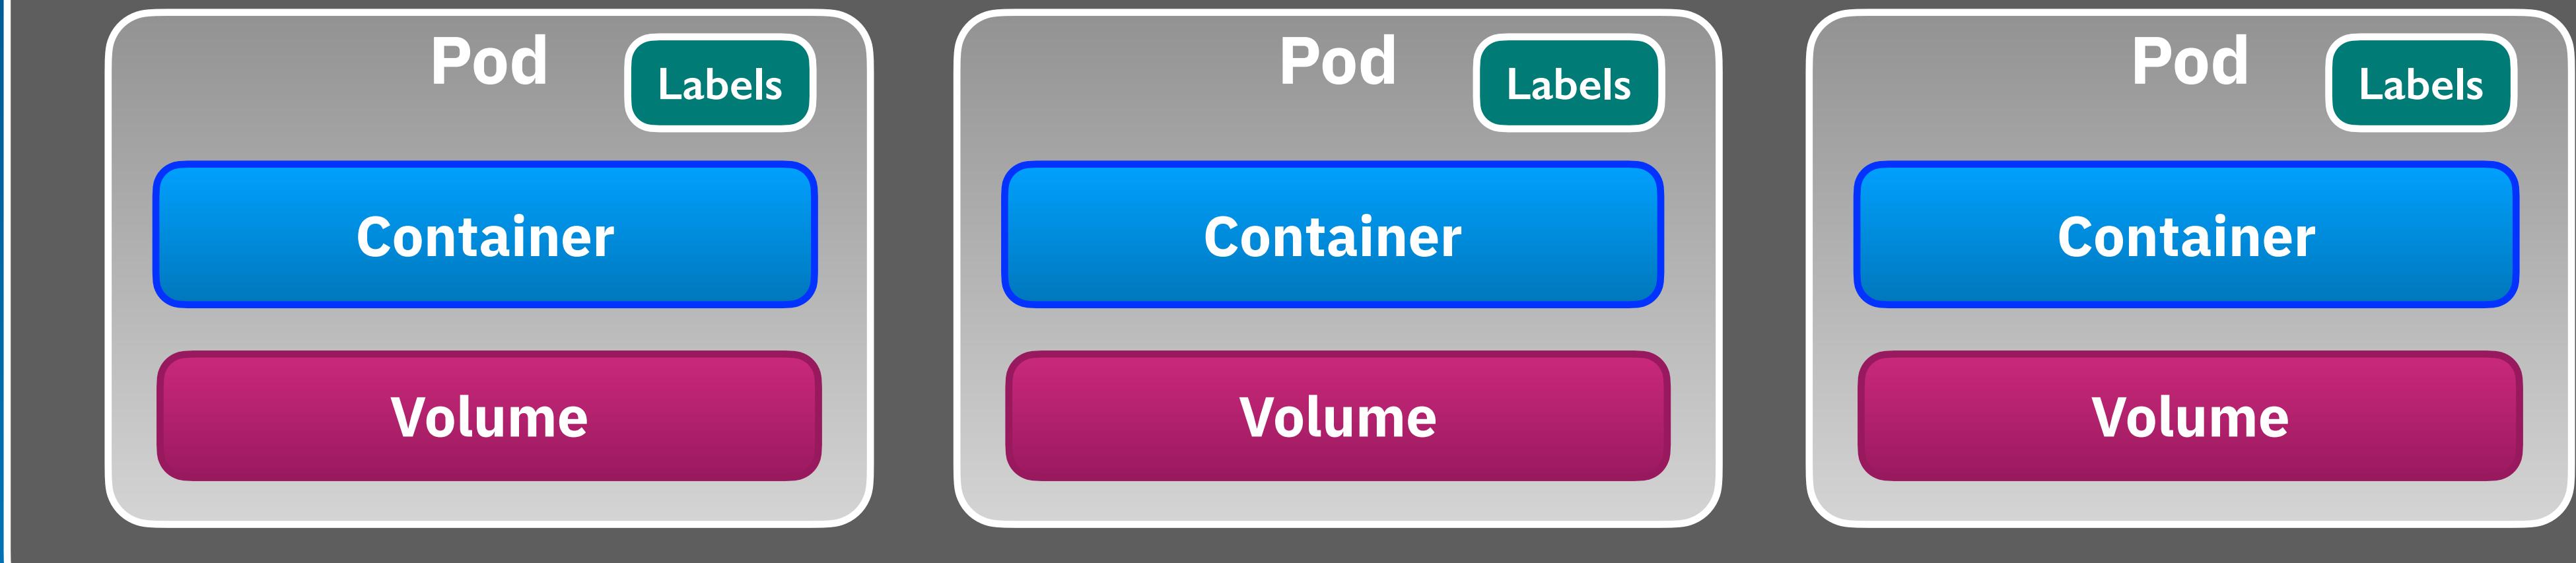
# Services Map to Pods via Labels



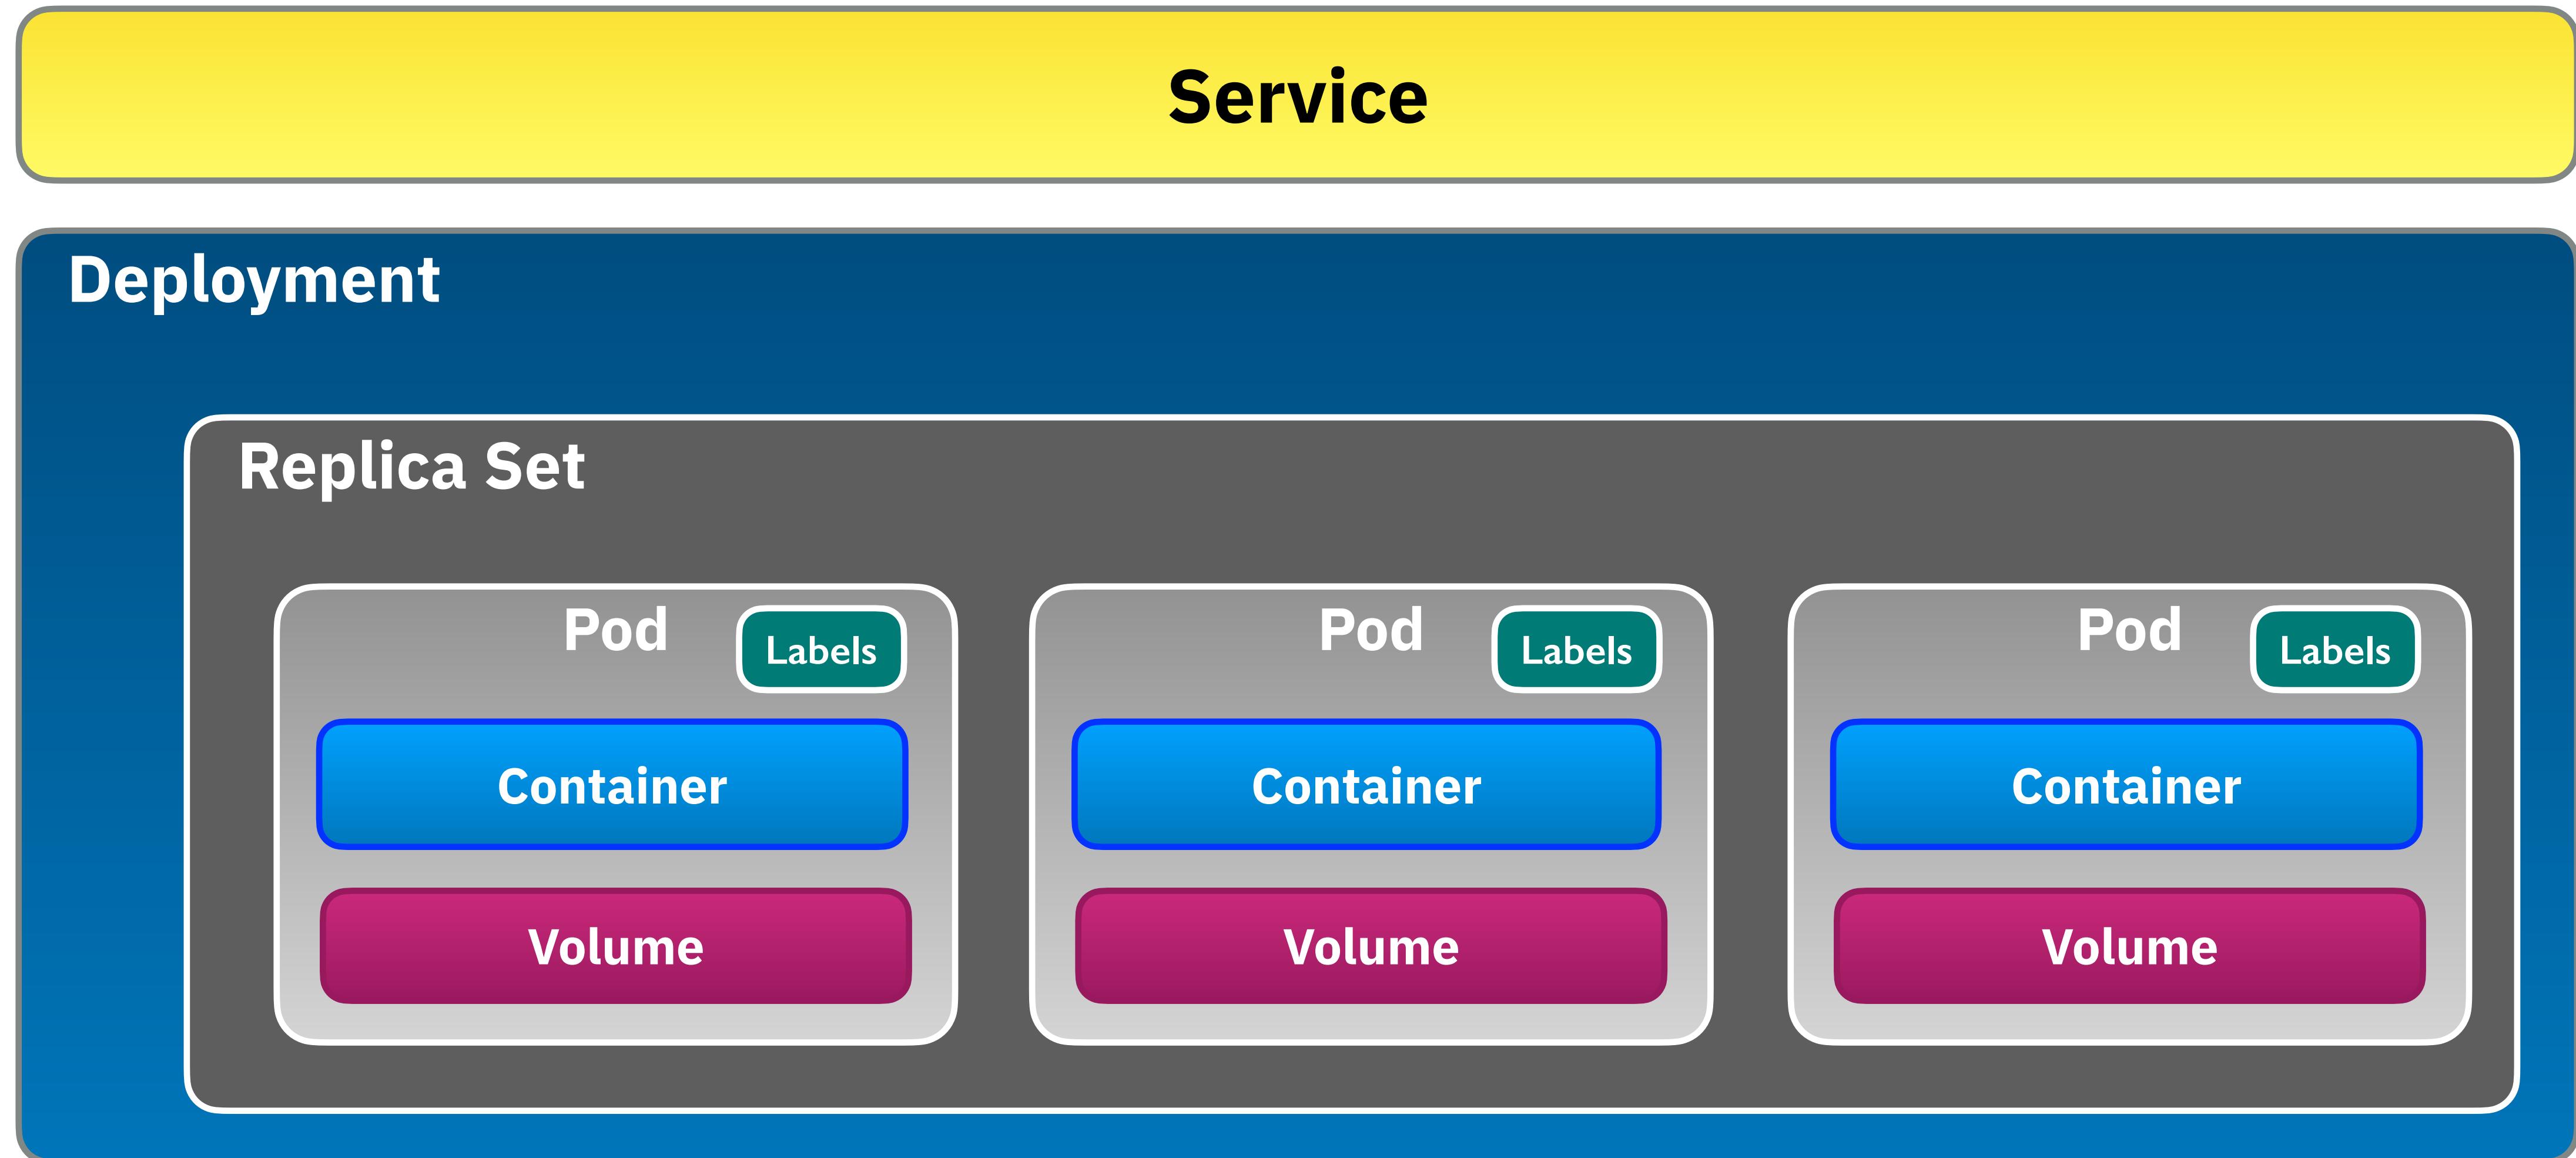
# Services Map to Pods via Labels

## Deployment

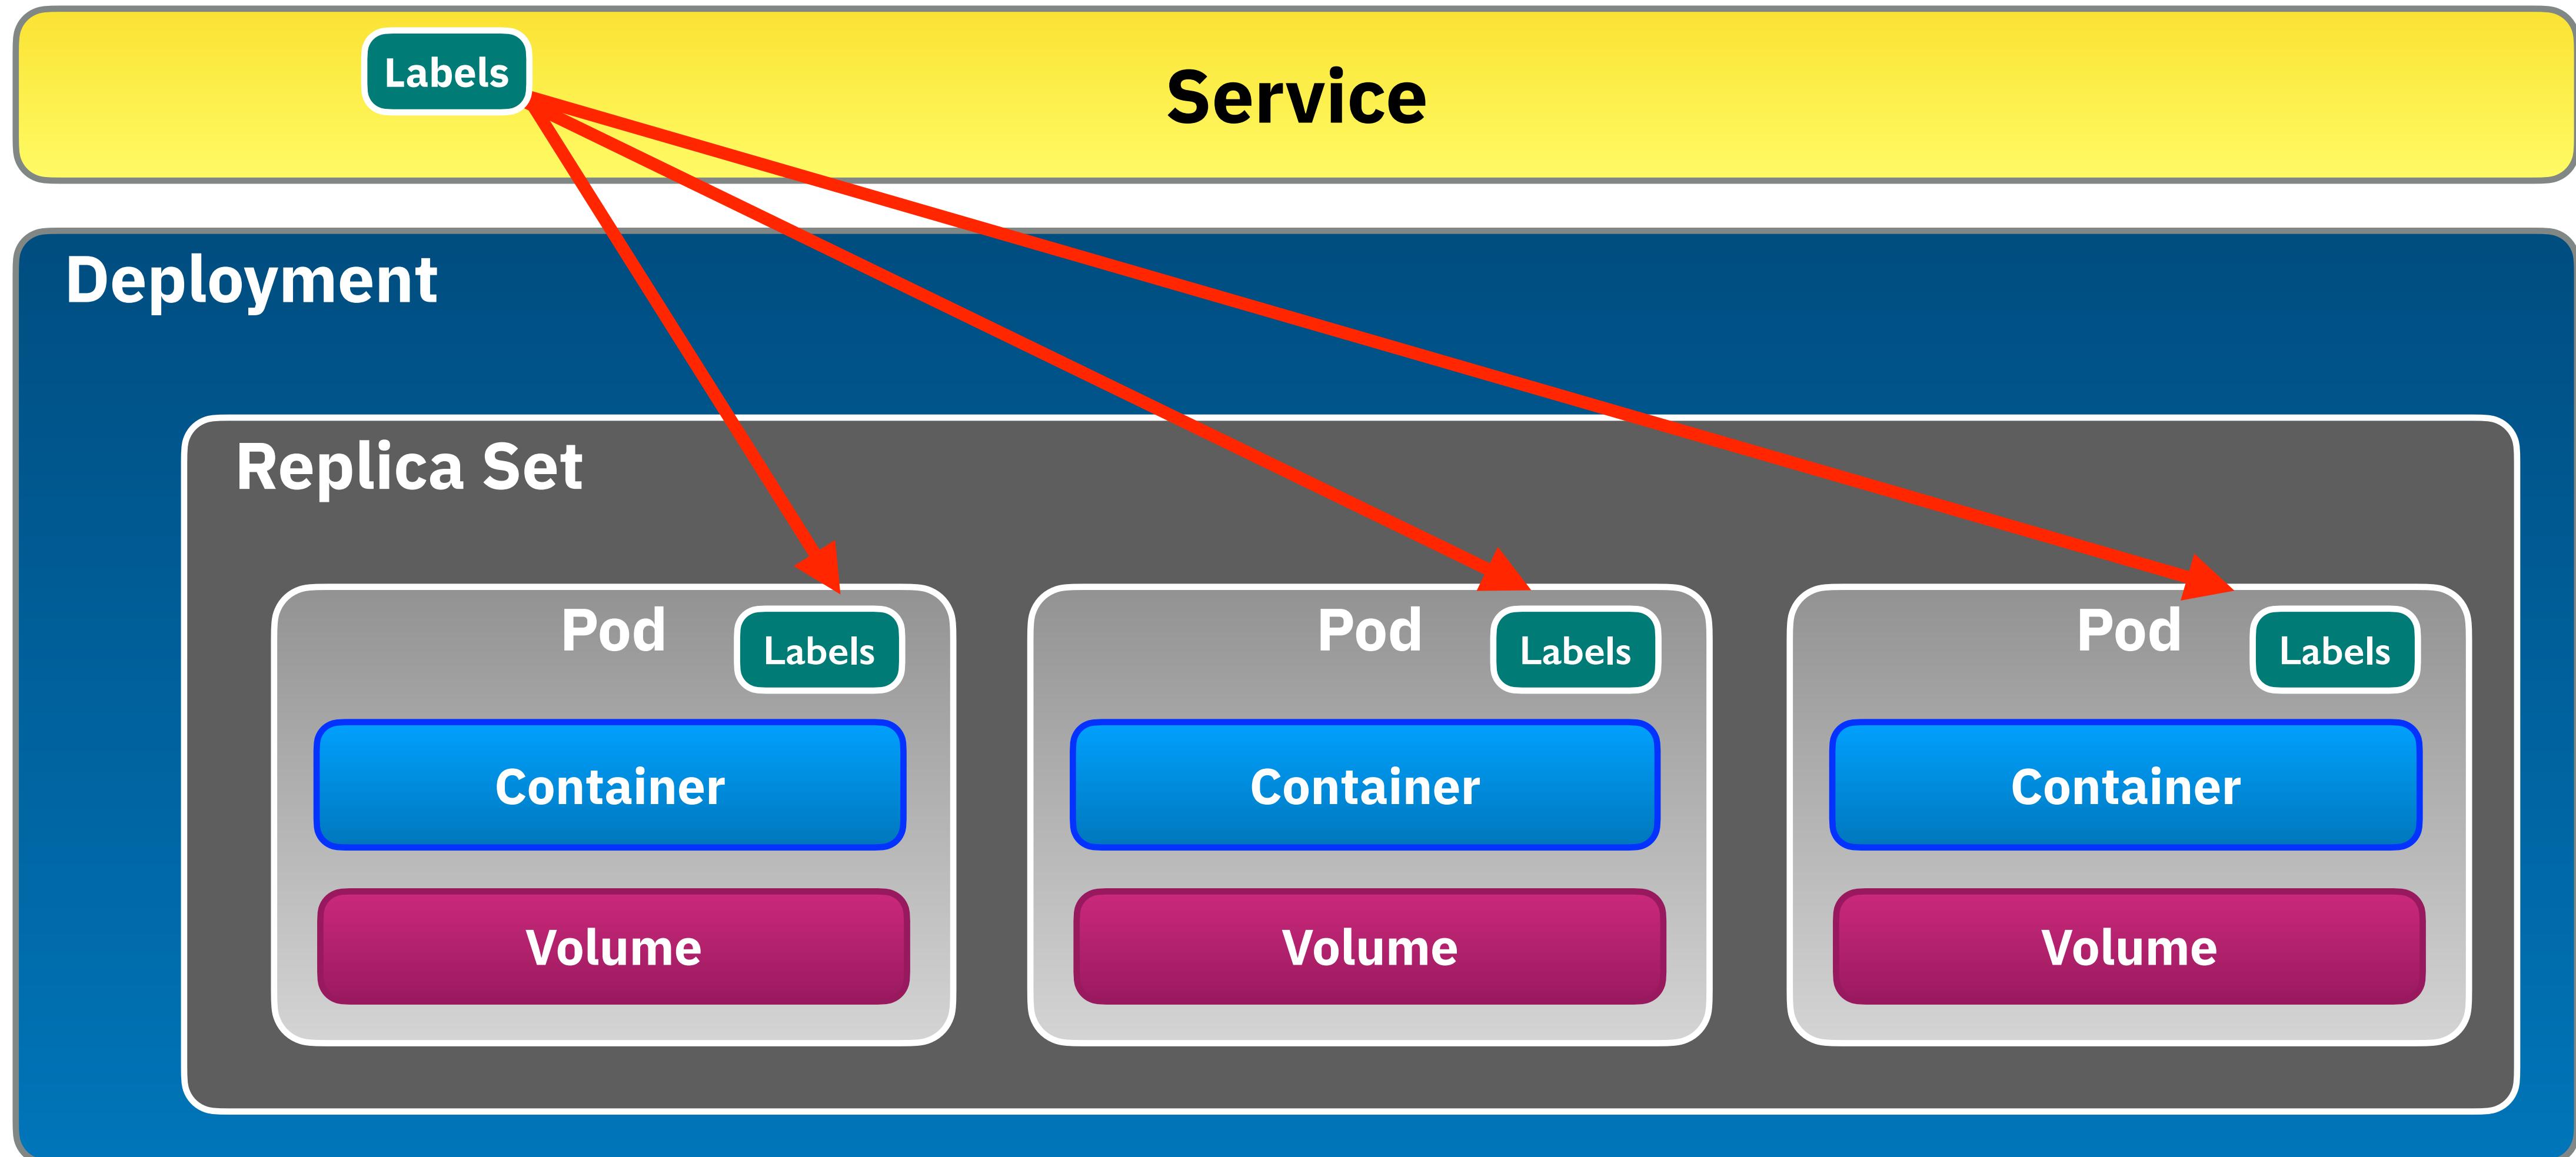
### Replica Set



# Services Map to Pods via Labels



# Services Map to Pods via Labels



# Use Case for Multiple Containers in a Pod

- Pods can be used to host vertically integrated application stacks (e.g. LAMP), but their primary motivation is to support co-located, co-managed helper programs, such as:
  - Content management systems, file and data loaders, local cache managers, etc.
  - Log and checkpoint backup, compression, rotation, snapshotting, etc.
  - Data change watchers, log tailers, logging and monitoring adapters, event publishers, etc.
  - Proxies, bridges, and adapters
  - Controllers, managers, configurators, and updaters

# Deployment using kubectl

- A containerized application can be deployed on Kubernetes using a deployment definition by executing a simple CLI command as follows:

```
$ kubectl create deployment <application-name> --image=<container-image>
```

# Create a Service

- You can expose a deployment as a service using `kubectl expose`:

```
$ kubectl expose deployment <application-name> --type=<service-type> \
--port=<port_number>
```

Valid <service-type>s are:

- LoadBalancer
- NodePort
- ClusterIP

# Deploy using YAML

- You can create a deployment and a service using yaml files

```
$ kubectl create -f deployment.yaml  
$ kubectl create -f service.yaml
```

- Update them

```
$ kubectl apply -f deployment.yaml  
$ kubectl apply -f service.yaml
```

- And just as easily delete them

```
$ kubectl delete -f deployment.yaml  
$ kubectl delete -f service.yaml
```

# Deployment YAML

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: hitcounter
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: hitcounter
10  template:
11    metadata:
12      labels:
13        app: hitcounter
14  spec:
15    containers:
16      - image: hitcounter:1.0
17        imagePullPolicy: IfNotPresent
18        name: hitcounter
19        ports:
20          - containerPort: 8080
21            protocol: TCP
22        restartPolicy: Always
```

# Service YAML

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: hitcounter-service
5 spec:
6   type: NodePort
7   selector:
8     app: hitcounter
9   ports:
10    - name: primary
11      protocol: TCP
12      port: 8080
```

~

# Linking Service to Pods

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: hitcounter-service
5 spec:
6   type: NodePort
7   selector:
8     app: hitcounter
9   ports:
10    - name: primary
11      protocol: TCP
12      port: 8080
```

~

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: hitcounter
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: hitcounter
10  template:
11    metadata:
12      labels:
13        app: hitcounter
14  spec:
15    containers:
16      - image: hitcounter:1.0
17        imagePullPolicy: IfNotPresent
18        name: hitcounter
19        ports:
20          - containerPort: 8080
21            protocol: TCP
22        restartPolicy: Always
```

# Linking Service to Pods

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: hitcounter-service
5 spec:
6   type: NodePort
7   selector:
8     app: hitcounter
9   ports:
10    - name: primary
11      protocol: TCP
12      port: 8080
~
```

Look for Pods with this Label

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: hitcounter
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: hitcounter
10  template:
11    metadata:
12      labels:
13        app: hitcounter
14  spec:
15    containers:
16      - image: hitcounter:1.0
17        imagePullPolicy: IfNotPresent
18        name: hitcounter
19        ports:
20          - containerPort: 8080
21            protocol: TCP
22        restartPolicy: Always
```

# Linking Service to Pods

Will match this Pod

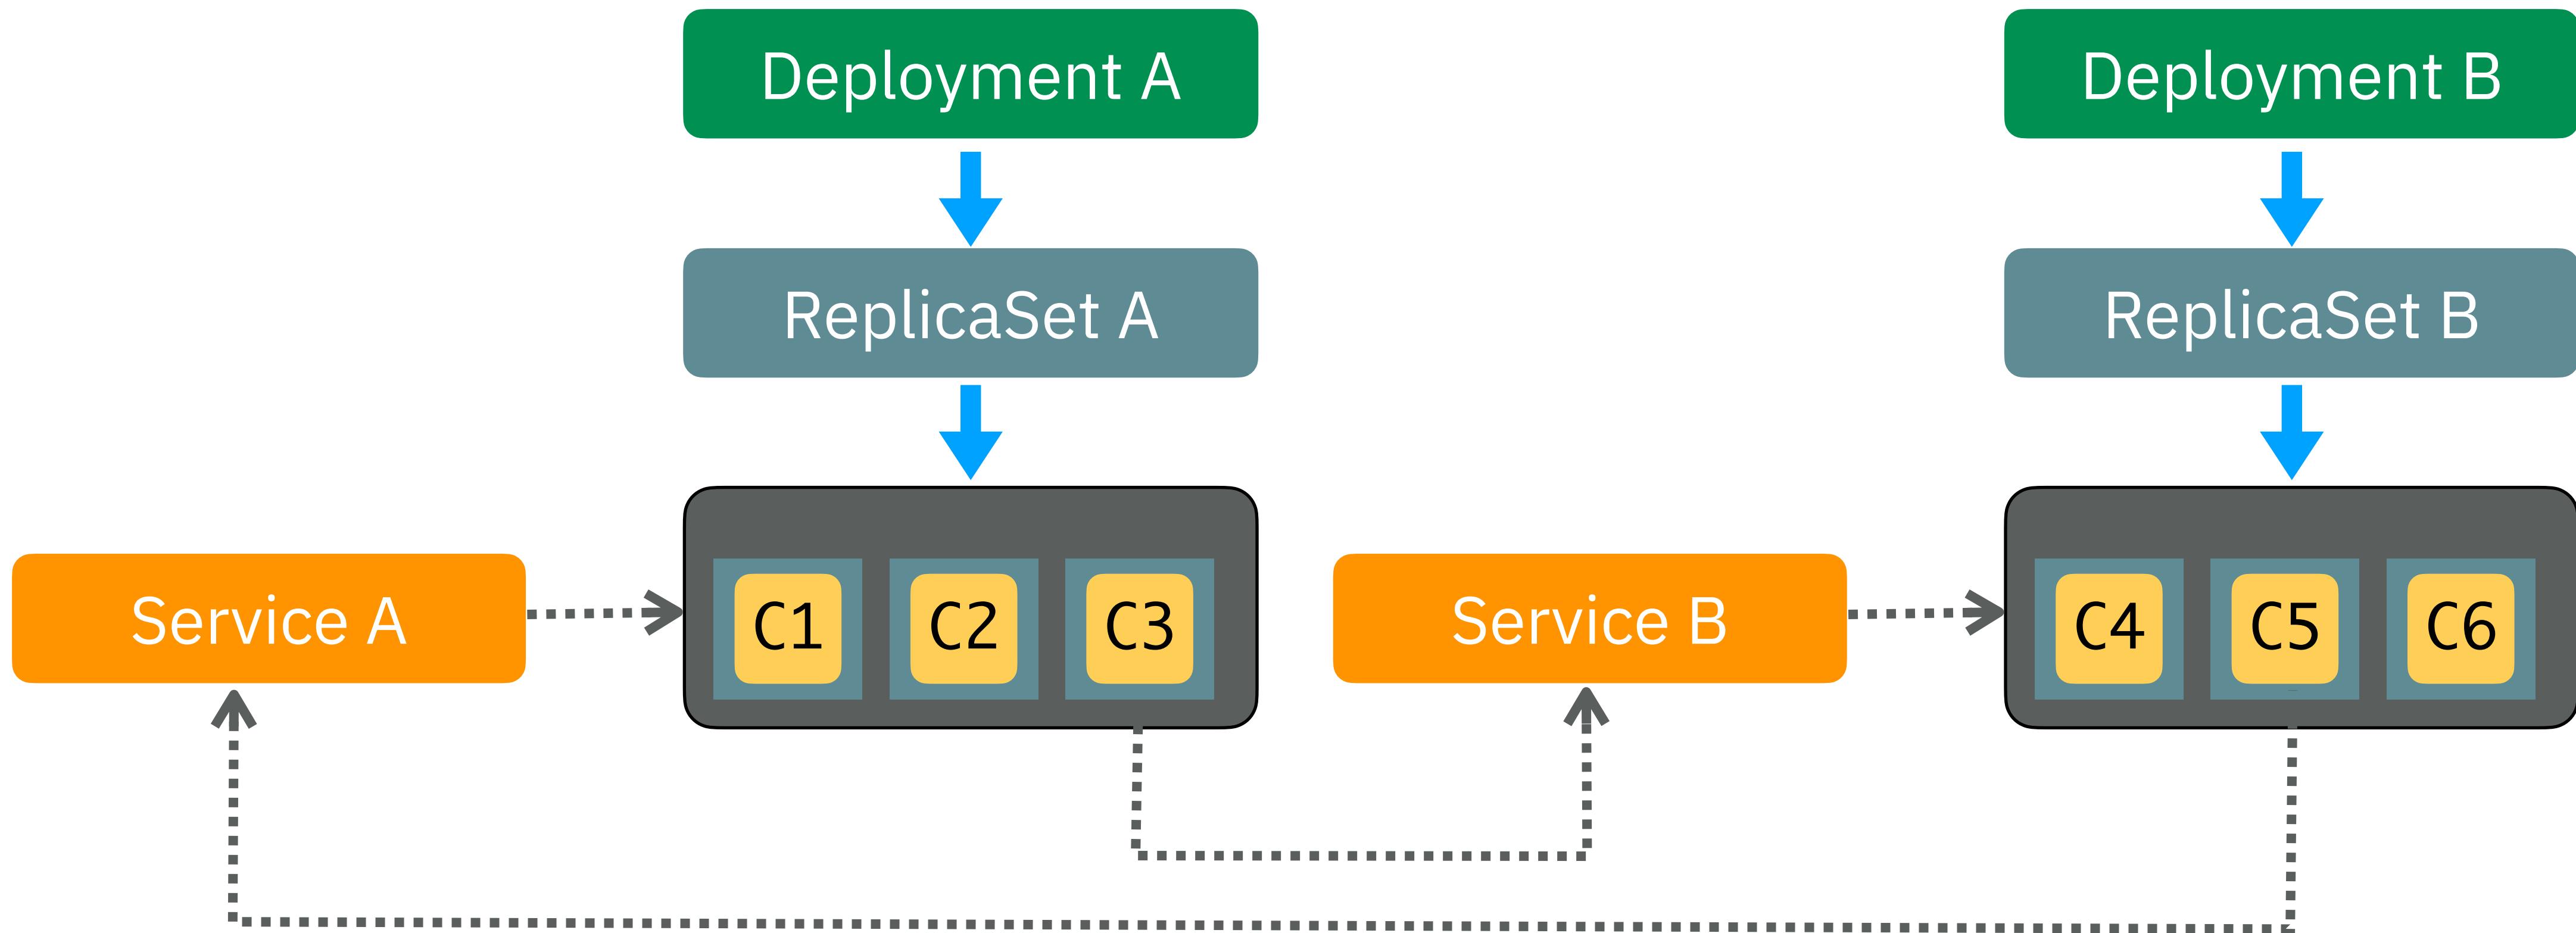
```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: hitcounter-service
5 spec:
6   type: NodePort
7   selector:
8     app: hitcounter
9   ports:
10    - name: primary
11      protocol: TCP
12      port: 8080
~
```

Look for Pods with this Label

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: hitcounter
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: hitcounter
10  template:
11    metadata:
12      labels:
13        app: hitcounter
14  spec:
15    containers:
16      - image: hitcounter:1.0
17        imagePullPolicy: IfNotPresent
18        name: hitcounter
19        ports:
20          - containerPort: 8080
21            protocol: TCP
22        restartPolicy: Always
```

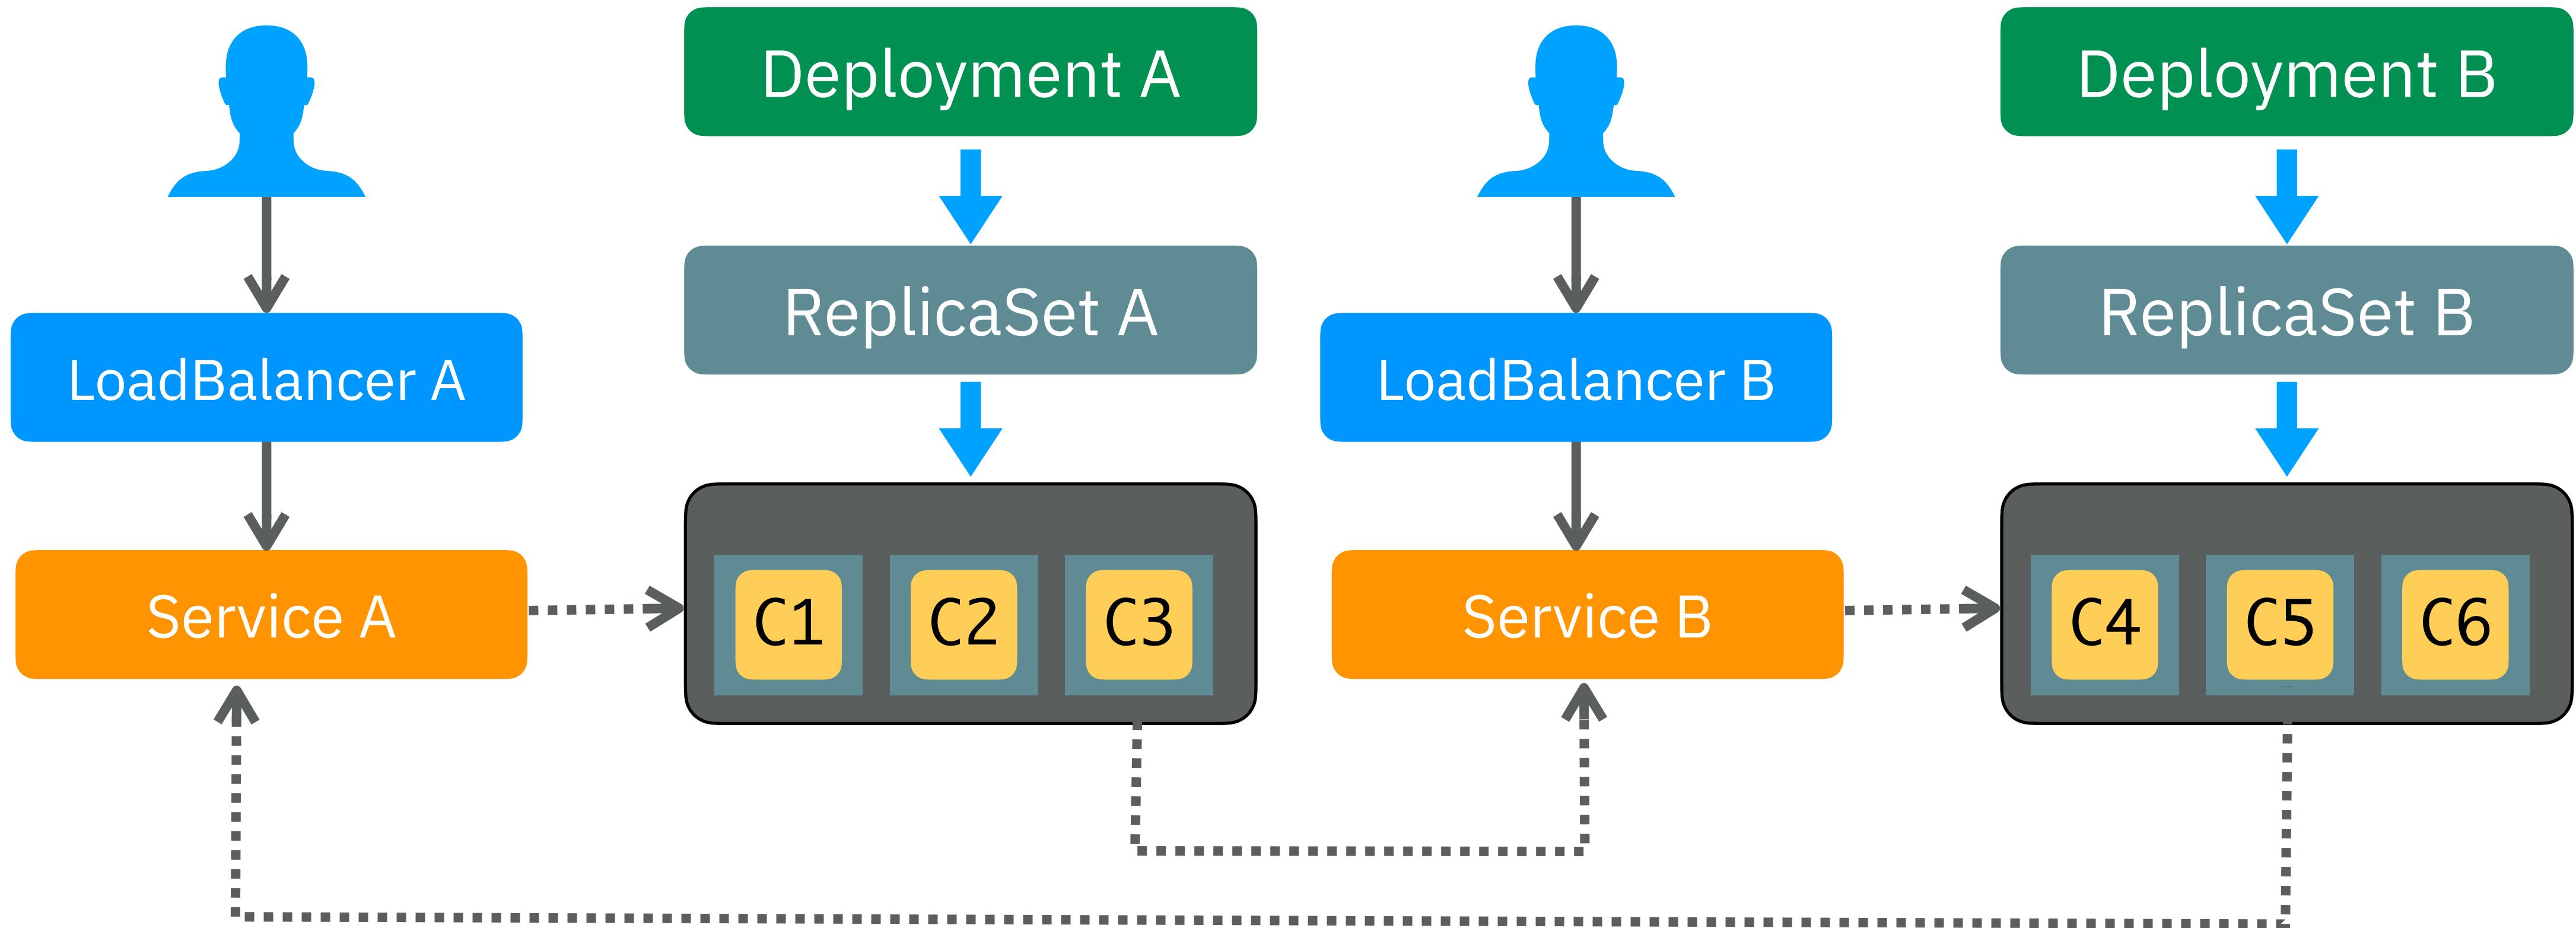
# Service Discovery

- Kubernetes provides internal routing so services can find each other



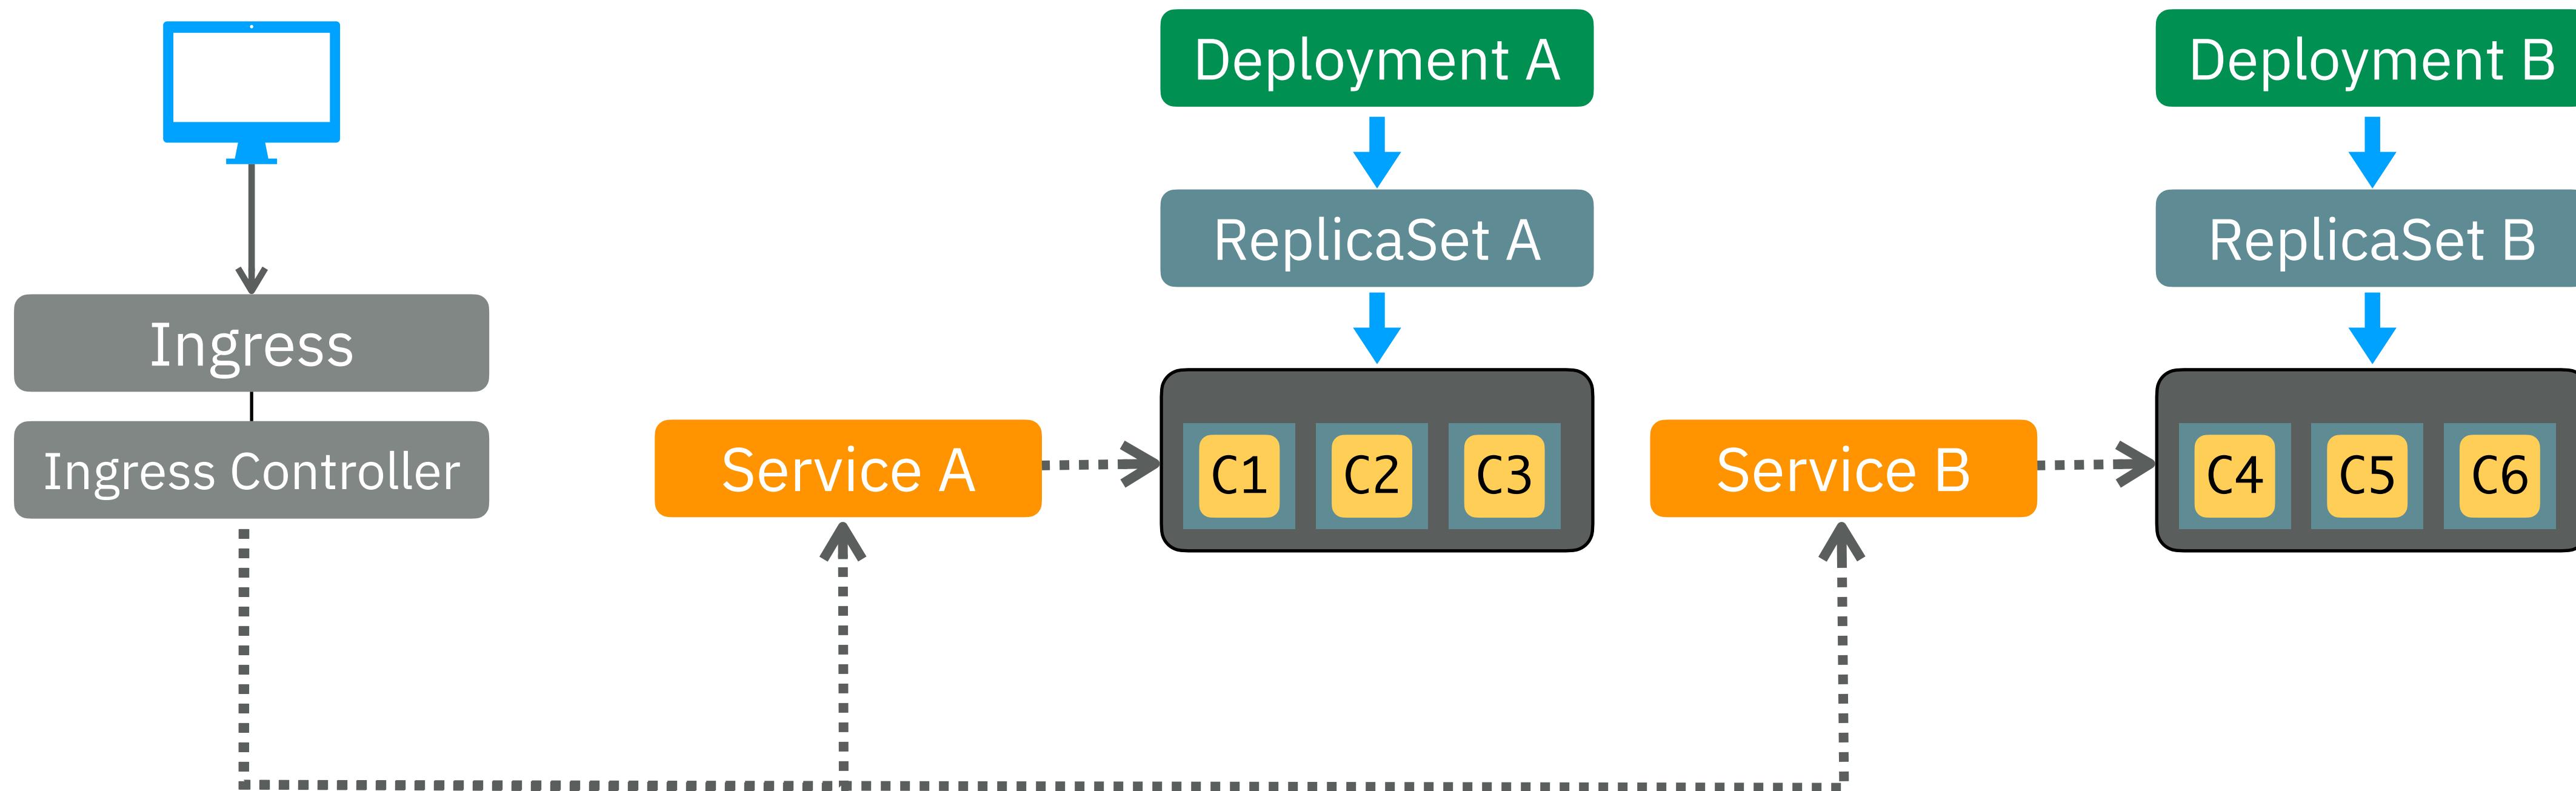
# External Service Access

- LoadBalancers provide external access for a single service



# External Routing

- Kubernetes provides an Ingress Controller to allow external network access or you can use NodePorts



# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.containers.mybluemix.net
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.containers.mybluemix.net
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Requests coming in on this URL

# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.containers.mybluemix.net
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Based on these PATHs

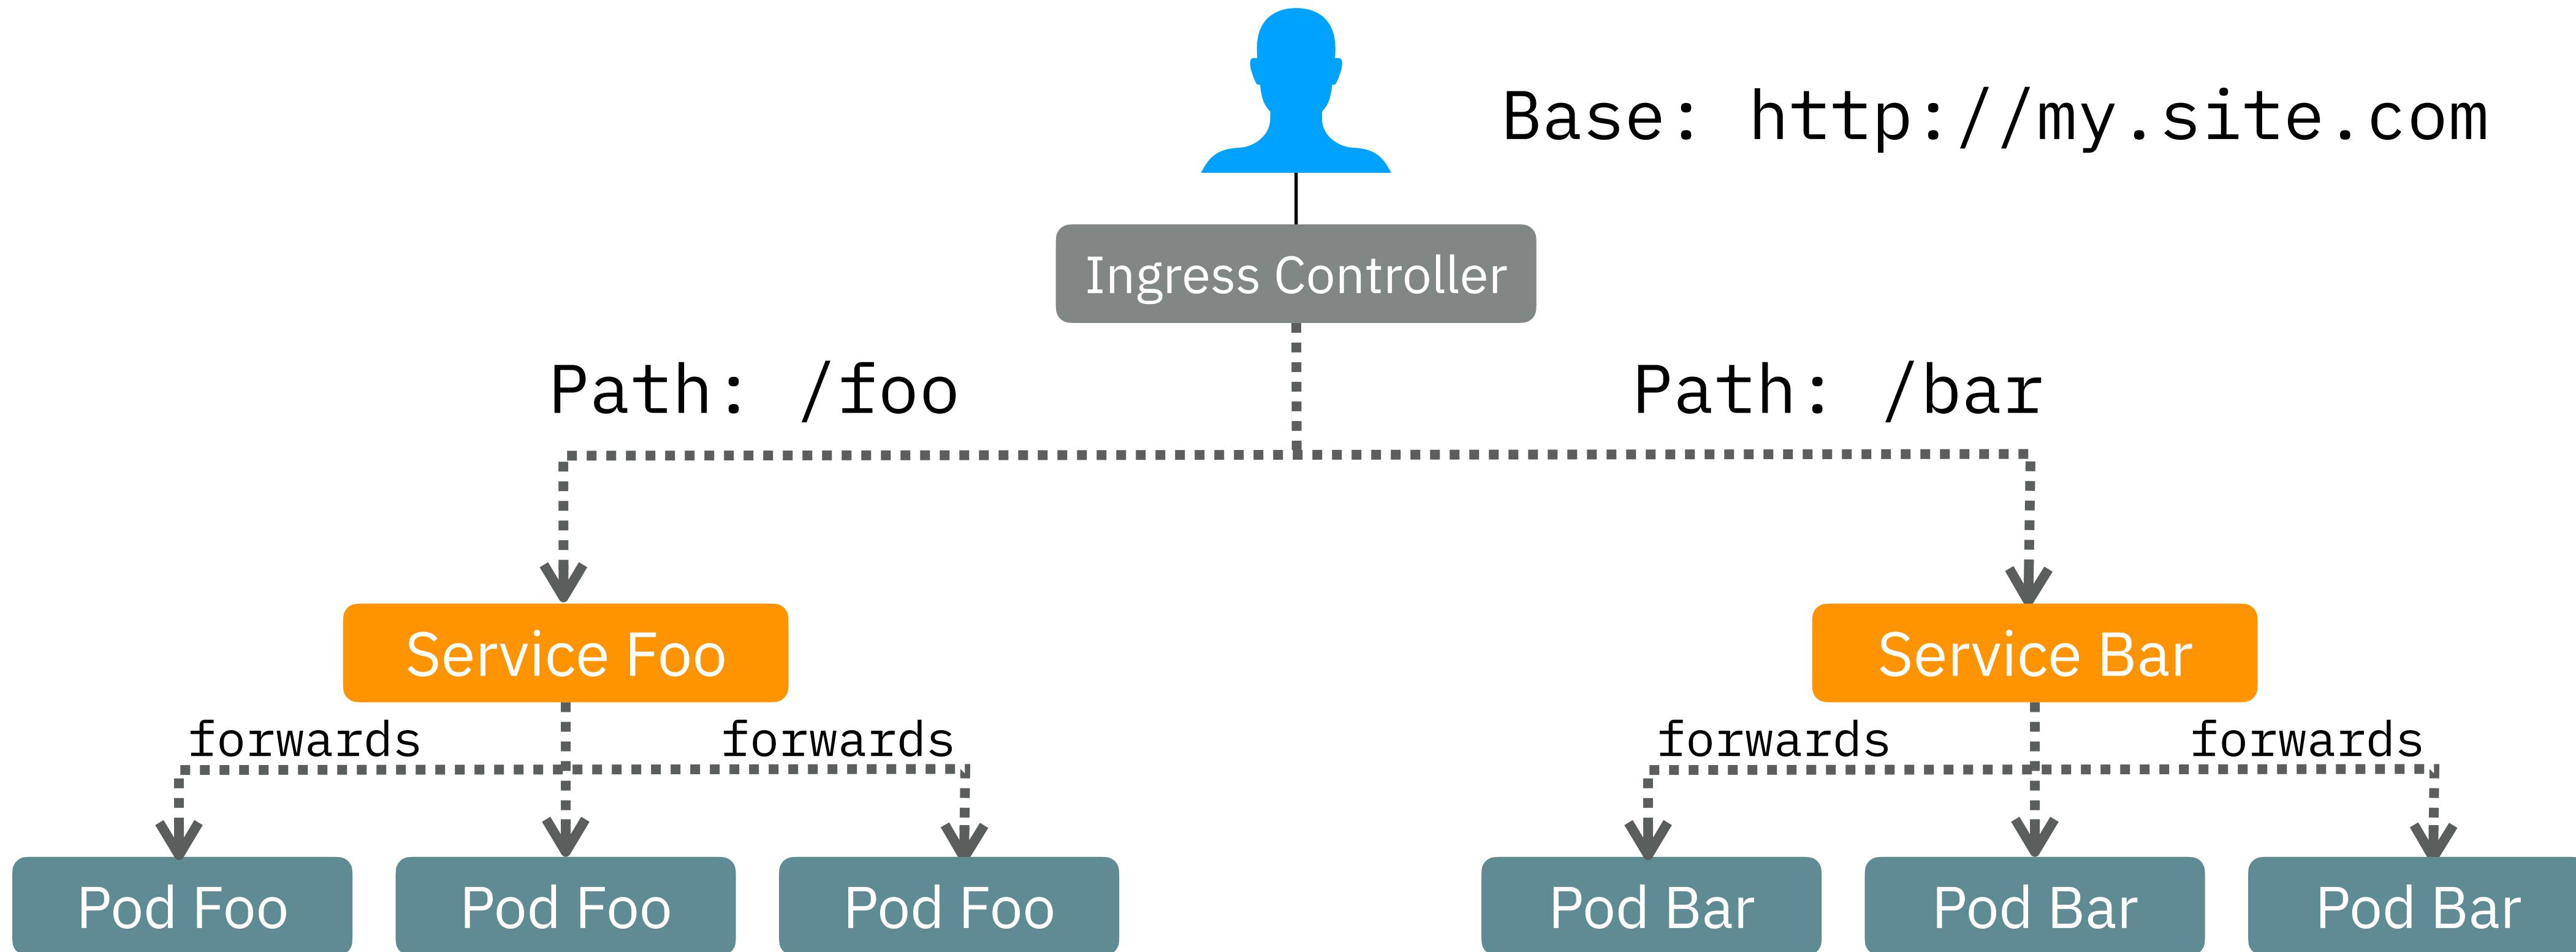
# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.containers.mybluemix.net
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Will be routed to these microservices

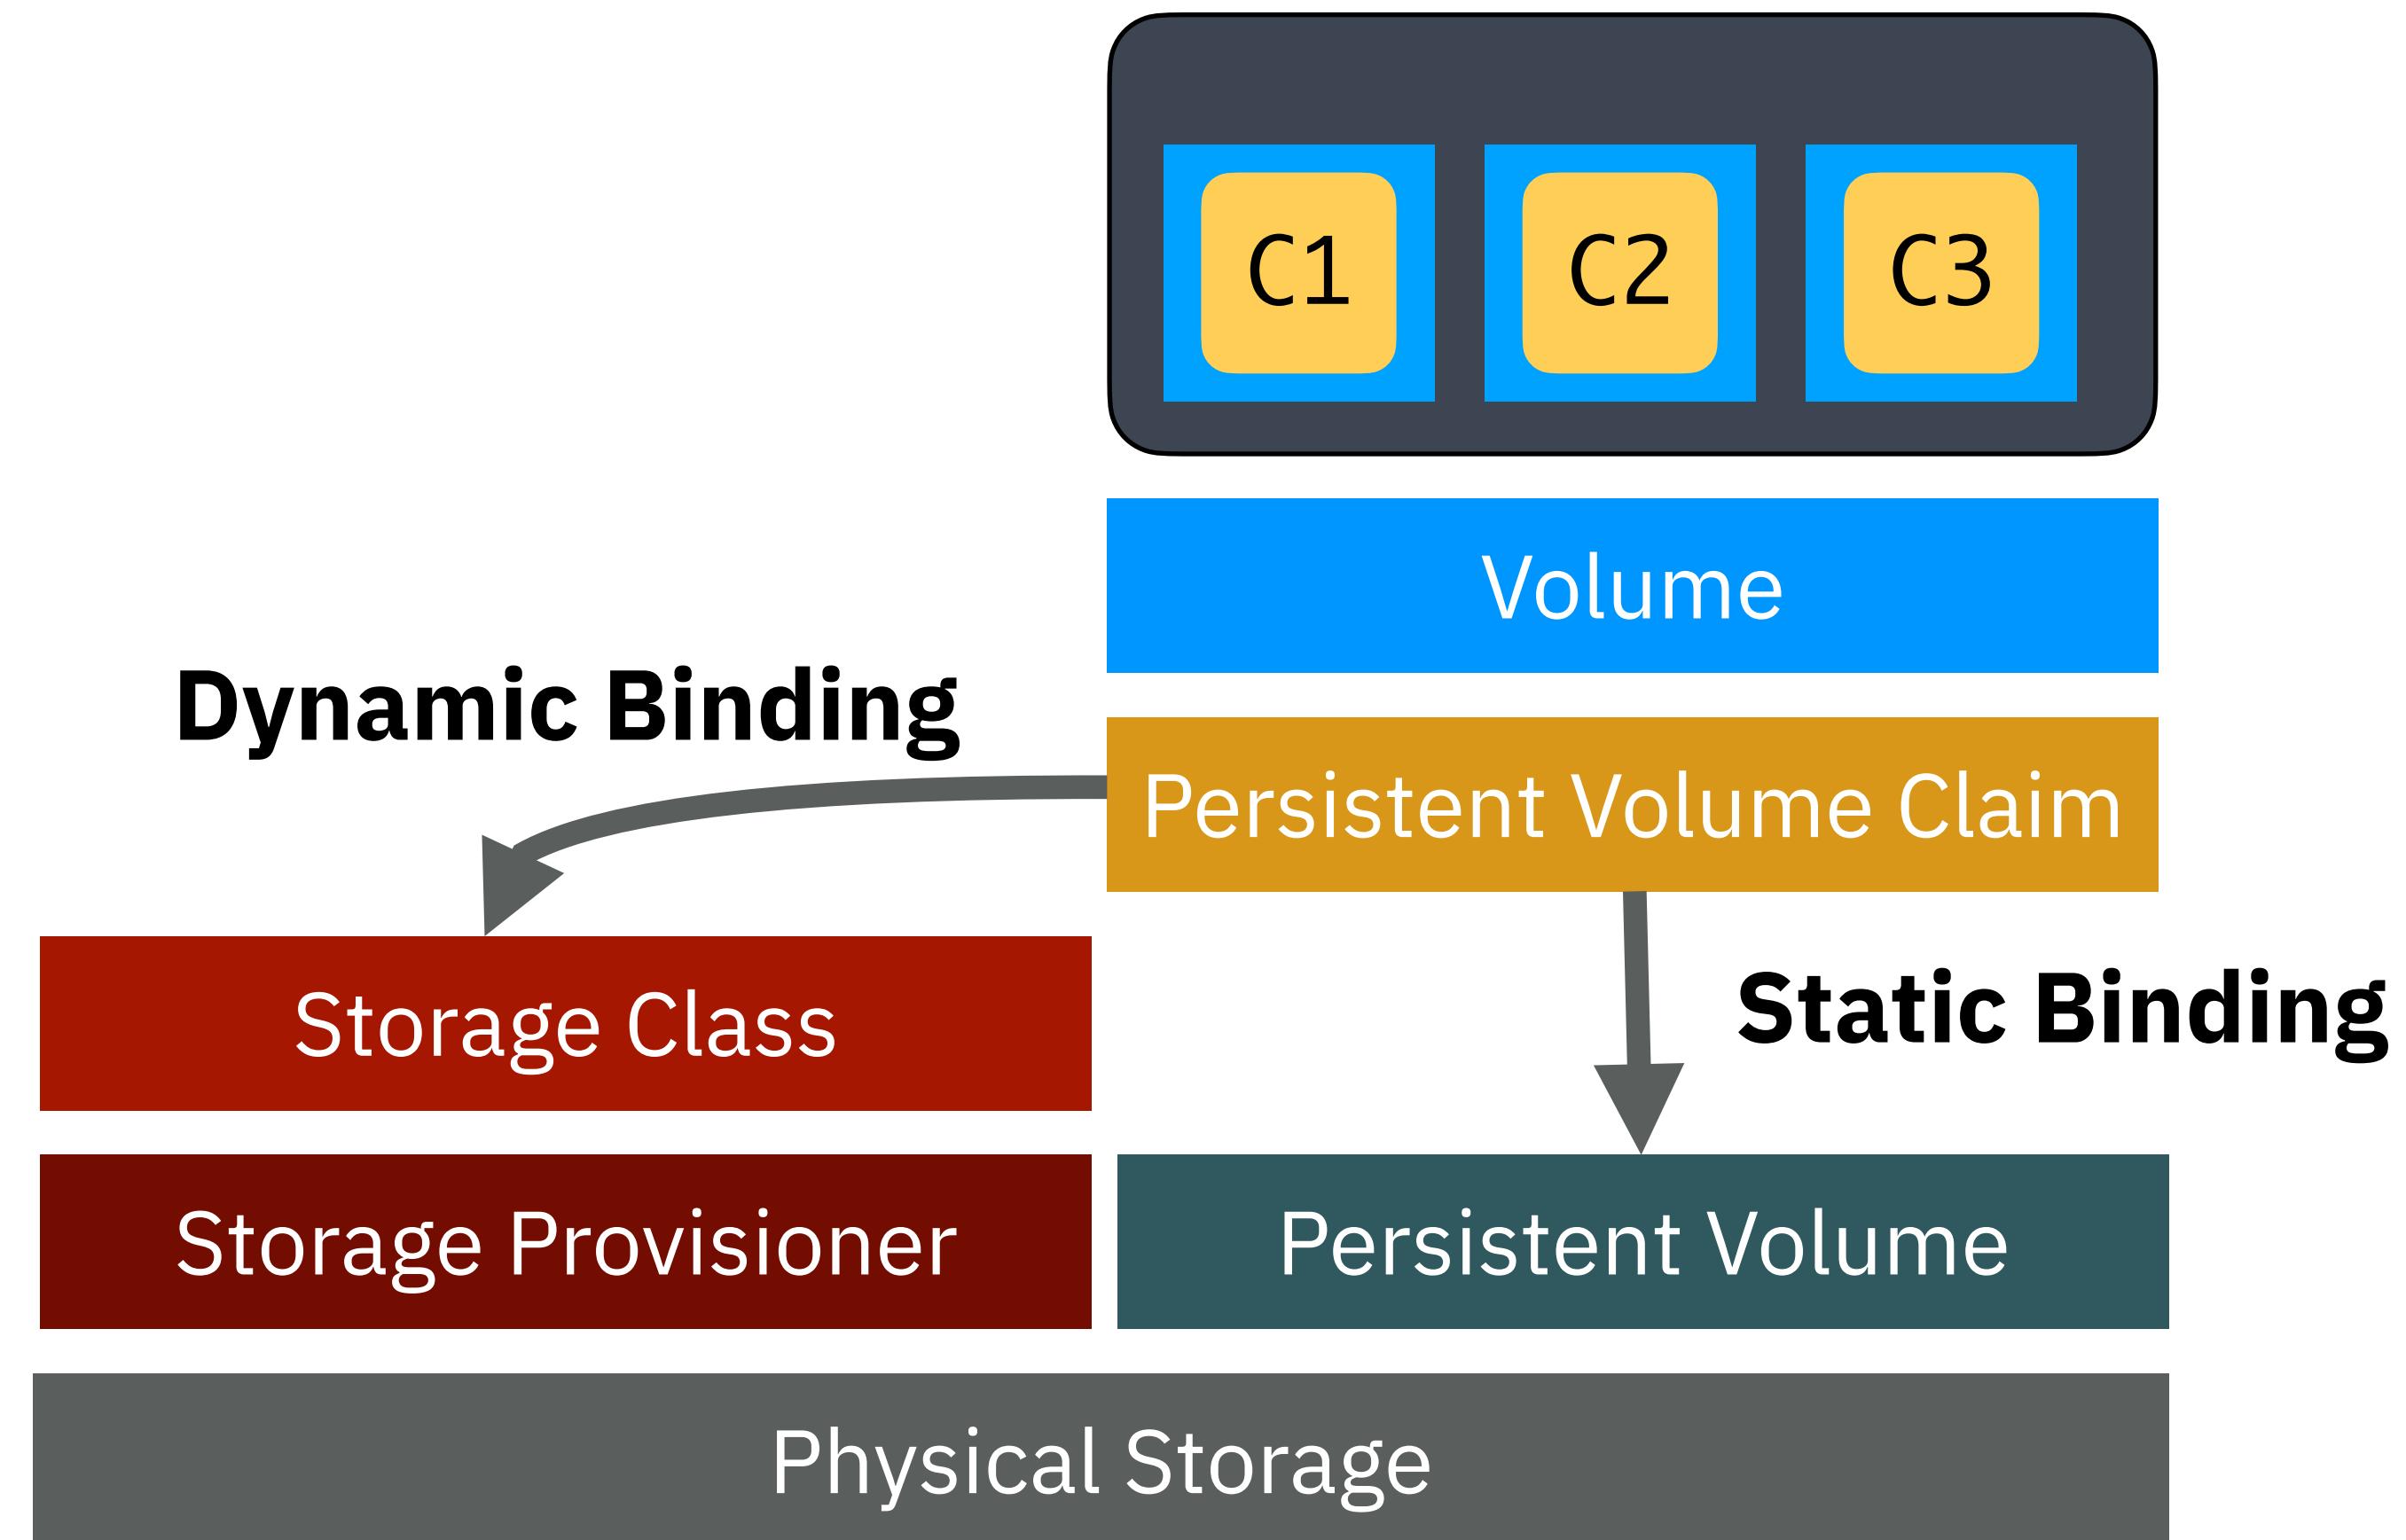
# Ingress Controller

- Single entry point into multiple kubernetes services



# Persistent Volumes

- Kubernetes loosely couples physical storage devices with containers by introducing an intermediate resource called persistent volume claims (PVCs).
- A PVC defines the disk size, disk type (ReadWriteOnce, ReadOnlyMany, ReadWriteMany) and dynamically links a storage device to a volume defined against a pod
- The binding process can either be done in a static way using PVs or dynamically be using a persistent storage provider



# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

~

# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

Volume named `redis-storage`

# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

Mounted at: /data/redis

Volume named redis-storage

# Configurations Management

(ConfigMap)

- Containers generally use environment variables for parameterizing their runtime configurations
- Kubernetes provides a way of managing more complex configuration files using a simple resource called ConfigMaps
- ConfigMaps can be created using directories, files or literal values using following CLI command:

```
$ kubectl create configmap <map-name> <data-source>  
  
# map-name: name of the config map  
# data-source: directory, file or literal value
```

# Credentials Management

(Secrets)

- Similar to ConfigMaps, Kubernetes provides another valuable resource called Secrets for managing sensitive information such as passwords, OAuth tokens, and ssh keys.
- A secret can be created for managing basic auth credentials using the following way:

```
# write credentials to two files
$ echo -n 'admin' > ./username.txt
$ echo -n '1f2d1e2e67df' > ./password.txt

# create a secret
$ kubectl create secret generic app-credentials --from-file=./username.txt --from-file=./password.txt
```

# Credentials From Environment Variables (Secrets)

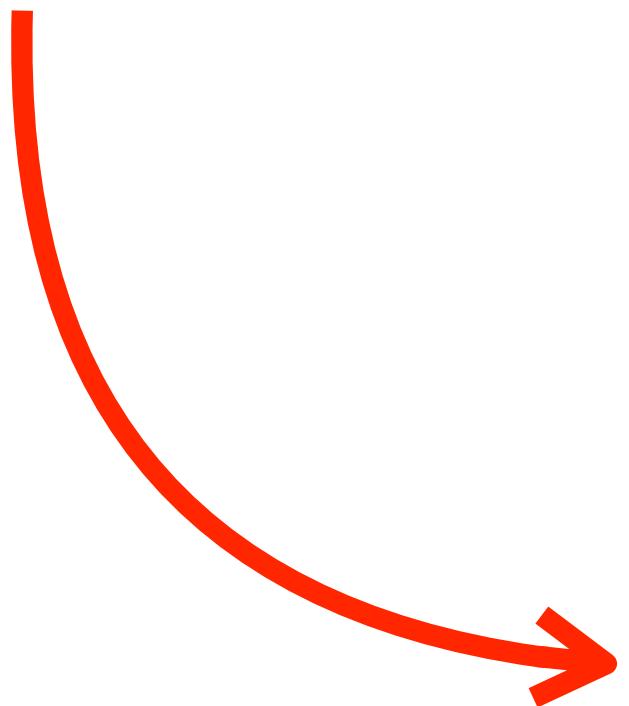
```
# export DATABASE_URI='postgres://admin:s3cr3t@postgres:5432/postgres'  
  
# kubectl create secret generic db-creds --from-literal=database-uri=$DATABASE_URI
```

- Creates the following

```
apiVersion: v1  
kind: Secret  
data:  
  database-uri: cG9zdGdyZXM6Ly9hZG1pbjpzM2NyM3RAcG9zdGdyZXM6NTQzMj9wb3N0Z3Jlcw==
```

# Create Secrets from Literals

```
kubectl create secret generic dev-db-secret \
--from-literal=username=devuser \
--from-literal=password='s3cr3t'
```



```
apiVersion: v1
kind: Secret
metadata:
  name: dev-db-secret
type: Opaque
data:
  username: ZGV2dXNlcgo=
  password: czNjcjN0Cg==
```

# Secret Yaml Example

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: ecommerce-apikey
5   namespace: default
6 data:
7   secret: <place base64 encoded secret here>
```

```
$ echo -n "this is my secret" | base64
dGhpcyBpcyBteSBzZWNyZXQ=
```

# Secret Yaml Example

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: ecommerce-apikey
5   namespace: default
6 data:
7 secret: <place base64 encoded secret here>
```

Place the base64 encode string in secret

```
$ echo -n "this is my secret" | base64
dGhpcyBpcyBteSBzZWNyZXQ=
```

# Using Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: pet-creds
  namespace: default
data:
  binding: dGhpcyBpcyBteSBzZWNyZXQ=
```

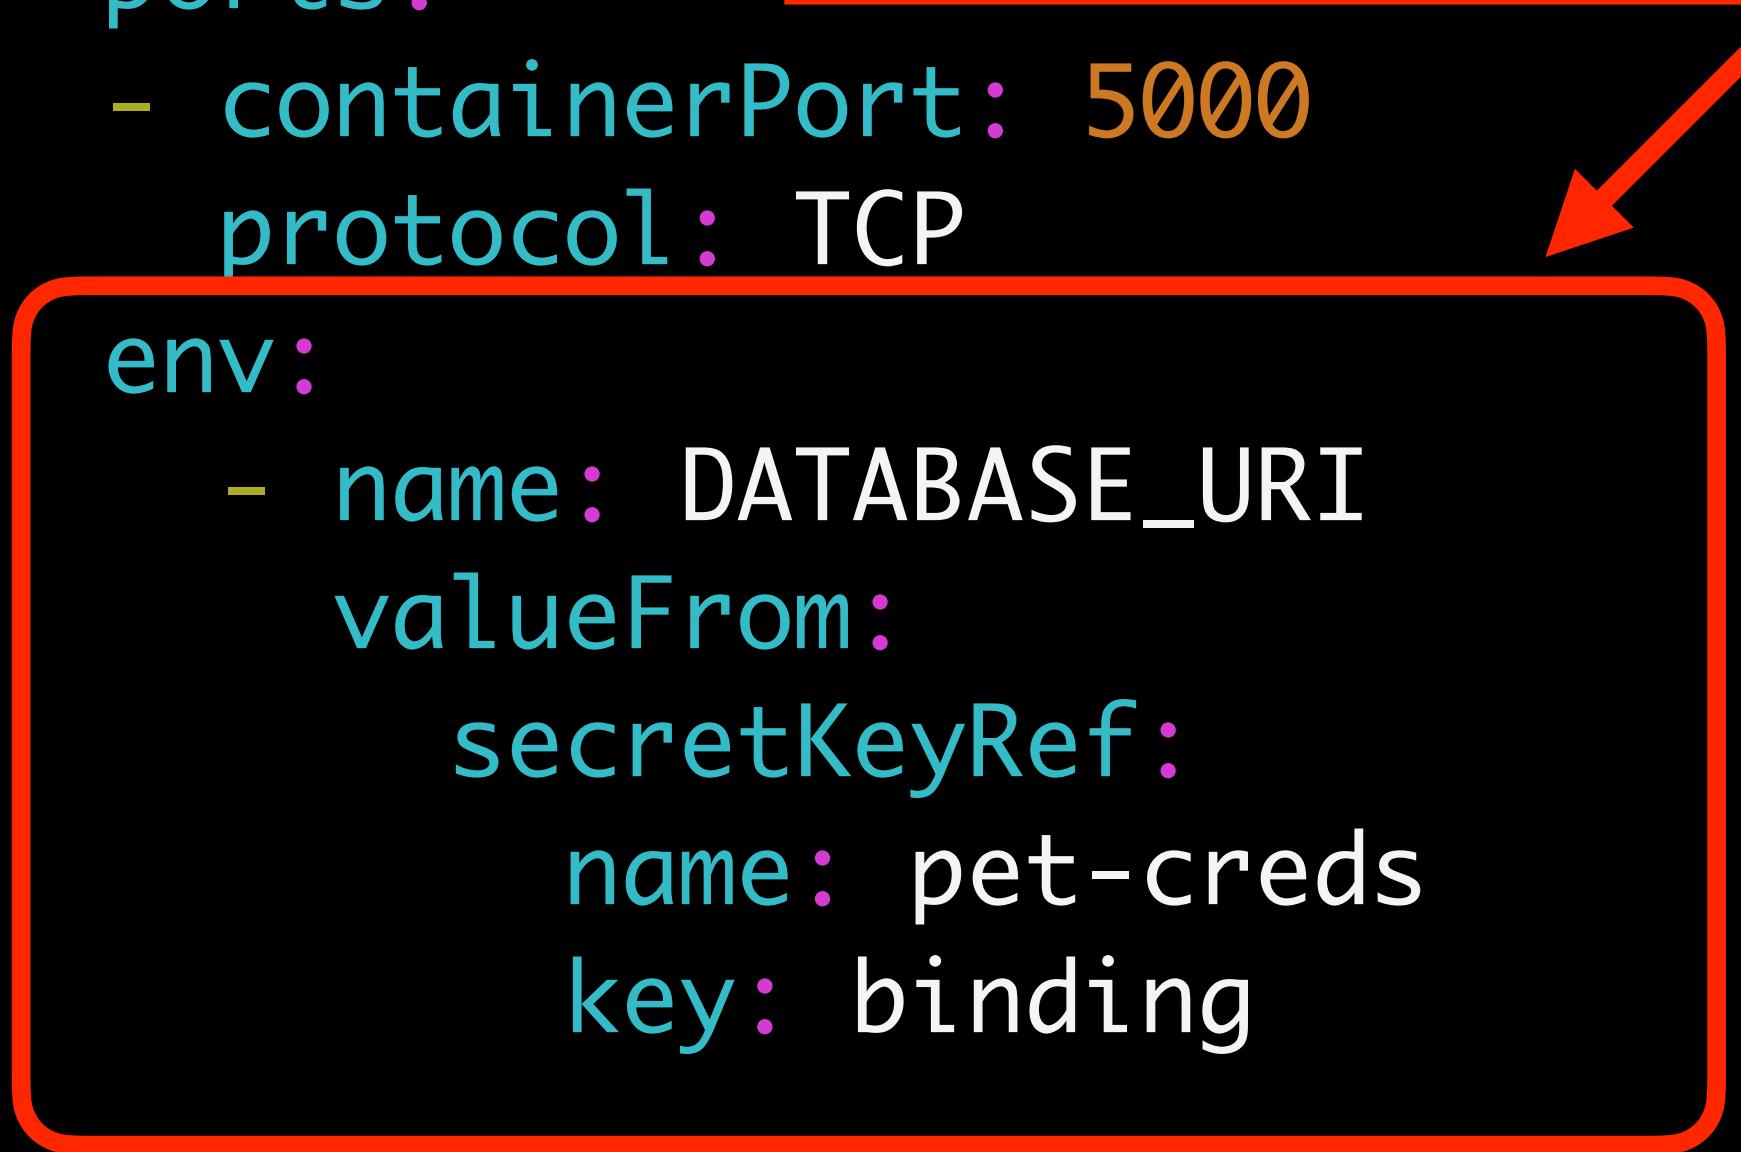
```
apiVersion: apps/v1
kind: Deployment
...
spec:
  containers:
    - name: pet-demo
      image: pet-demo:v1
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5000
          protocol: TCP
      env:
        - name: DATABASE_URI
          valueFrom:
            secretKeyRef:
              name: pet-creds
              key: binding
```

# Using Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: pet-creds
  namespace: default
data:
  binding: dGhpcyBpcyBteSBzZWNyZXQ=
```

```
apiVersion: apps/v1
kind: Deployment
...
spec:
  containers:
    - name: pet-demo
      image: pet-demo:v1
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5000
          protocol: TCP
      env:
        - name: DATABASE_URI
          valueFrom:
            secretKeyRef:
              name: pet-creds
              key: binding
```

Secrets exposed as environment variables



# Kubernetes Rolling Updates (Zero Downtime Deployments)

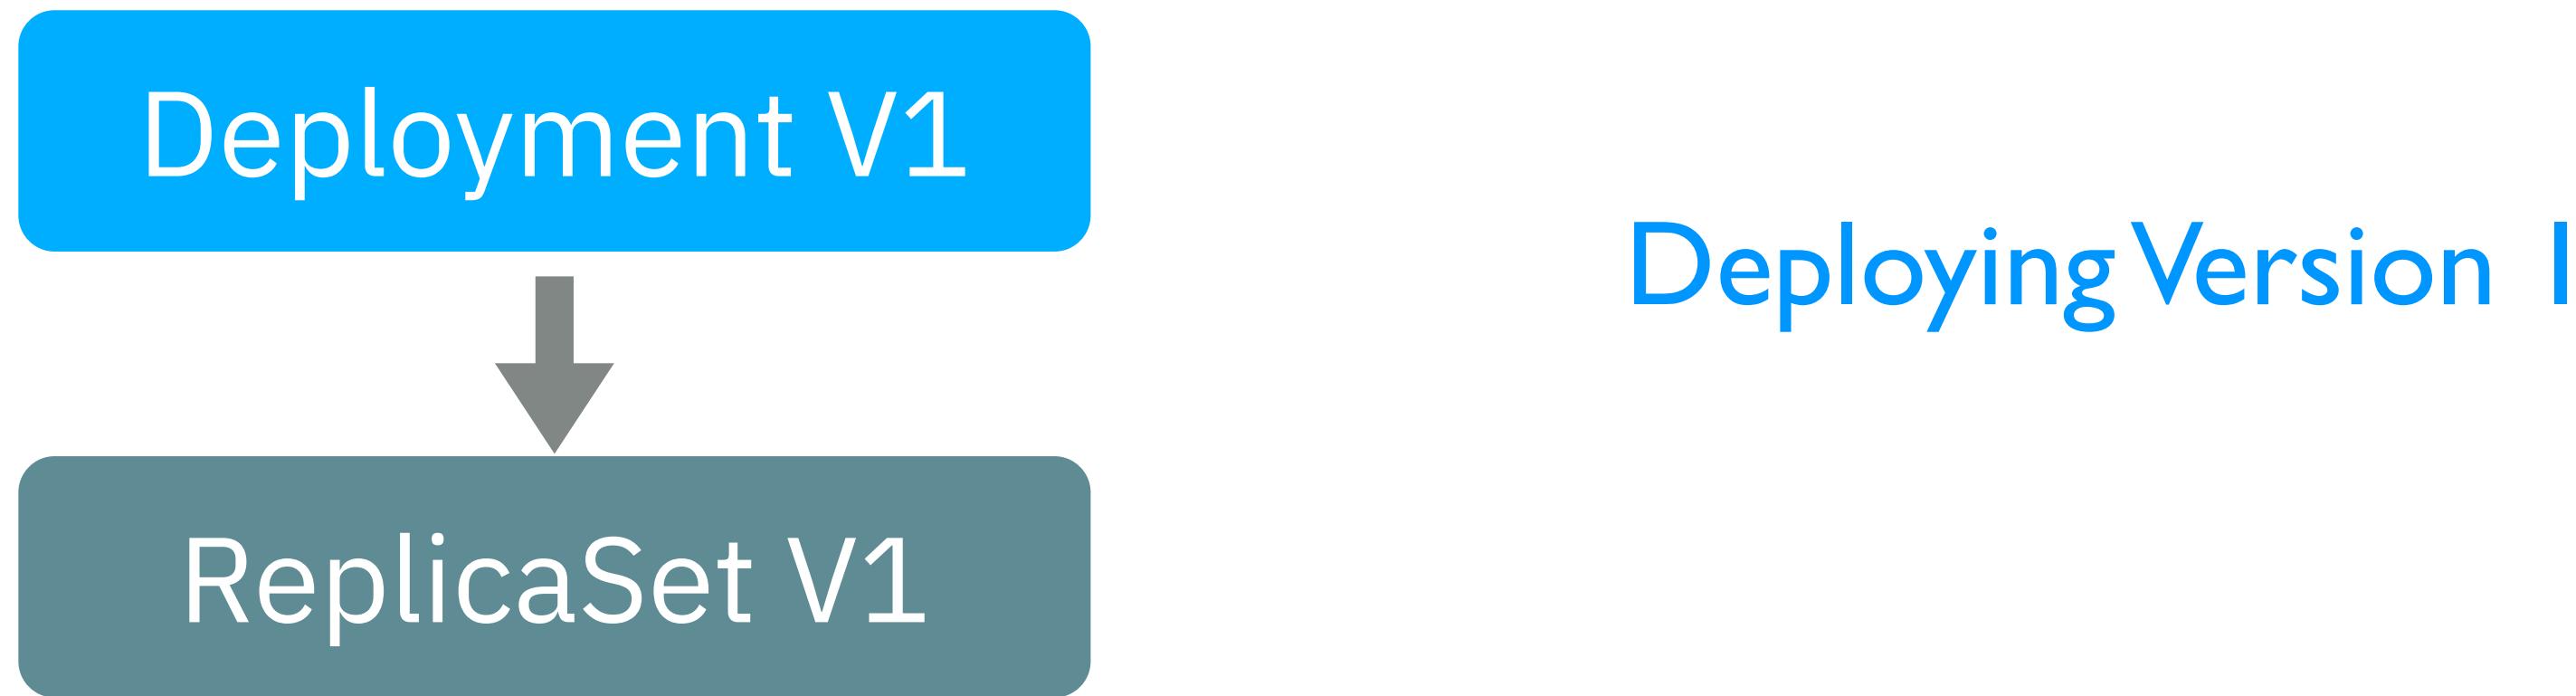
```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```

Deployment V1

Deploying Version I

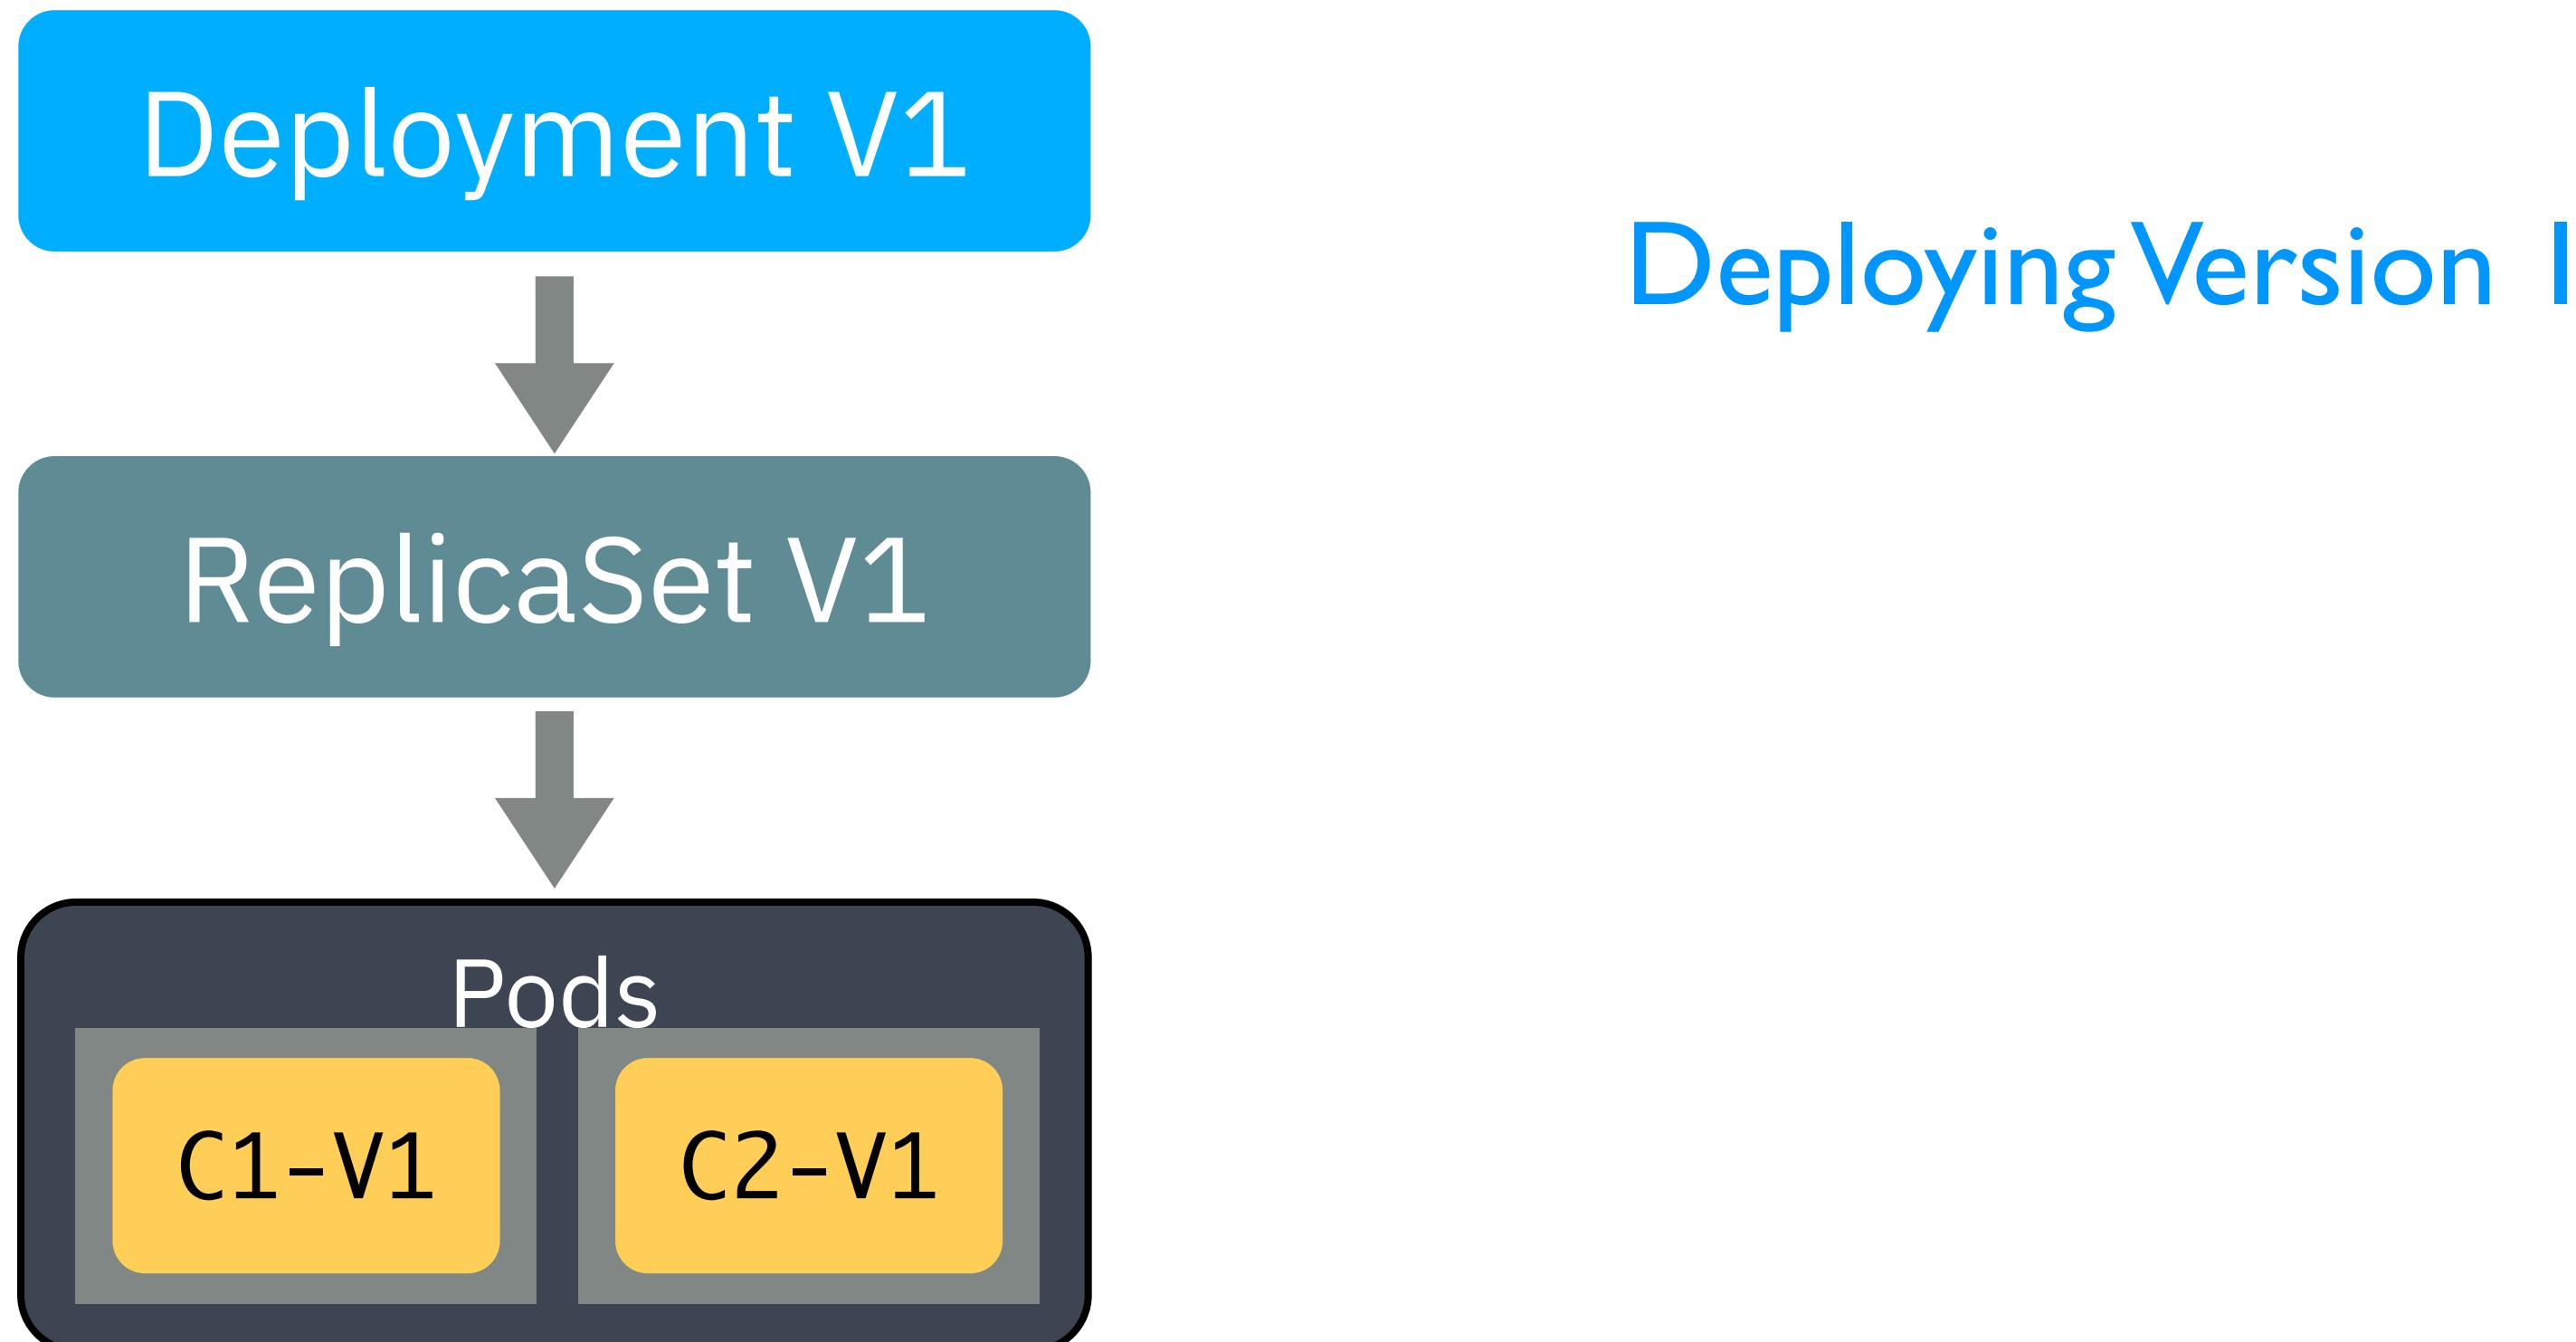
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



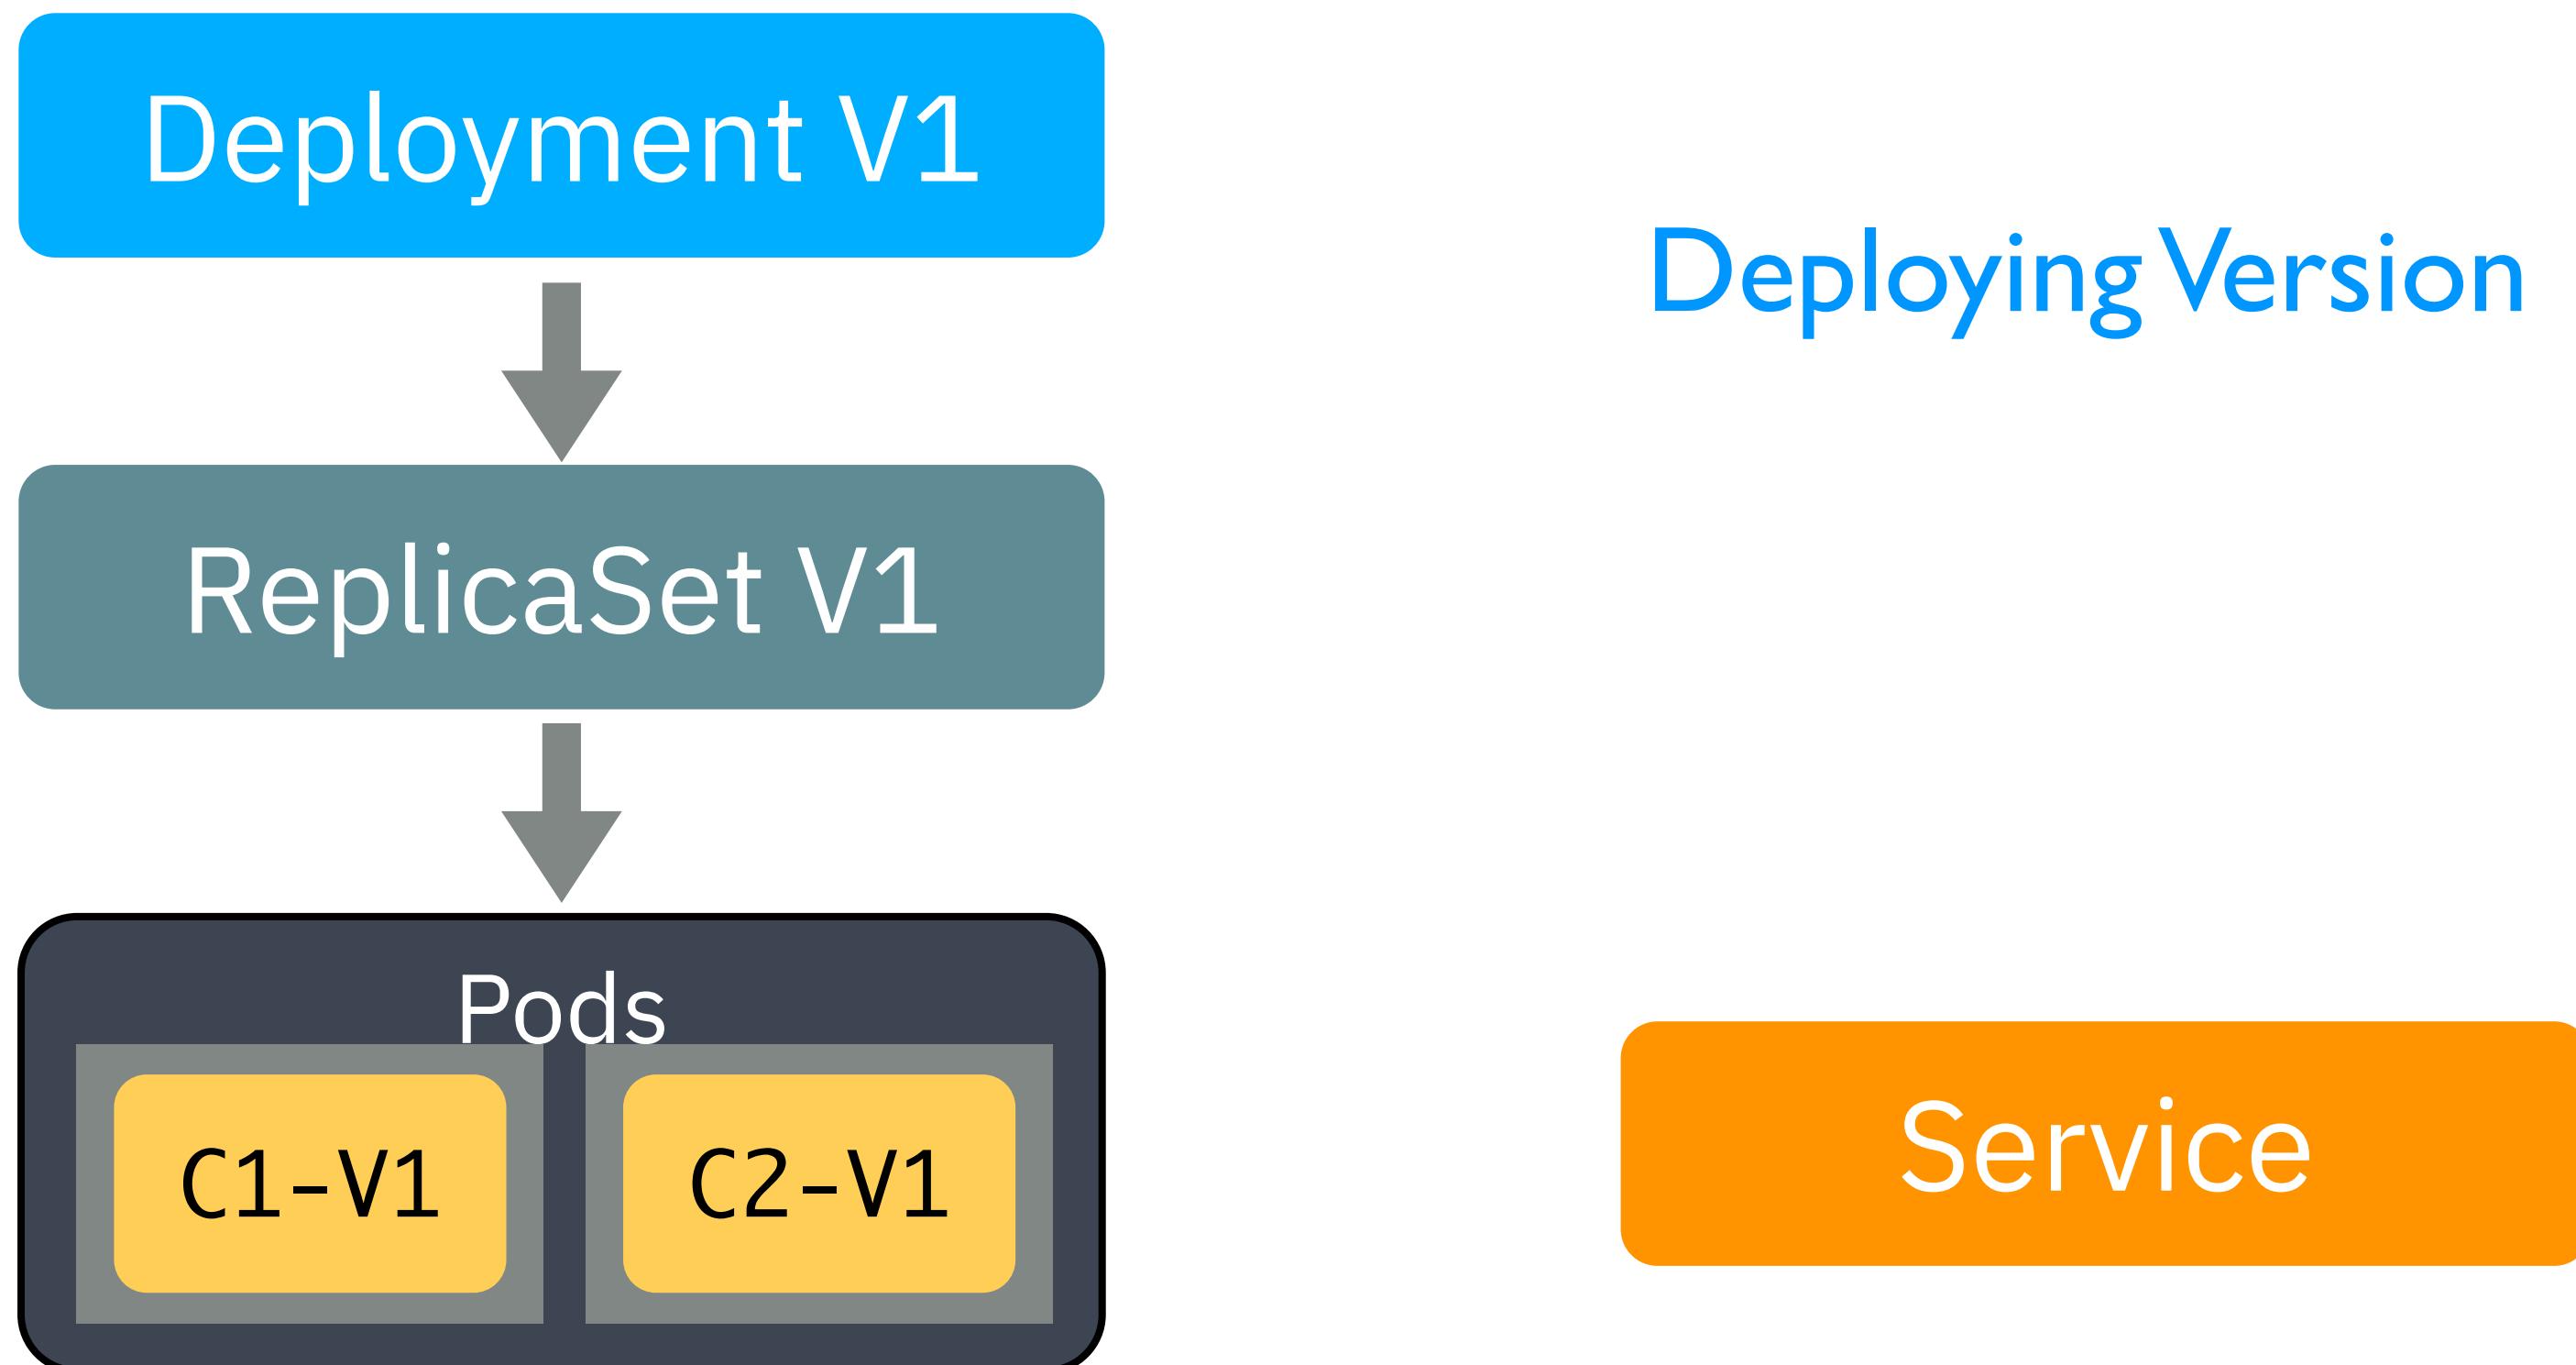
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



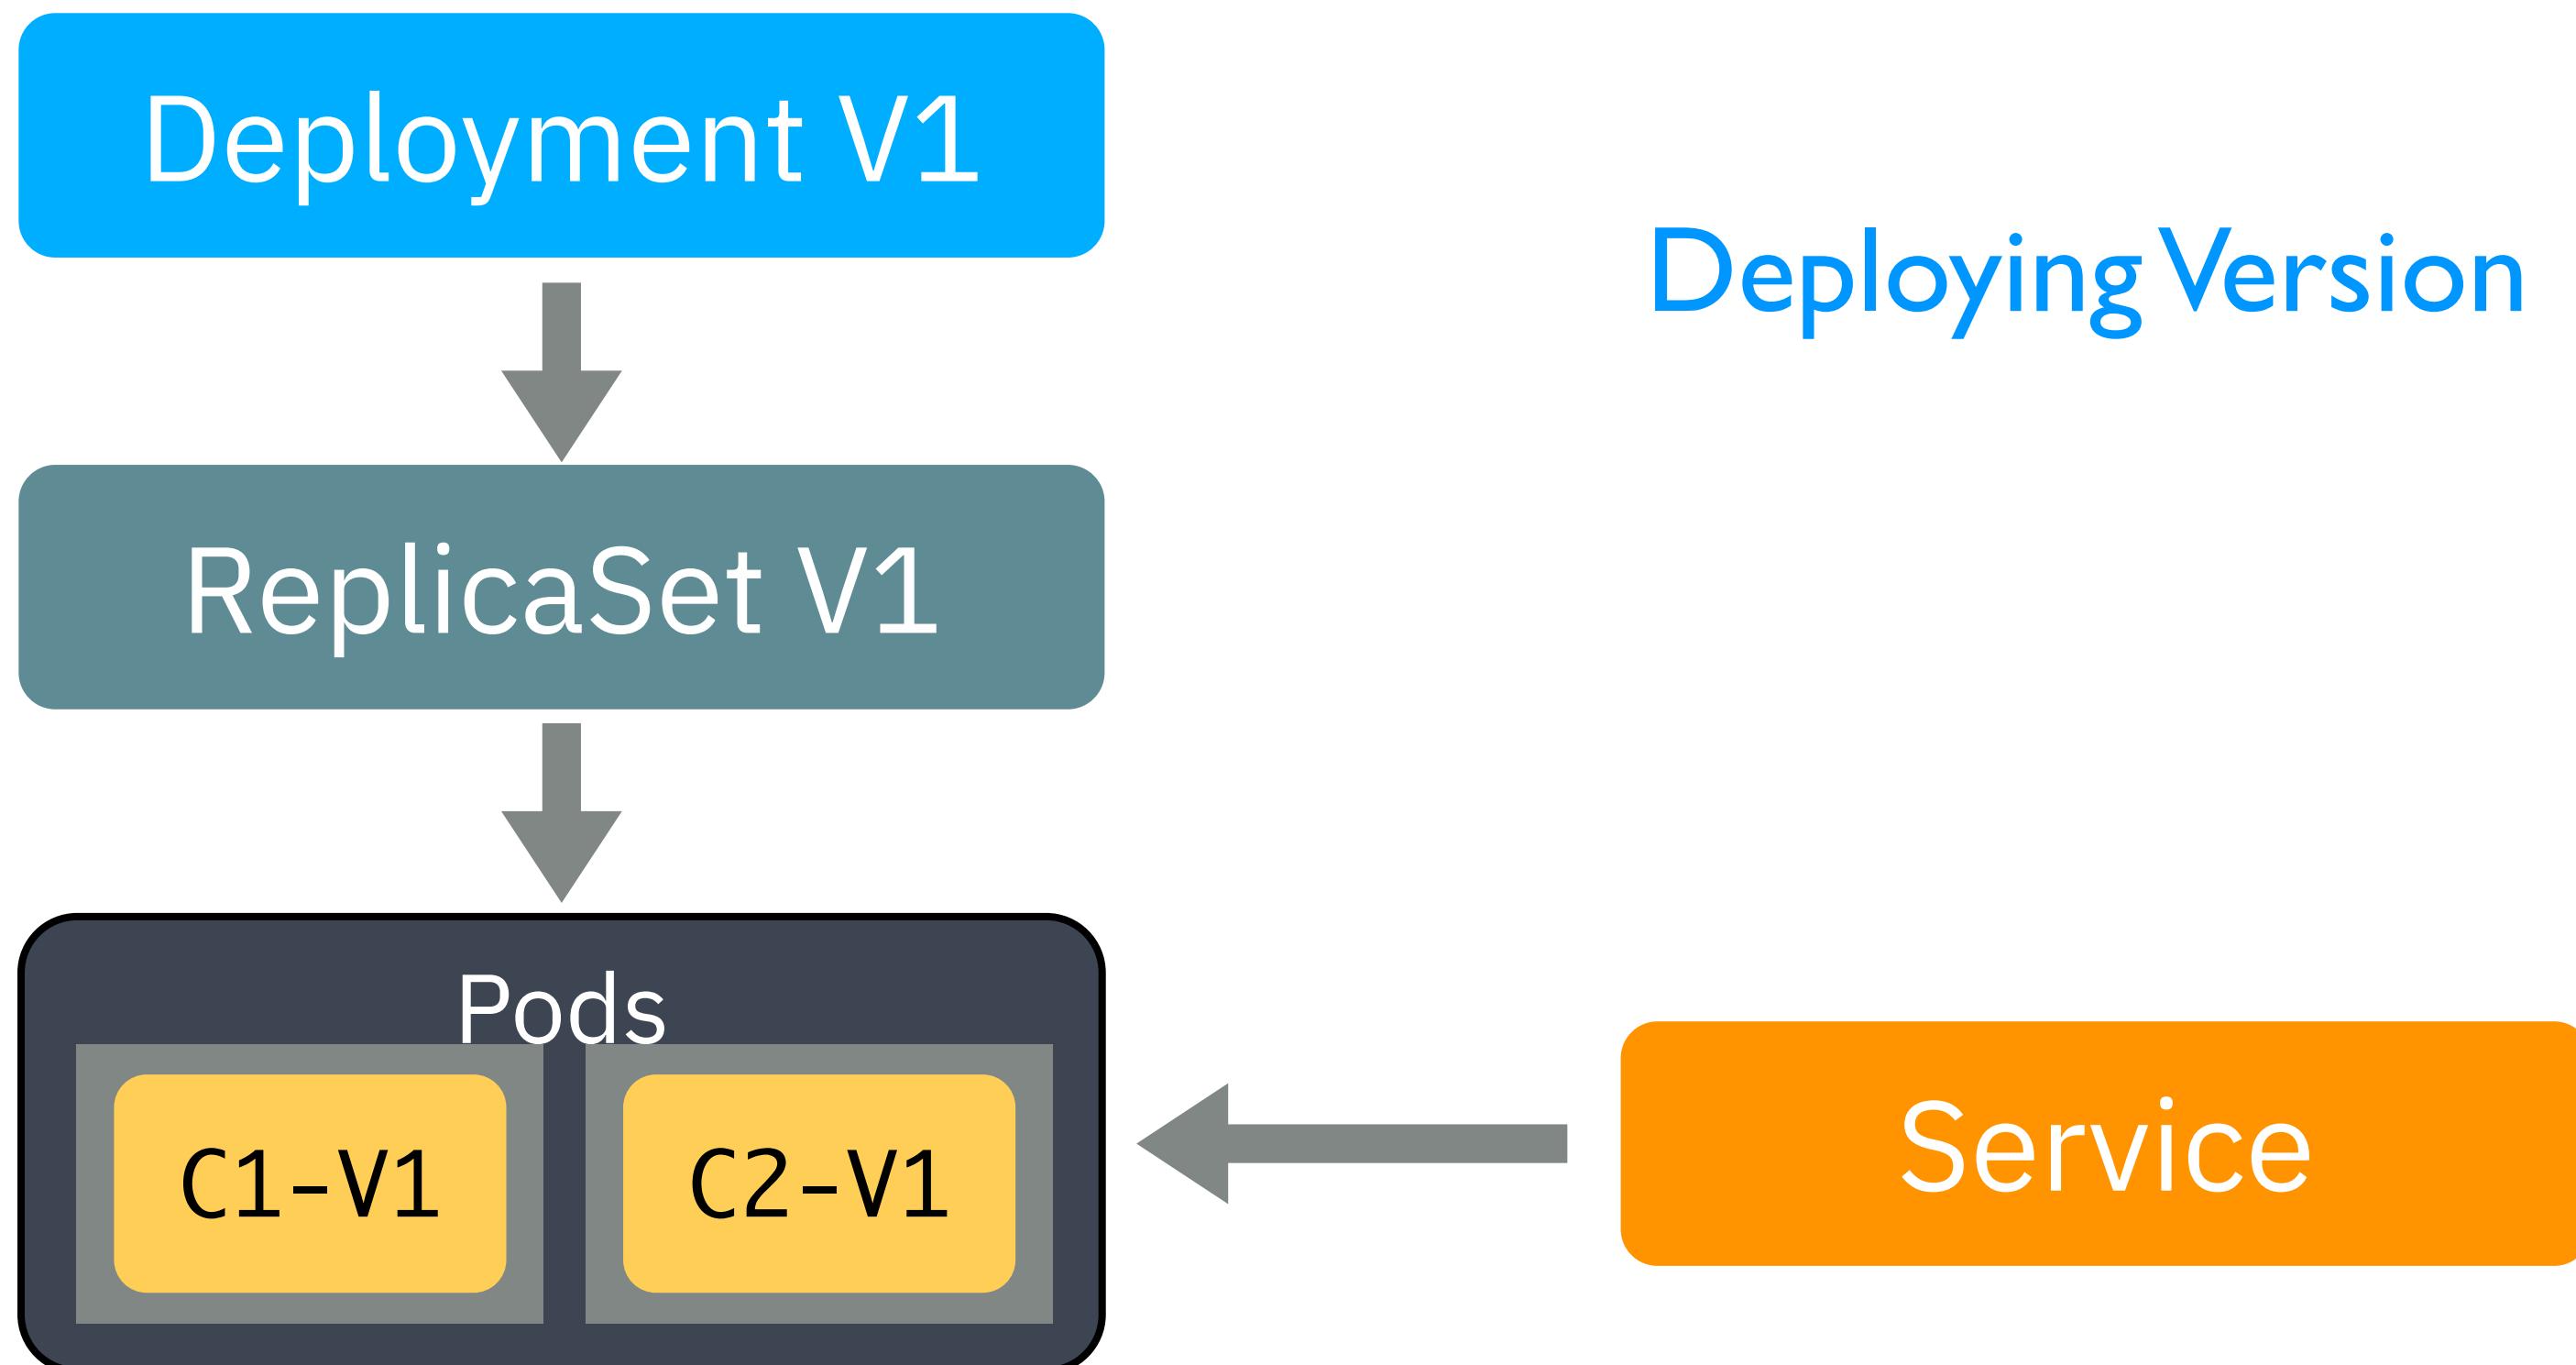
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



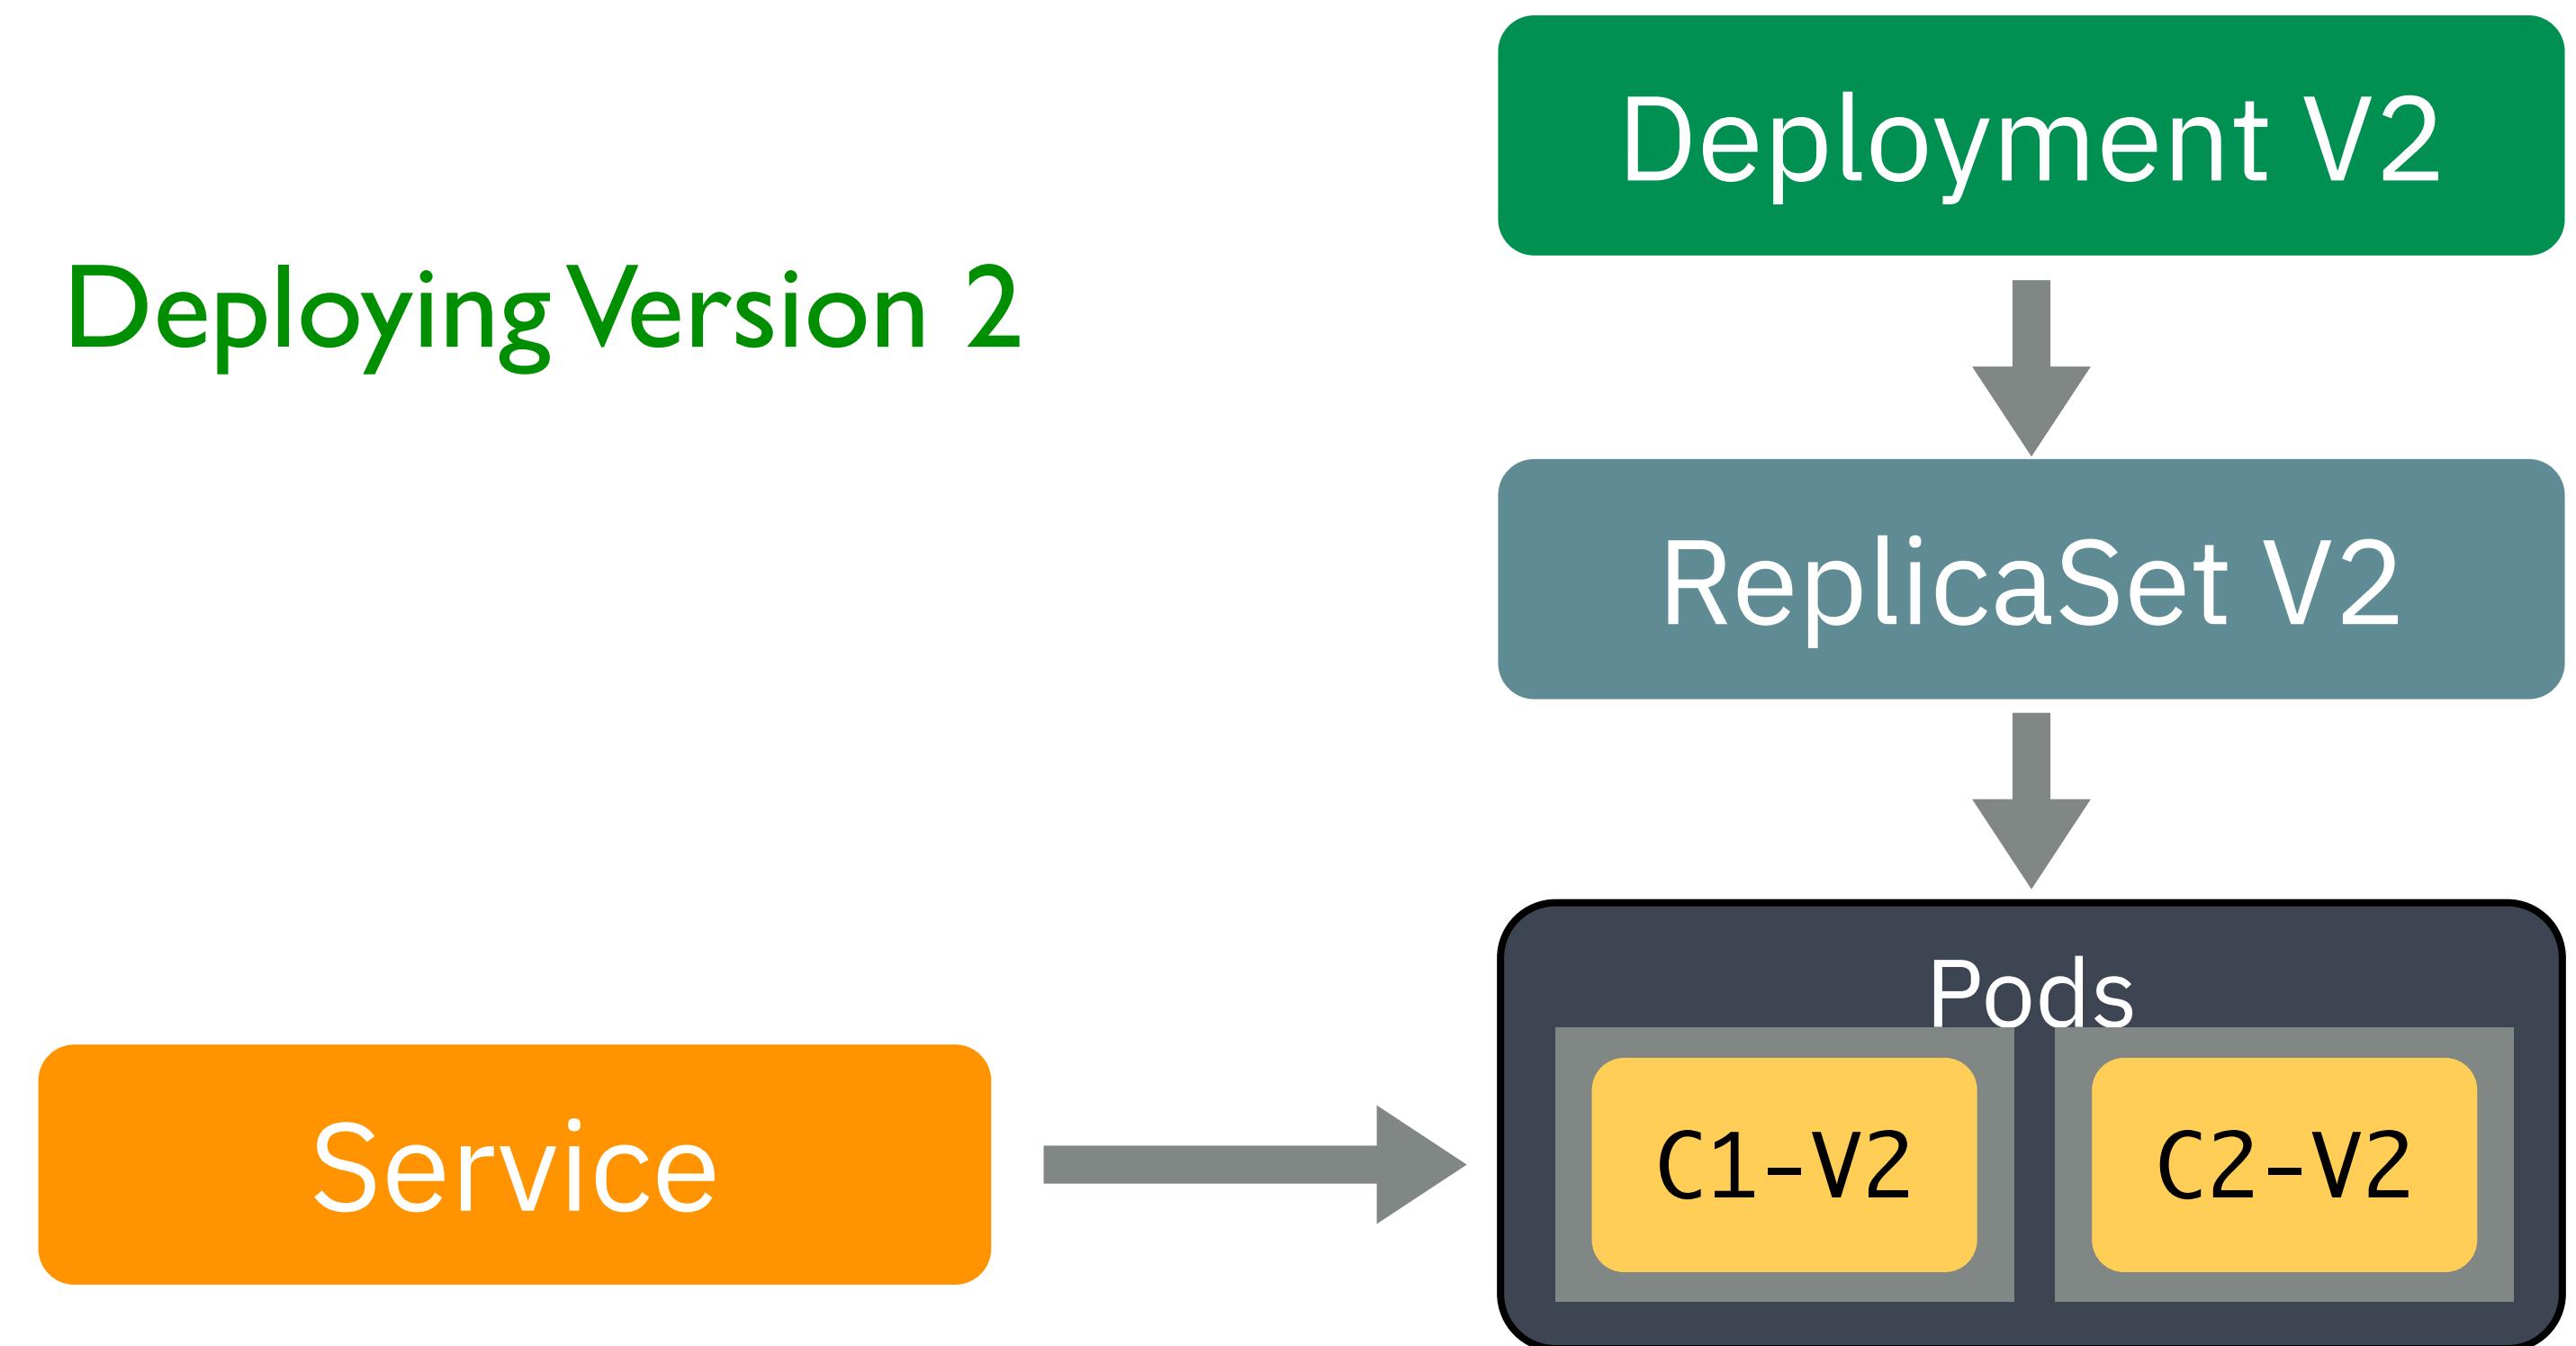
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

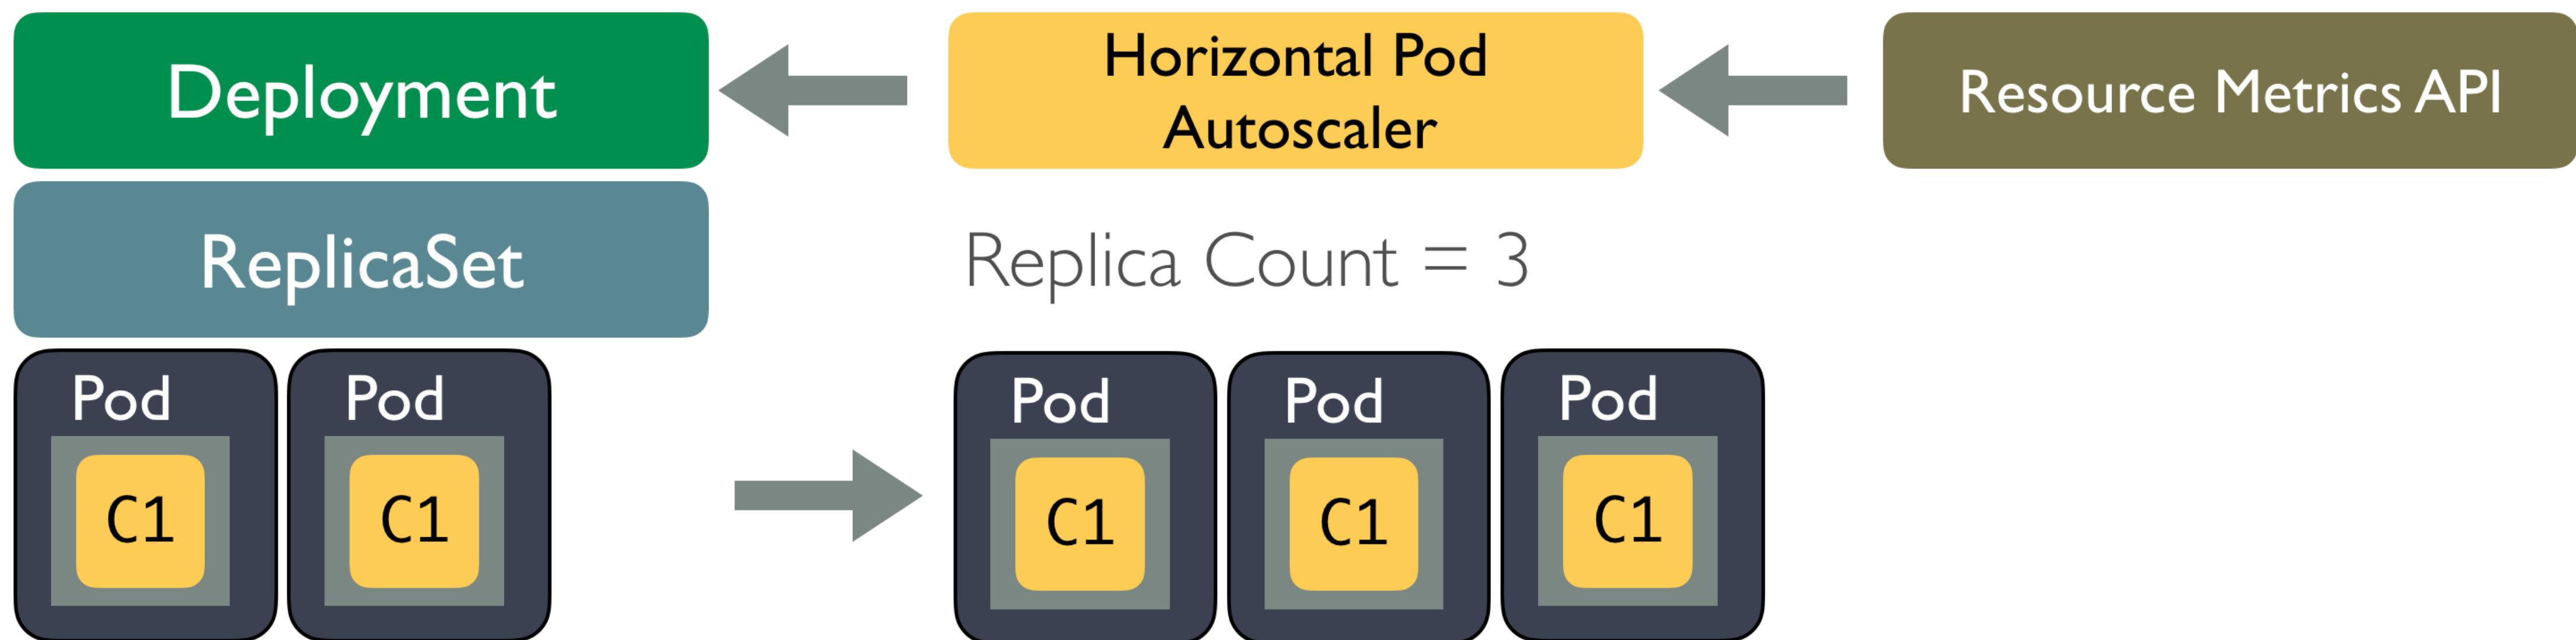
```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Autoscaling

Kubernetes allows pods to be manually scaled either using ReplicaSets or Deployments. This can be achieved using the following CLI command:

```
$ kubectl scale --replicas=<desired-instance-count> deployment/<application-name>
```



# Kubernetes Distributions

kubernetes, minikube, openshift, minishift, oco, minikube, microk8s, k3s, k3d

The screenshot shows the official Kubernetes website's 'Getting started' section. Under the 'Installing kubeadm' heading, it provides instructions for installing the kubeadm toolbox. It includes a 'Before you begin' section with several bullet points and a 'kubeadm' logo.

The screenshot shows the Kubernetes 'Tutorials' section. Under the 'Hello Minikube' heading, it provides instructions for running a simple Hello World Node.js app on Kubernetes using Minikube and Katacoda. It includes a 'Note' section and a 'Objectives' section with several bullet points.

The screenshot shows the Canonical MicroK8s page. It features a large heading 'Autonomous low-ops Kubernetes for clusters, workstations, edge and IoT'. Below it is a 'certified kubernetes' badge. The 'Hello Minikube' section is identical to the one on the official Kubernetes site.

The screenshot shows the Red Hat OpenShift homepage. It highlights 'OPENSHIFT 4.3 IS HERE' and 'The Kubernetes platform for big ideas'. It includes a 'Get started' button and a 'What's new in OpenShift 4.3?' section.

The screenshot shows the OKD Minishift landing page. It features a large 'okd' logo and the text 'Minishift: Develop Applications Locally in a Containerized OKD Cluster'. It includes an 'ABOUT', 'GET STARTED', and 'RESOURCES' navigation bar.

The screenshot shows the Rancher K3s page. It features a large 'Lightweight Kubernetes' heading and a 'Great For' section with categories: Edge, IoT, CI, and ARM. It includes a code snippet for curling to get.k3s.io and a 'For detailed installation, refer to the docs' link.

# Hands-On

# Confirm that your cluster Works

- Make sure that your cluster is configured correctly

```
$ kubectl get all
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3m44s

# Deploy Nginx

- Create a Deployment

```
$ kubectl create deployment my-nginx --image=nginx:alpine
deployment.apps/my-nginx created
```

- Check that the deployment and pods were created

```
$ kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
my-nginx   1/1     1           1           3h42m

$ kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
my-nginx-b49cf867-q7mjn      1/1     Running   0          3h42m
```

# Deploy Nginx

- Create a Service

```
$ kubectl expose deploy my-nginx --type=NodePort --port=80
service/my-nginx exposed
```

- Check that the service was created

\$ kubectl get service						
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	
kubernetes	ClusterIP	172.21.0.1	<none>	443/TCP	317d	
my-nginx	NodePort	172.21.12.226	<none>	80:30371/TCP	3h42m	

# Kubectl get all

```
vagrant@kubernetes:~$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/my-nginx-64f47865f-2z5w8	1/1	Running	0	2m47s
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP
service/my-nginx	NodePort	10.104.182.40	<none>	80:30965/TCP
NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/my-nginx	1/1	1	1	2m47s
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/my-nginx-64f47865f	1	1	1	2m47s

# Naming Convention

- We made a deployment called my-nginx and it created:

**Deployment:** my-nginx

**ReplicaSet:** my-nginx-b49cf867

**Pod:** my-nginx-b49cf867-q7mjn

# Naming Convention

- We made a deployment called my-nginx and it created:



# Naming Convention

- We made a deployment called my-nginx and it created:

**Deployment:** my-nginx

**ReplicaSet:** my-nginx-b49cf867

Pod name starts with ReplicaSet Name

**Pod:** my-nginx-b49cf867-q7mjn

# Get the NodePort

\$ kubectl get service my-nginx					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	NodePort	172.21.9.54	<none>	80:30187/TCP	37d

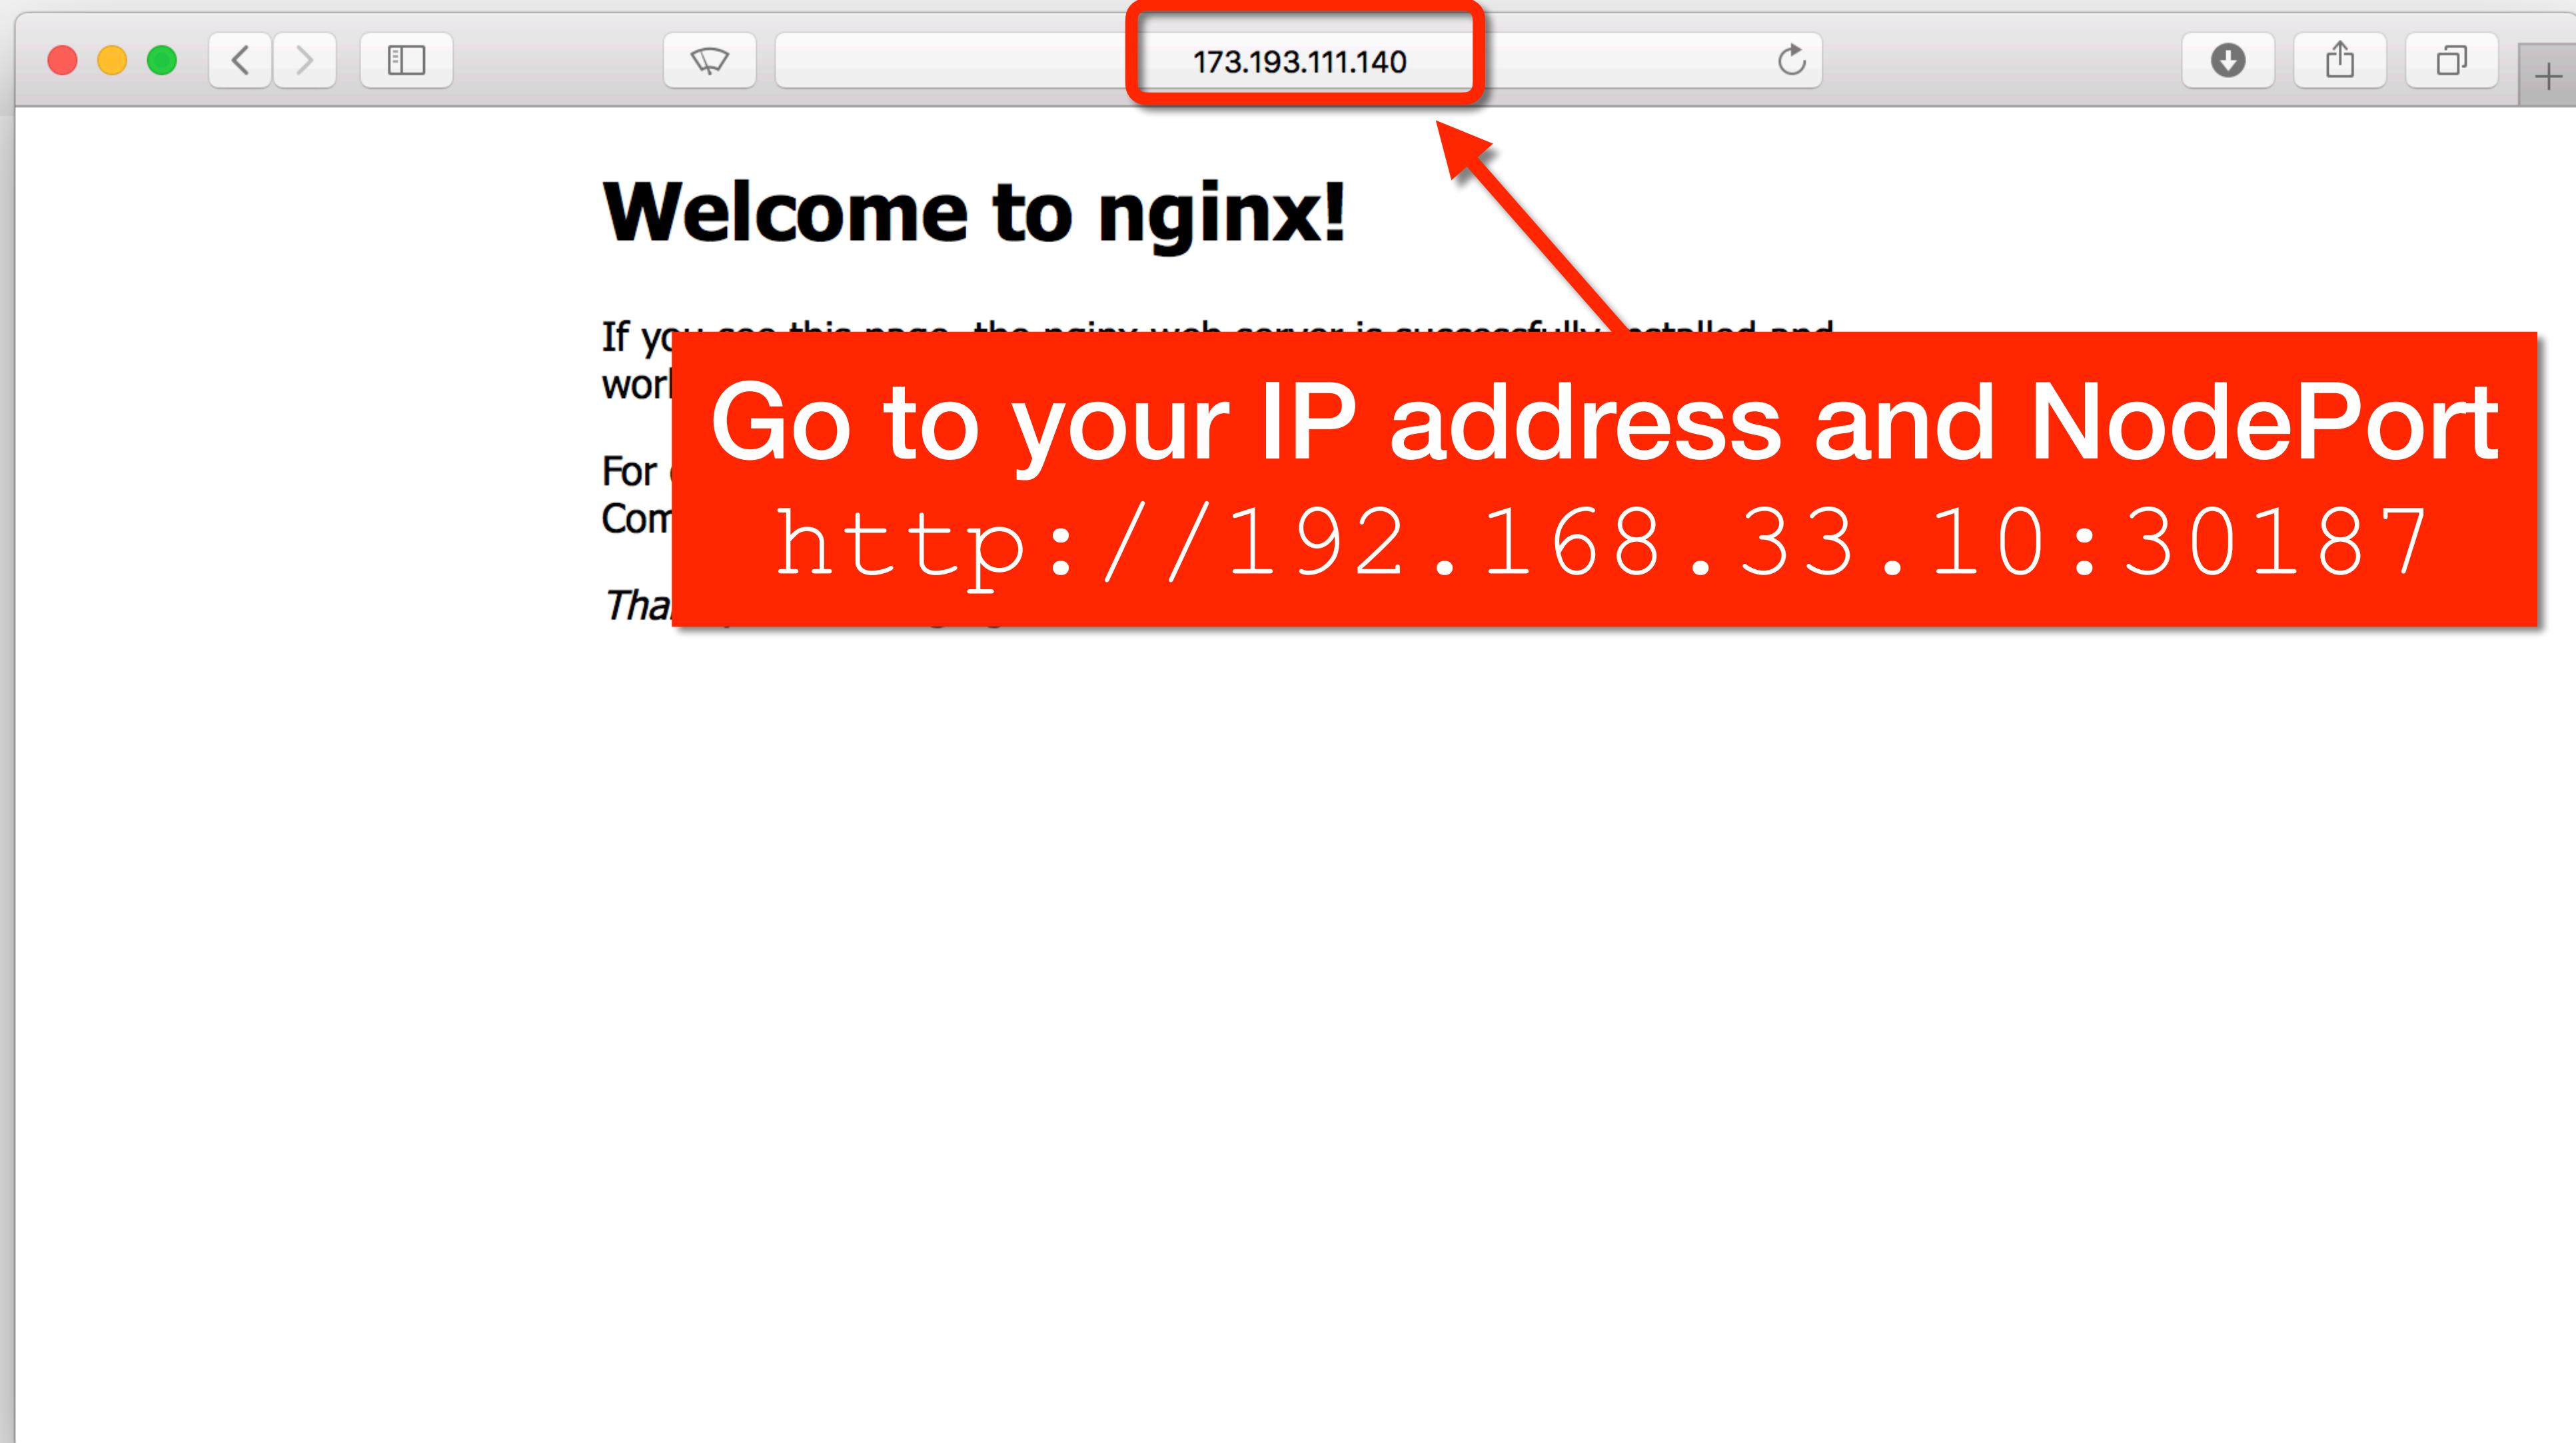
This is the NodePort that was assigned

# Get the Public IP address

The IP address of the Vagrant Virtual Machine is in the Vagrantfile

**192.168.33.10**

# Access NodePort



# Destroy Nginx

- Delete a Service

```
$ kubectl delete svc my-nginx  
service "my-nginx" deleted
```

- Delete the deployment

```
$ kubectl delete deploy my-nginx  
deployment.apps "my-nginx" deleted
```

# Let's Deploy our Project to Kubernetes



# Build our Project as a Docker Image

- First let's build our project as a Docker image

```
$ cd /vagrant  
$ docker build -t hitcounter:1.0 .
```

# Push our Image to the Microk8s registry

- First let's tag the Docker image we just built

```
$ docker tag hitcounter:1.0 localhost:32000(hitcounter:1.0
```

- Then we push it to the remote registry

```
$ docker push localhost:32000(hitcounter:1.0
```

# Create A Redis Service\*

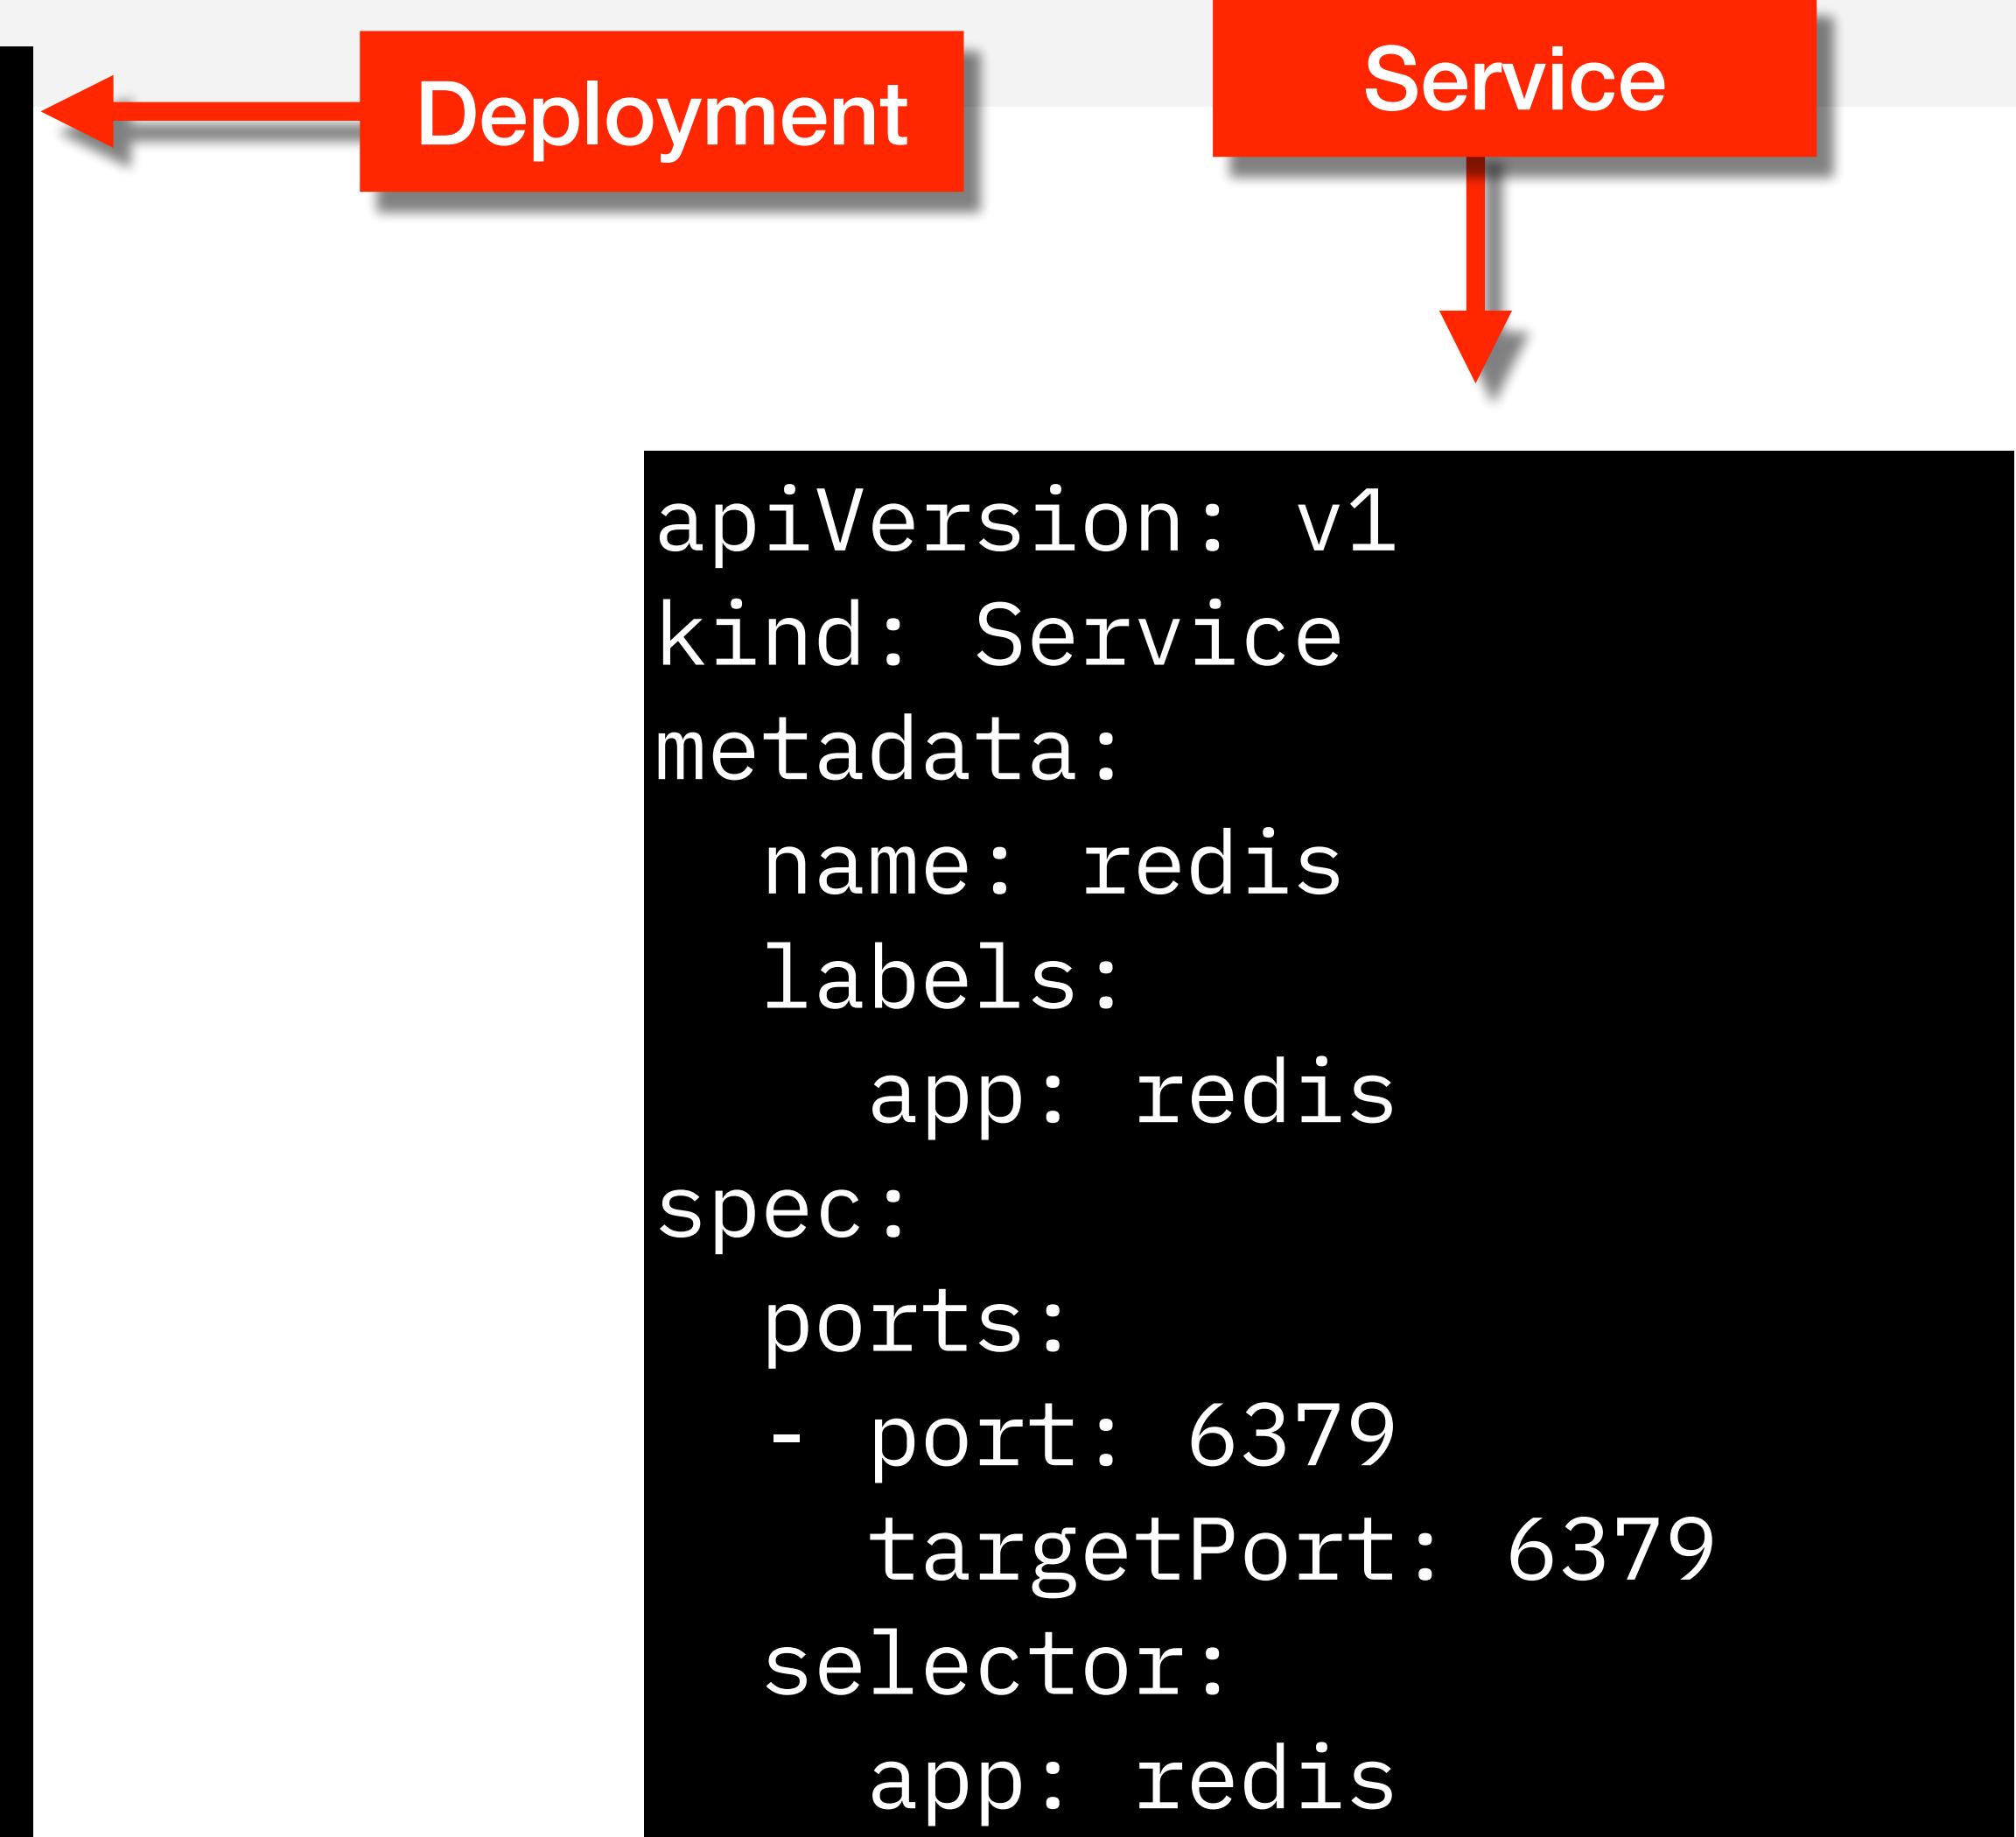
- The demo application requires Redis to store it's state
- We can deploy a Redis container from the `redis.yaml` file in the `./deploy` folder of the current repository

```
$ kubectl apply -f deploy/redis.yaml
deployment.apps/redis created
service/redis created
```

\* Note: in a "real" deployment we would attach a volume to the Redis container to persist it's data beyond the life of the container

# Redis.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis
  labels:
    app: redis
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: redis
          image: "redis:alpine"
          ports:
            - containerPort: 6379
              protocol: TCP
```



# View the Logs from Redis

- You can view the logs of the pod to see how the service is doing

```
$ kubectl logs redis-784d69fbfd-qp5p5

1:C 01 Dec 2019 20:54:39.303 # o000o000o000o Redis is starting o000o000o000o
1:C 01 Dec 2019 20:54:39.303 # Redis version=5.0.6, bits=64, commit=00000000, modified=0, pid=1, just started
1:C 01 Dec 2019 20:54:39.303 # Warning: no config file specified, using the default config. In order to specify a
config file use redis-server /path/to/redis.conf
1:M 01 Dec 2019 20:54:39.305 * Running mode=standalone, port=6379.
1:M 01 Dec 2019 20:54:39.305 # WARNING: The TCP backlog setting of 511 cannot be enforced because /proc/sys/net/core/
somaxconn is set to the lower value of 128.
1:M 01 Dec 2019 20:54:39.305 # Server initialized
1:M 01 Dec 2019 20:54:39.305 # WARNING you have Transparent Huge Pages (THP) support enabled in your kernel. This will
create latency and memory usage issues with Redis. To fix this issue run the command 'echo never > /sys/kernel/mm/
transparent_hugepage/enabled' as root, and add it to your /etc/rc.local in order to retain the setting after a reboot.
Redis must be restarted after THP is disabled.
1:M 01 Dec 2019 20:54:39.305 * Ready to accept connections
```

# Create your Credentials Secret

- Let's use the command line method of creating a secret for DATABASE\_URI

```
$ export DATABASE_URI="redis://redis:6379"  
  
$ kubectl create secret generic redis-creds \  
--from-literal=uri=$DATABASE_URI  
  
secret/redis-creds created
```

# Check your Credentials Secret

- Let's see what was created for **There is your base64 encoded secret**

```
$ kubectl get secret redis-creds -o yaml
apiVersion: v1
data:
  uri: cmVkaXM6Ly86QHJlZGlz0jYzNzkvMA==
kind: Secret
metadata:
  creationTimestamp: "2019-11-20T03:53:22Z"
  name: redis-creds
  namespace: default
  resourceVersion: "3437"
  selfLink: /api/v1/namespaces/default/secrets/redis-creds
  uid: aee28dcf-f1f4-40ba-83fc-8b775c208264
type: Opaque
```

# Now we can Deploy Our Image

- You can use kubectl apply with the files under /vagrant/deploy

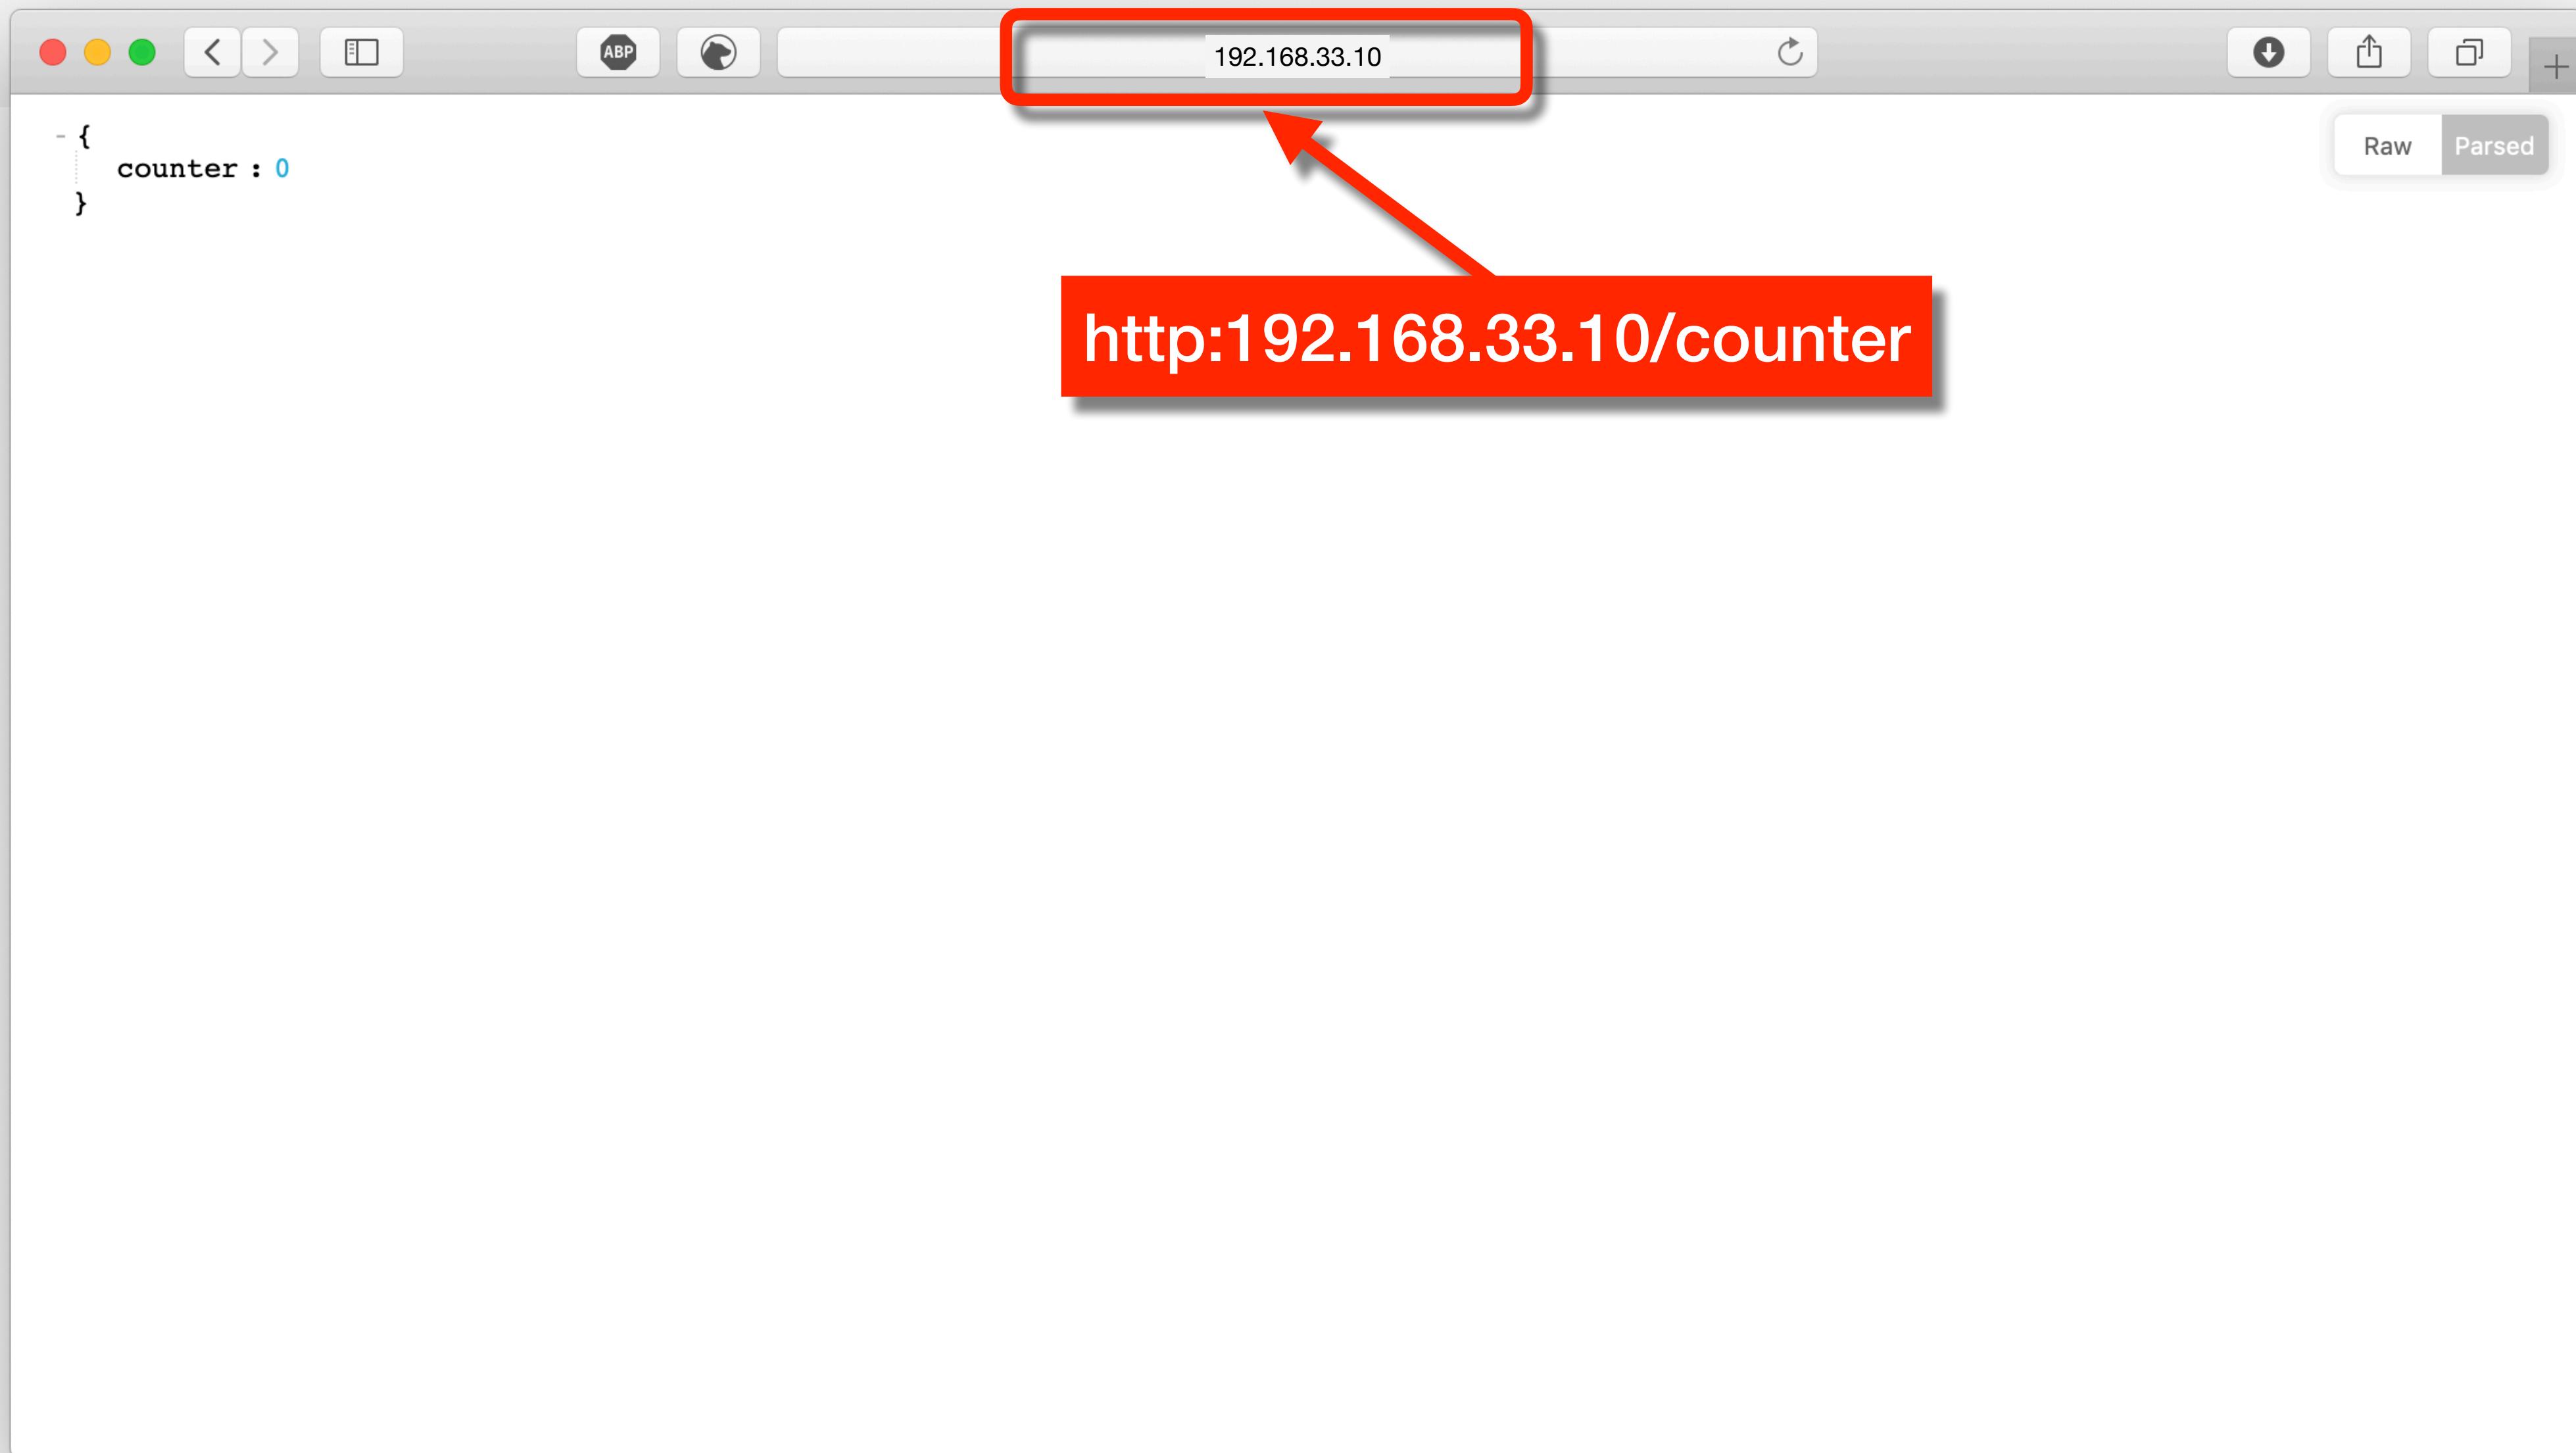
```
$ cd /vagrant/deploy
```

```
$ kubectl apply -f deployment.yaml
deployment.apps/hitcounter created
```

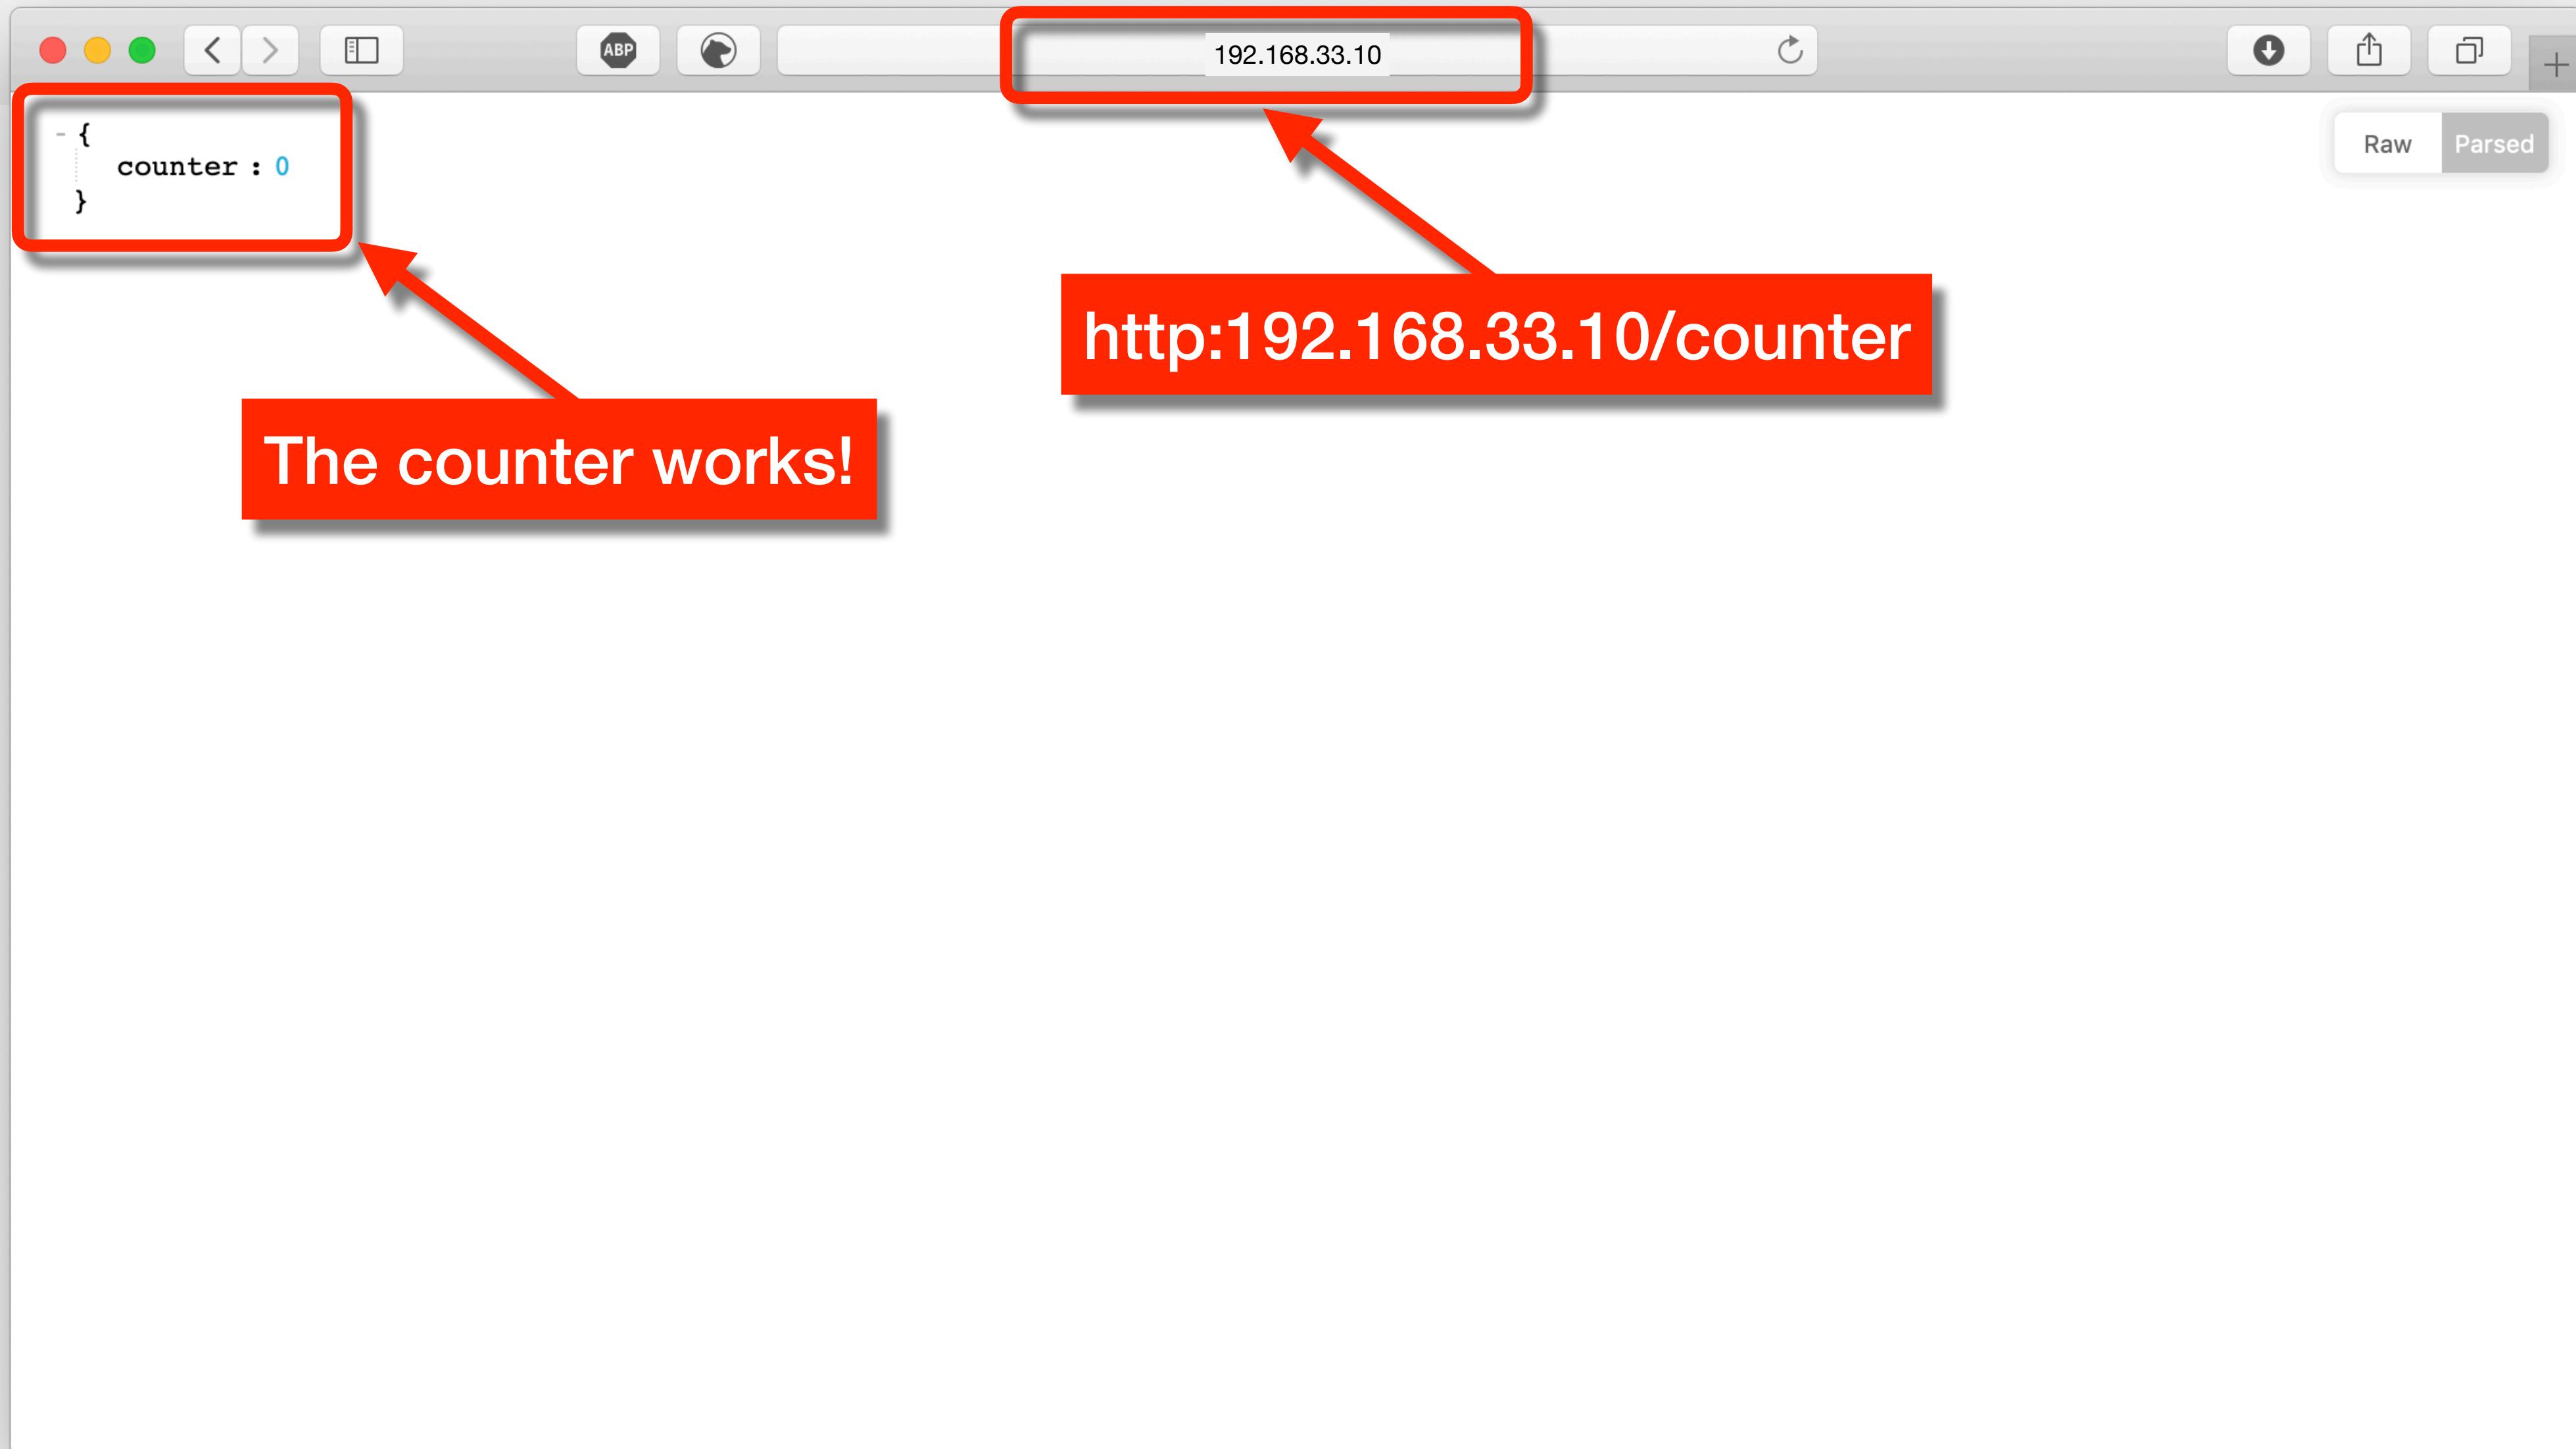
```
$ kubectl apply -f service.yaml
service/hitcounter-service created
```

```
$ kubectl apply -f ingress.yaml
ingress.extensions/hitcounter-ingress created
```

# Your App on Kubernetes



# Your App on Kubernetes

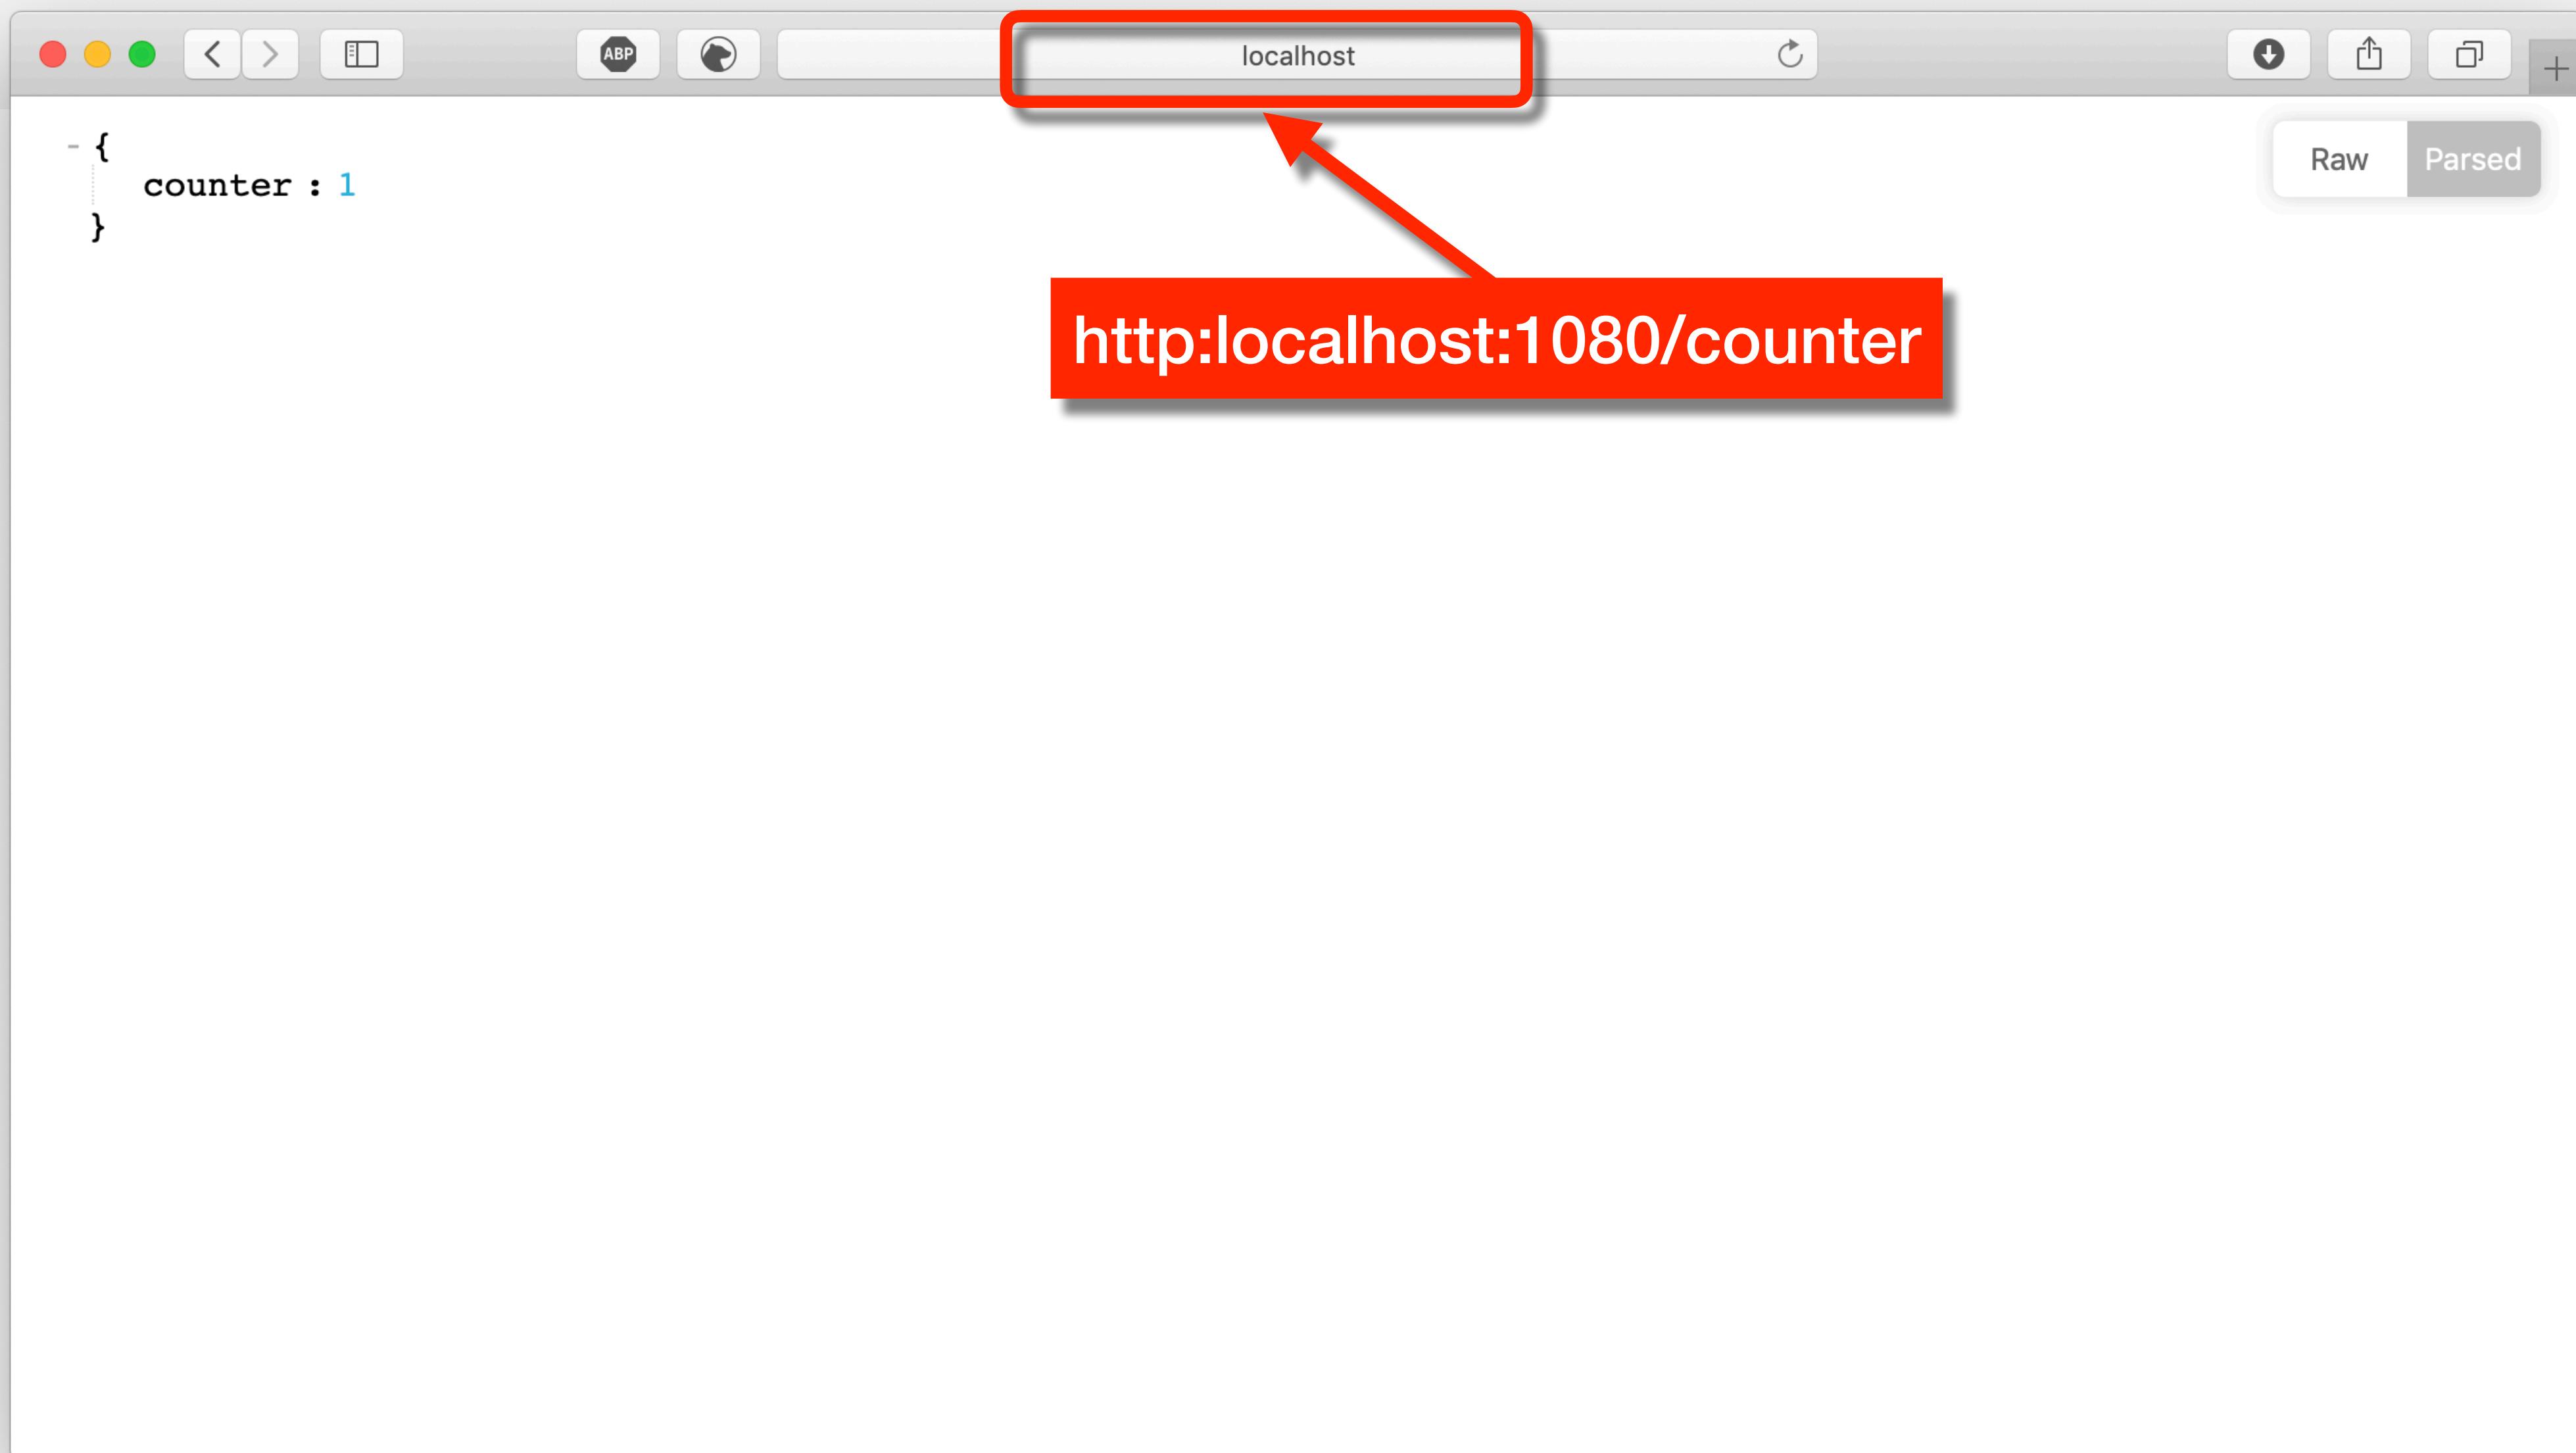


# Use HTTP command to increment the counter

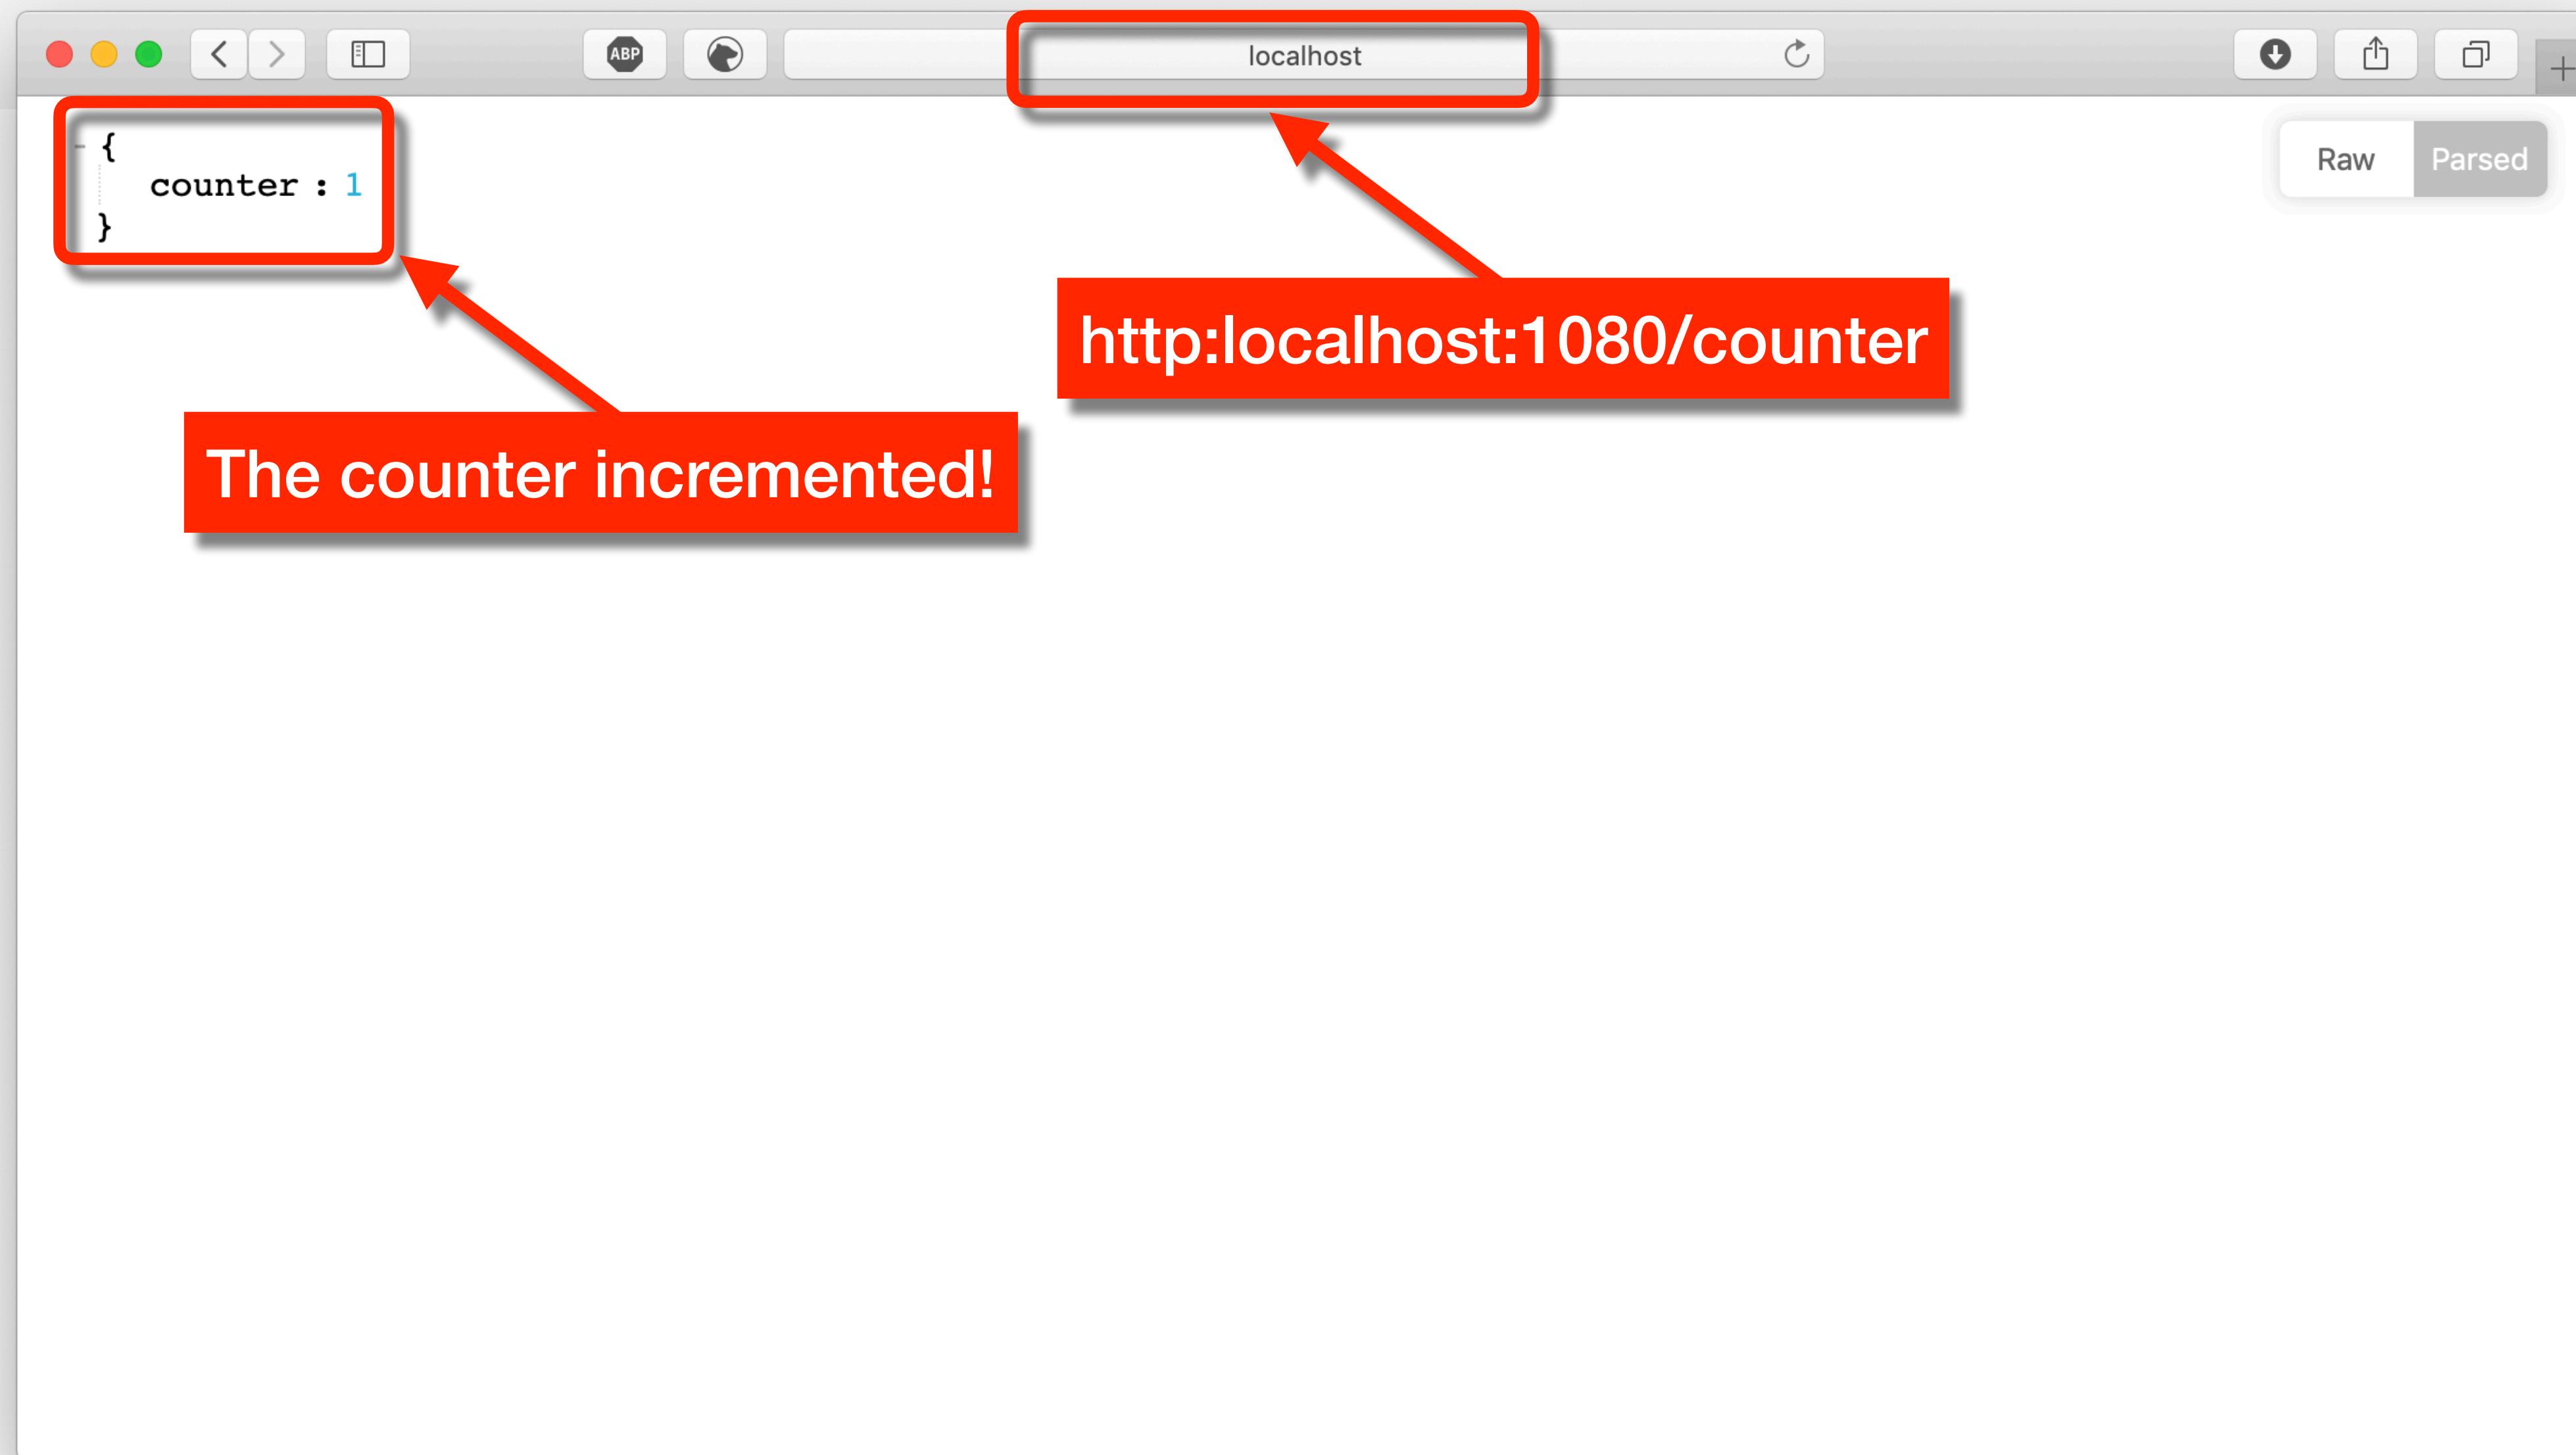
```
$ http post localhost/counters/foo
HTTP/1.1 201 CREATED
Connection: keep-alive
Content-Length: 14
Content-Type: application/json
Date: Wed, 20 Nov 2019 04:14:51 GMT
Server: openresty/1.15.8.2

{
  "counter": 1
  "Name": "foo"
}
```

# Check again



# Check again



# Summary

You just created your first Kubernetes deployment

You learned how to create deployments

You learned how to expose deployments as services

You learned how to scale deployments

You can now deploy your microservices on any Kubernetes cloud

