

# Introduction to Kubernetes



Instructor:

**John J Rofrano**

Senior Technical Staff Member, DevOps Champion

IBM T.J. Watson Research Center

[rofrano@us.ibm.com](mailto:rofrano@us.ibm.com) (@JohnRofrano)

# What will you learn?

- What is Kubernetes
- Why would you use a Kubernetes for deployment
- Kubernetes Architecture Overview
- How to deploy your own containers in Kubernetes





# Create your Kubernetes Cluster

- We will use Minikube to create a Kubernetes Cluster in our VM

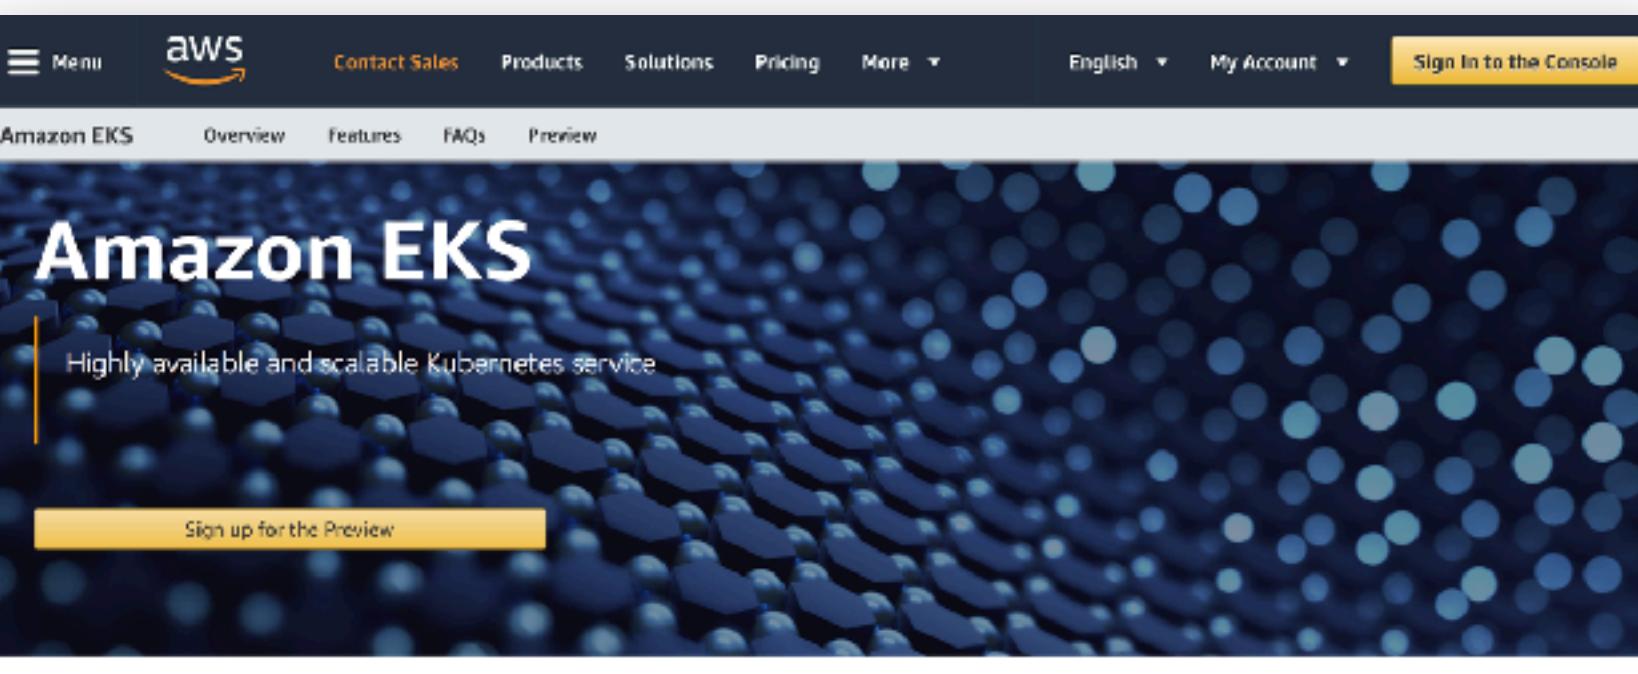
```
$ export CHANGE_MINIKUBE_NONE_USER=true  
  
$ sudo minikube start --vm-driver=none
```

# What is Kubernetes?



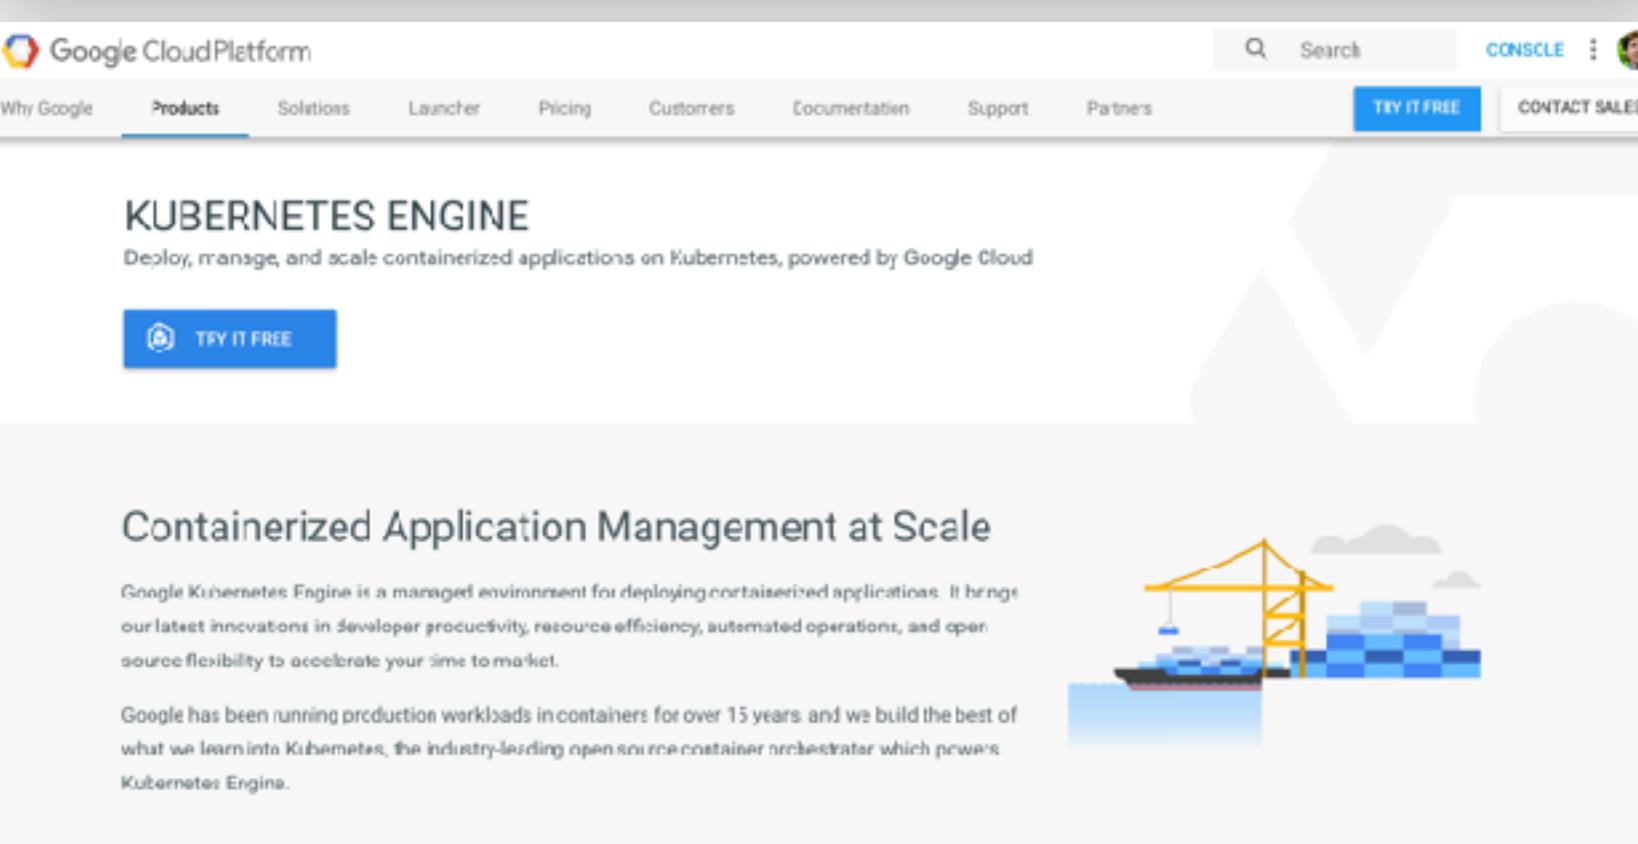
# Kubernetes is the de facto Orchestrator for Containers

Every major cloud provider support Kubernetes

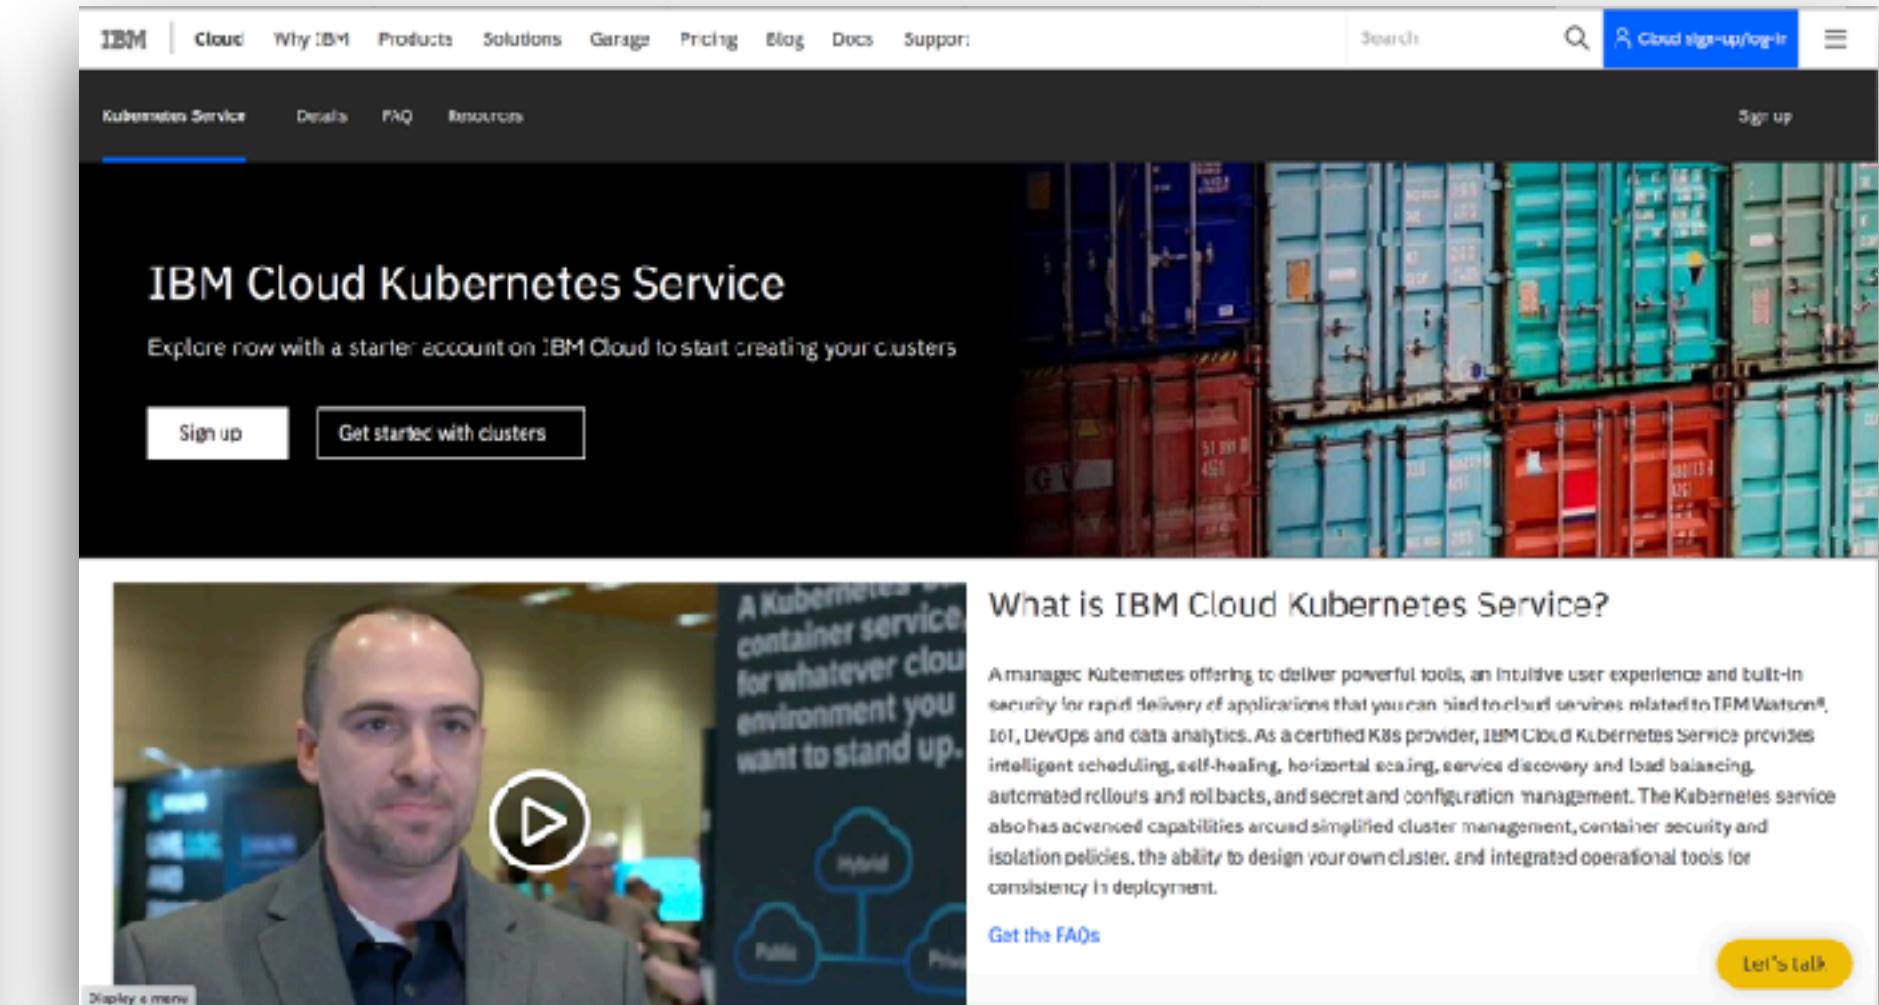


The screenshot shows the Amazon EKS (Amazon Elastic Container Service for Kubernetes) homepage. It features a dark blue background with a grid of blue dots. The title "Amazon EKS" is prominently displayed. Below it, the text "Highly available and scalable Kubernetes service" is visible. A yellow button labeled "Sign up for the Preview" is at the bottom.

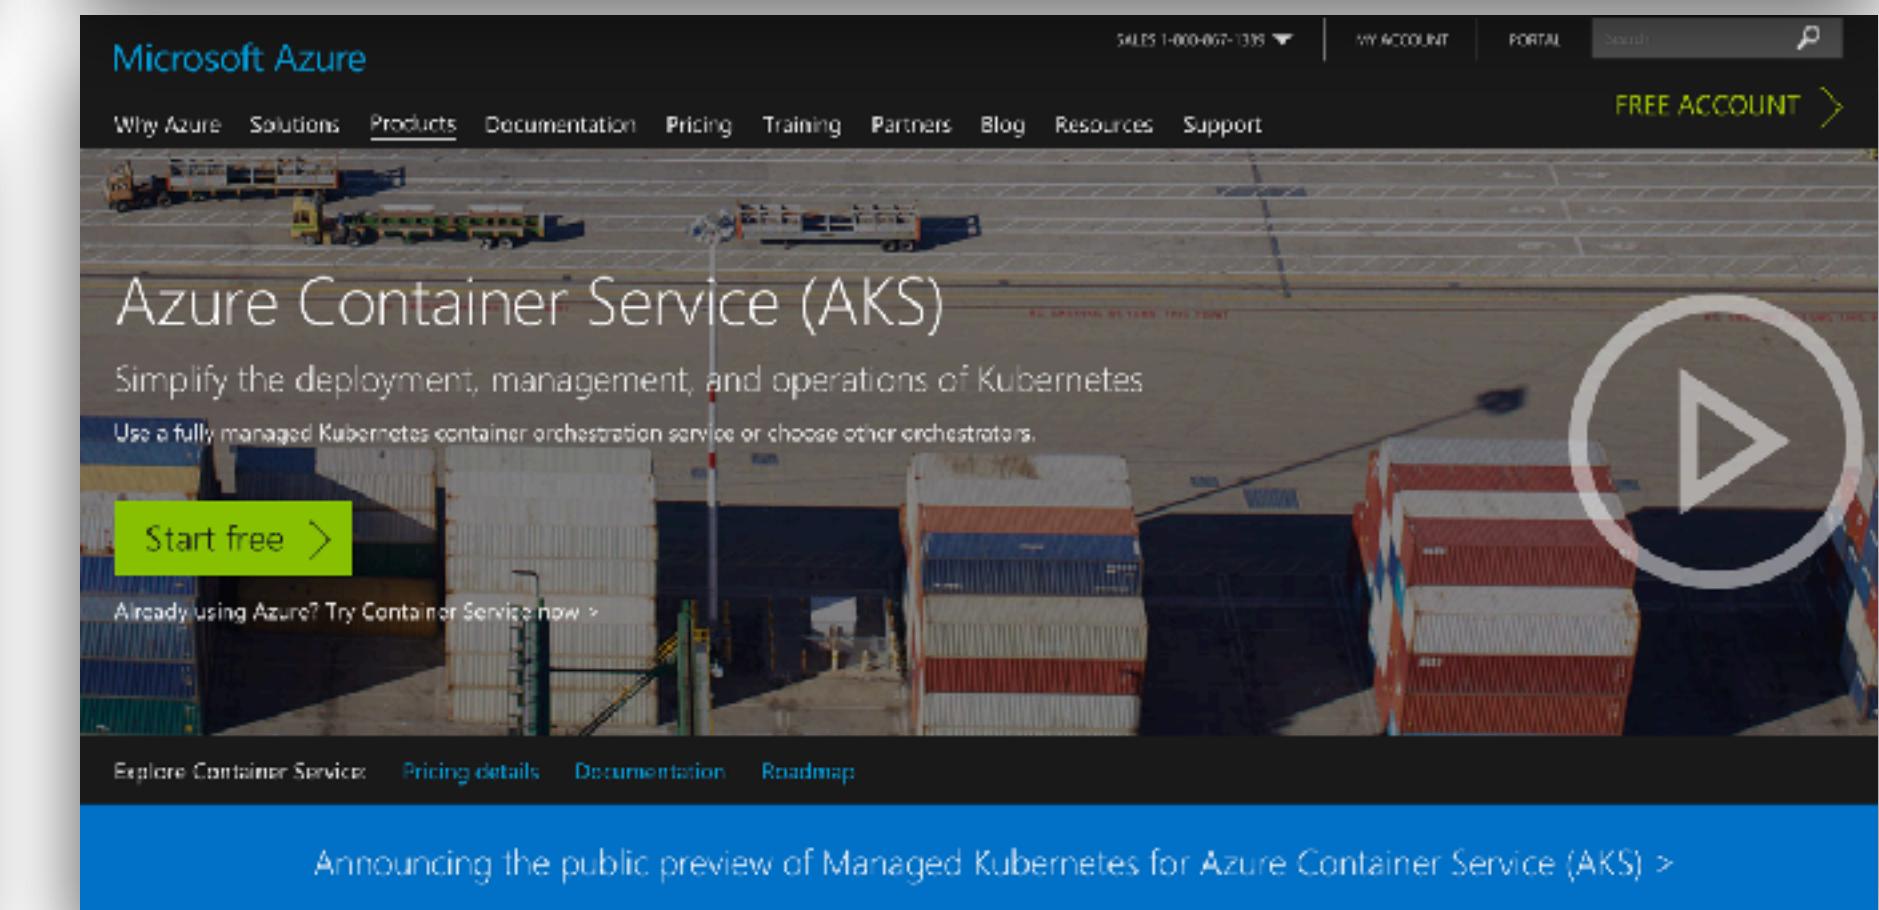
Amazon Elastic Container Service for Kubernetes (Amazon EKS) is a managed service that makes it easy for you to run Kubernetes on AWS without needing to install and operate your own Kubernetes clusters. Kubernetes is an open-source system for automating the deployment, scaling, and management of containerized applications. Operating Kubernetes for production applications presents a number of challenges. You need to manage the scaling and availability of your Kubernetes masters and persistence layer by ensuring that you have chosen appropriate instance types, running them across multiple Availability Zones, monitoring their health, and replacing unhealthy nodes. You need to patch and upgrade your masters and worker nodes to ensure that you are running the latest version of Kubernetes. This all requires expertise and a lot of manual work. With Amazon EKS, upgrades and high availability are managed for you by AWS. Amazon EKS runs three Kubernetes masters across three Availability Zones in order to ensure high availability. Amazon EKS automatically detects and replaces unhealthy masters, and it provides automated version upgrades and patching for the masters.



The screenshot shows the Google Cloud Platform Kubernetes Engine page. The title "KUBERNETES ENGINE" is at the top, followed by the subtext "Deploy, manage, and scale containerized applications on Kubernetes, powered by Google Cloud". A "TRY IT FREE" button is present. Below this, the section "Containerized Application Management at Scale" is shown, featuring a small illustration of a construction crane over clouds.



The screenshot shows the IBM Cloud Kubernetes Service page. The title "IBM Cloud Kubernetes Service" is at the top. Below it, a video player shows a man speaking. To the right of the video, the text "What is IBM Cloud Kubernetes Service?" is followed by a detailed description of the service's features, including intelligent scheduling, self-healing, horizontal scaling, service discovery, and load balancing. A "Get the FAQs" button is at the bottom right.



The screenshot shows the Microsoft Azure Container Service (AKS) page. The title "Azure Container Service (AKS)" is at the top. Below it, the text "Simplify the deployment, management, and operations of Kubernetes" is followed by a "Start free" button. A large play button icon is on the right side of the page.

# Why choose Kubernetes?

- An open-source system for automating deployment, scaling, and management of containerized applications.
- No Vendor Lock-In
- Large Community of support
- Robust platform for container orchestration
- Based on 15 years of experience of running production workloads at Google, combined with best-of-breed ideas and practices from the community



# Kubernetes is an Orchestration Platform

- **Scheduling**  
decide where containers run
  - **Lifecycle and Health**  
keep containers running and restart them if they fail
  - **Scaling**  
grow and shrink deployments as needed
  - **Naming and Discovery**  
help containers find each other
  - **Load Balancing**  
distribute traffic across containers
- ...and a whole lot more



# Why Do You Need Container Orchestration?

- Deploy applications to servers without worrying about specific servers
- Scale the application horizontally up and down
- Restore the application if the server on which it worked fails
  - This is called container auto-healing or rescheduling



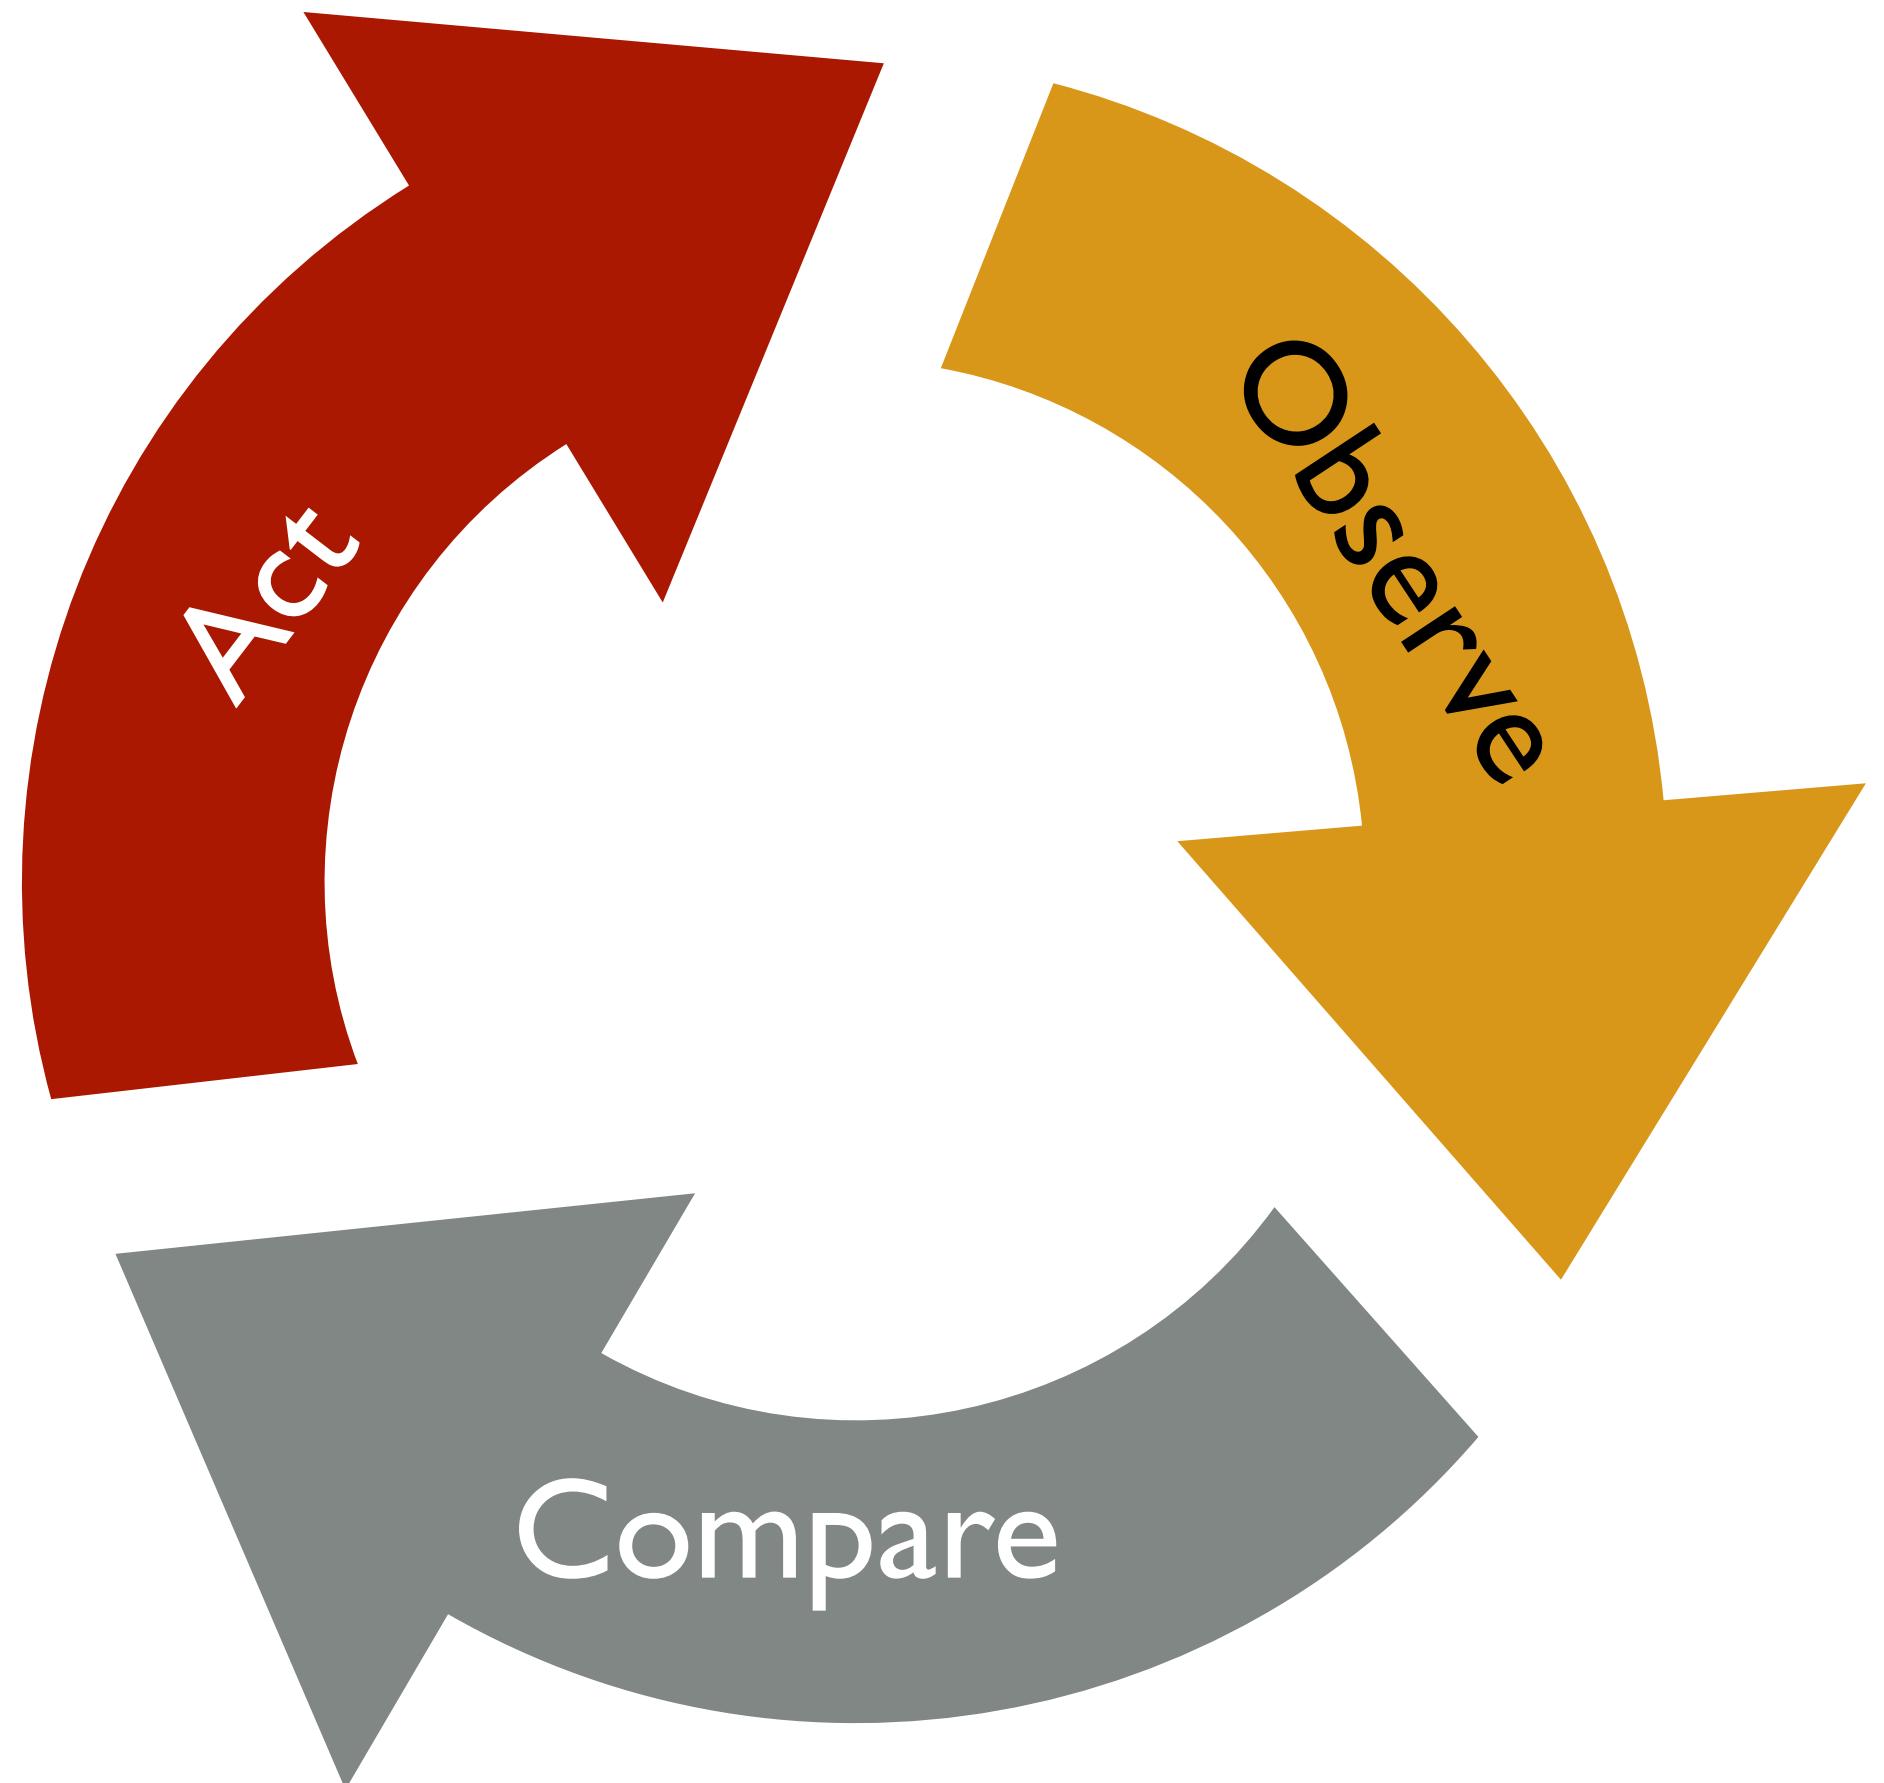
# Kubernetes has a Declarative API

Kubernetes is a Declarative Model

You express the desired state

Kubernetes maintains it

What could be simpler?

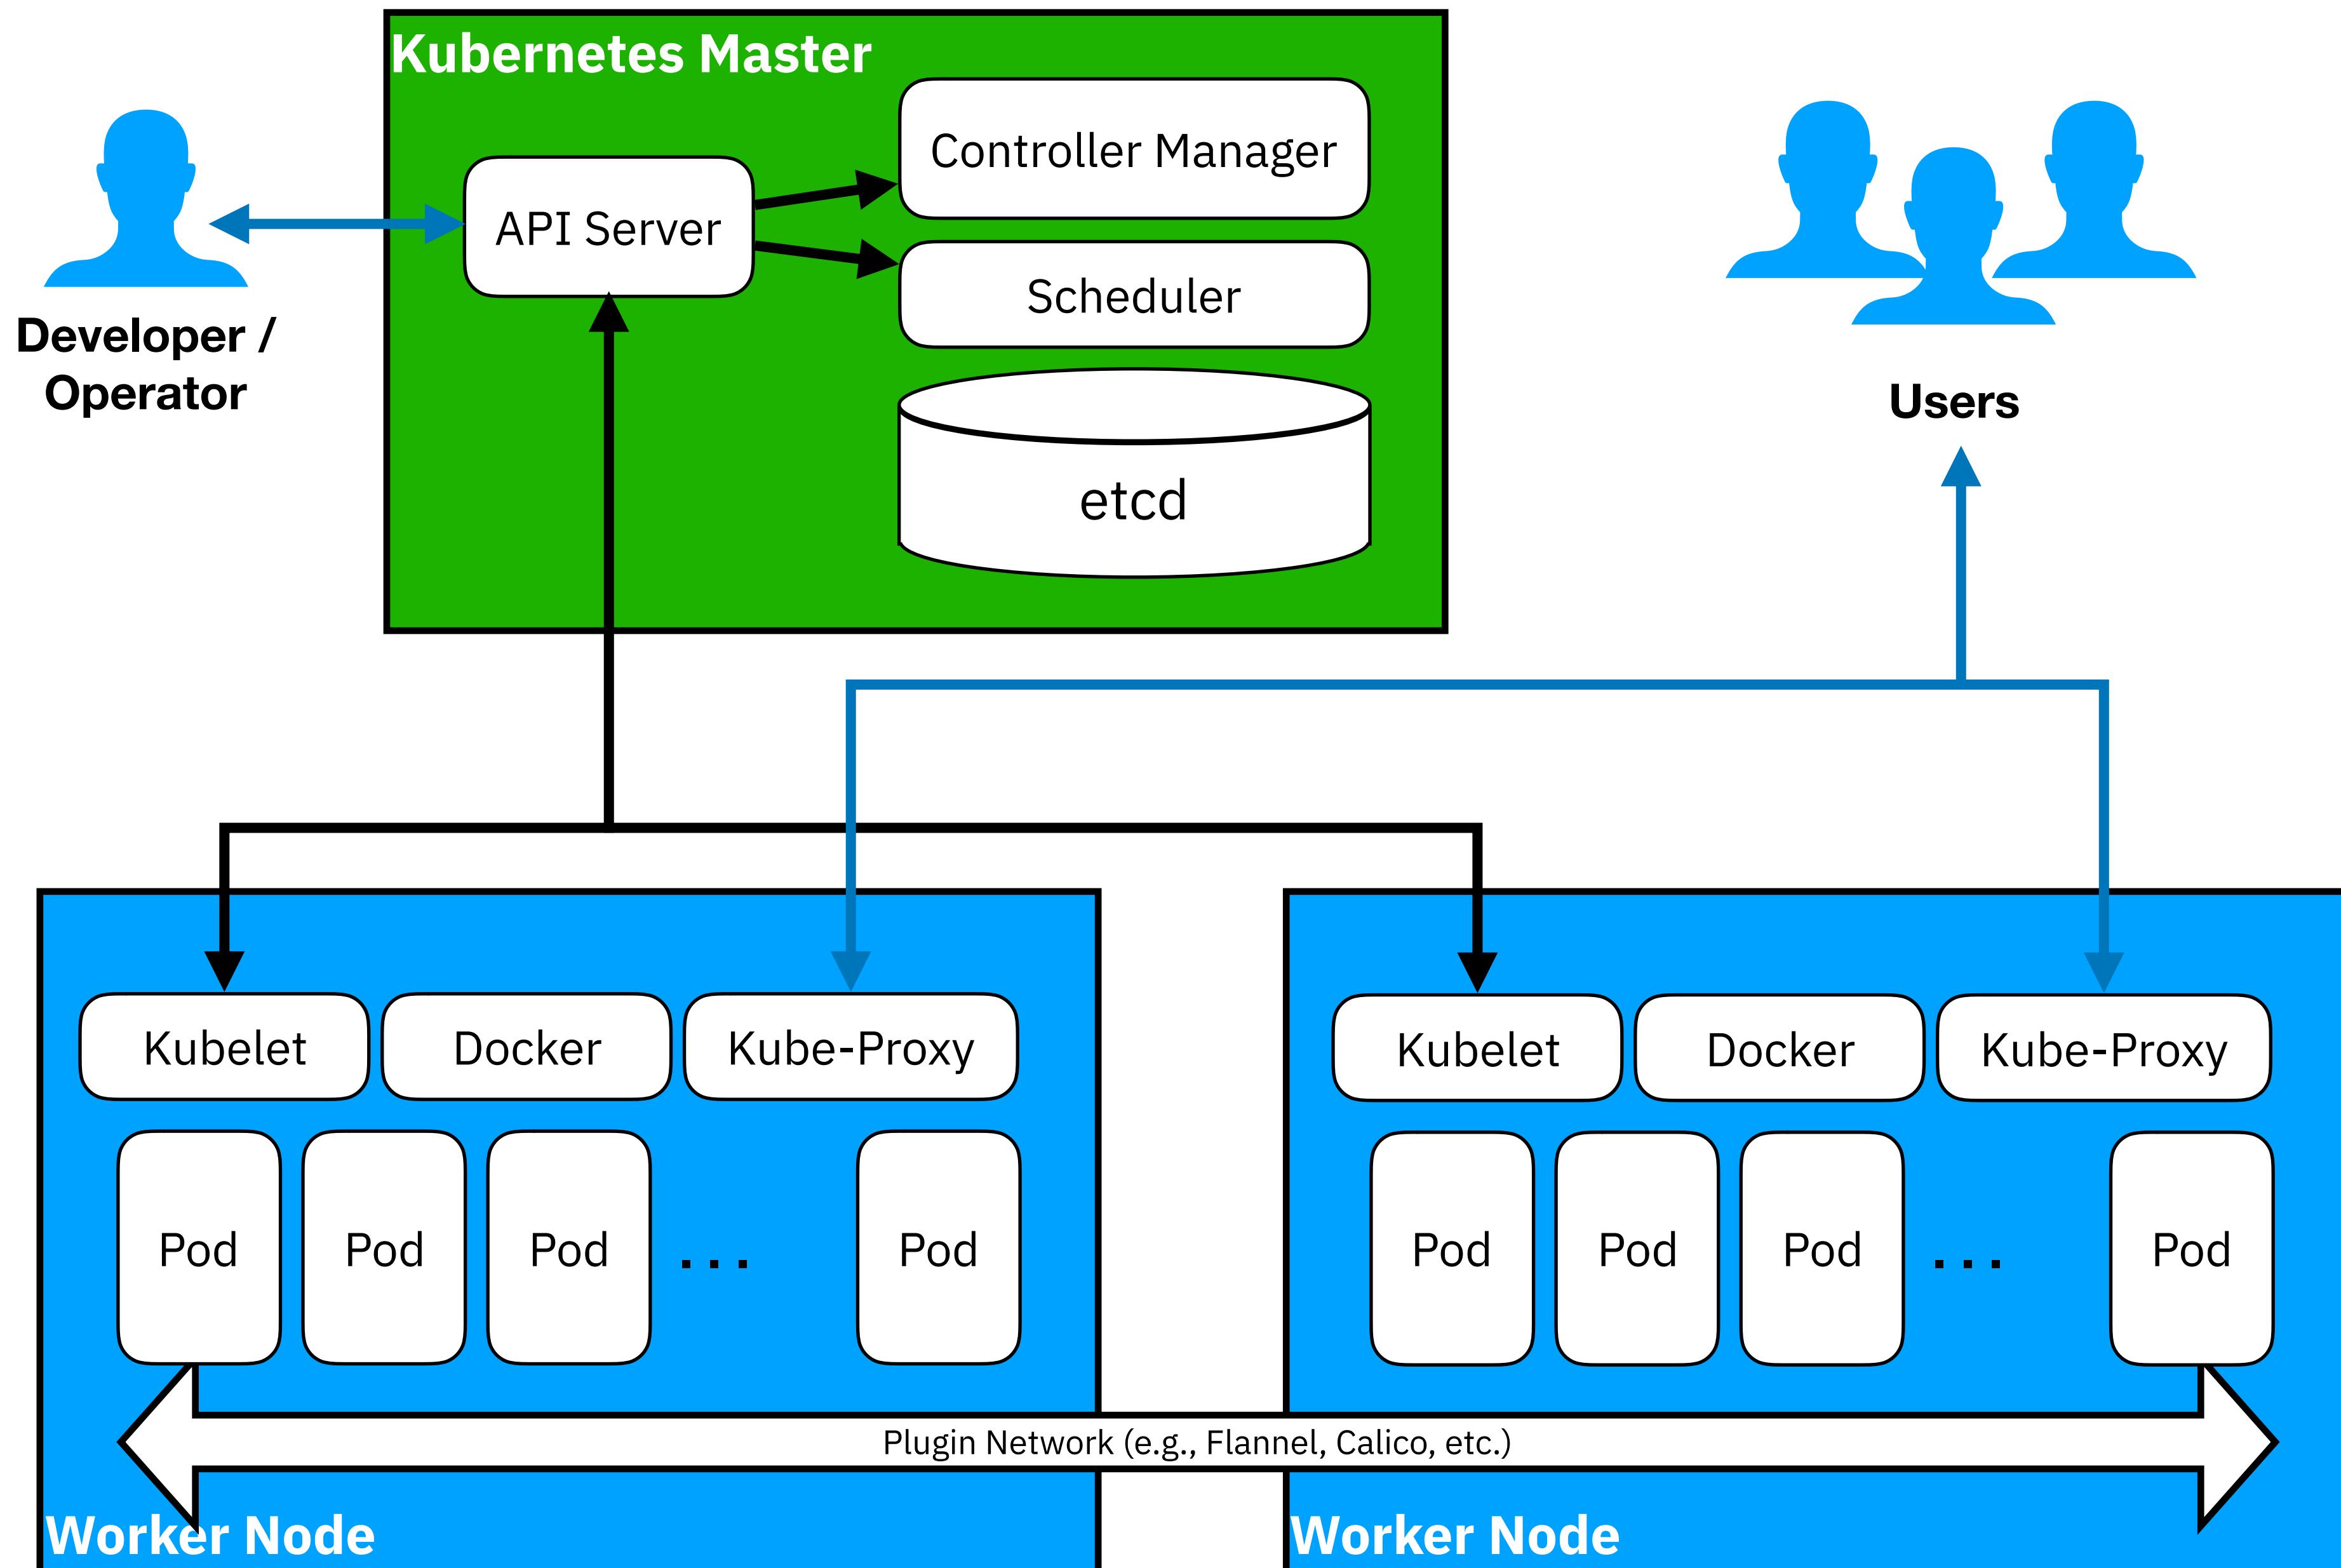


# Kubernetes Architecture



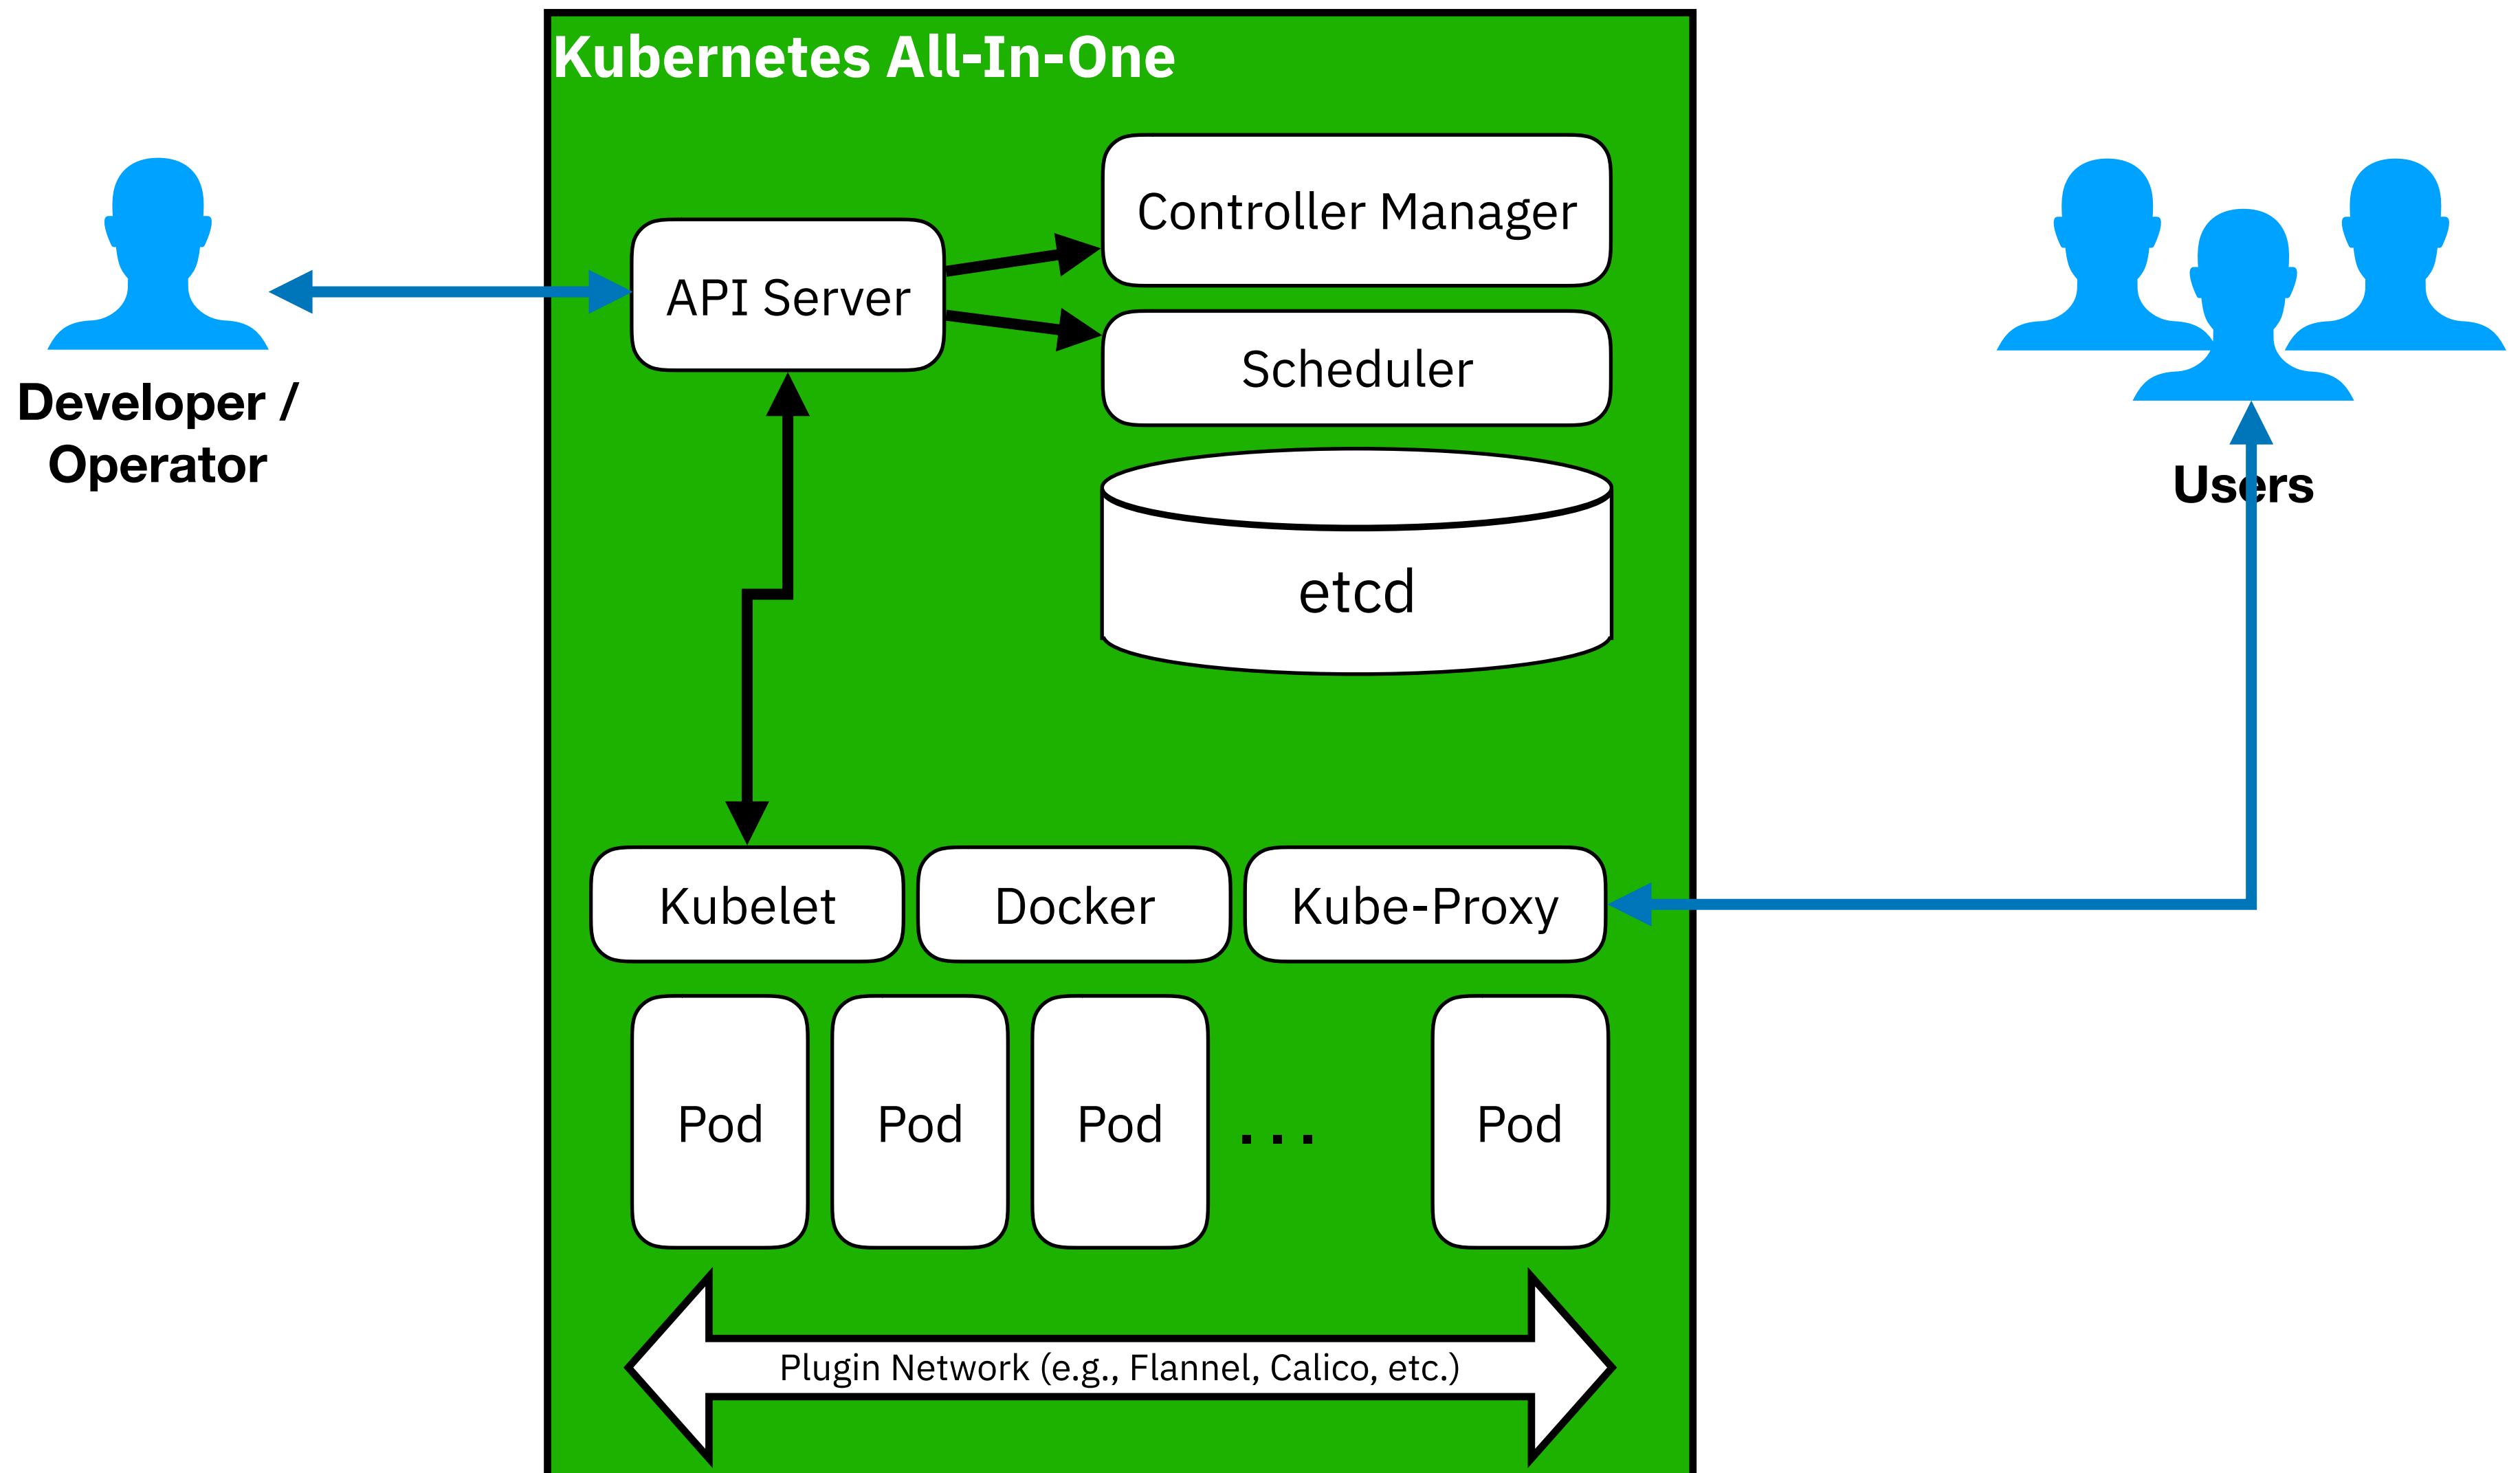
# Kubernetes Architecture

- There can be one or more Master Nodes (HA)
- There can be zero or more Worker Nodes
- Everyone communicates via the API Server
- Kubelet on each Worker Node acting as an agent
- Everything can be on one node for development use



# Kubernetes All-In-One

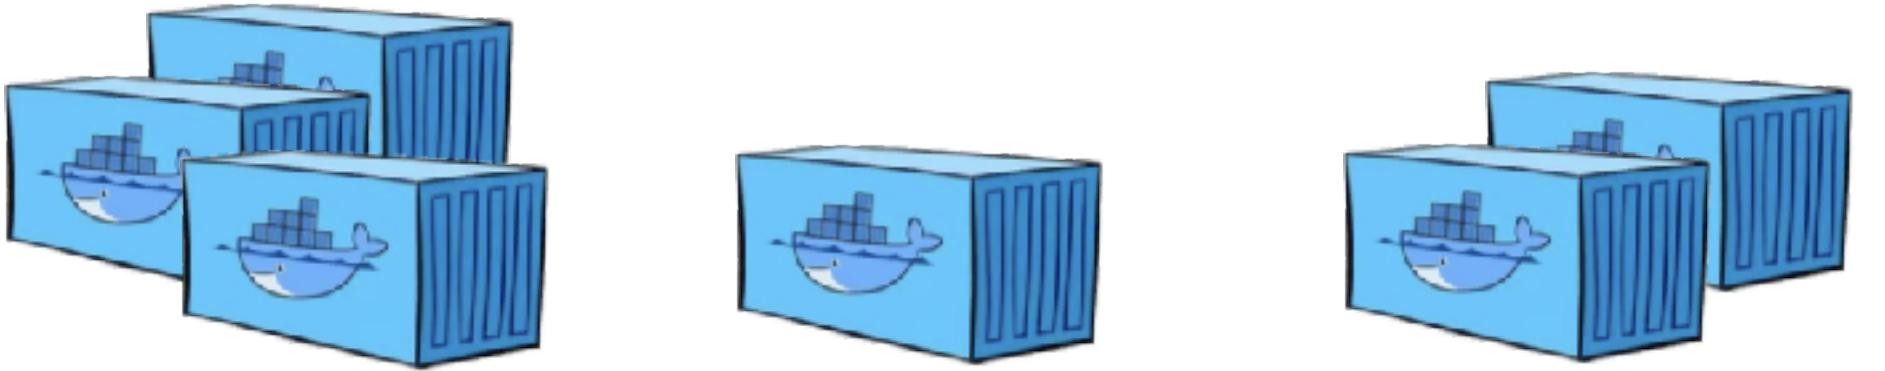
- You can run Kubernetes all in one VM for development work (not for production!)
- **MiniKube** is great for setting this up
- **Minishift** will deploy a development version of RedHat OpenShift (Origin) which uses Kubernetes



# Kubernetes is the new "Cloud OS"

## Managing Containers with VMs

- SysAdmins must decide where to place containers
- Workload balancing is manual



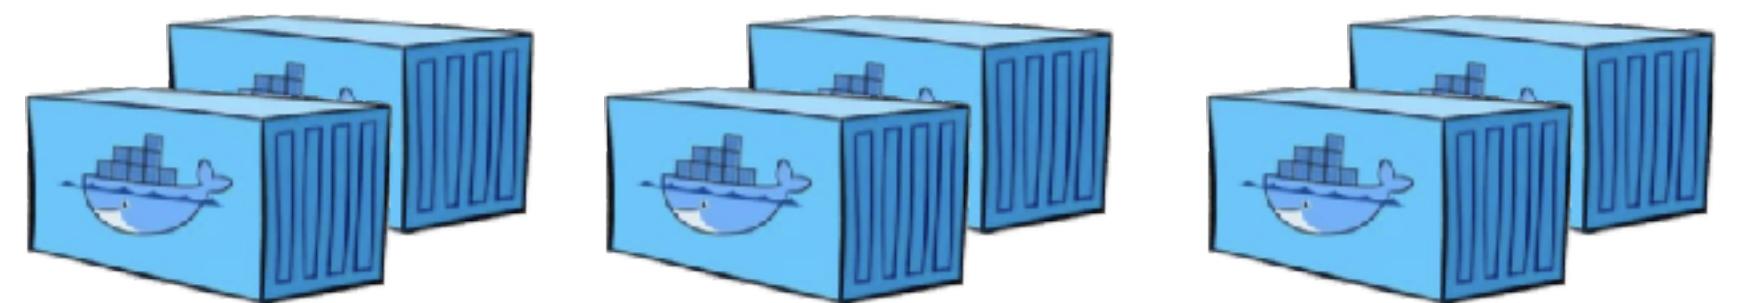
Virtual  
Machine

Virtual  
Machine

Virtual  
Machine

## Managing Containers with Kubernetes

- Kubernetes schedules and optimizes workloads automatically



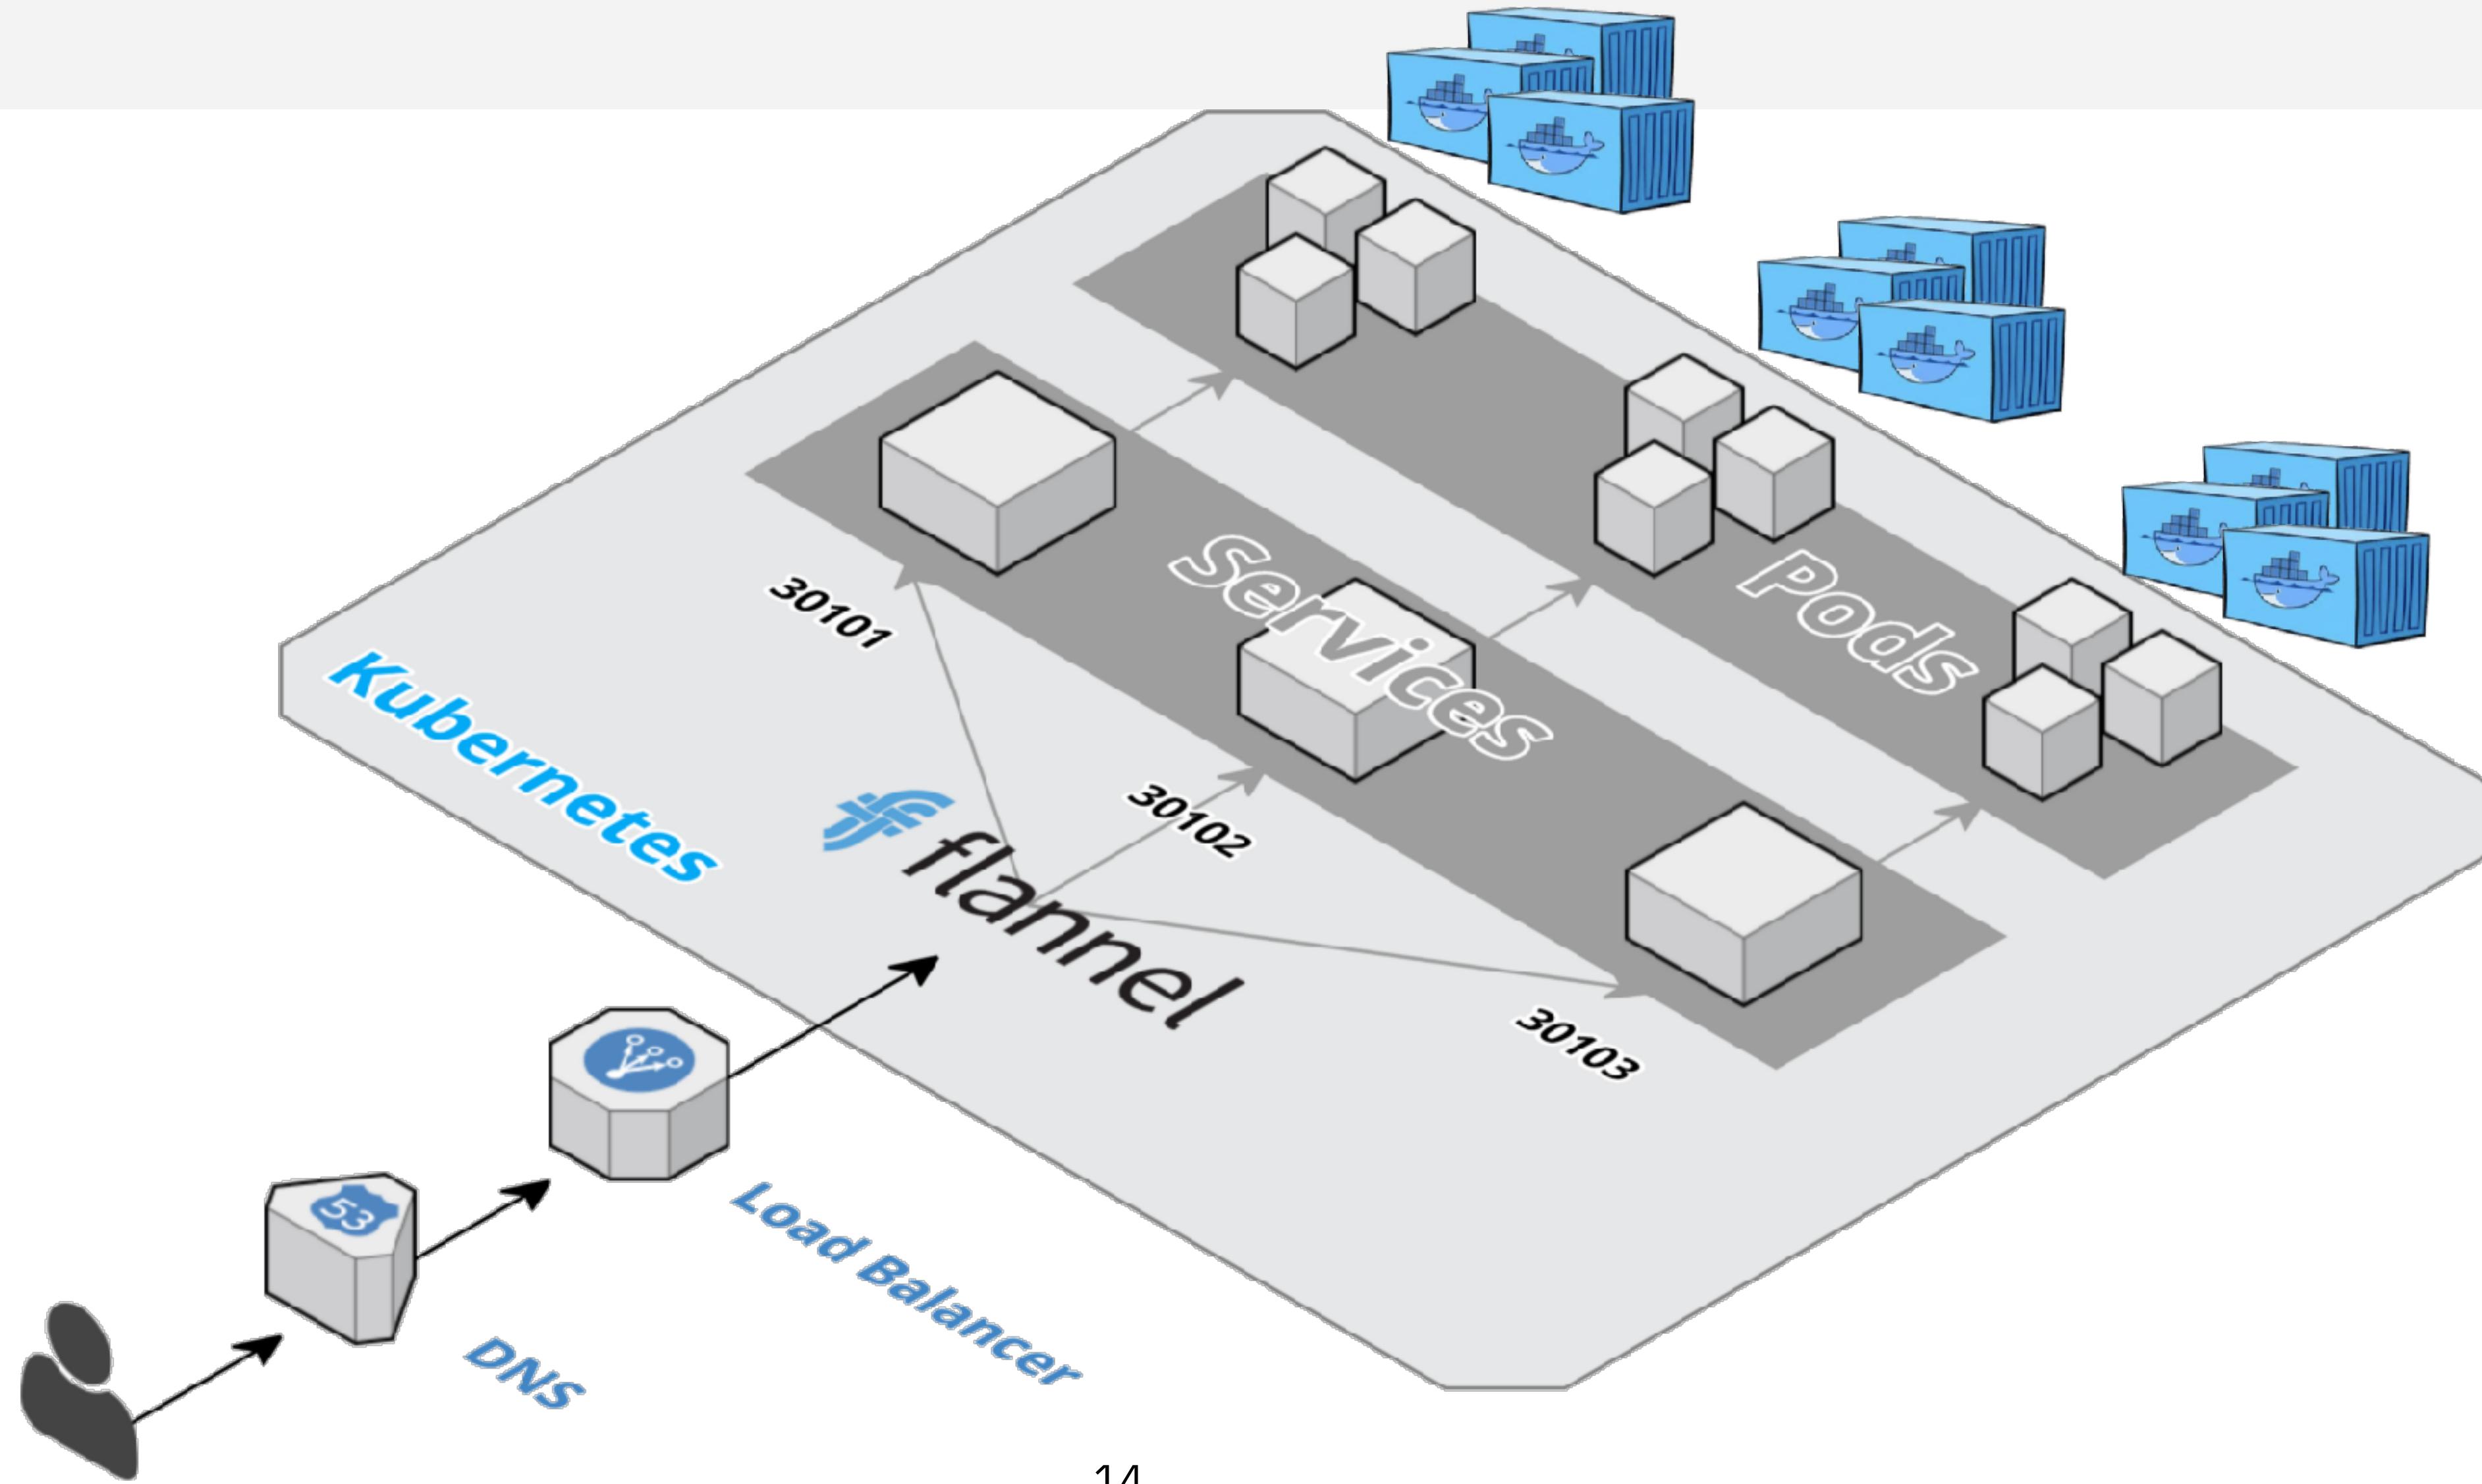
Kubernetes

Master VM

Node VM

Node VM

# Docker with Kubernetes



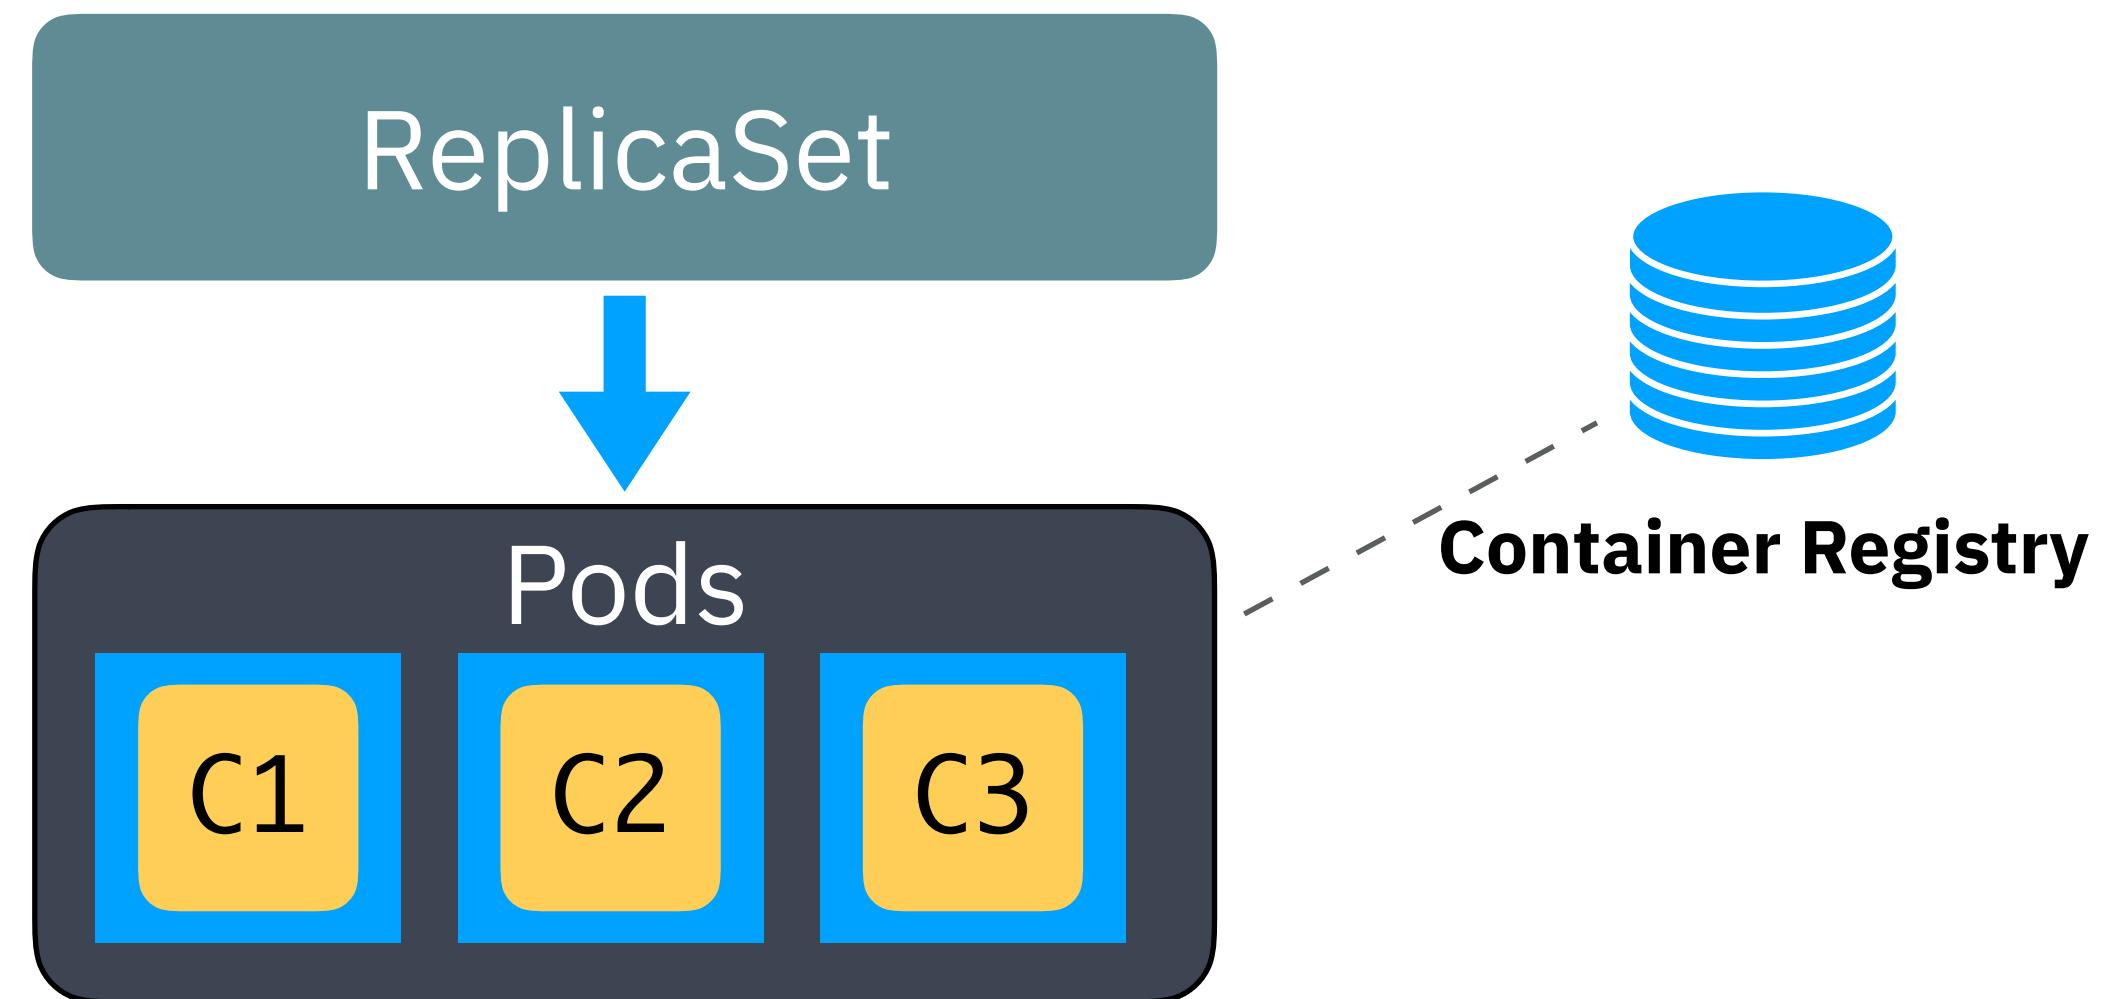
# Kubernetes Pods

- Pods
  - Containers run in Pods
  - The Pod is the smallest deployment possible
  - Containers are deployed from a container registry (public or private)



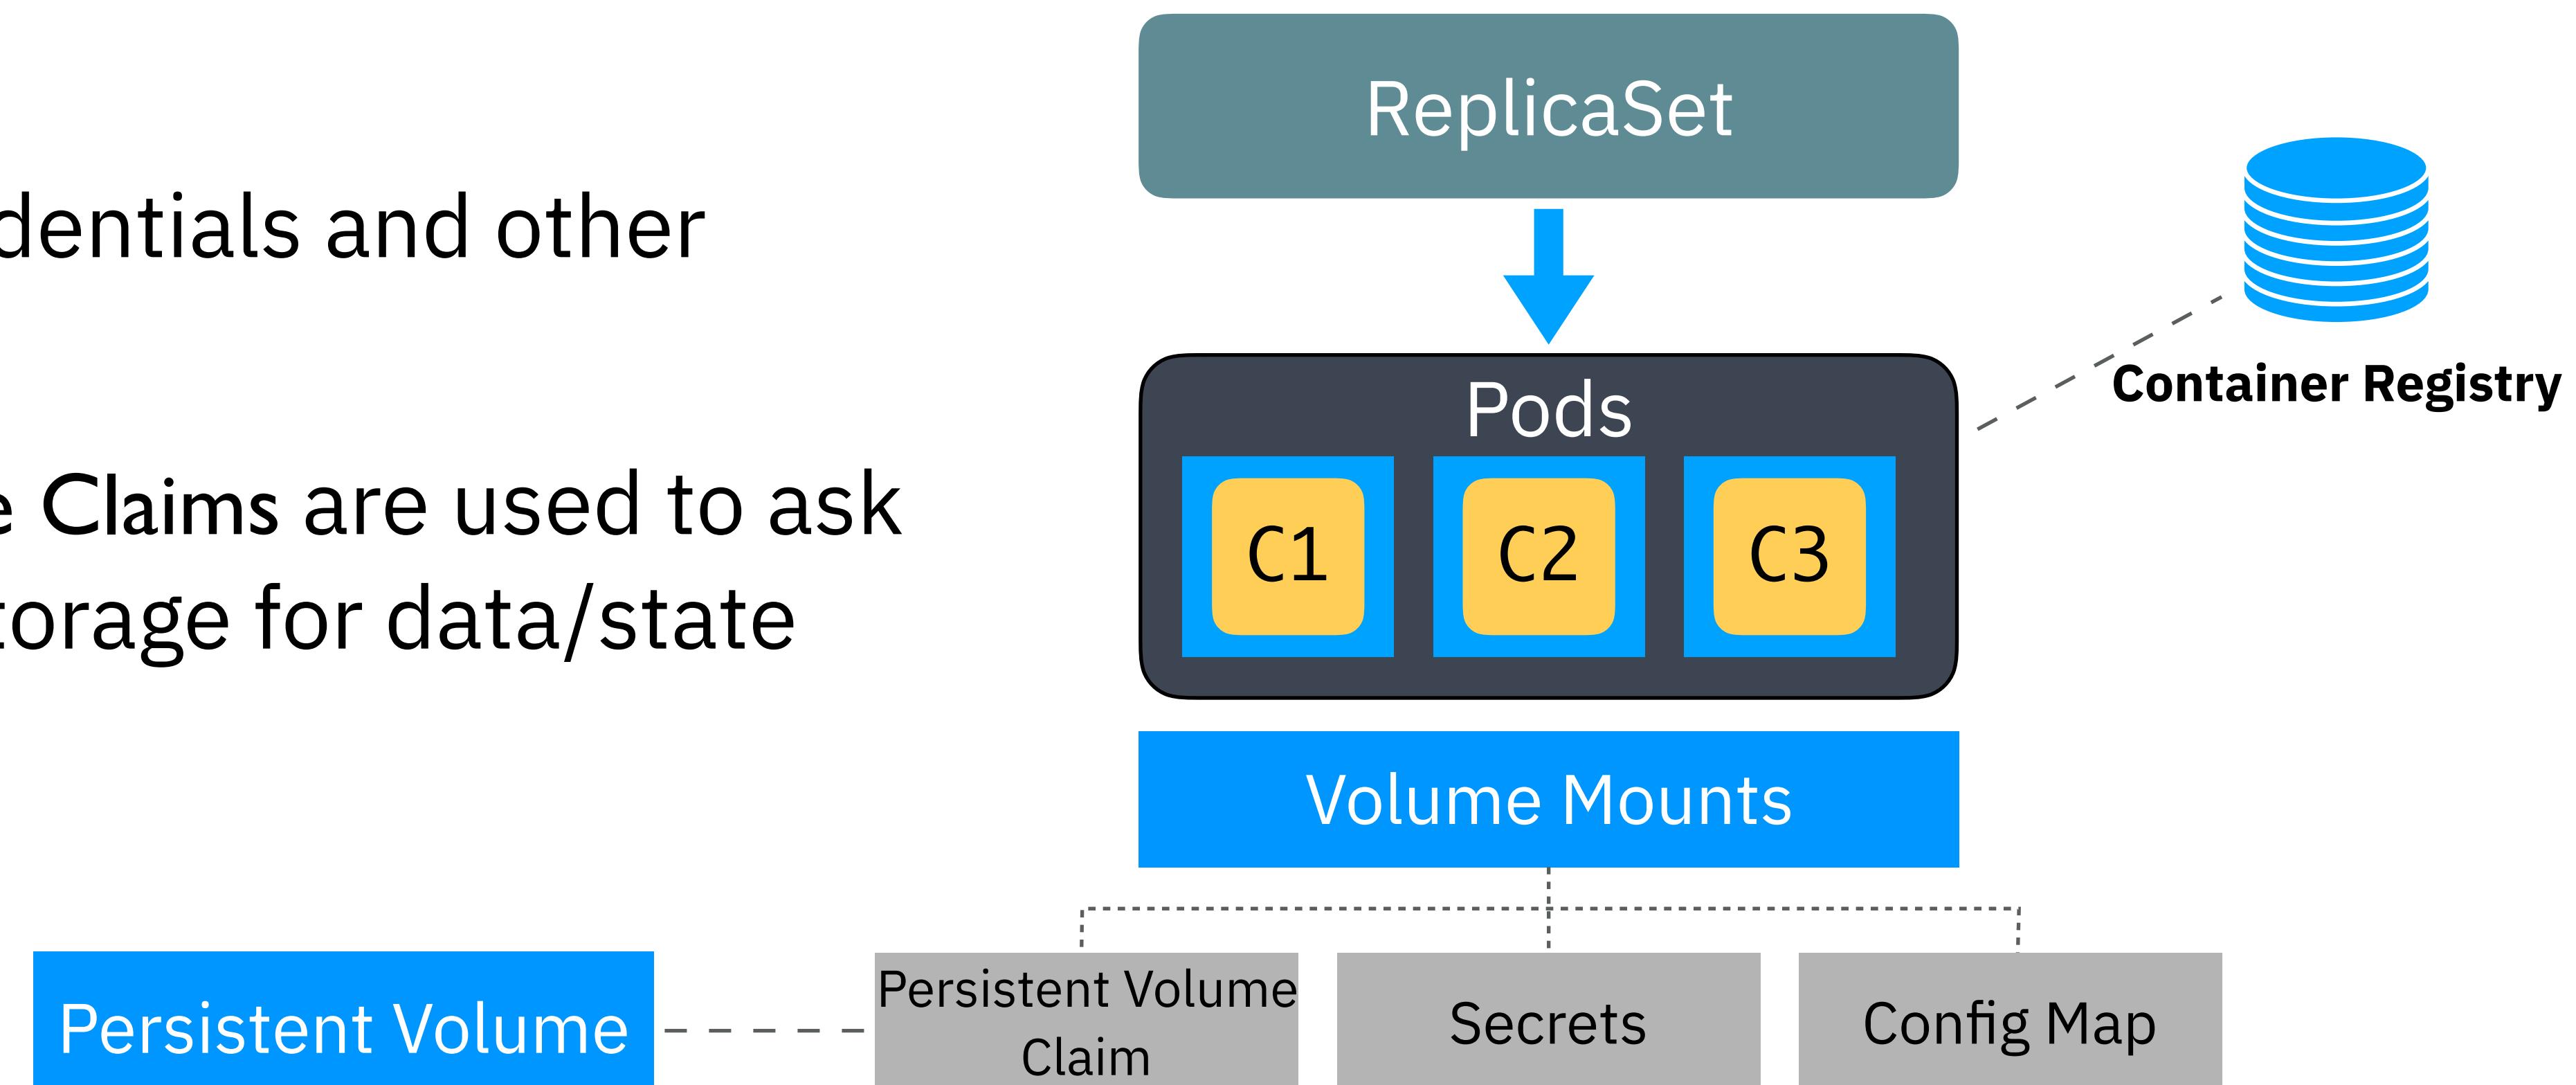
# Kubernetes ReplicaSets

- ReplicaSet
  - ReplicaSets allow multiple copies of containers to be deployed
  - Each one in its own Pod
  - If a container dies, the ReplicaSet will spawn a new one



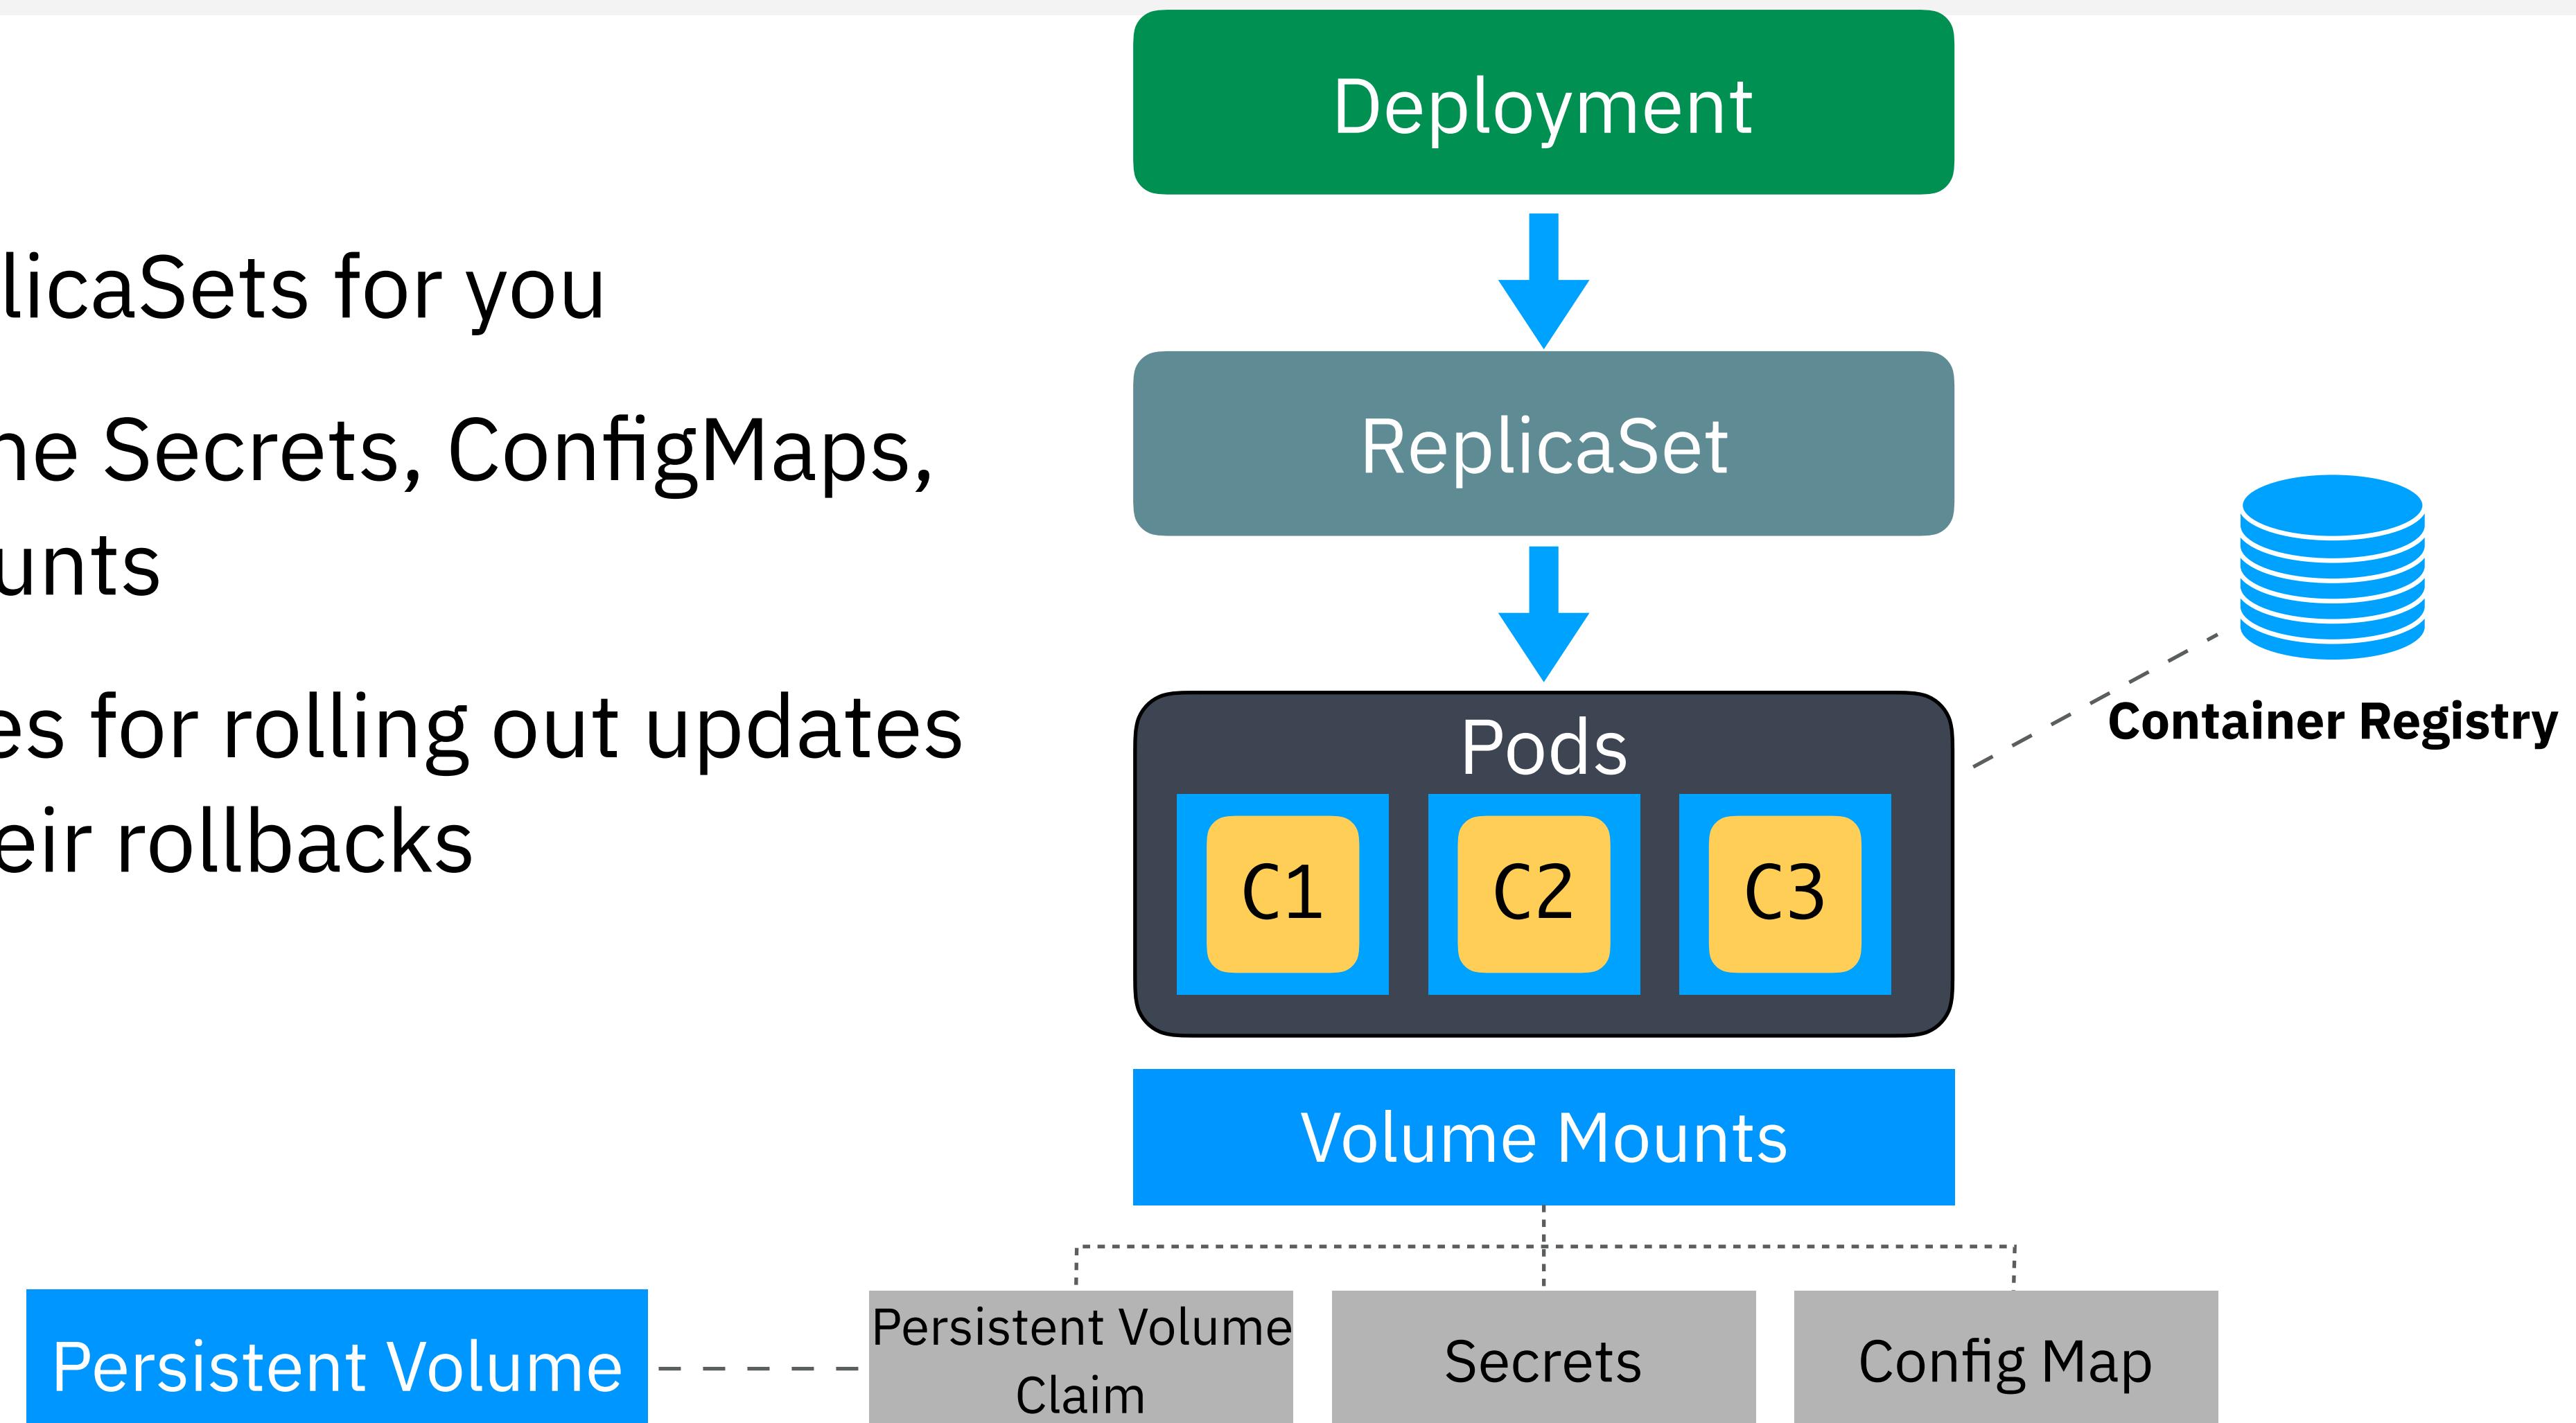
# Kubernetes Volume Mounts

- Volumes
  - ConfigMaps hold configuration parameters
  - Secrets hold credentials and other secrets
  - Persistent Volume Claims are used to ask for persistent storage for data/state



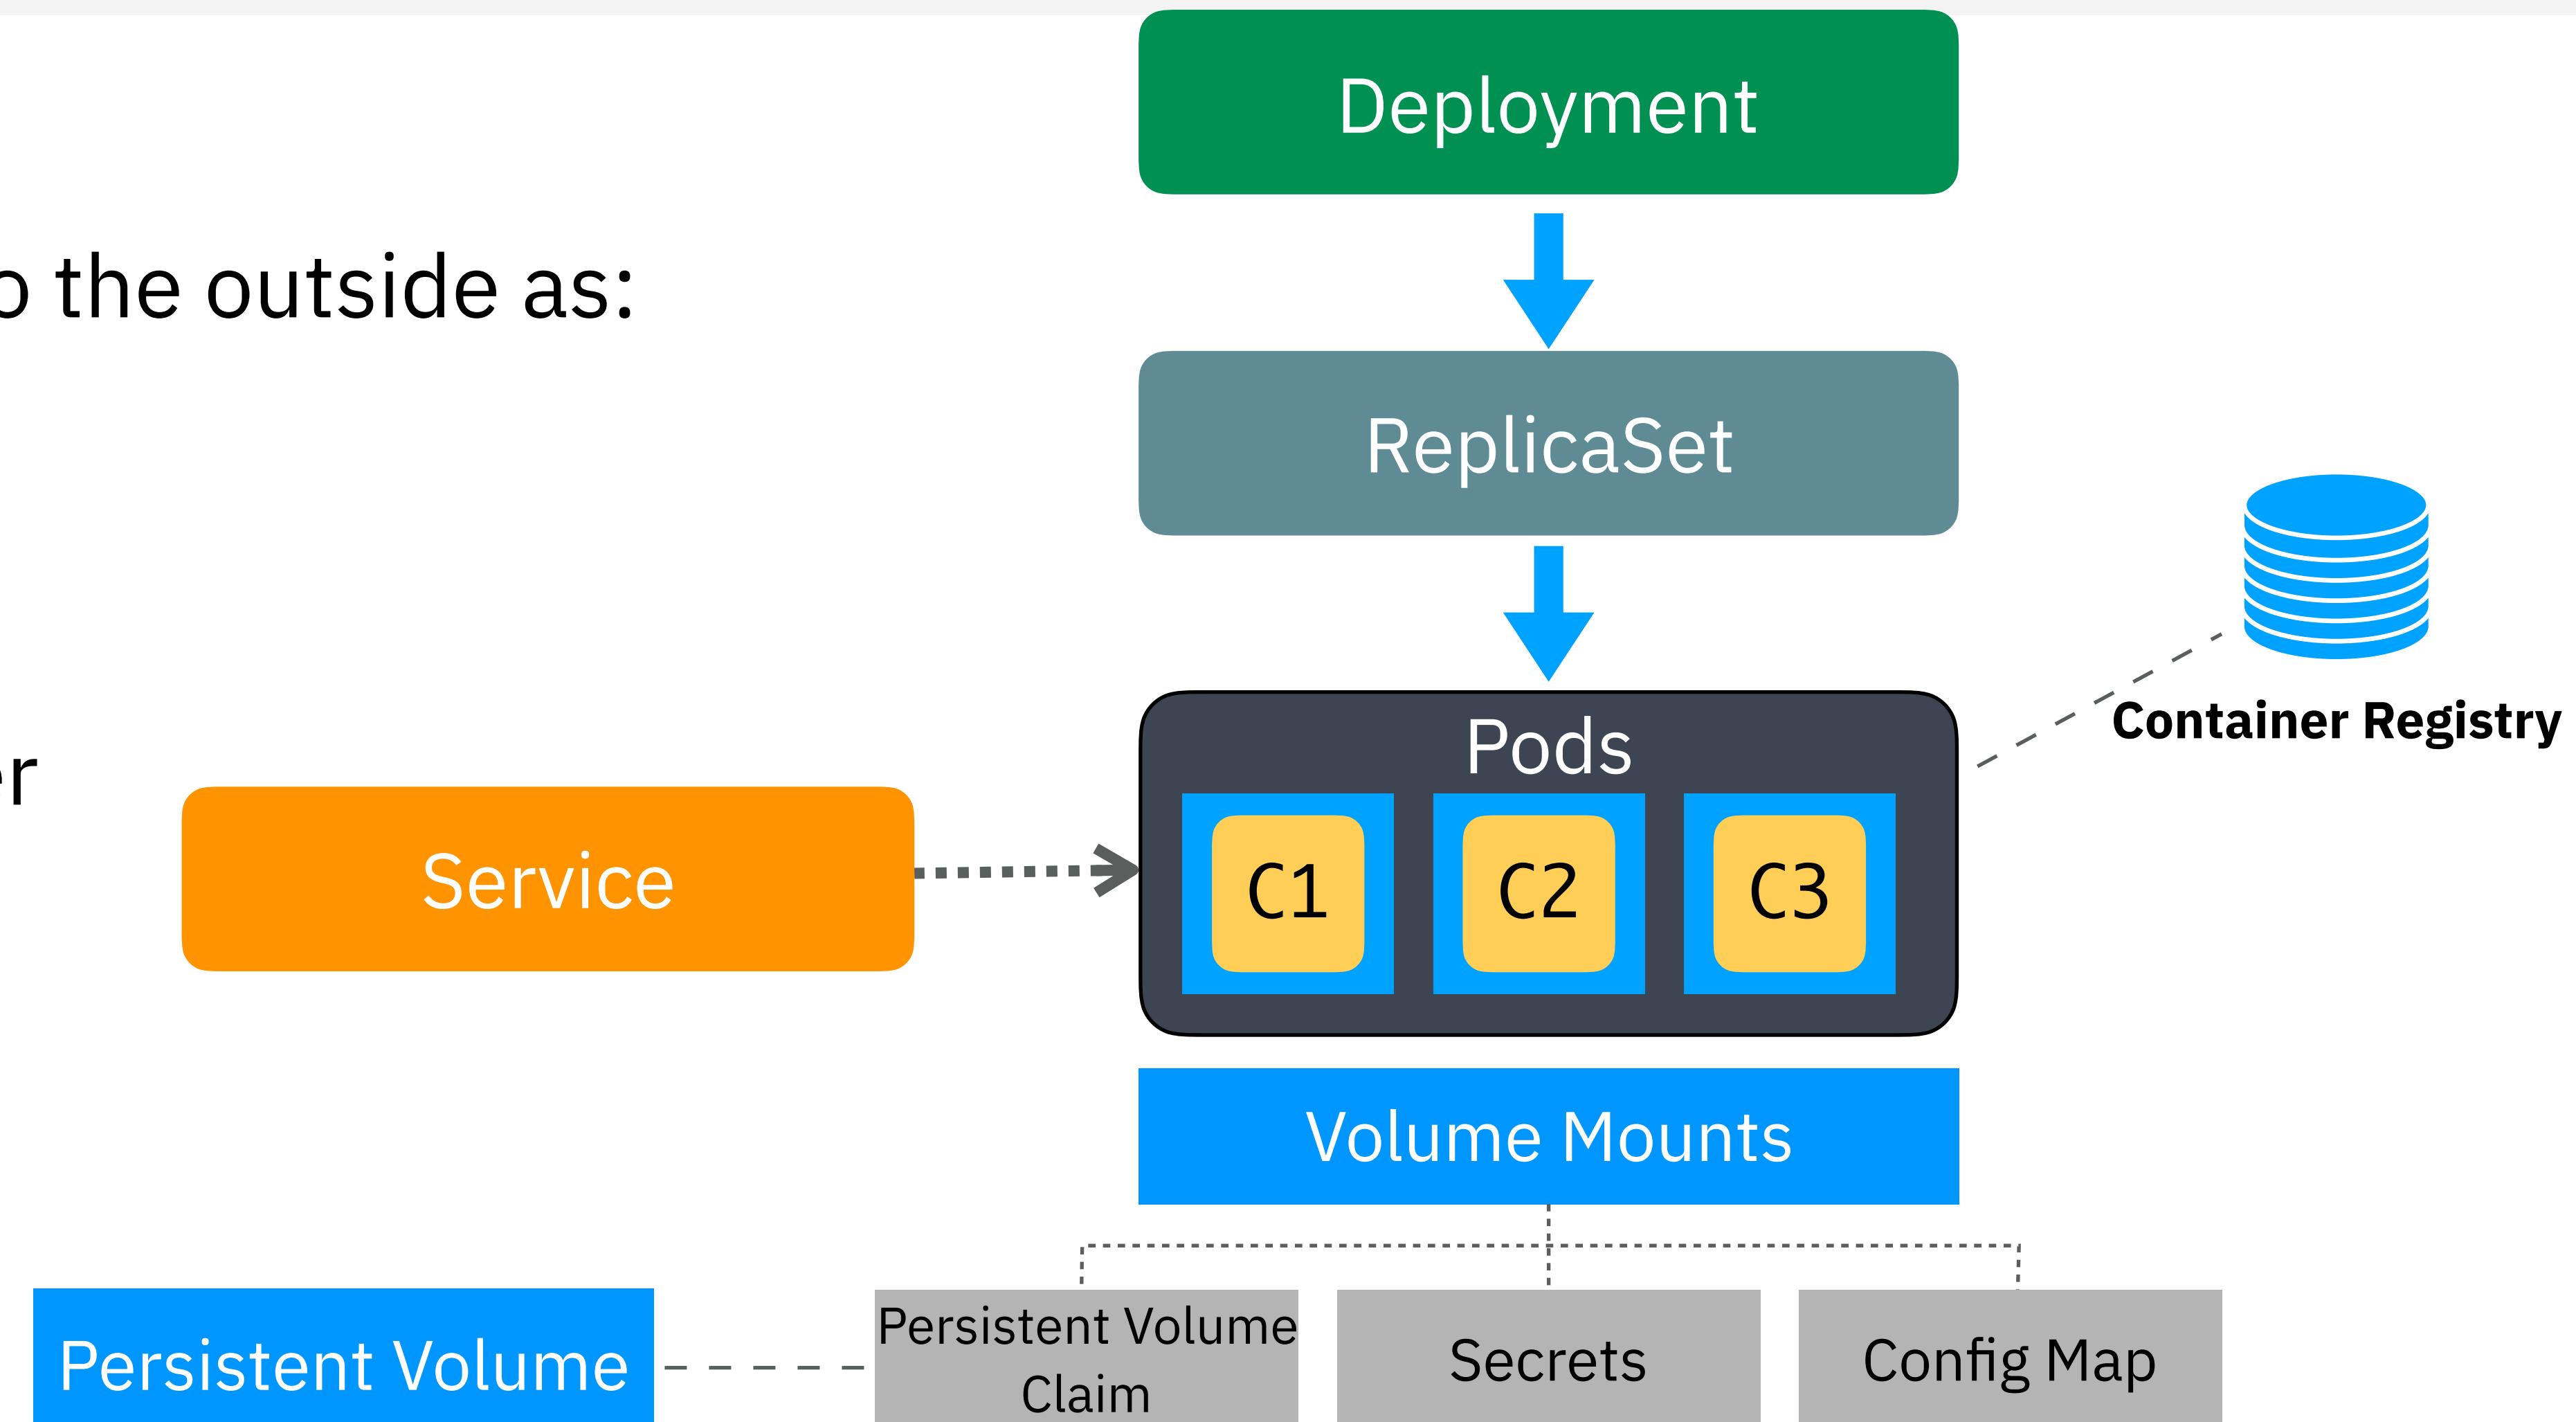
# Kubernetes Deployments

- Deployment
  - Sets up the ReplicaSets for you
  - Also specified the Secrets, ConfigMaps, and Volume Mounts
  - Provides features for rolling out updates and handling their rollbacks



# Kubernetes Service

- Service
  - Exposes Pods to the outside as:
    - ClusterIP
    - NodePort
    - Load Balancer

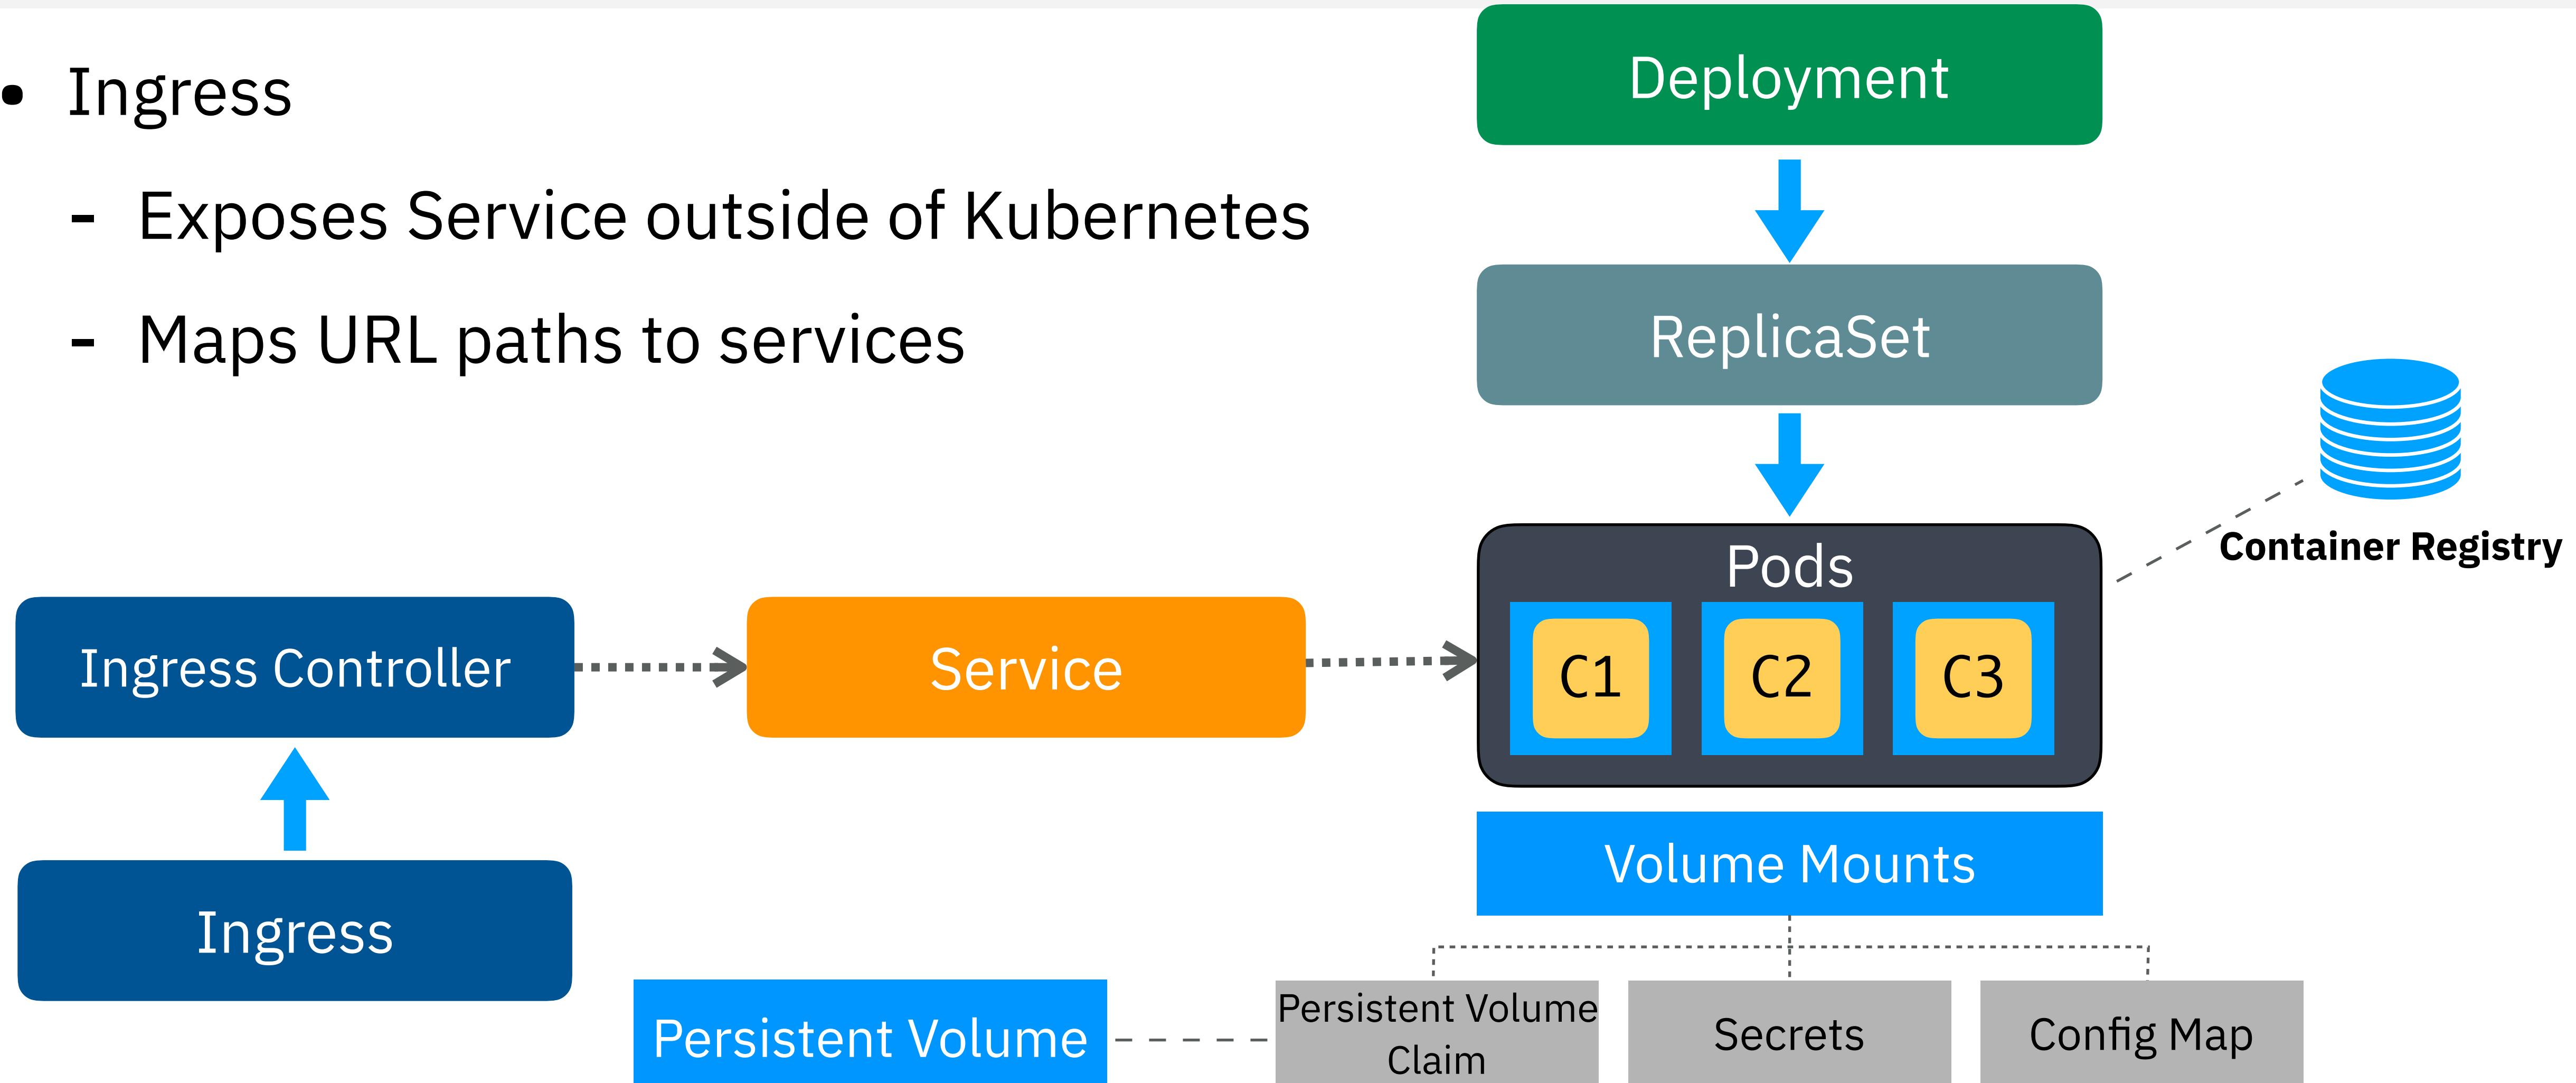


# Types of Services

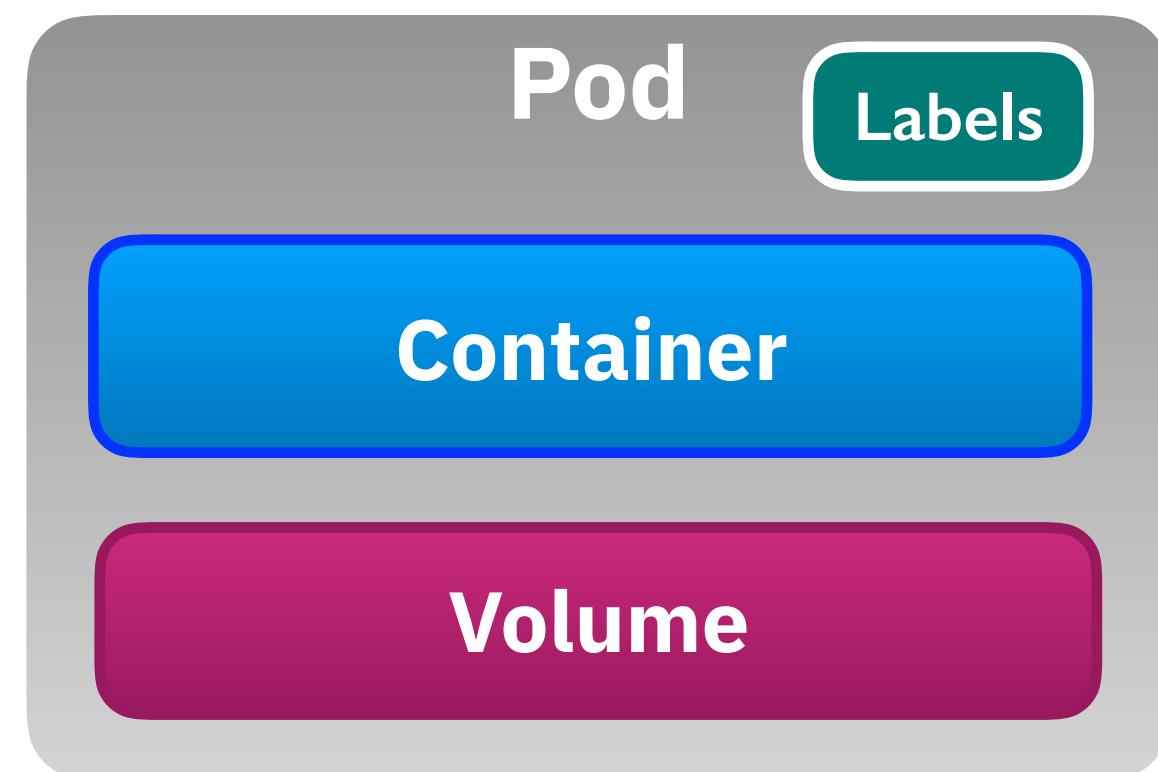
- **LoadBalancer** – Only available via cloud providers. Front end for service that balances the load across multiple backends from a single IP address
- **NodePort** – Exposes the service as an arbitrary port on every worker node in the cluster
- **ClusterIP** – Service is only accessible from other services within the cluster (no external exposure)

# Kubernetes Ingress Controller

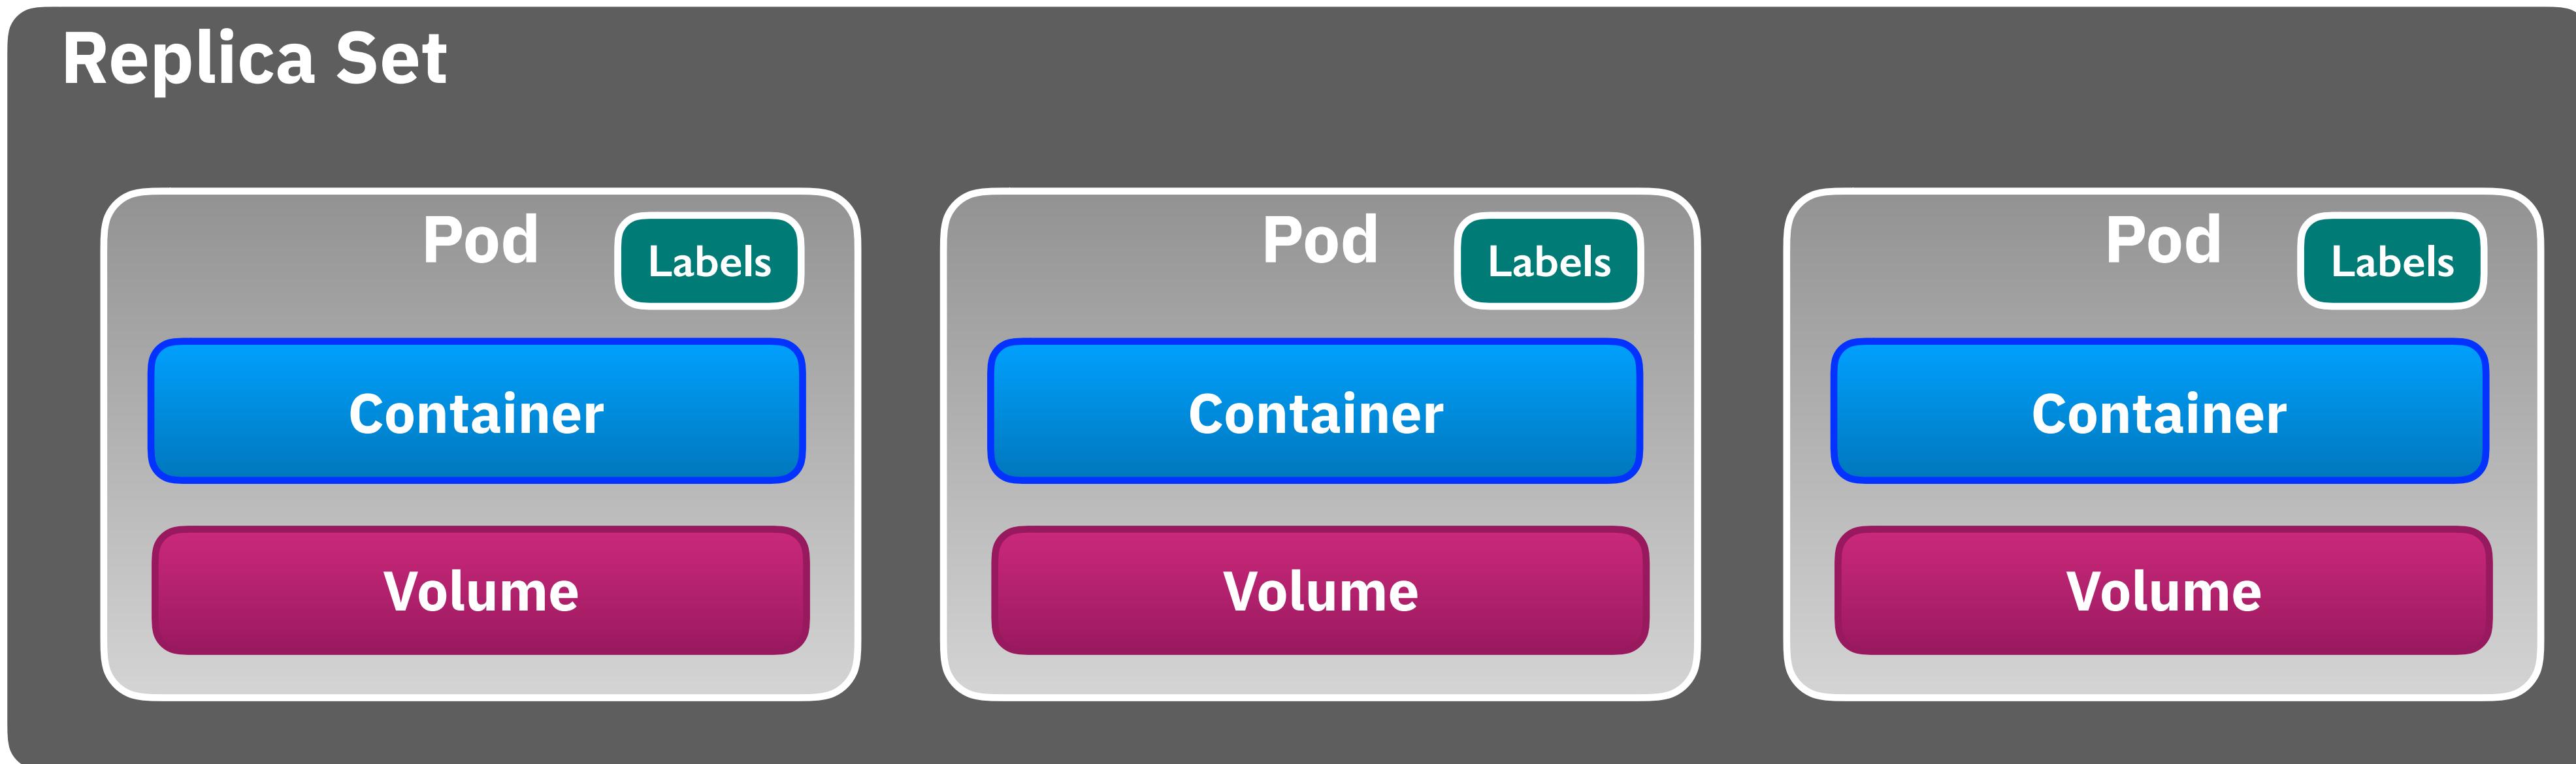
- Ingress
  - Exposes Service outside of Kubernetes
  - Maps URL paths to services



# Services Map to Pods via Labels



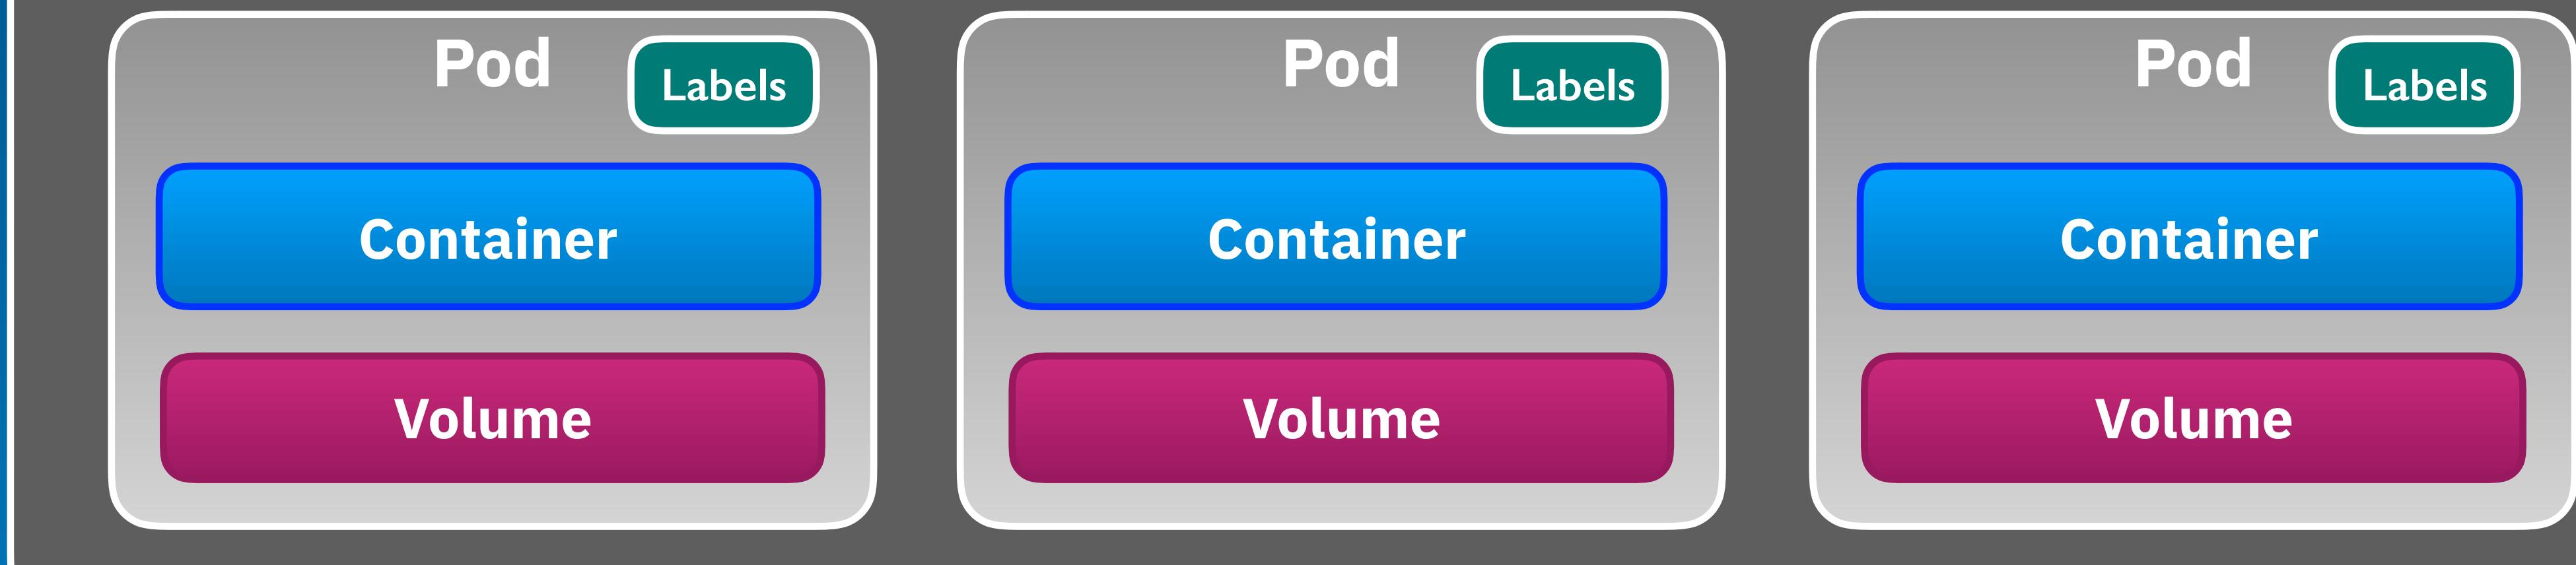
# Services Map to Pods via Labels



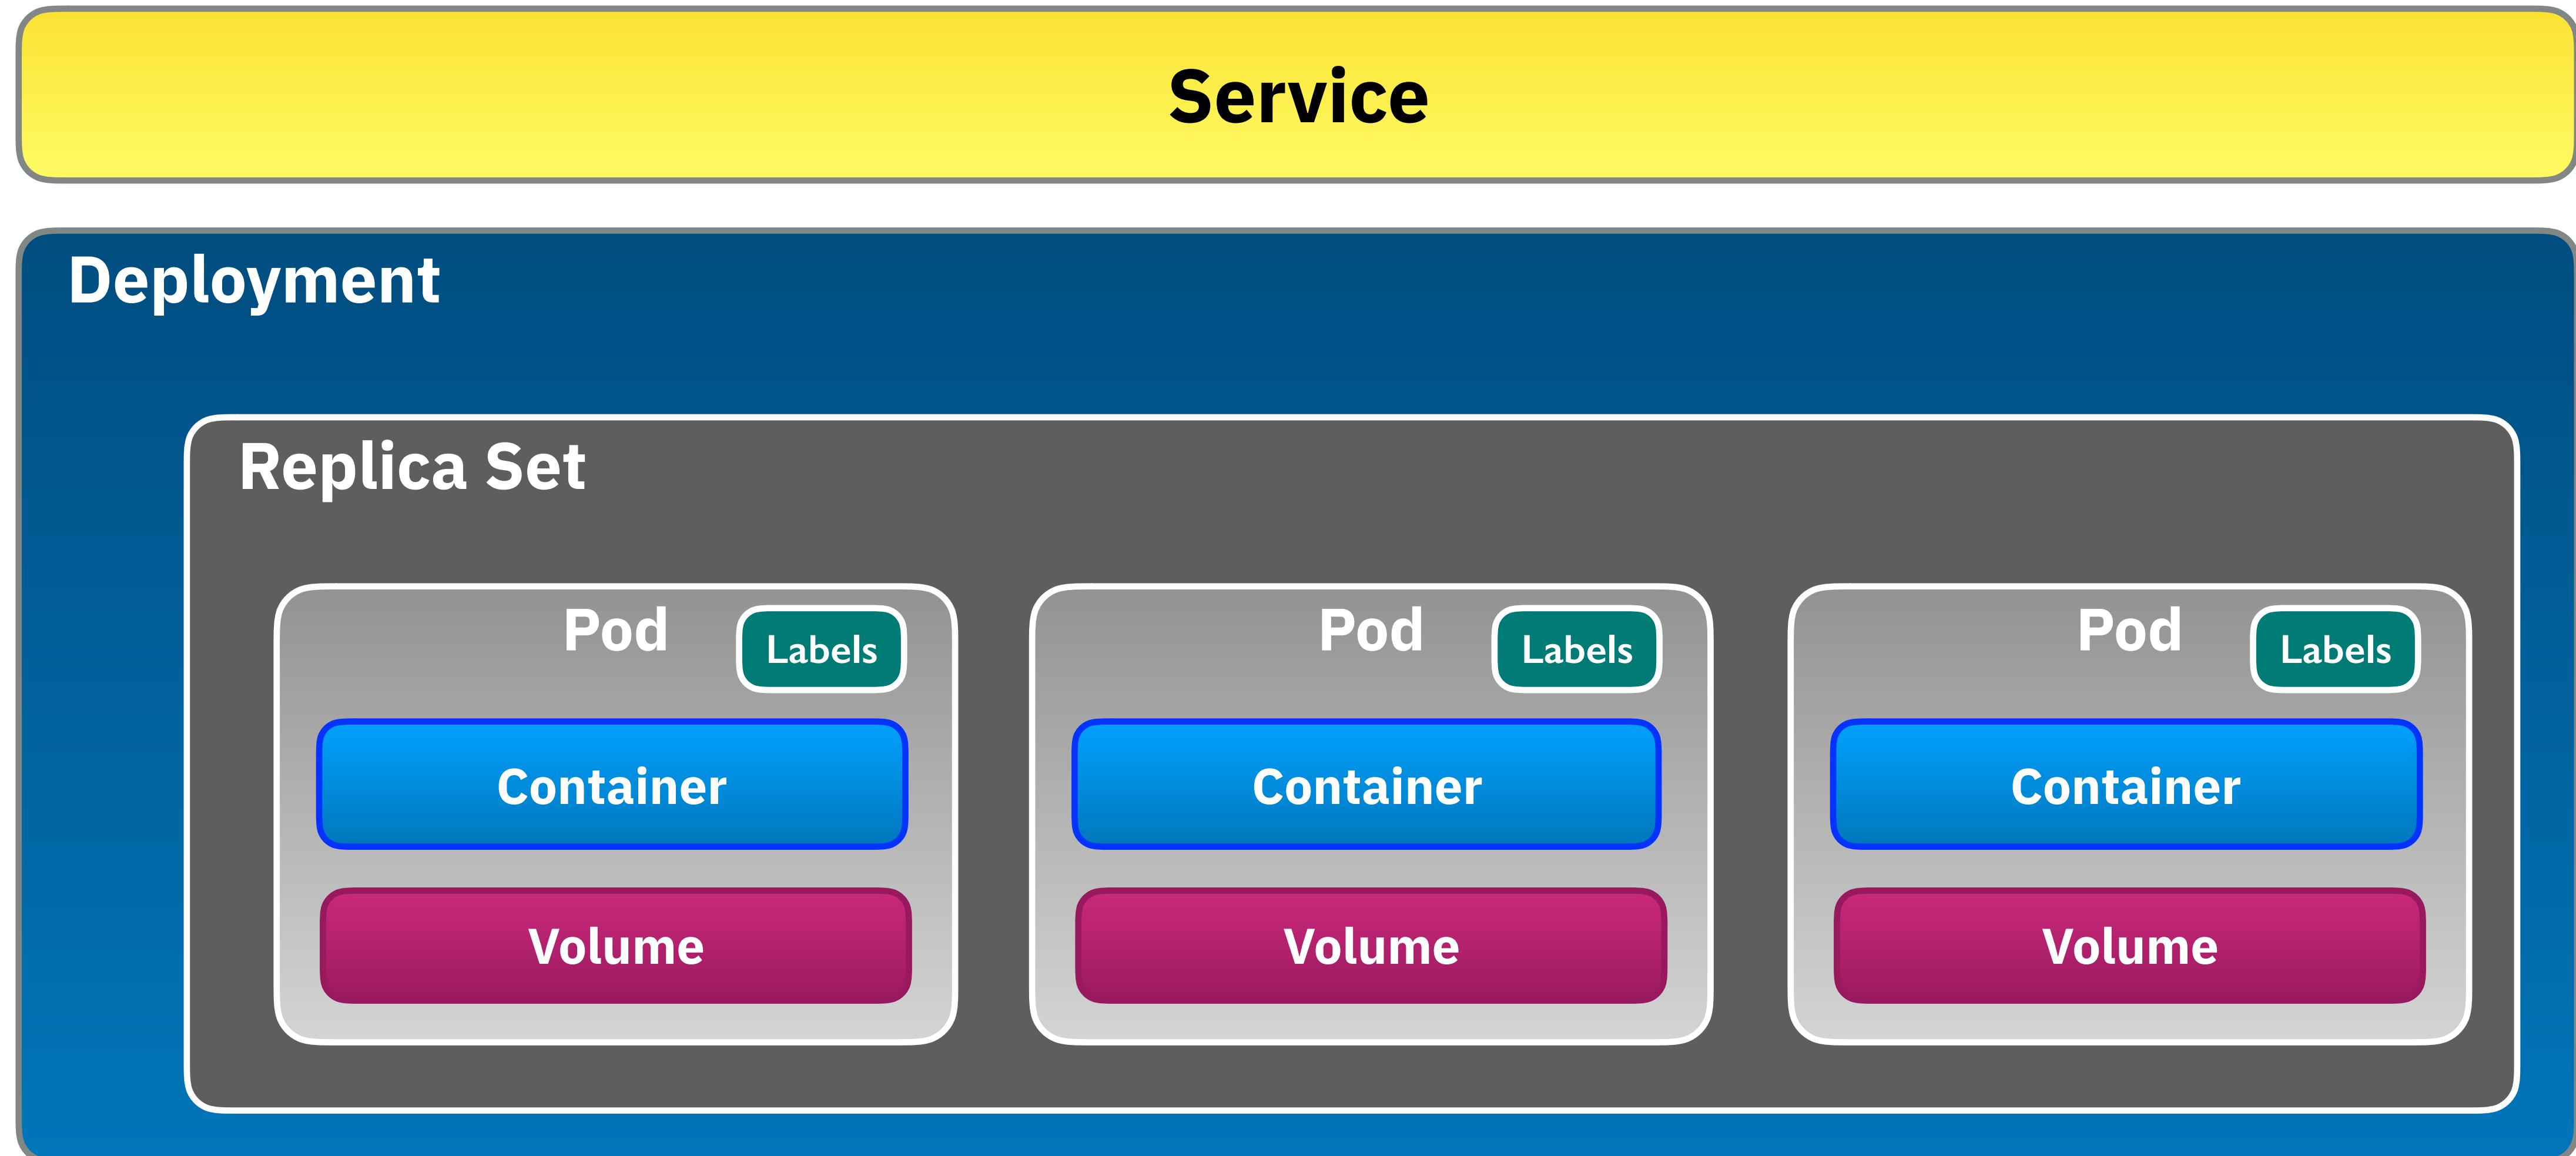
# Services Map to Pods via Labels

## Deployment

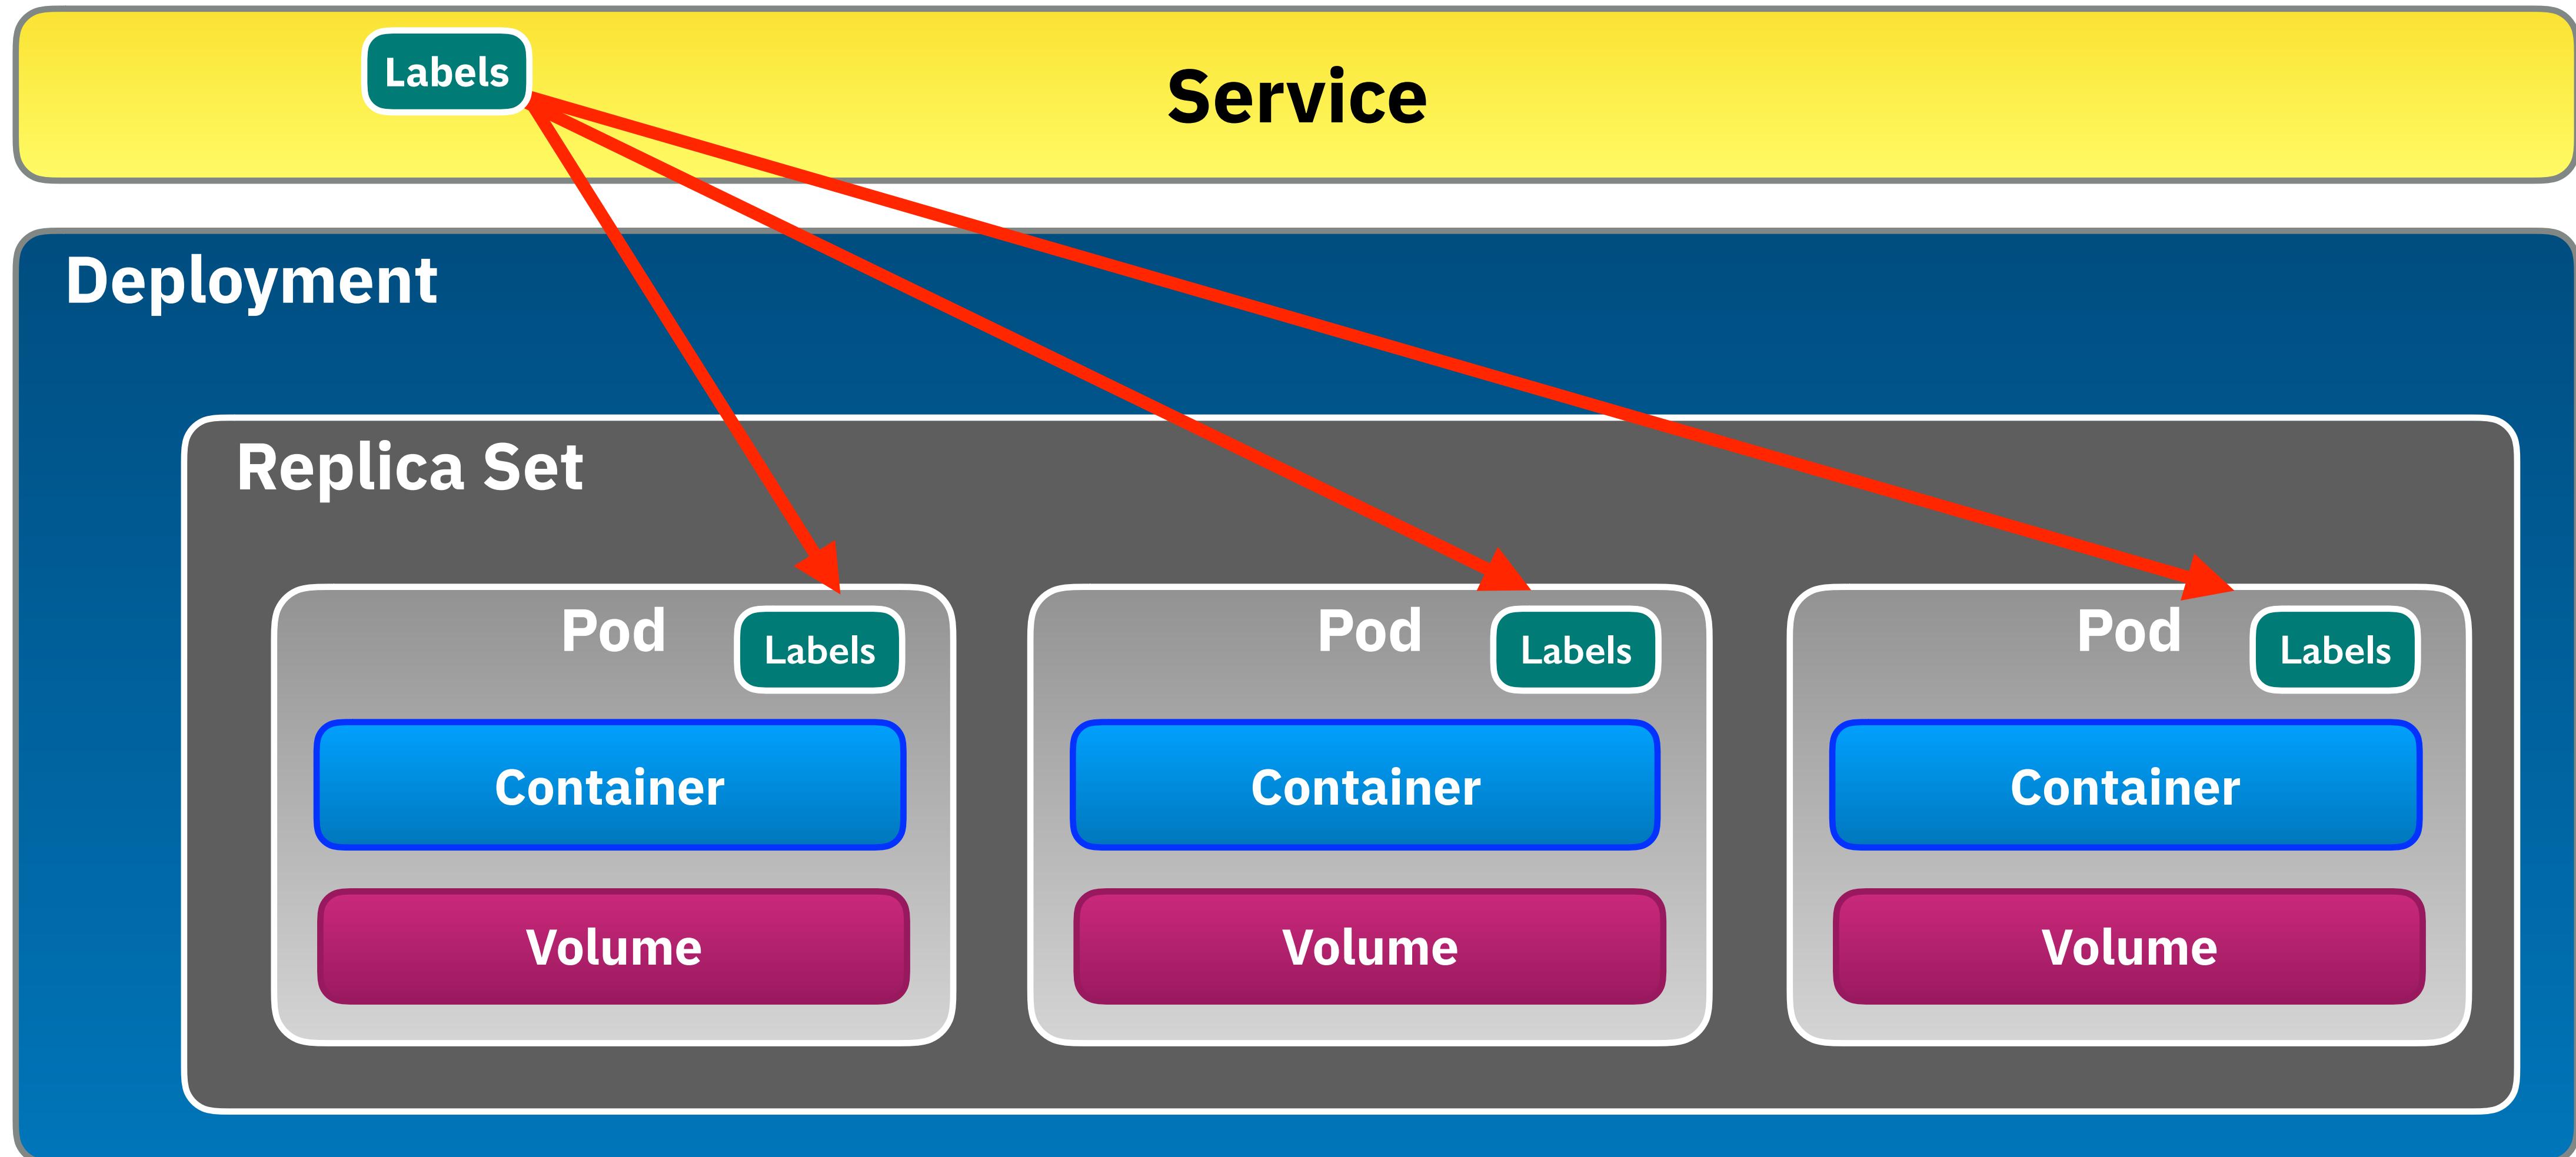
### Replica Set



# Services Map to Pods via Labels



# Services Map to Pods via Labels



# Use Case for Multiple Containers in a Pod

- Pods can be used to host vertically integrated application stacks (e.g. LAMP), but their primary motivation is to support co-located, co-managed helper programs, such as:
  - Content management systems, file and data loaders, local cache managers, etc.
  - Log and checkpoint backup, compression, rotation, snapshotting, etc.
  - Data change watchers, log tailers, logging and monitoring adapters, event publishers, etc.
  - Proxies, bridges, and adapters
  - Controllers, managers, configurators, and updaters

# Deployment using kubectl

- A containerized application can be deployed on Kubernetes using a deployment definition by executing a simple CLI command as follows:

```
kubectl create deployment <application-name> --image=<container-image>
```

```
$ kubectl create deployment hello-service --image=hello:v1
```

# Create a Service

- You can expose a deployment as a service using `kubectl expose`:

```
kubectl expose deployment <application-name> --type=<service-type> \  
--port=<port_number>
```

```
$ kubectl expose deployment hello-service --type=NodePort --port=3000
```

Valid <service-type>s are:

- LoadBalancer
- NodePort
- ClusterIP

# Deploy using YAML

- You can create a deployment and a service using yaml files

```
$ kubectl create -f deployment.yaml  
$ kubectl create -f service.yaml
```

- Update them

```
$ kubectl apply -f deployment.yaml  
$ kubectl apply -f service.yaml
```

- And just as easily delete them

```
$ kubectl delete -f deployment.yaml  
$ kubectl delete -f service.yaml
```

# Deployment YAML

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pet-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: pet-demo
  template:
    metadata:
      labels:
        app: pet-demo
    spec:
      containers:
        - name: pet-demo
          image: pet-demo:v1
          imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5000
          protocol: TCP
      restartPolicy: Always
```

# Service YAML

```
apiVersion: v1
kind: Service
metadata:
  name: pet-demo-service
spec:
  type: NodePort
  selector:
    App: pet-demo
  ports:
    - name: primary
      protocol: TCP
      port: 5000
```

# Linking Service to Pods

```
apiVersion: v1
kind: Service
metadata:
  name: pet-demo-service
spec:
  type: NodePort
  selector:
    app: pet-demo
  ports:
    - name: primary
      protocol: TCP
      port: 5000
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pet-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: pet-demo
  template:
    metadata:
      labels:
        app: pet-demo
    spec:
      containers:
        - name: pet-demo
          image: pet-demo:v1
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 5000
              protocol: TCP
          restartPolicy: Always
```

# Linking Service to Pods

```
apiVersion: v1
kind: Service
metadata:
  name: pet-demo-service
spec:
  type: NodePort
  selector:
    app: pet-demo
  ports:
    - name: primary
      protocol: TCP
      port: 5000
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pet-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: pet-demo
  template:
    metadata:
      labels:
        app: pet-demo
    spec:
      containers:
        - name: pet-demo
          image: pet-demo:v1
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 5000
              protocol: TCP
          restartPolicy: Always
```

# Linking Service to Pods

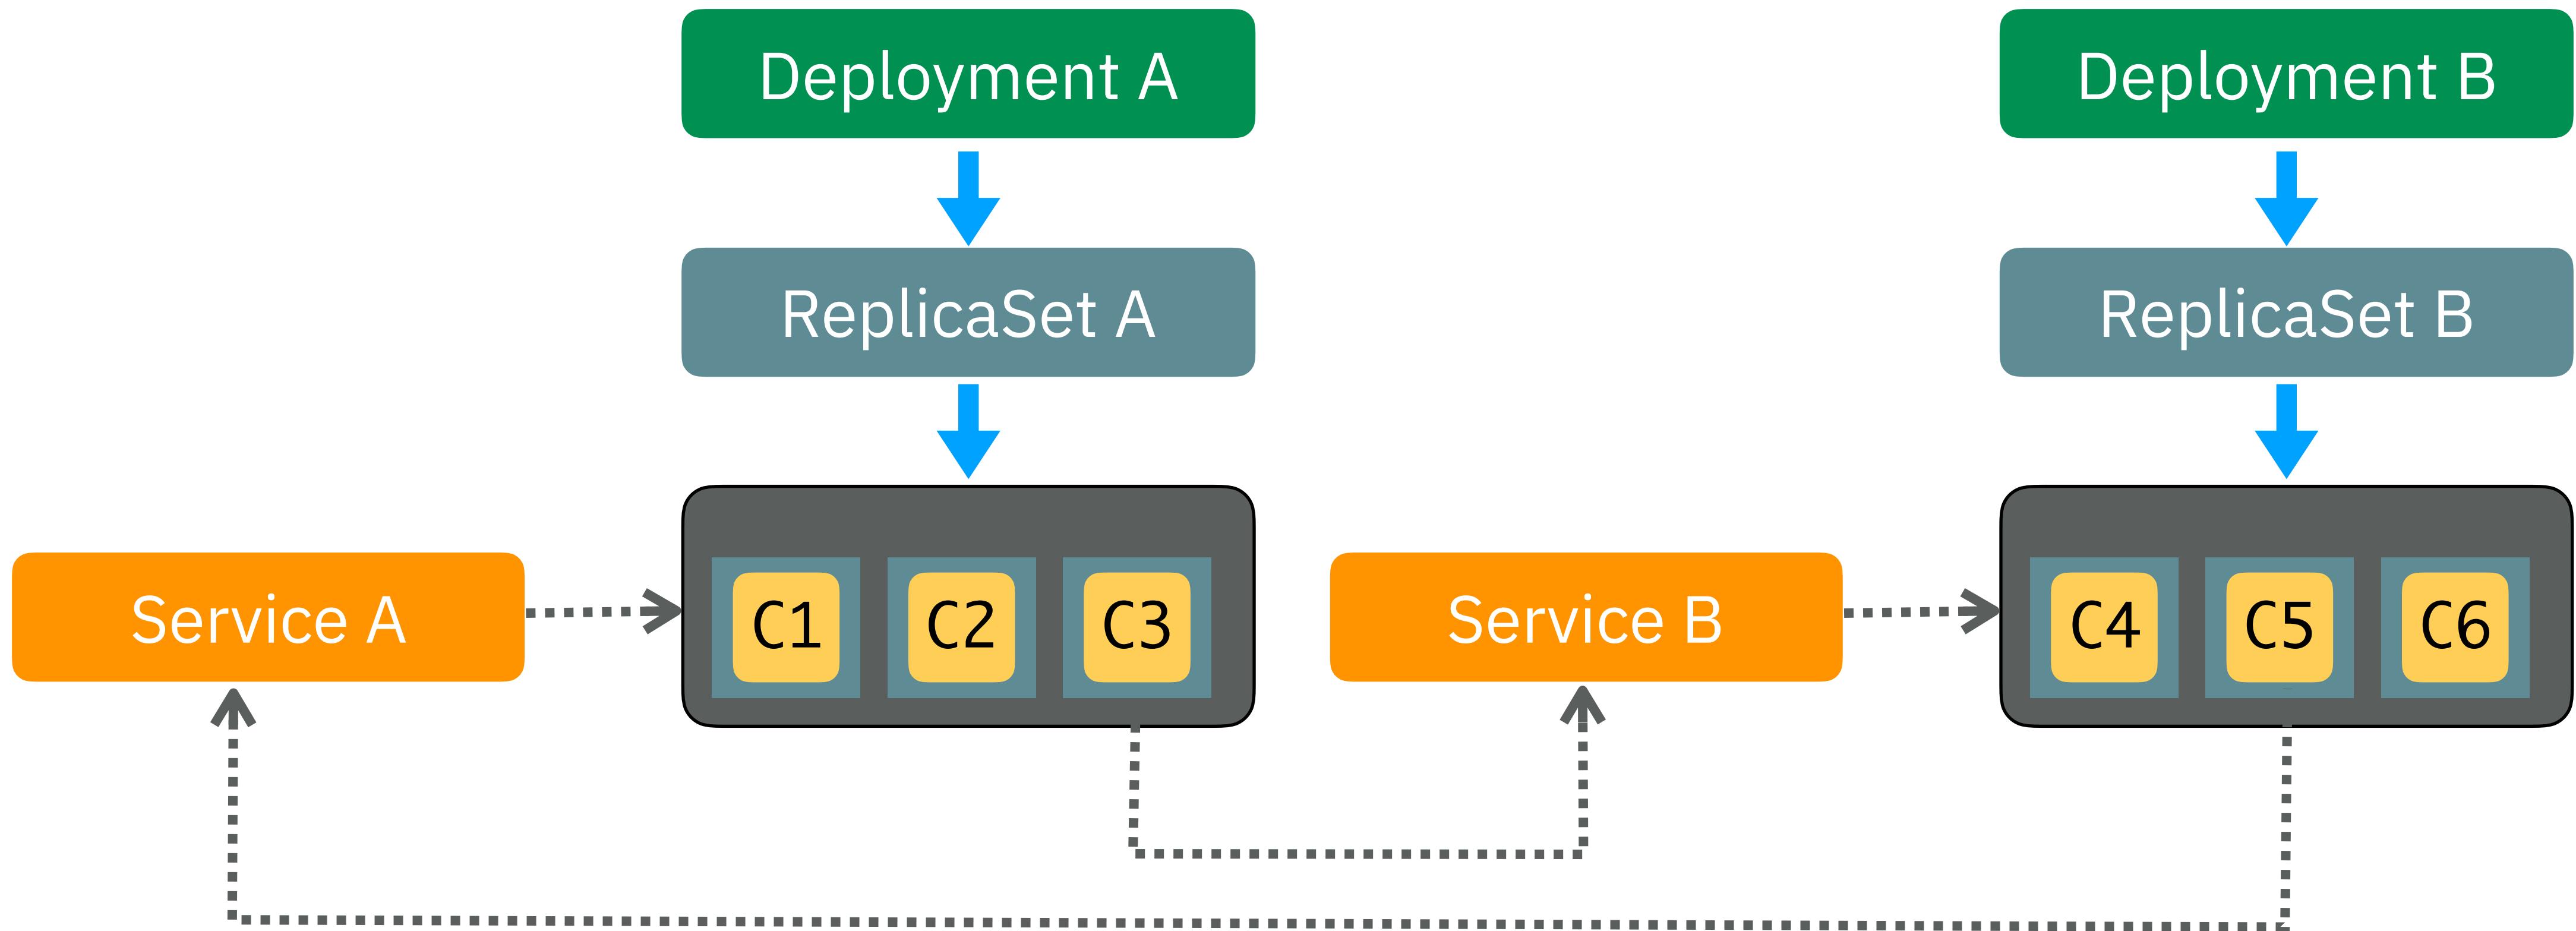
```
apiVersion: v1
kind: Service
metadata:
  name: pet-demo-service
spec:
  type: NodePort
  selector:
    app: pet-demo
  ports:
    - name: primary
      protocol: TCP
      port: 5000
```

Will match this Pod

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pet-demo
spec:
  replicas: 2
  selector:
    matchLabels:
      app: pet-demo
  template:
    metadata:
      labels:
        app: pet-demo
    spec:
      containers:
        - name: pet-demo
          image: pet-demo:v1
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 5000
              protocol: TCP
      restartPolicy: Always
```

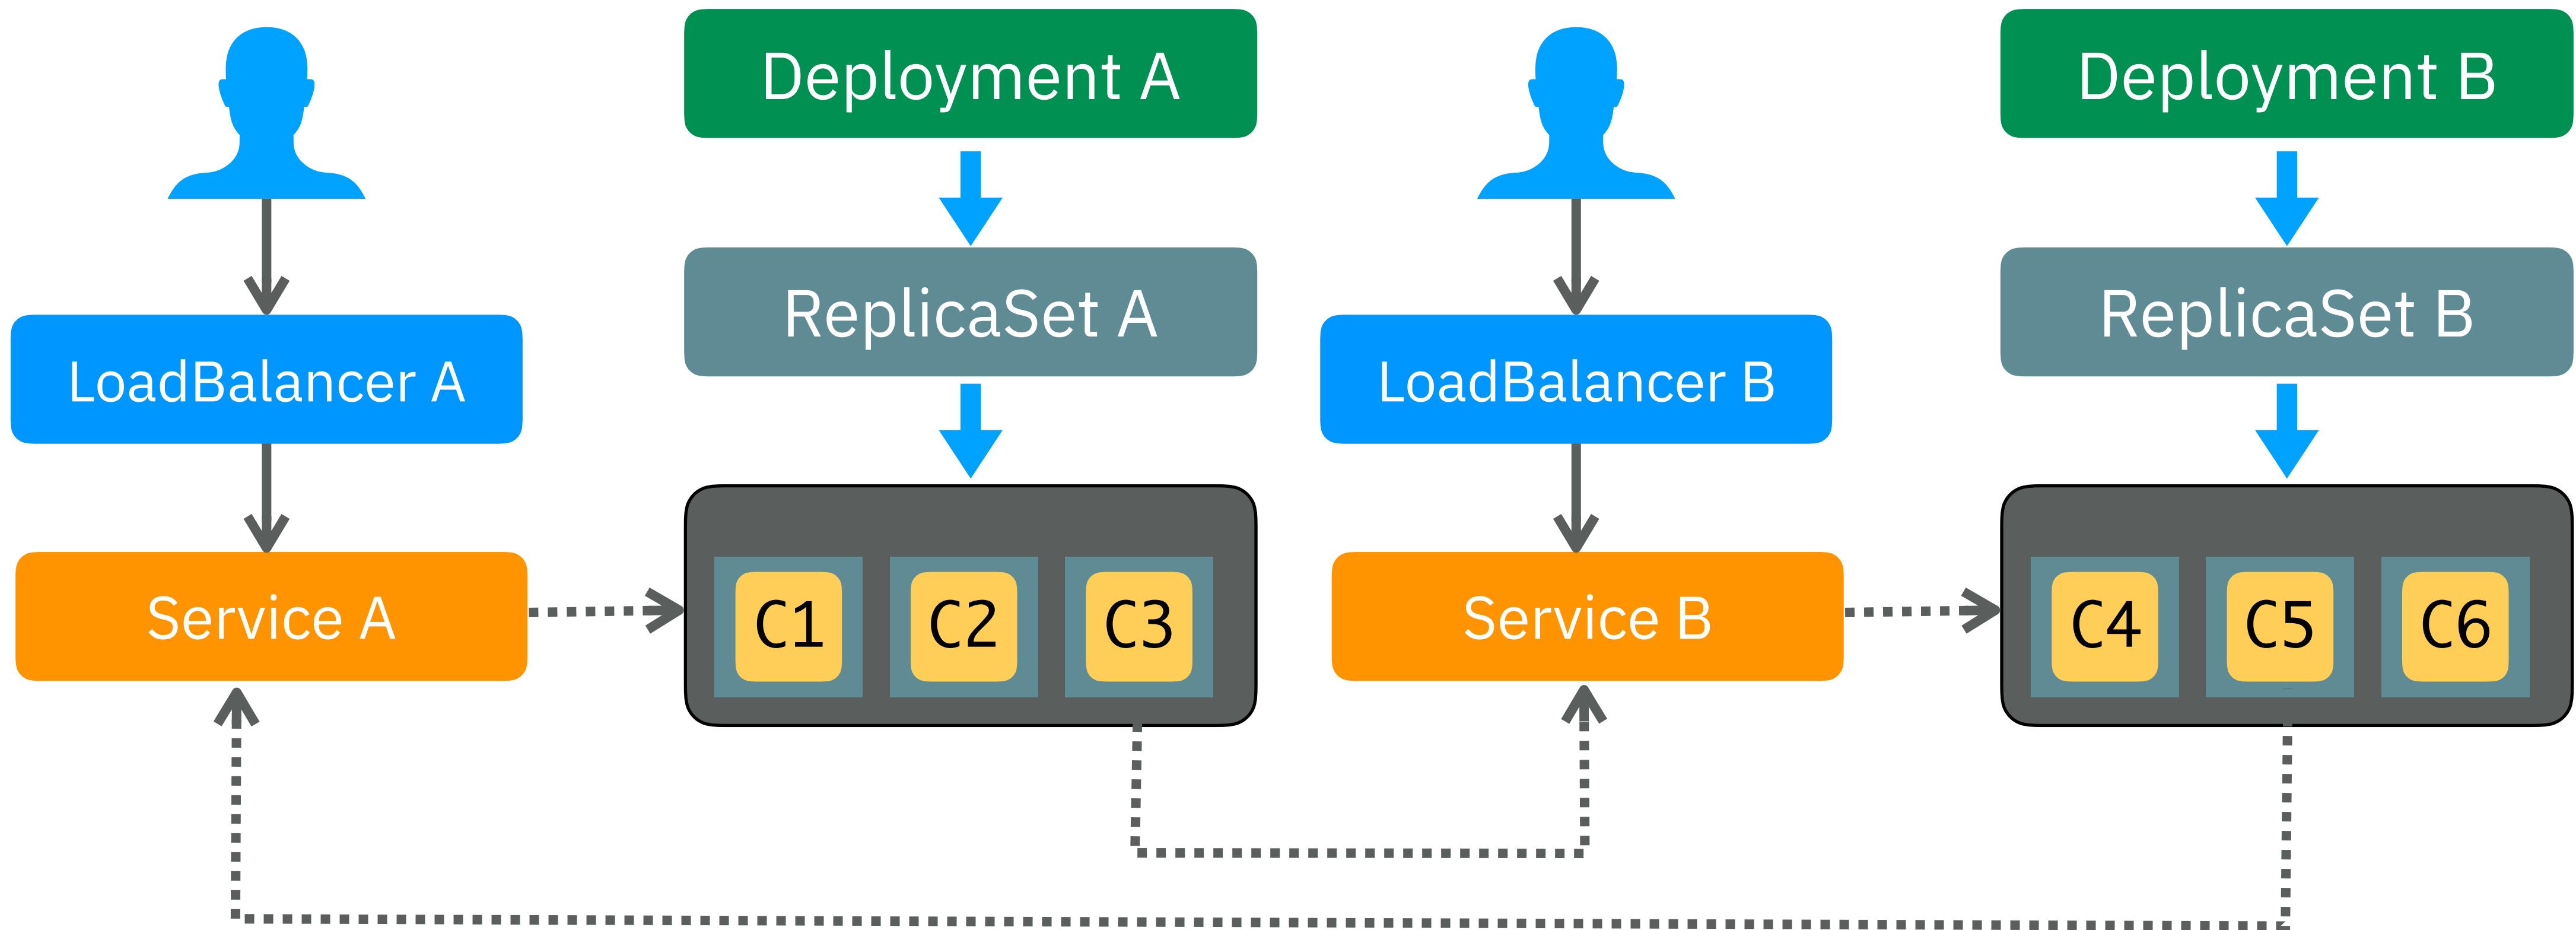
# Service Discovery

- Kubernetes provides internal routing so services can find each other



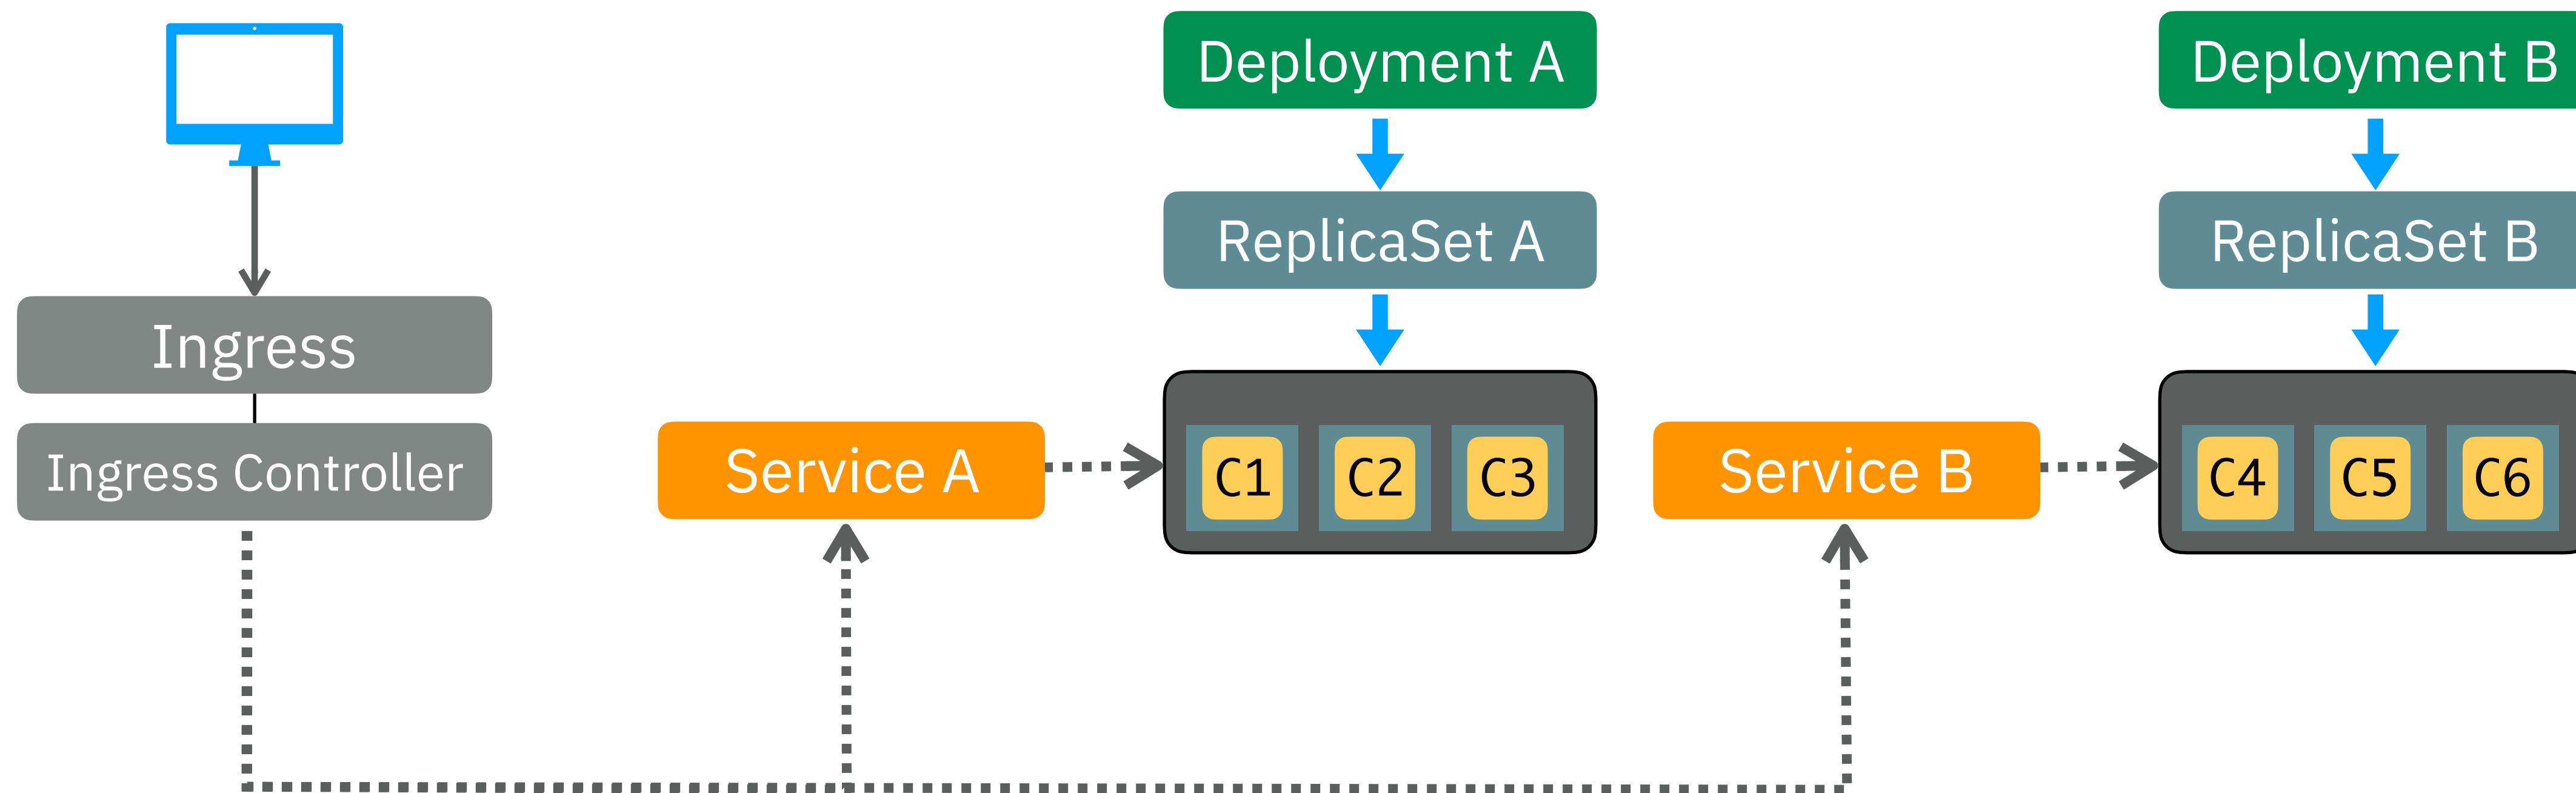
# External Service Access

- LoadBalancers provide external access for a single service



# External Routing

- Kubernetes provides an Ingress Controller to allow external network access or you can use NodePorts



# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.mycompany.com
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.mycompany.com
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Requests coming in on this URL

# Ingress Example

```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.mycompany.com
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Requests coming in on this URL

Will be routed to these microservices

# Ingress Example

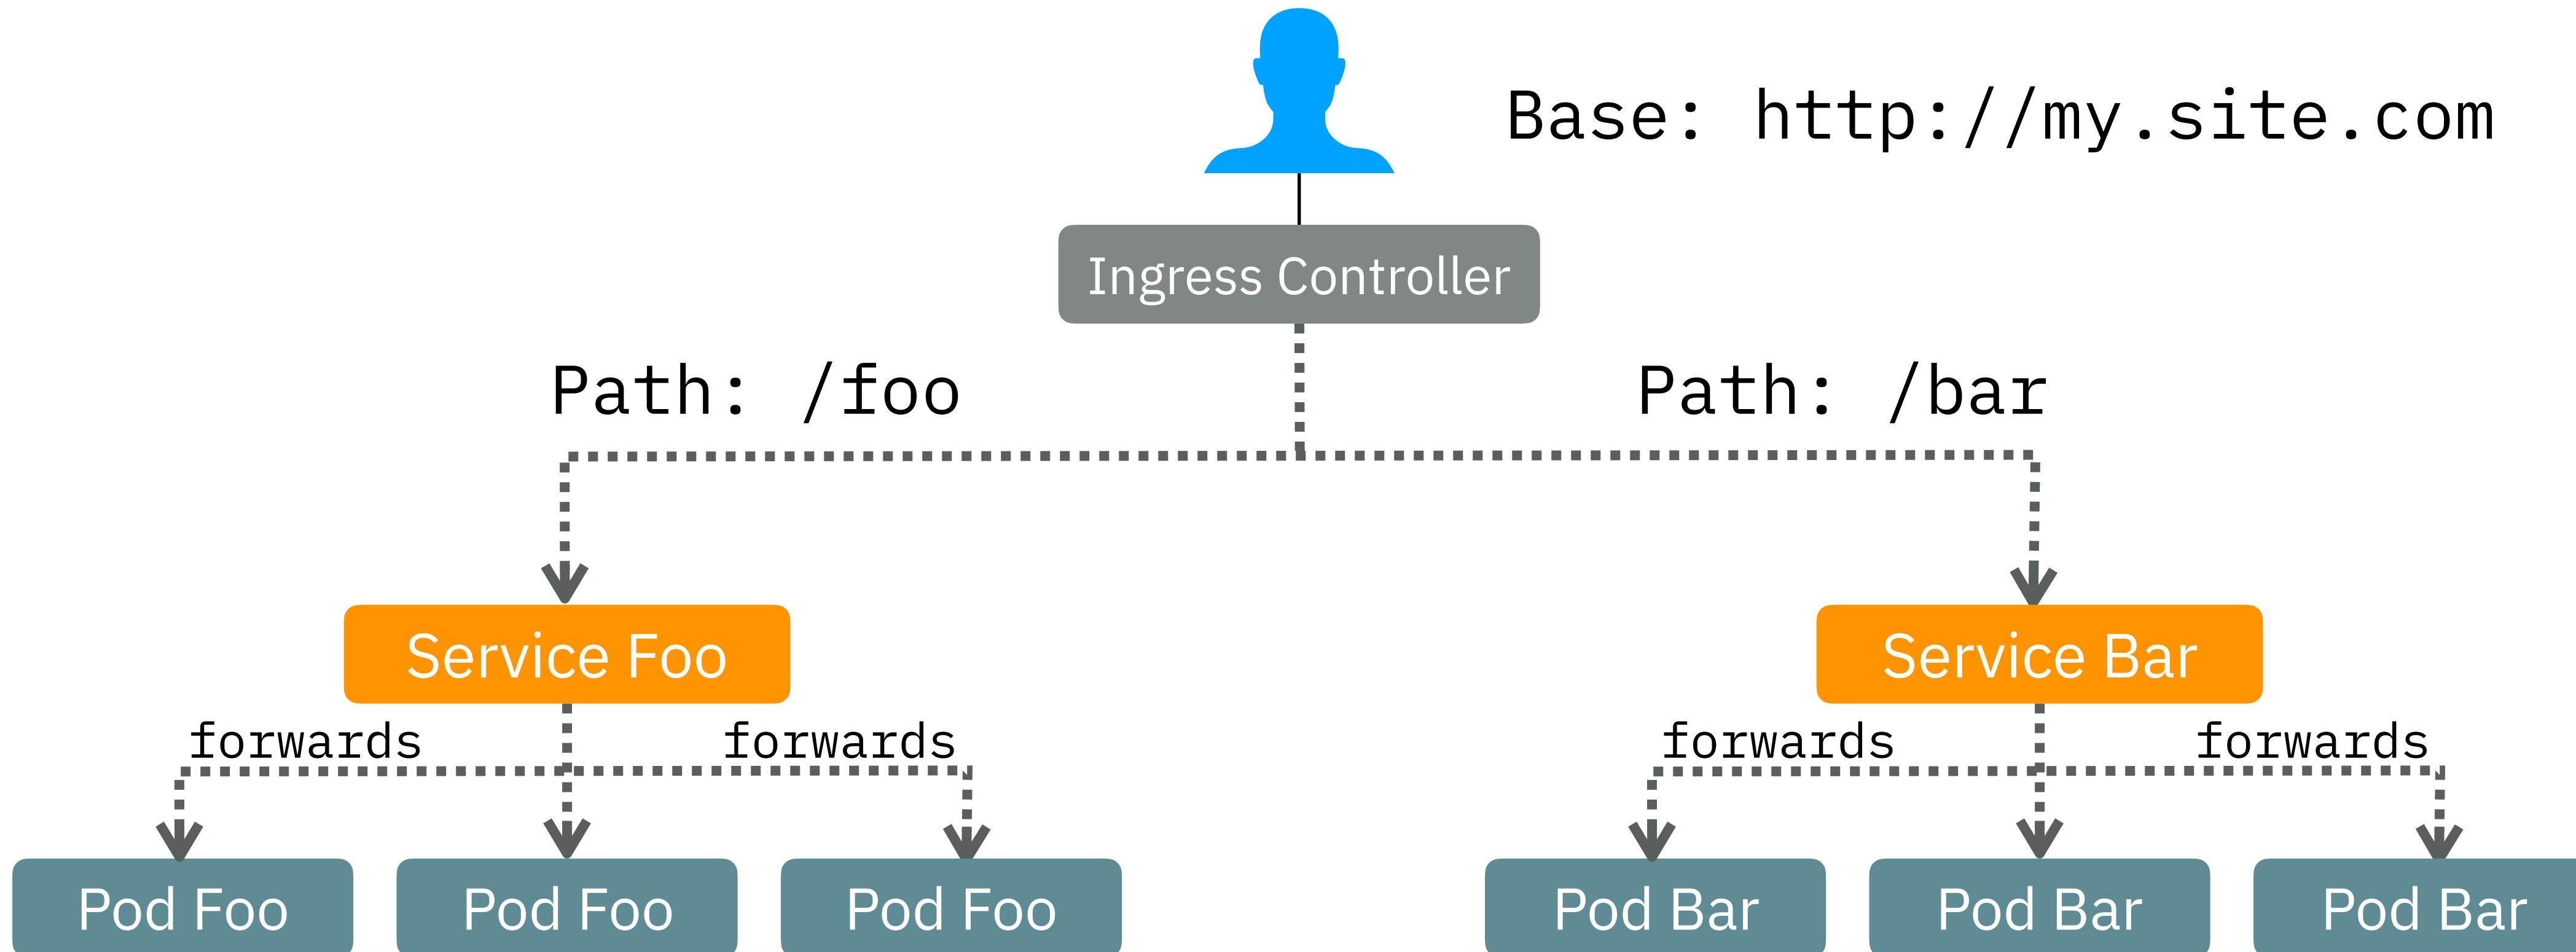
```
1 apiVersion: extensions/v1beta1
2 kind: Ingress
3 metadata:
4   name: ecommerce
5 spec:
6   # tls:
7   # - secretName: tls
8   rules:
9     - host: ecommerce.mycompany.com
10    http:
11      paths:
12        - path: /shopcarts
13          backend:
14            serviceName: shopcart-service
15            servicePort: 5000
16        - path: /catalog
17          backend:
18            serviceName: catalog-service
19            servicePort: 5000
20        - path: /orders
21          backend:
22            serviceName: order-service
23            servicePort: 5000
24        - path: /recommendations
25          backend:
26            serviceName: recommendation-service
27            servicePort: 5000
```

Requests coming in on this URL

Based on these PATHs

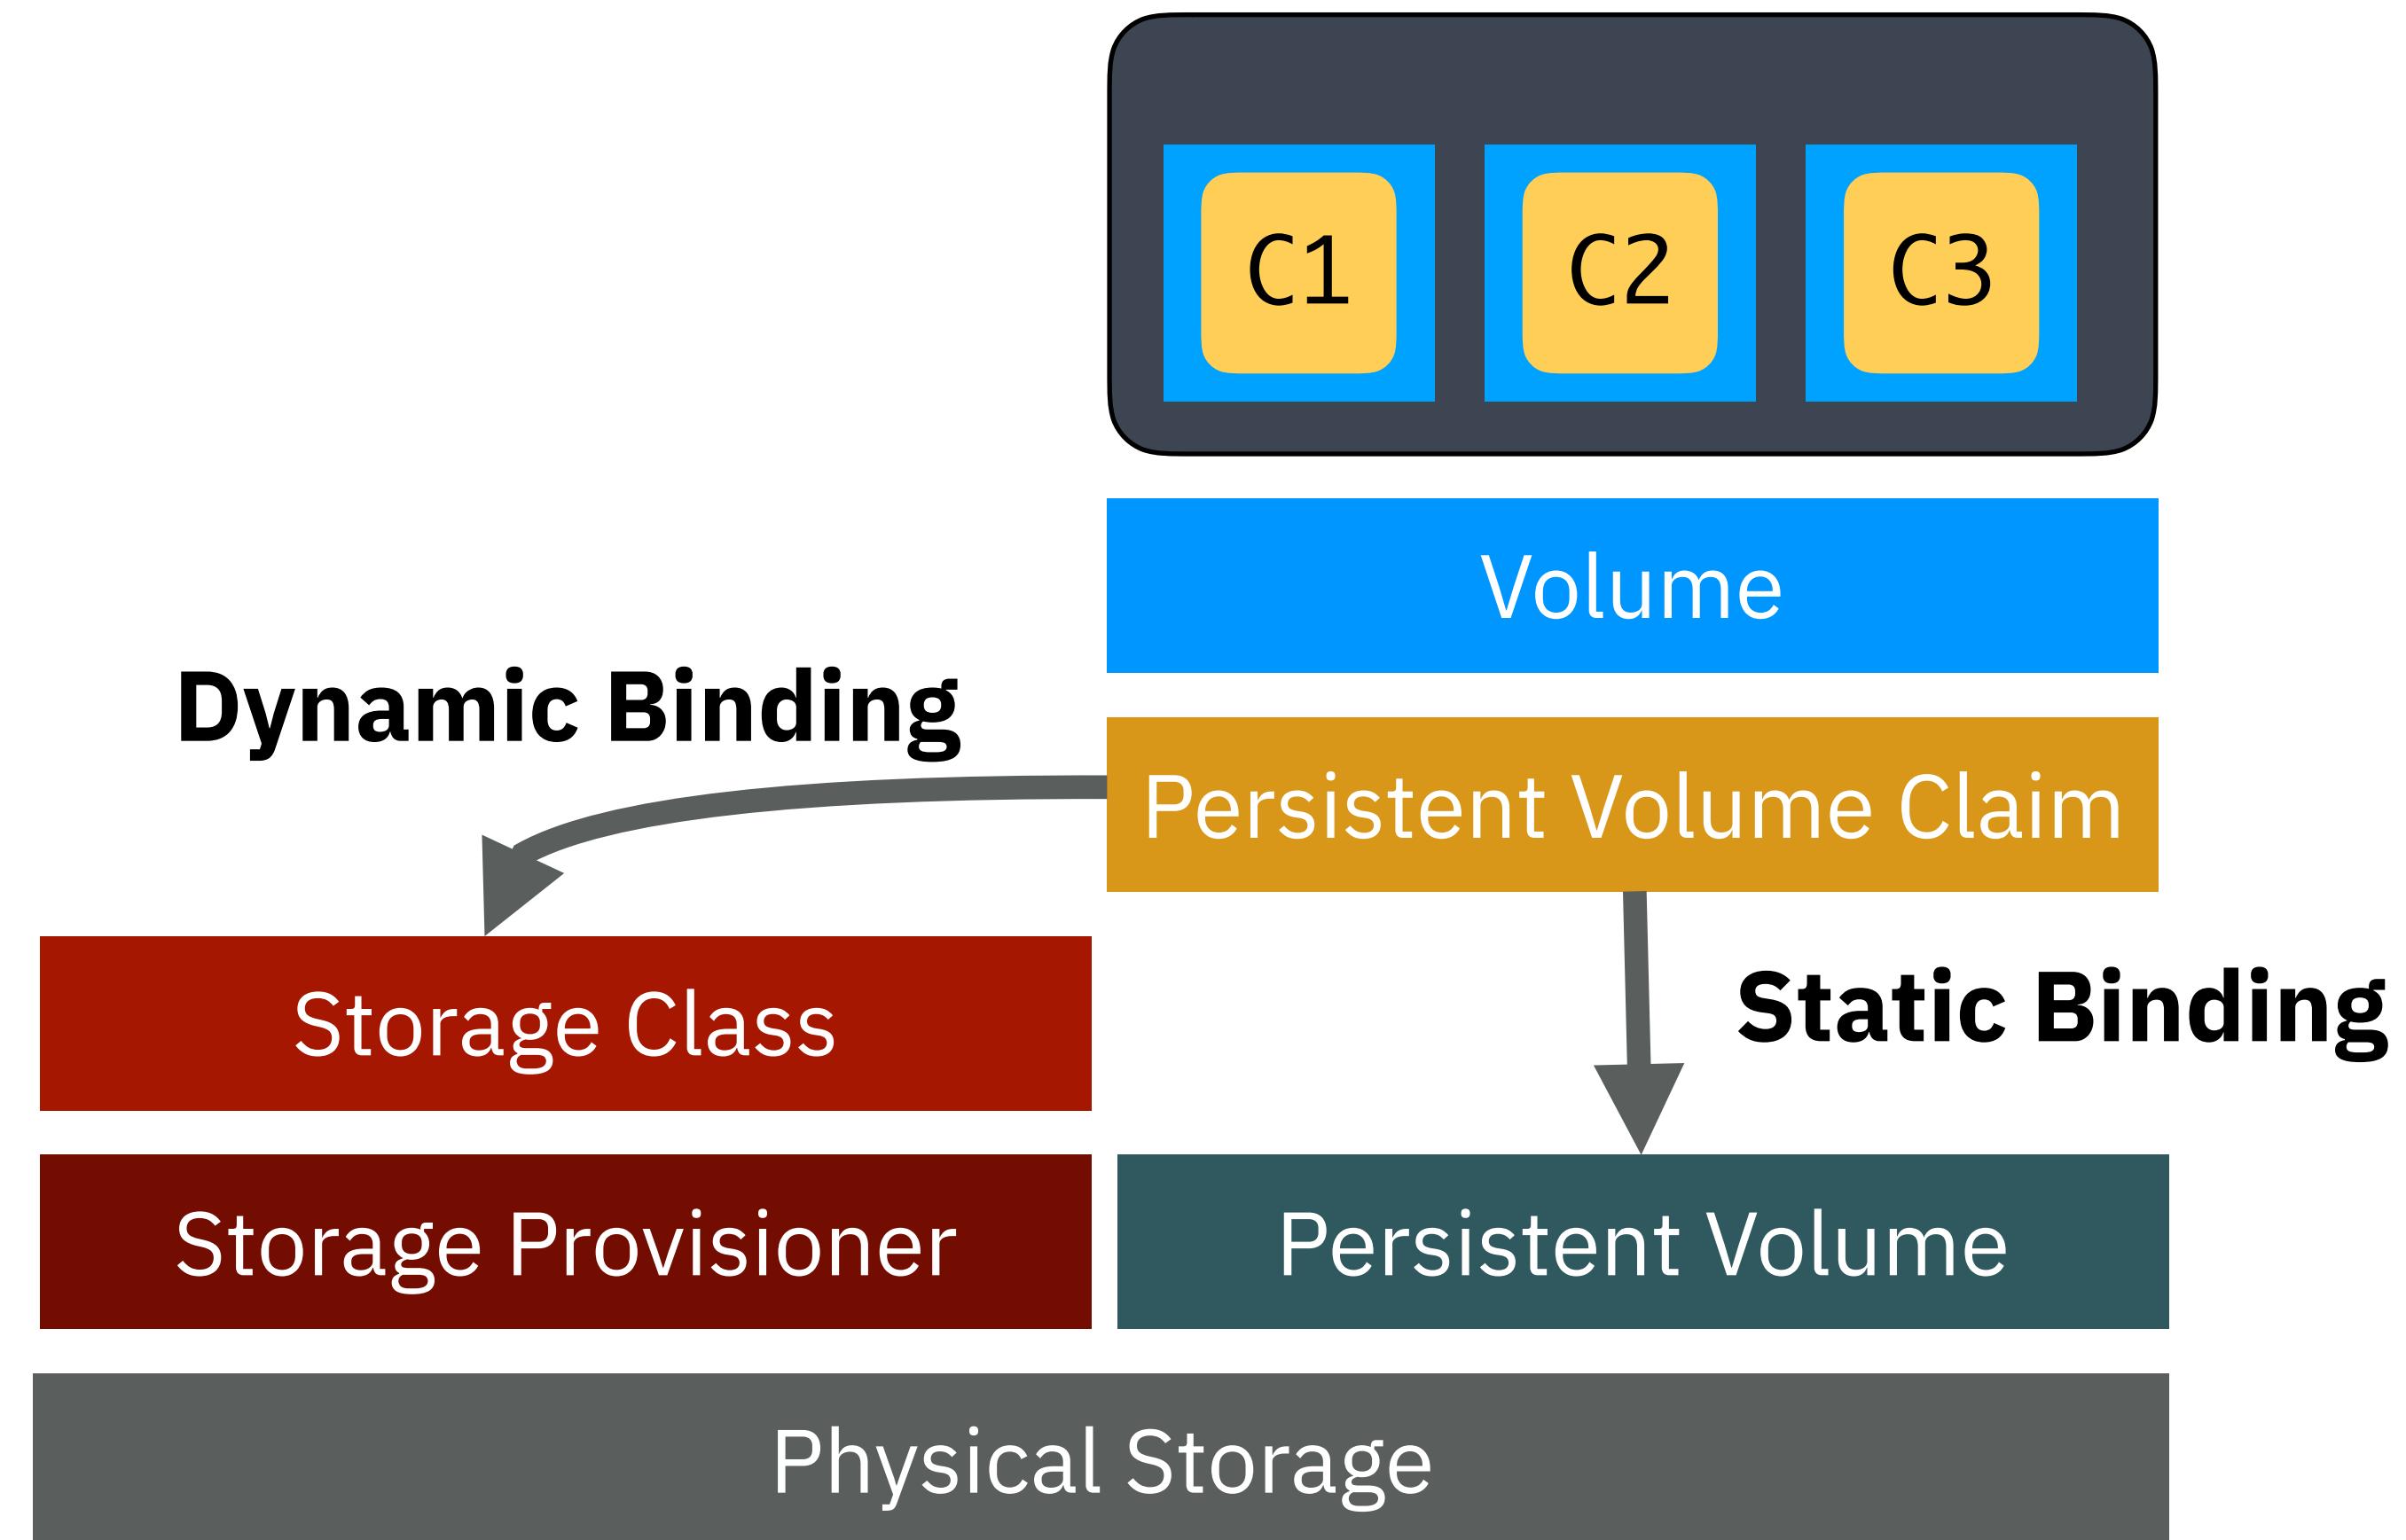
# Ingress Controller

- Single entry point into multiple kubernetes services



# Persistent Volumes

- Kubernetes loosely couples physical storage devices with containers by introducing an intermediate resource called persistent volume claims (PVCs).
- A PVC defines the disk size, disk type (ReadWriteOnce, ReadOnlyMany, ReadWriteMany) and dynamically links a storage device to a volume defined against a pod
- The binding process can either be done in a static way using PVs or dynamically be using a persistent storage provider



# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

~

# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

Volume named `redis-storage`

# Example Volume Mount

- There is a volume named `redis-storage` and that is connected to the `redis` container via the `VolumeMount: /data/redis`

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

Mounted at: /data/redis

Volume named redis-storage

# Configurations Management

(ConfigMap)

- Containers generally use environment variables for parameterizing their runtime configurations
- Kubernetes provides a way of managing more complex configuration files using a simple resource called ConfigMaps
- ConfigMaps can be created using directories, files or literal values using following CLI command:

```
$ kubectl create configmap <map-name> <data-source>  
  
# map-name: name of the config map  
# data-source: directory, file or literal value
```

# Credentials Management

(Secrets)

- Similar to ConfigMaps, Kubernetes provides another valuable resource called Secrets for managing sensitive information such as passwords, OAuth tokens, and ssh keys.
- A secret can be created for managing basic auth credentials using the following way:

```
# write credentials to two files
$ echo -n 'admin' > ./username.txt
$ echo -n '1f2d1e2e67df' > ./password.txt

# create a secret
$ kubectl create secret generic app-credentials --from-file=./username.txt --from-file=./password.txt
```

# Secret Yaml Example

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: ecommerce-apikey
5   namespace: default
6 data:
7   secret: <place base64 encoded secret here>
```

```
$ echo -n "this is my secret" | base64
dGhpcyBpcyBteSBzZWNyZXQ=
```

# Secret Yaml Example

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: ecommerce-apikey
5   namespace: default
6 data:
7 secret: <place base64 encoded secret here>
```

Place the base64 encode string in secret

```
$ echo -n "this is my secret" | base64
dGhpcyBpcyBteSBzZWNyZXQ=
```

# Using Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: pet-creds
  namespace: default
data:
  binding: dGhpcyBpcyBteSBzZWNyZXQ=
```

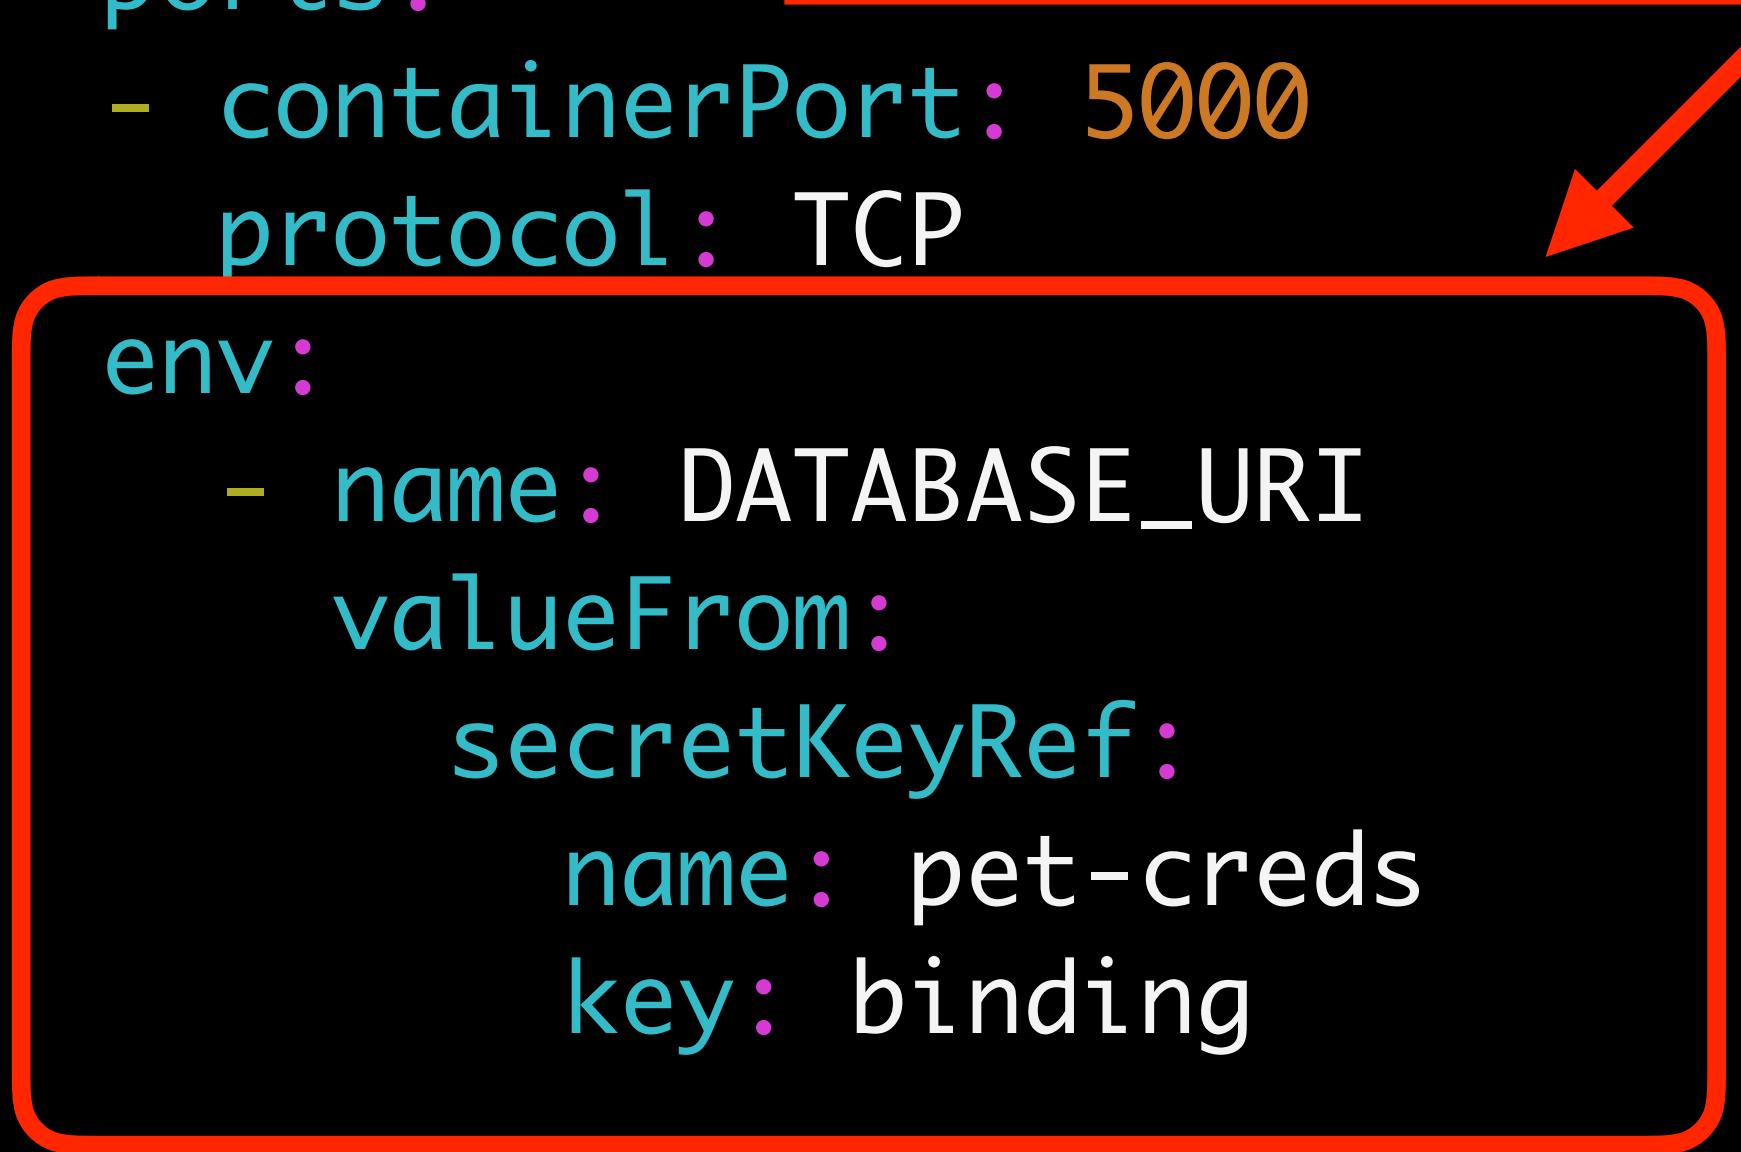
```
apiVersion: apps/v1
kind: Deployment
...
spec:
  containers:
    - name: pet-demo
      image: pet-demo:v1
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5000
          protocol: TCP
      env:
        - name: DATABASE_URI
          valueFrom:
            secretKeyRef:
              name: pet-creds
              key: binding
```

# Using Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: pet-creds
  namespace: default
data:
  binding: dGhpcyBpcyBteSBzZWNyZXQ=
```

```
apiVersion: apps/v1
kind: Deployment
...
spec:
  containers:
    - name: pet-demo
      image: pet-demo:v1
      imagePullPolicy: IfNotPresent
      ports:
        - containerPort: 5000
          protocol: TCP
      env:
        - name: DATABASE_URI
          valueFrom:
            secretKeyRef:
              name: pet-creds
              key: binding
```

Secrets exposed as environment variables



# Kubernetes Rolling Updates (Zero Downtime Deployments)

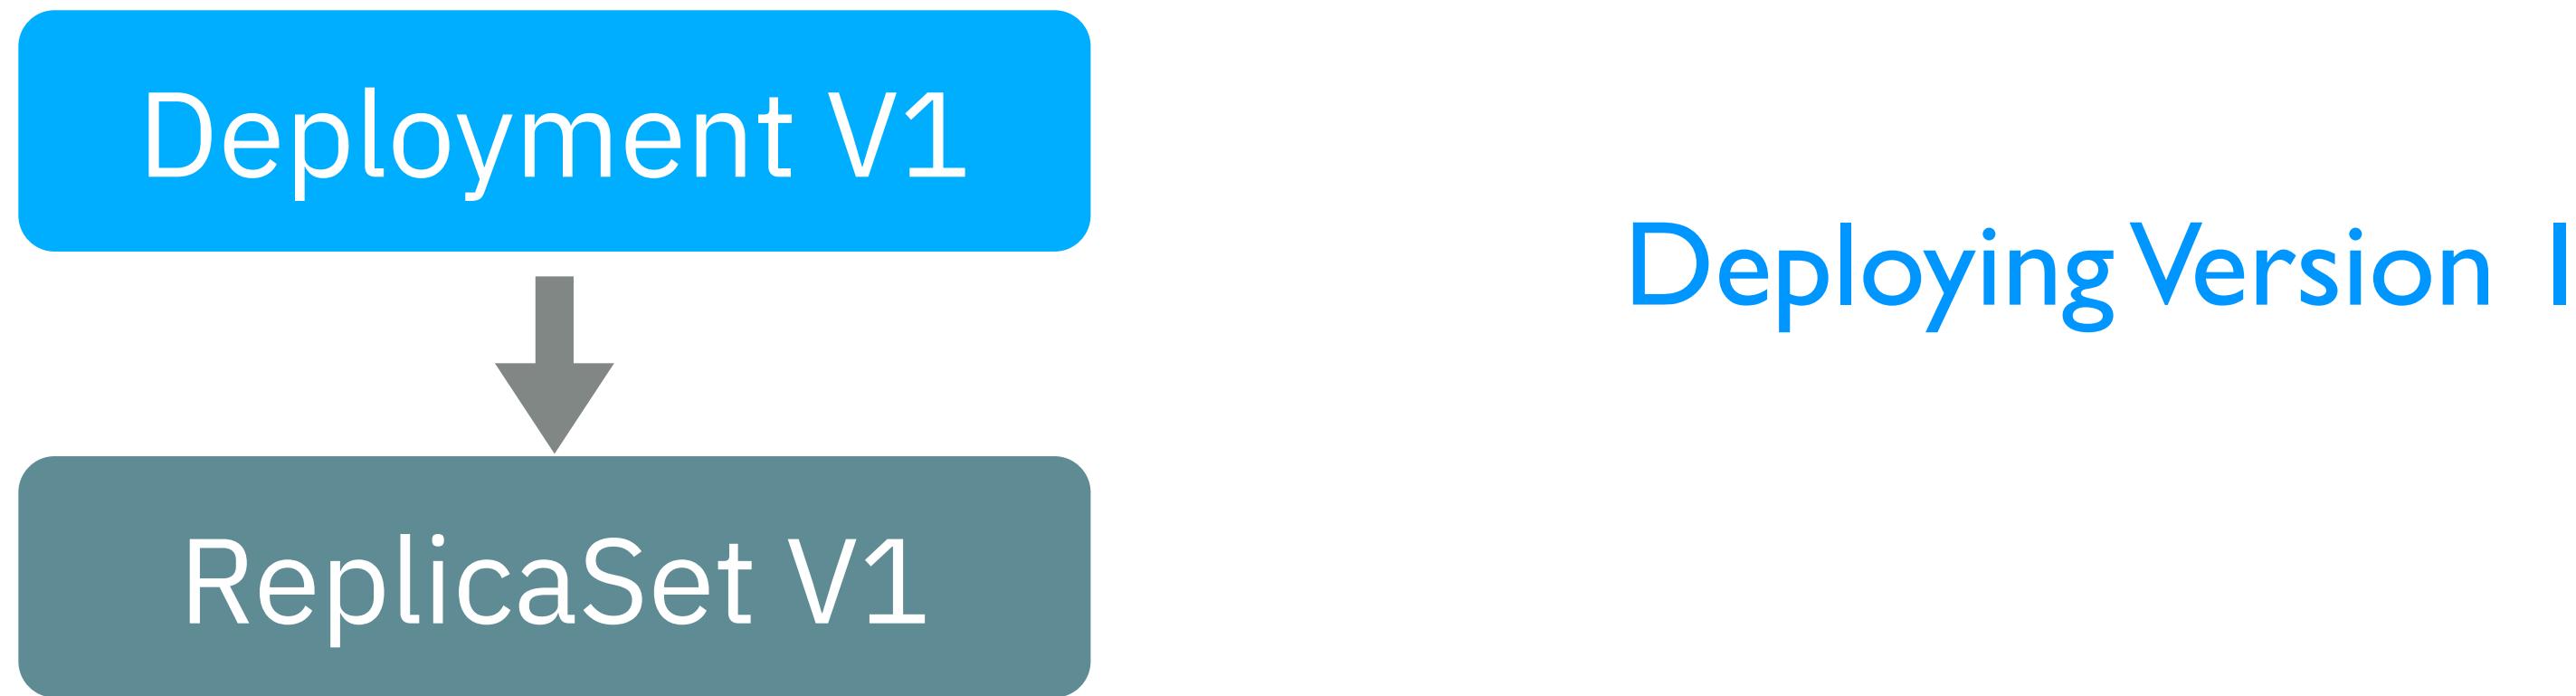
```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```

Deployment V1

Deploying Version I

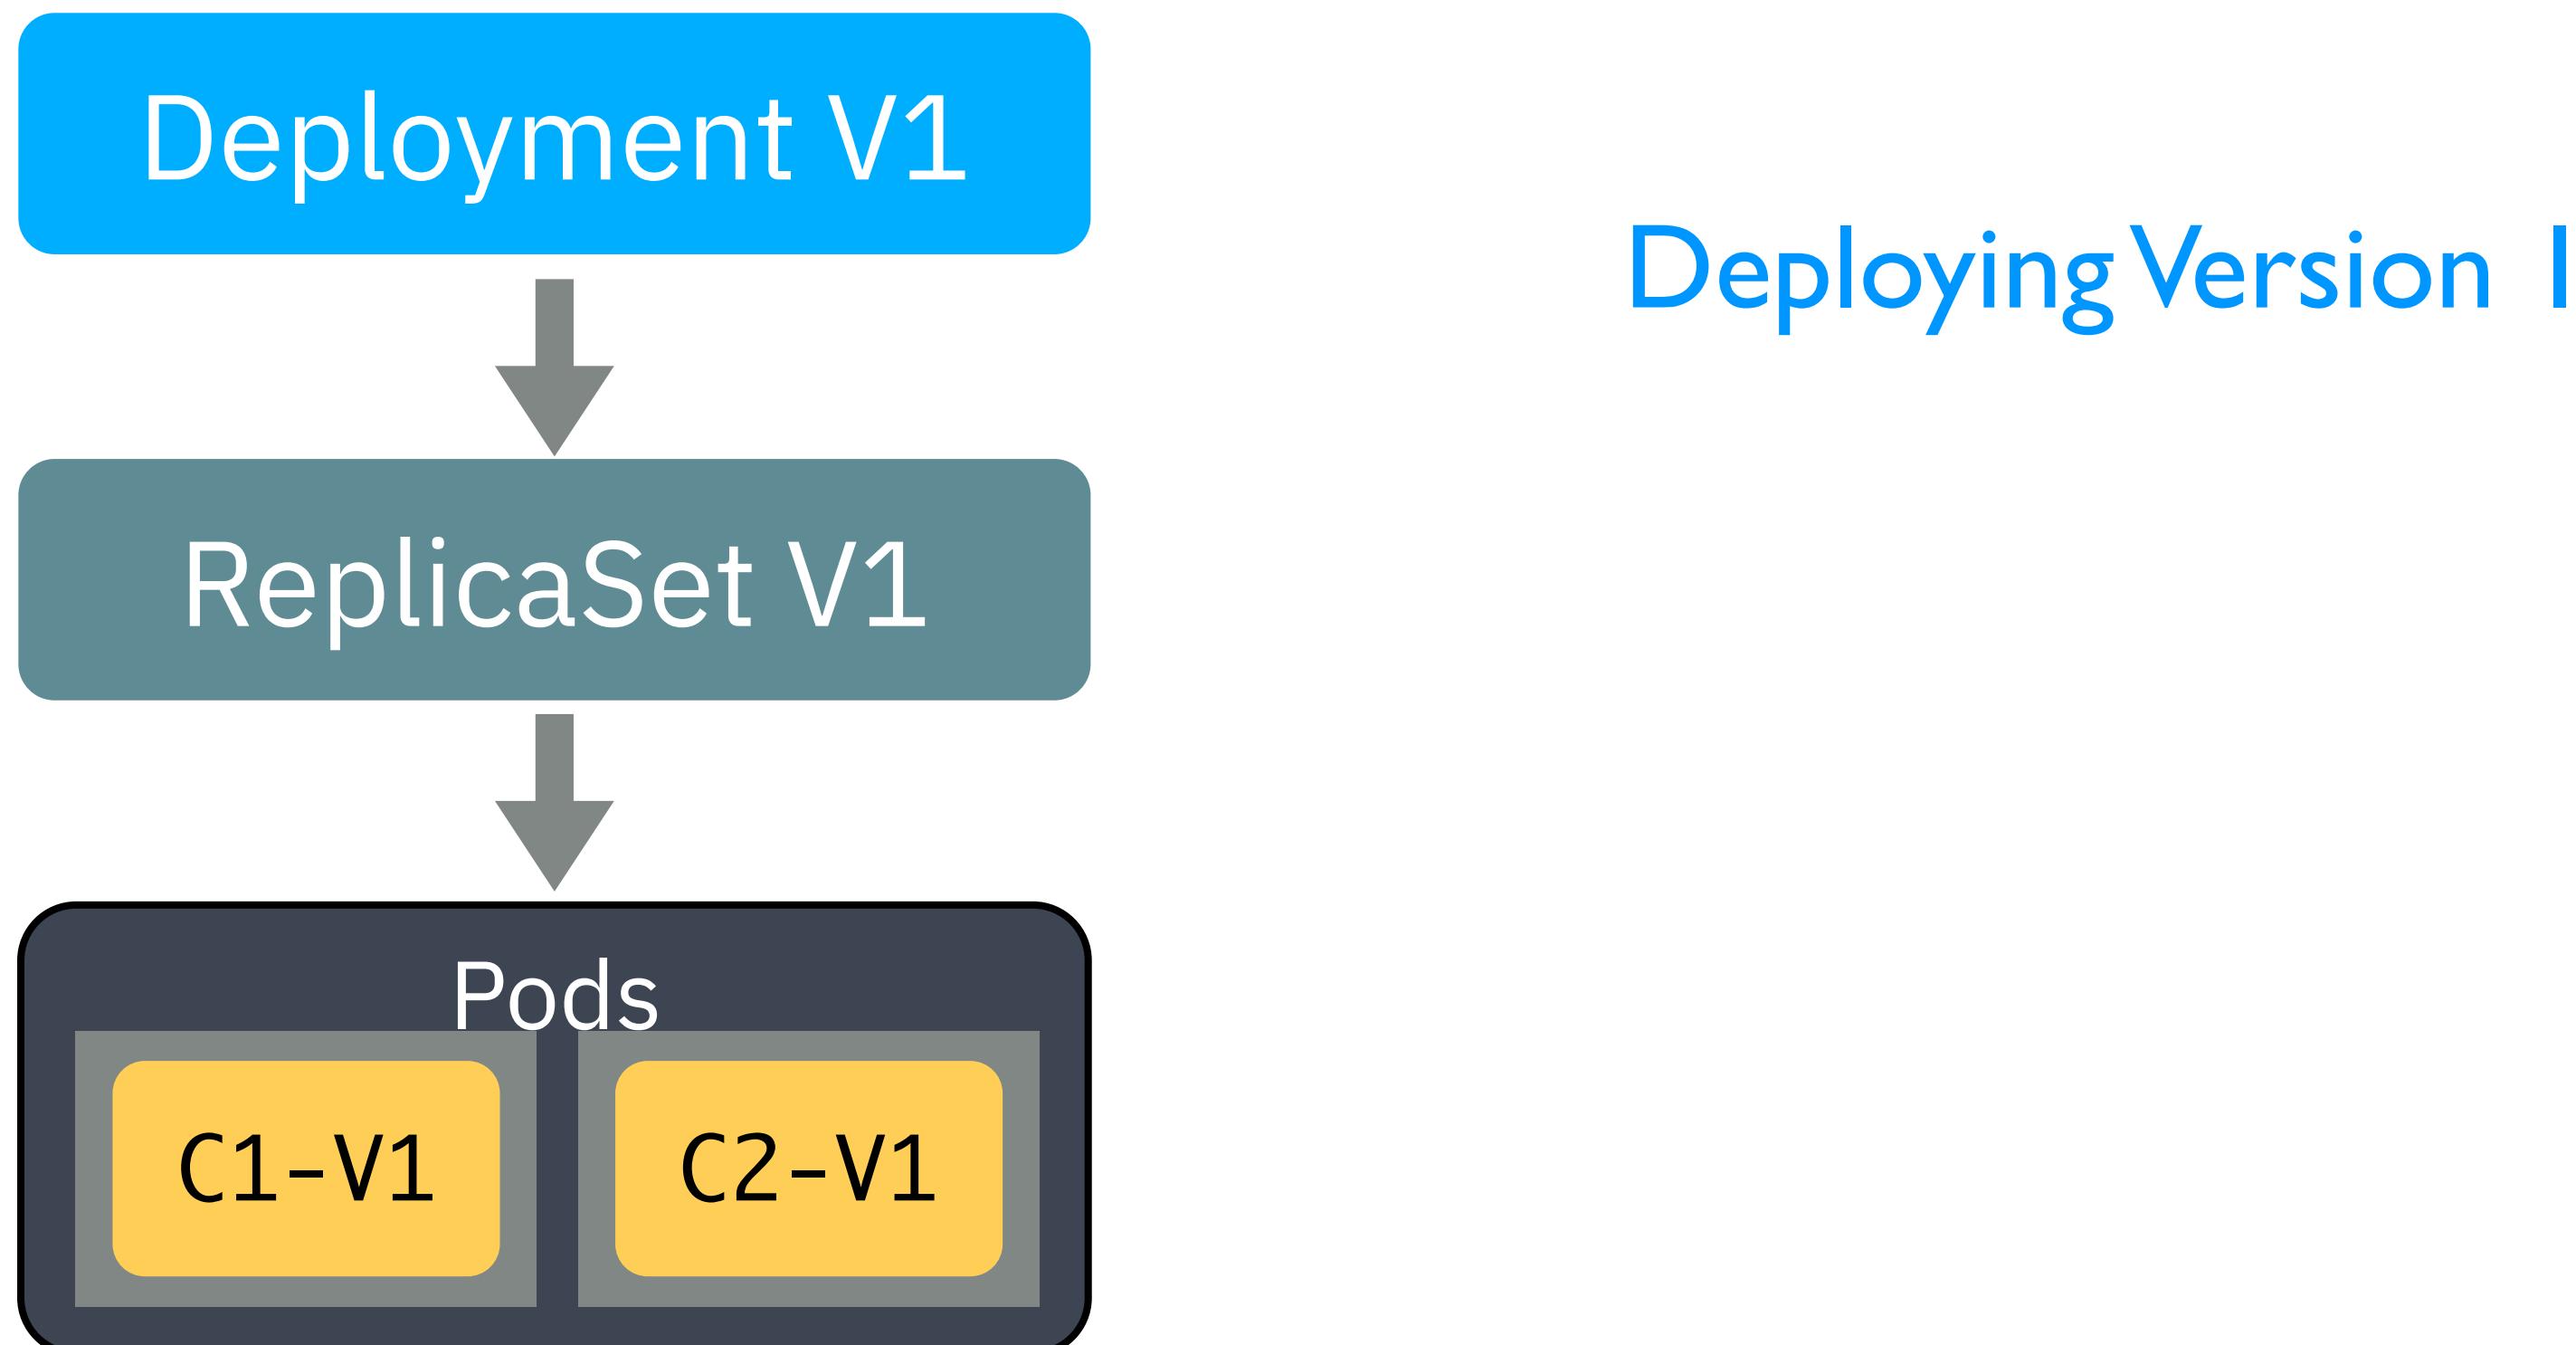
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



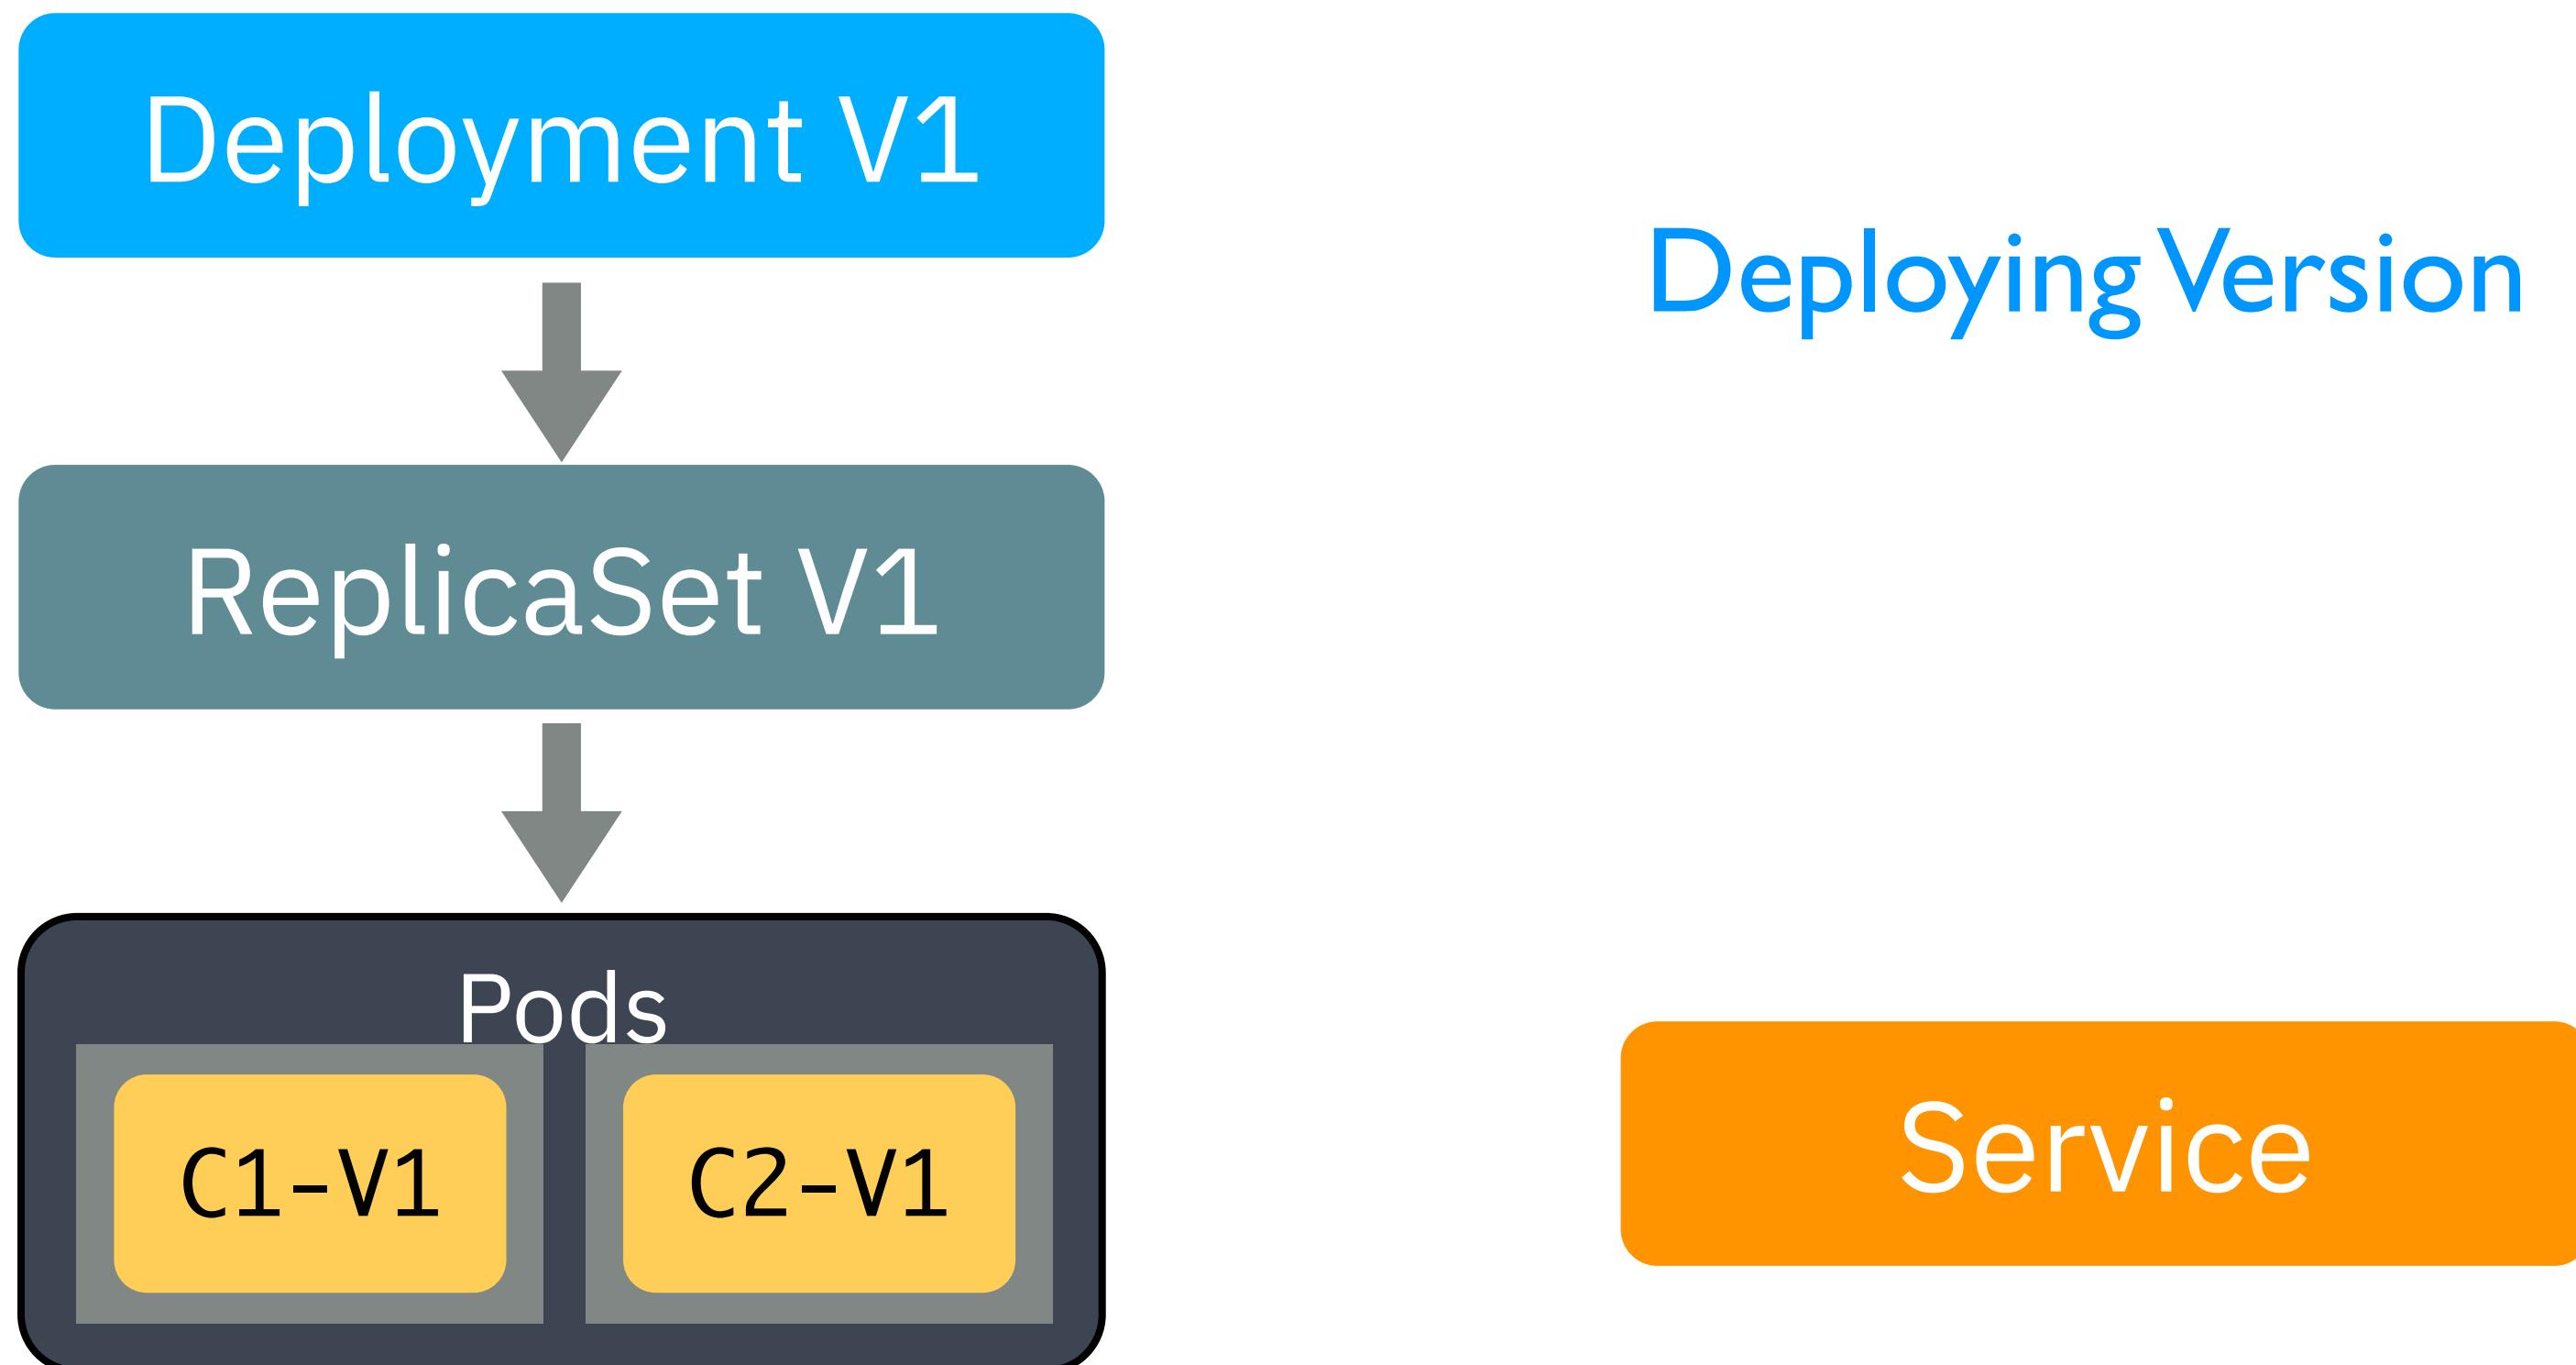
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



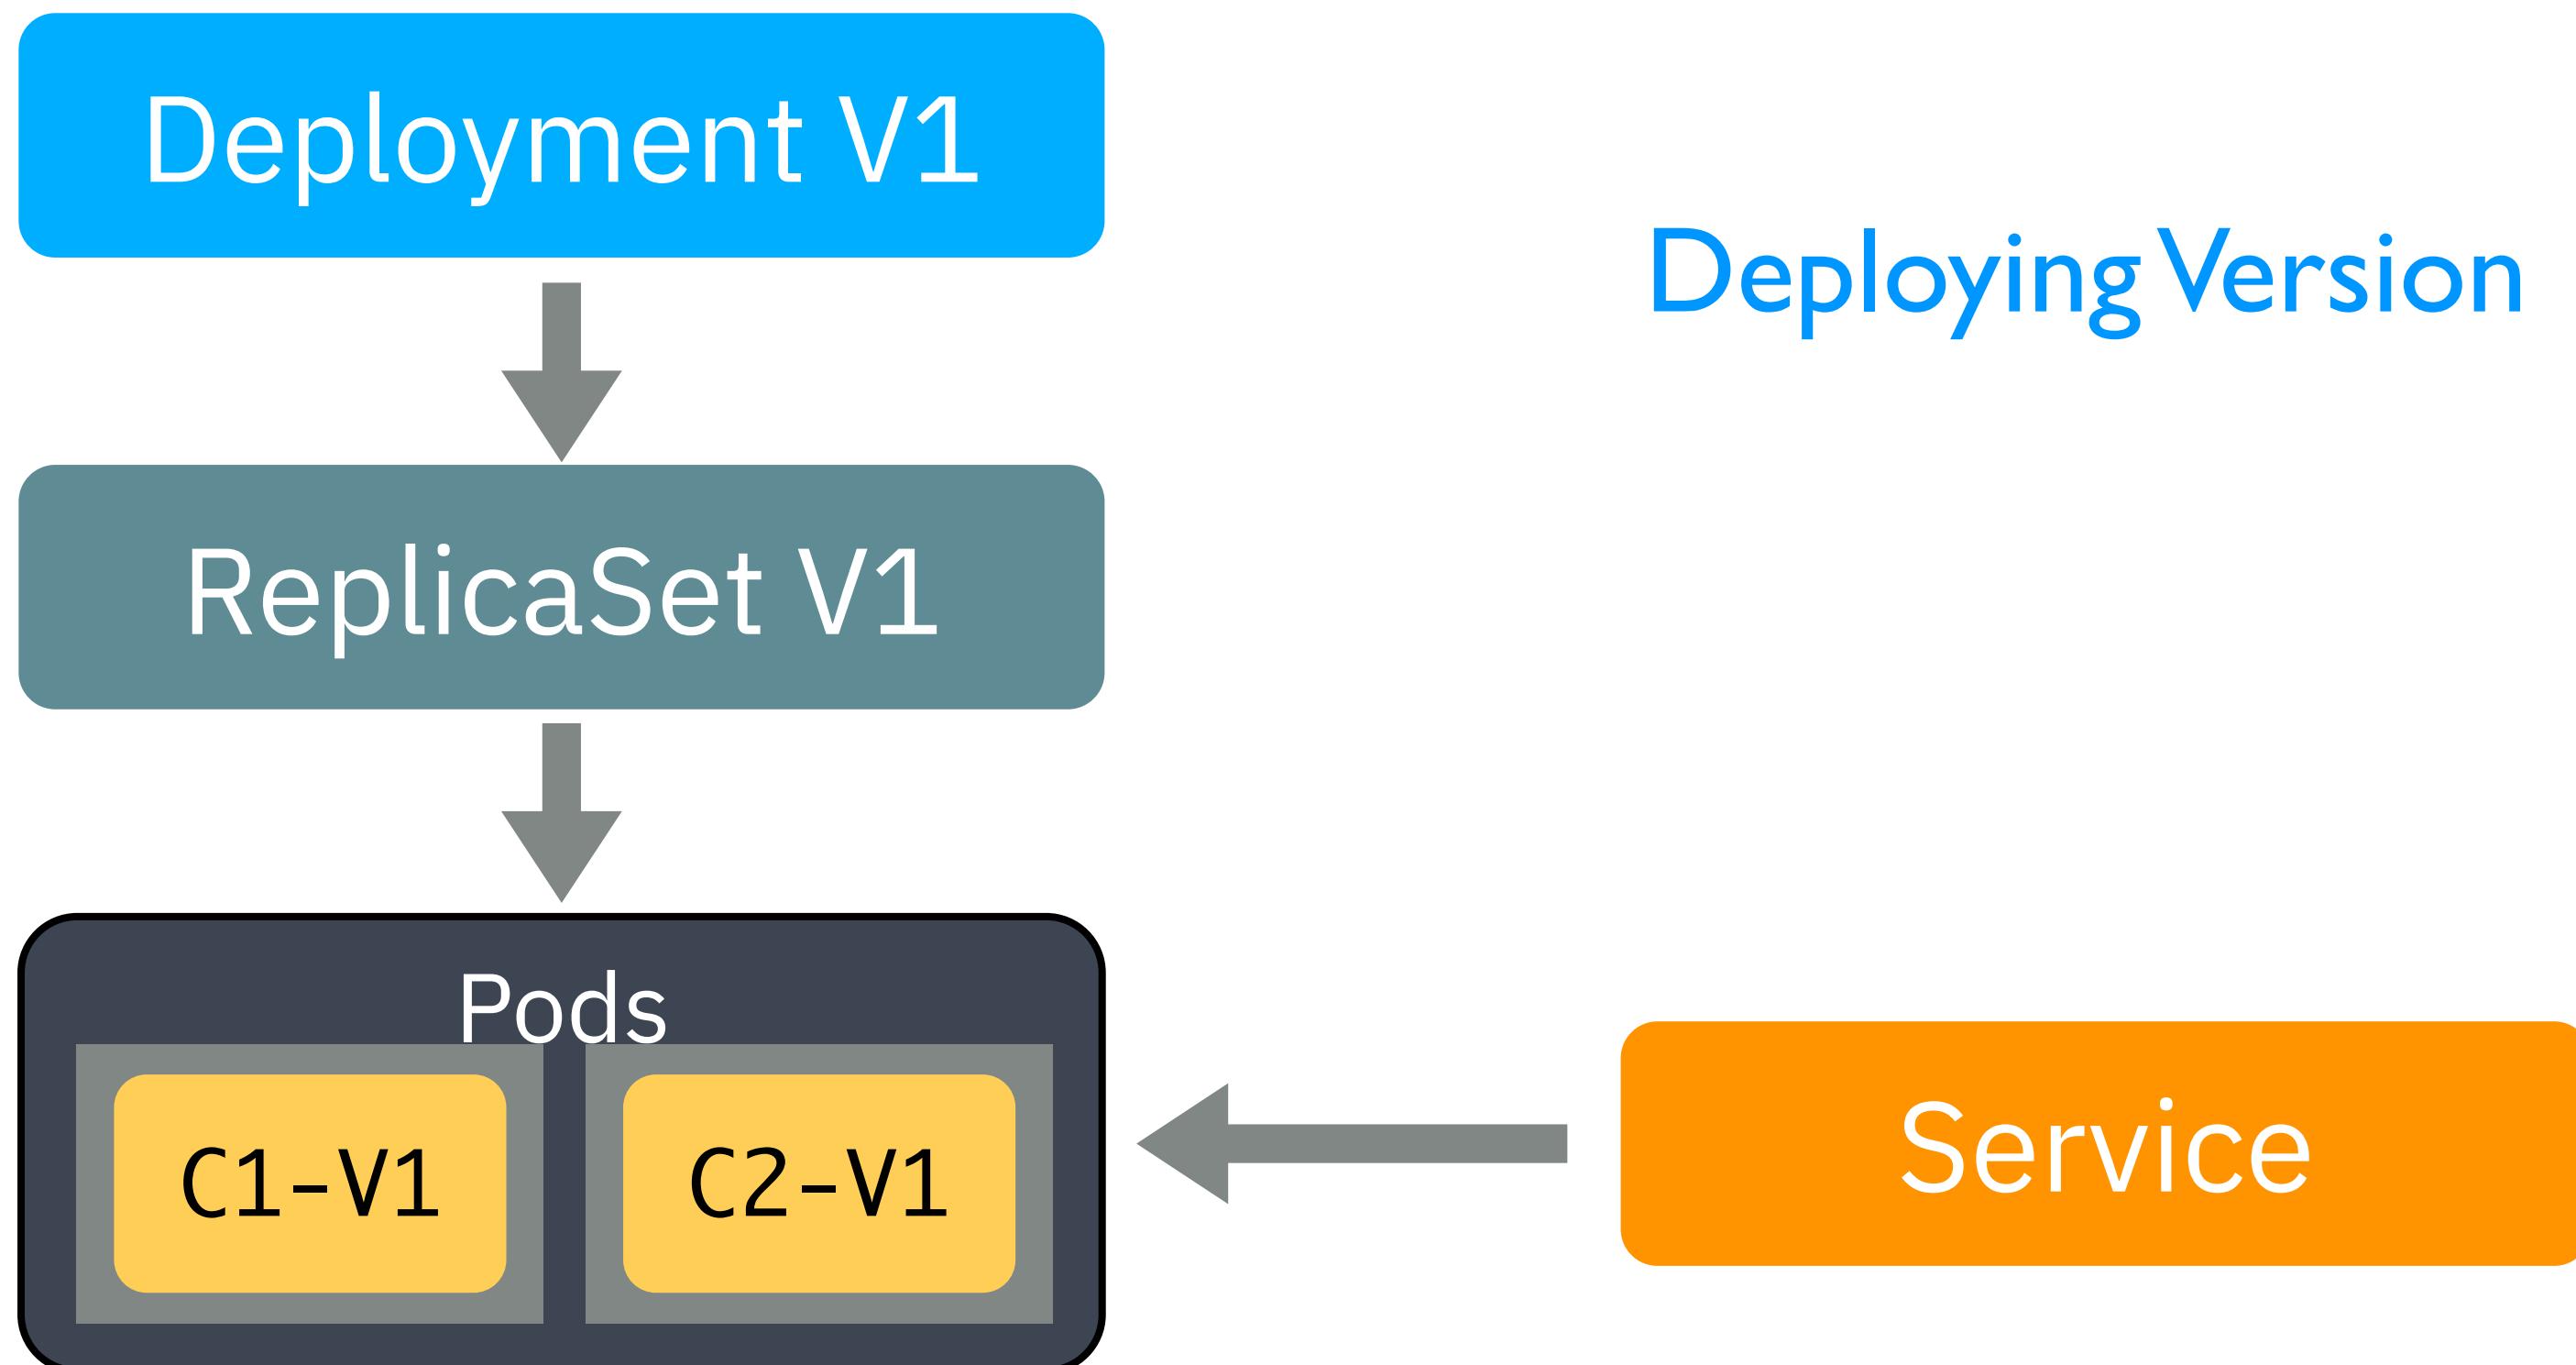
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



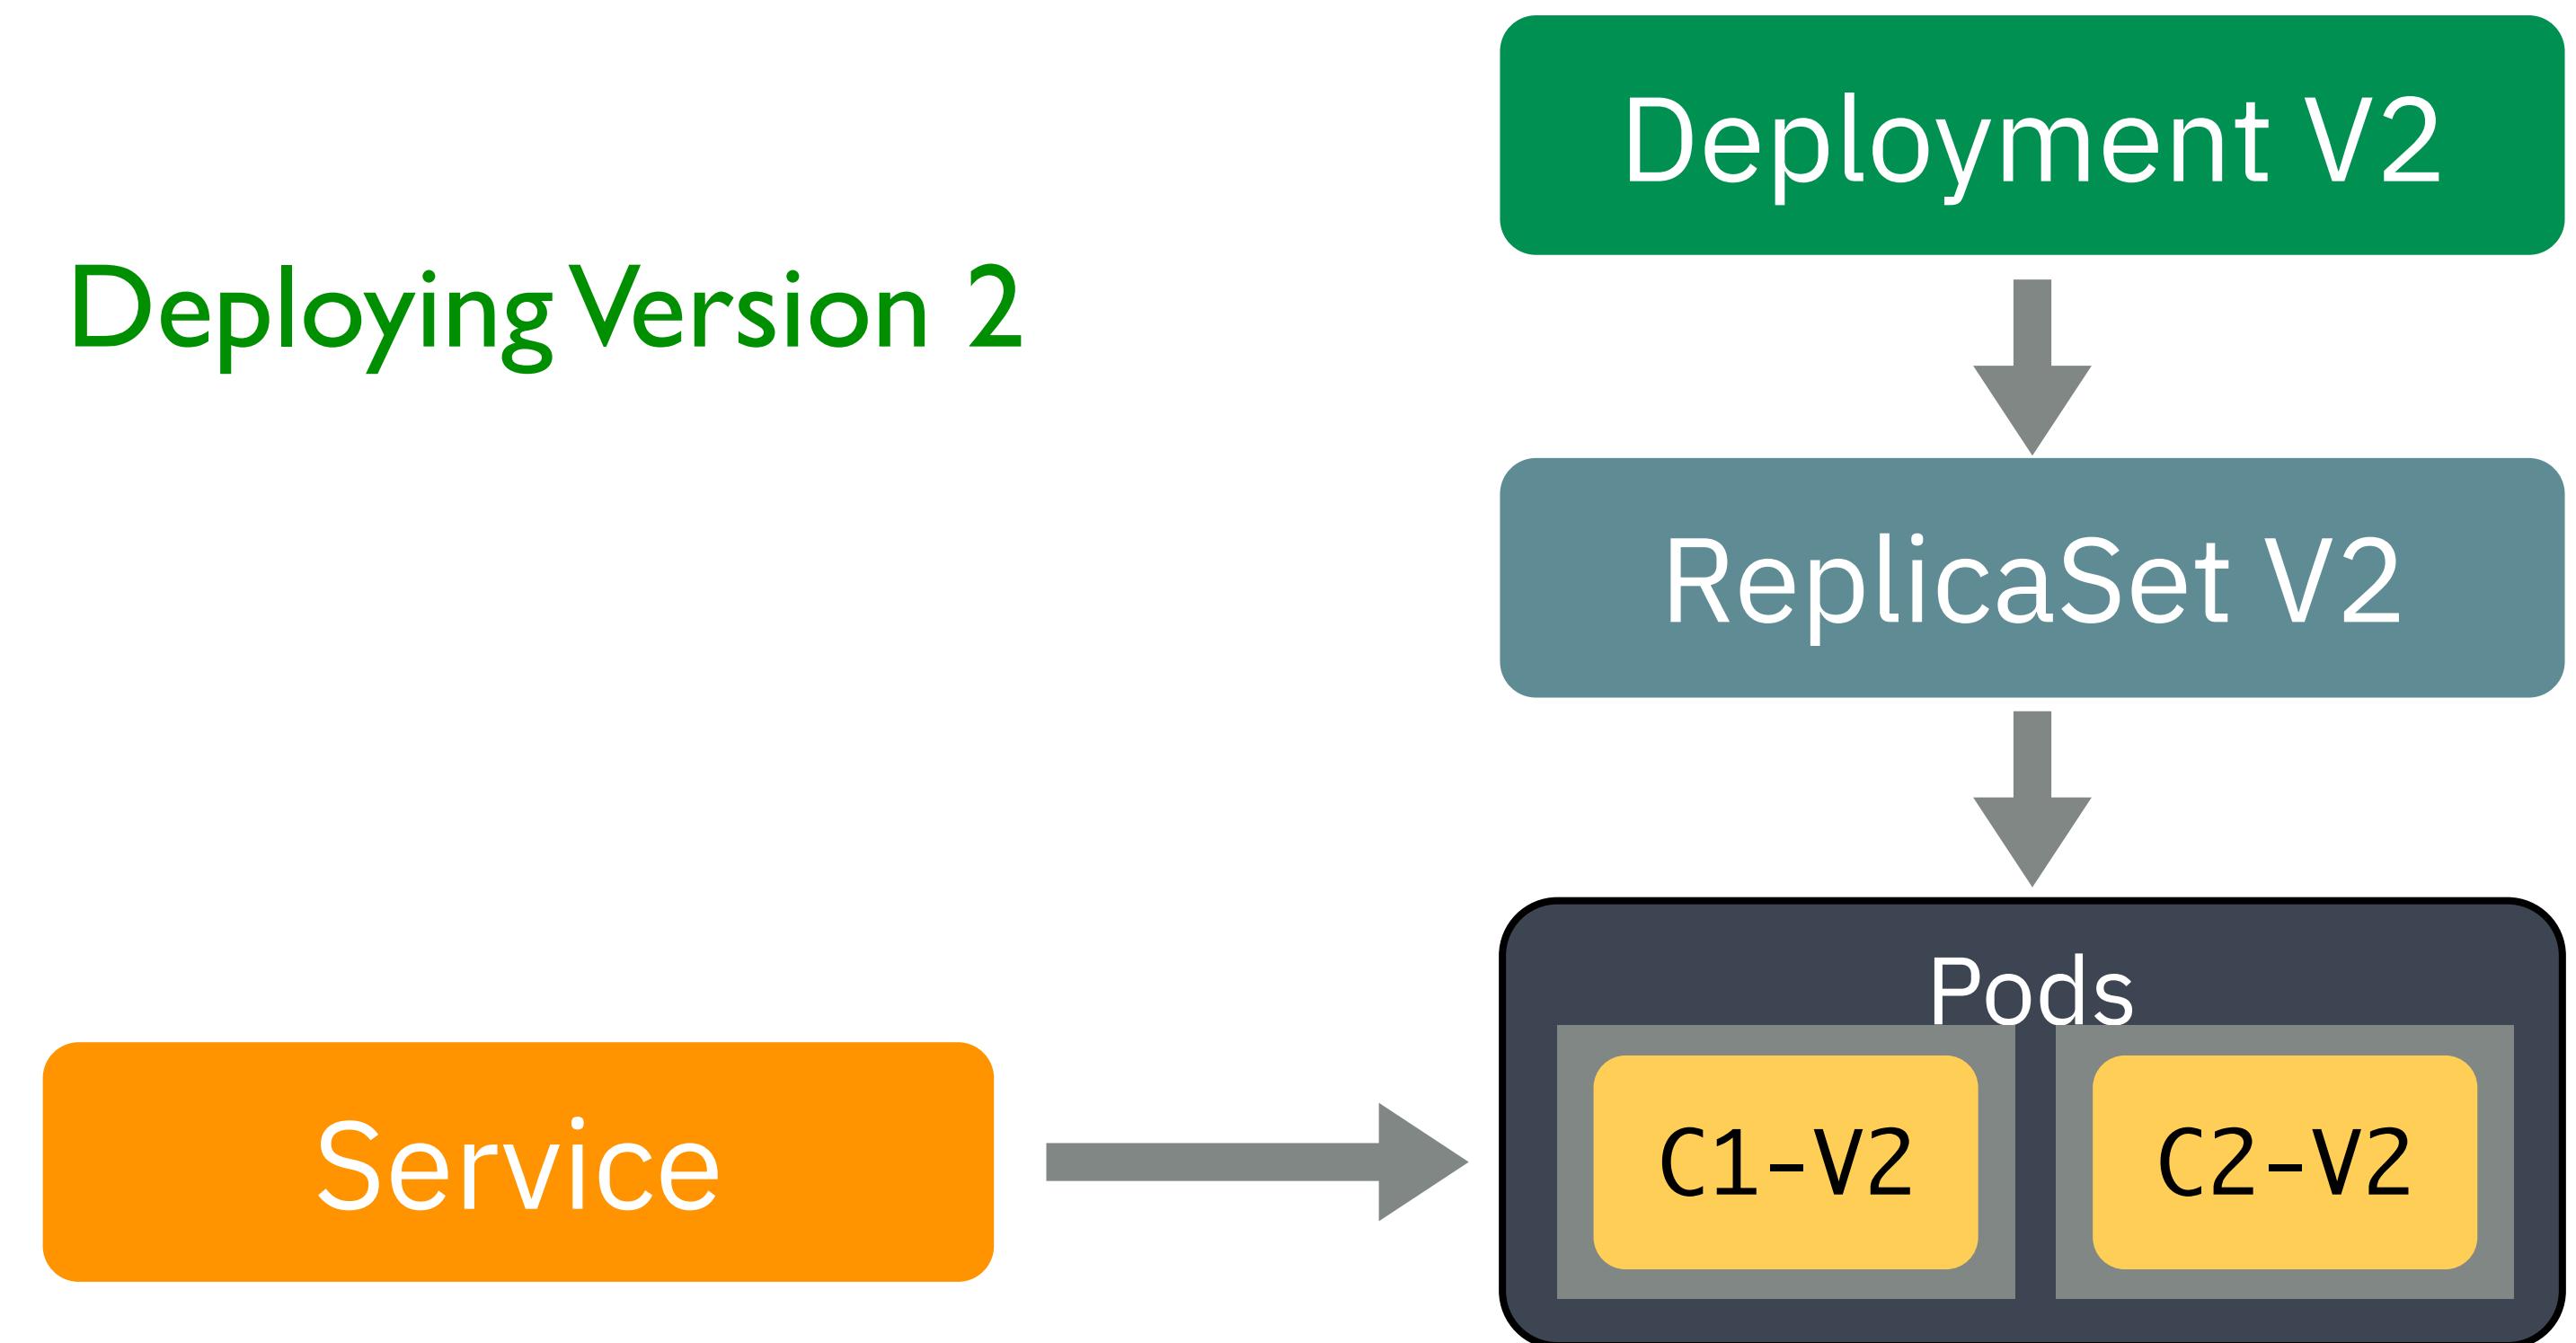
# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```



# Kubernetes Rolling Updates (Zero Downtime Deployments)

```
$ kubectl set image deployment/<application-name> <container-name>=<container-image-name>:<new-version>
```

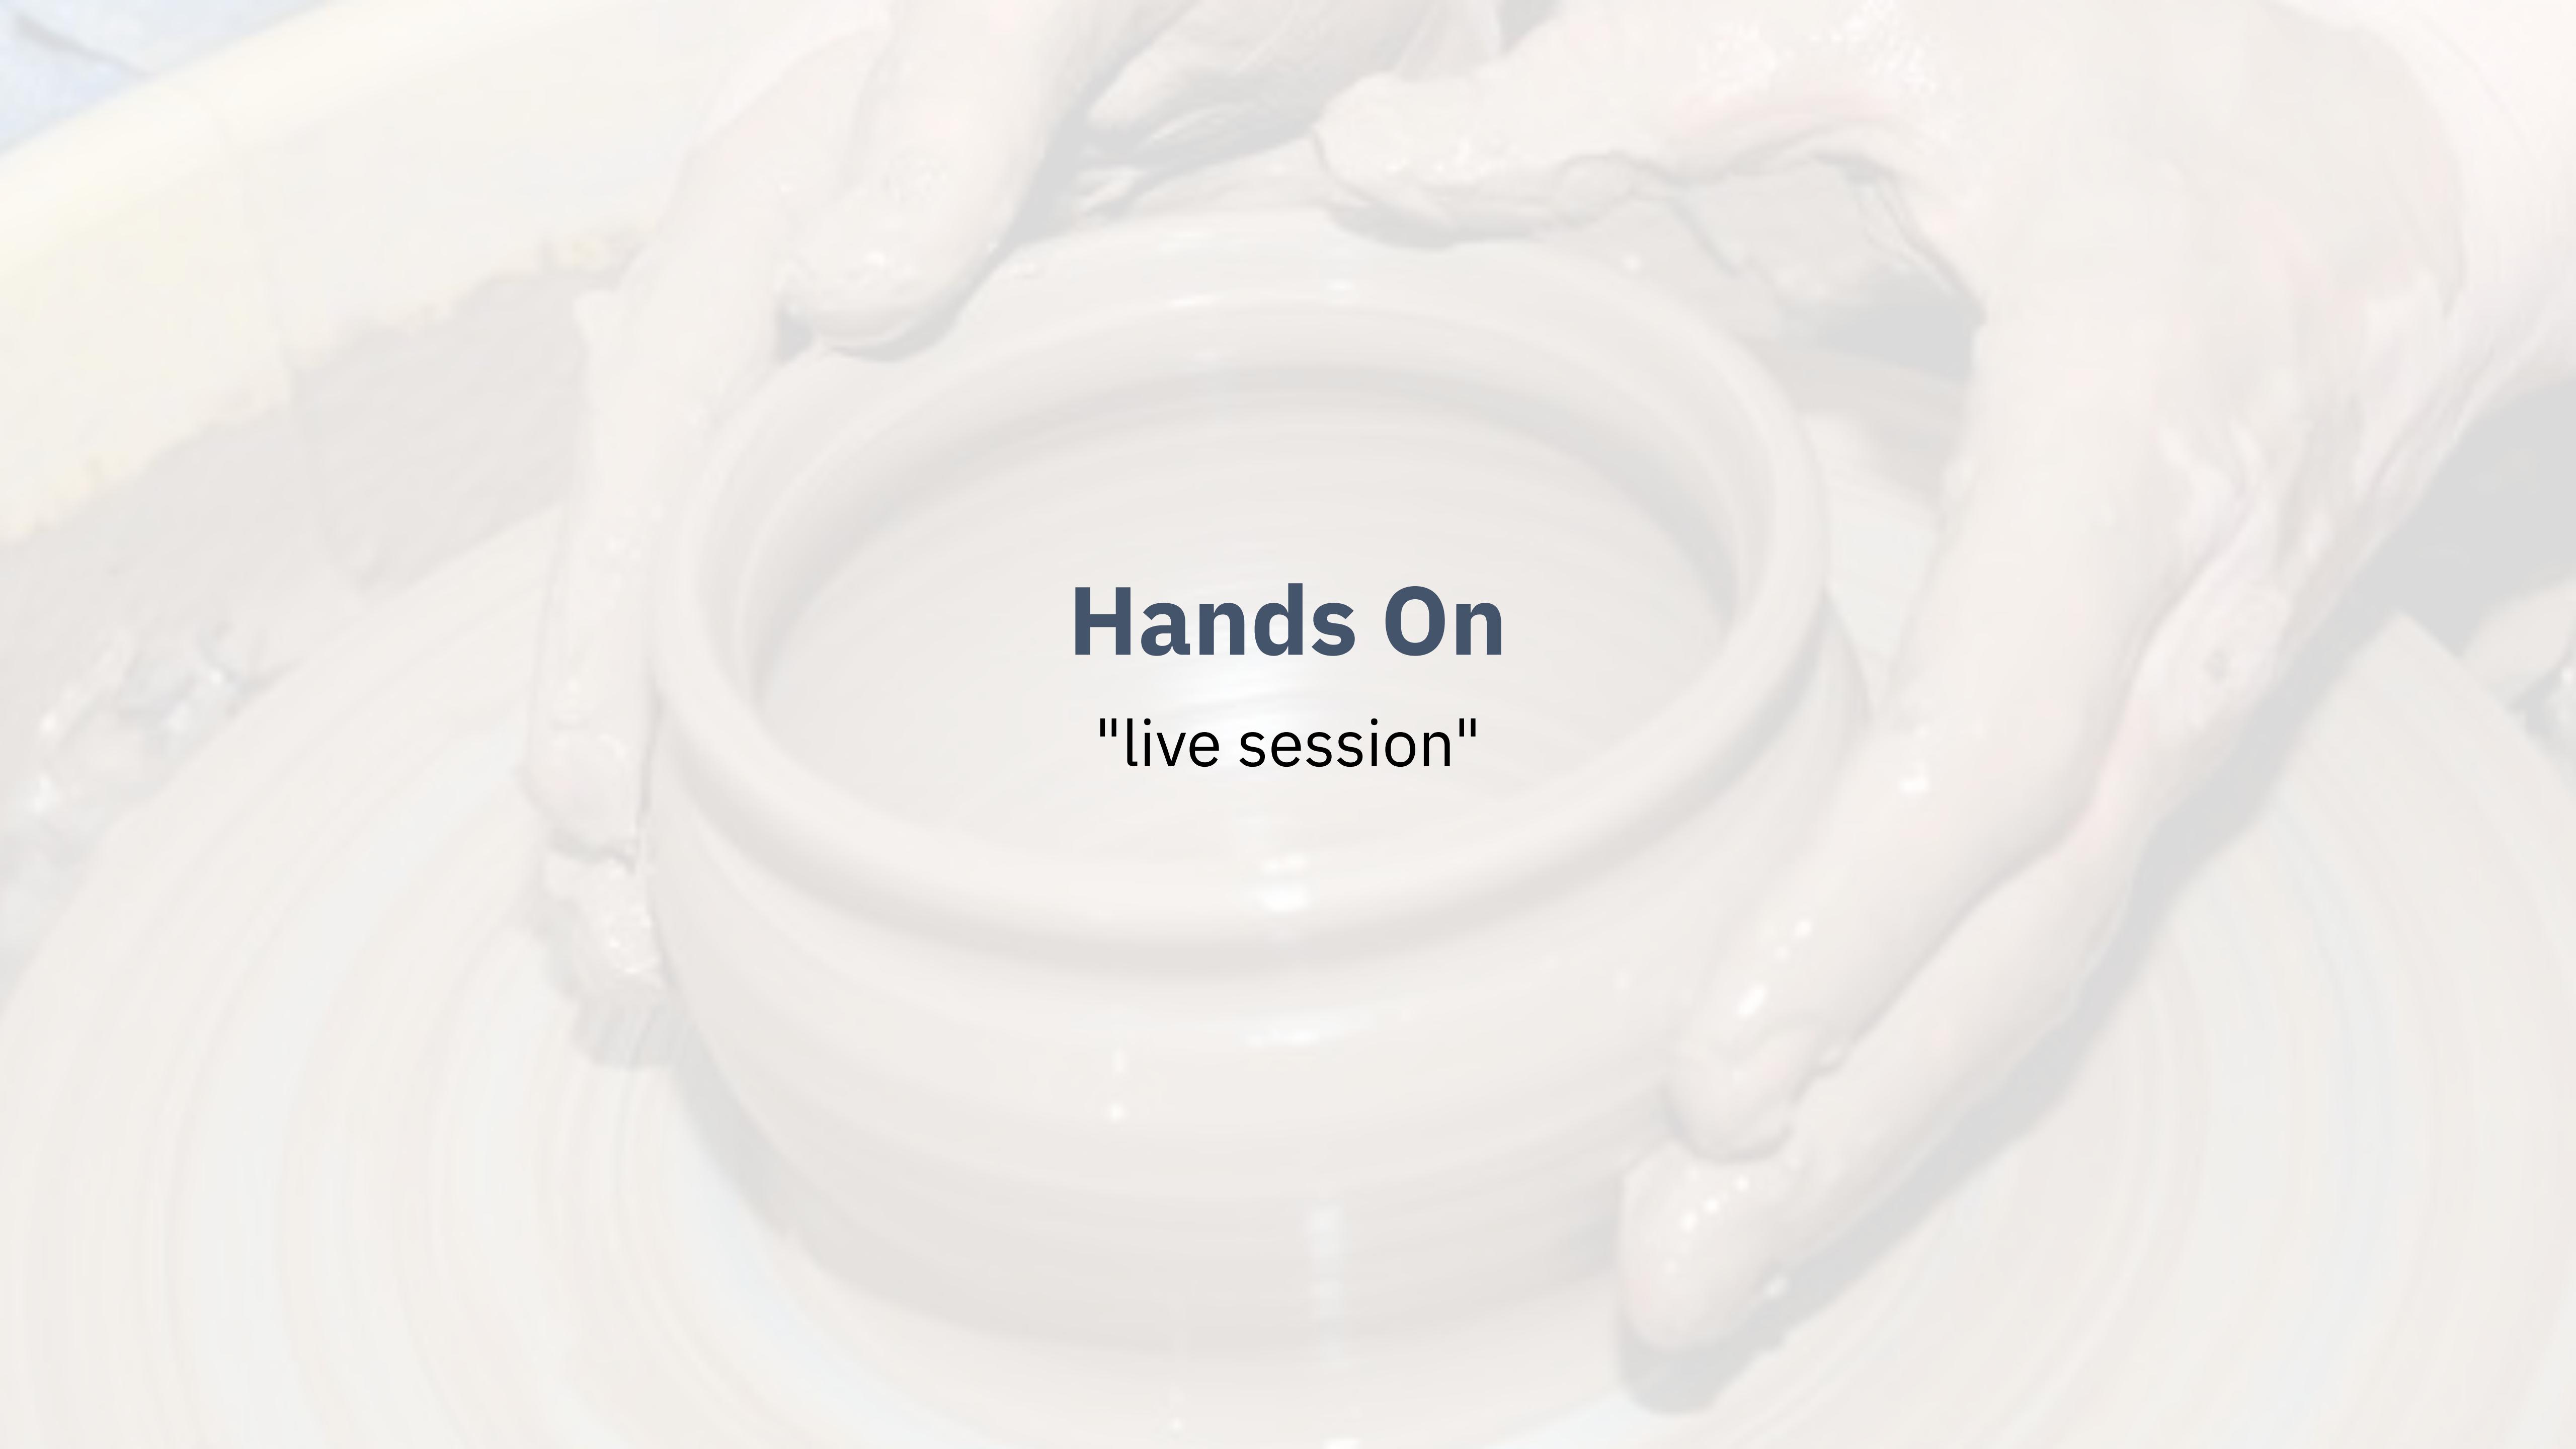


# Kubernetes Autoscaling

Kubernetes allows pods to be manually scaled either using ReplicaSets or Deployments. This can be achieved using the following CLI command:

```
$ kubectl scale --replicas=<desired-instance-count> deployment/<application-name>
```



A grayscale photograph showing a close-up of a person's hands. The hands are positioned over a laptop keyboard, with fingers resting on the keys. The background is slightly blurred, emphasizing the hands and the laptop.

# Hands On

"live session"



# Configure your Kubernetes Cluster

- We need to set permissions on the metadata that minikube needs:

```
$ sudo chown -R $USER $HOME/.kube $HOME/.minikube  
  
$ sudo minikube addons enable ingress
```

# Confirm that your cluster Works

- Make sure that your cluster is configured correctly

```
$ kubectl get all
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3m44s

# Deploy Nginx

- Create a Deployment

```
$ kubectl create deployment my-nginx --image=nginx:alpine
deployment.apps/my-nginx created
```

- Check that the deployment and pods were created

```
$ kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
my-nginx   1/1     1           1           3h42m
```

```
$ kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
my-nginx-b49cf867-q7mjn      1/1     Running   0          3h42m
```

# Deploy Nginx

- Create a Service

```
$ kubectl expose deploy my-nginx --type=NodePort --port=80
service/my-nginx exposed
```

- Check that the service was created

\$ kubectl get service				
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
kubernetes	ClusterIP	172.21.0.1	<none>	443/TCP
317d				
my-nginx	NodePort	172.21.12.226	<none>	80:30371/TCP
3h42m				

# Kubectl get all

```
vagrant@ubuntu-xenial:~$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/my-nginx-b49cf867-q7mjn	1/1	Running	0	3h33m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	172.21.0.1	<none>	443/TCP	317d
service/my-nginx	NodePort	172.21.12.226	<none>	80:30371/TCP	3h30m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/my-nginx	1/1	1	1	3h33m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/my-nginx-b49cf867	1	1	1	3h33m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/my-nginx	1/1	1	1	3h33m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/my-nginx-b49cf867	1	1	1	3h33m

# Naming Convention

- We made a deployment called my-nginx and it created:

**Deployment:** my-nginx

**ReplicaSet:** my-nginx-b49cf867

**Pod:** my-nginx-b49cf867-q7mjn

# Naming Convention

- We made a deployment called my-nginx and it created:



# Naming Convention

- We made a deployment called my-nginx and it created:

**Deployment:** my-nginx

**ReplicaSet:** my-nginx-b49cf867

Pod name starts with ReplicaSet Name

**Pod:** my-nginx-b49cf867-q7mjn

# Get the NodePort

\$ kubectl get service my-nginx					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	NodePort	172.21.9.54	<none>	80:30187/TCP	37d

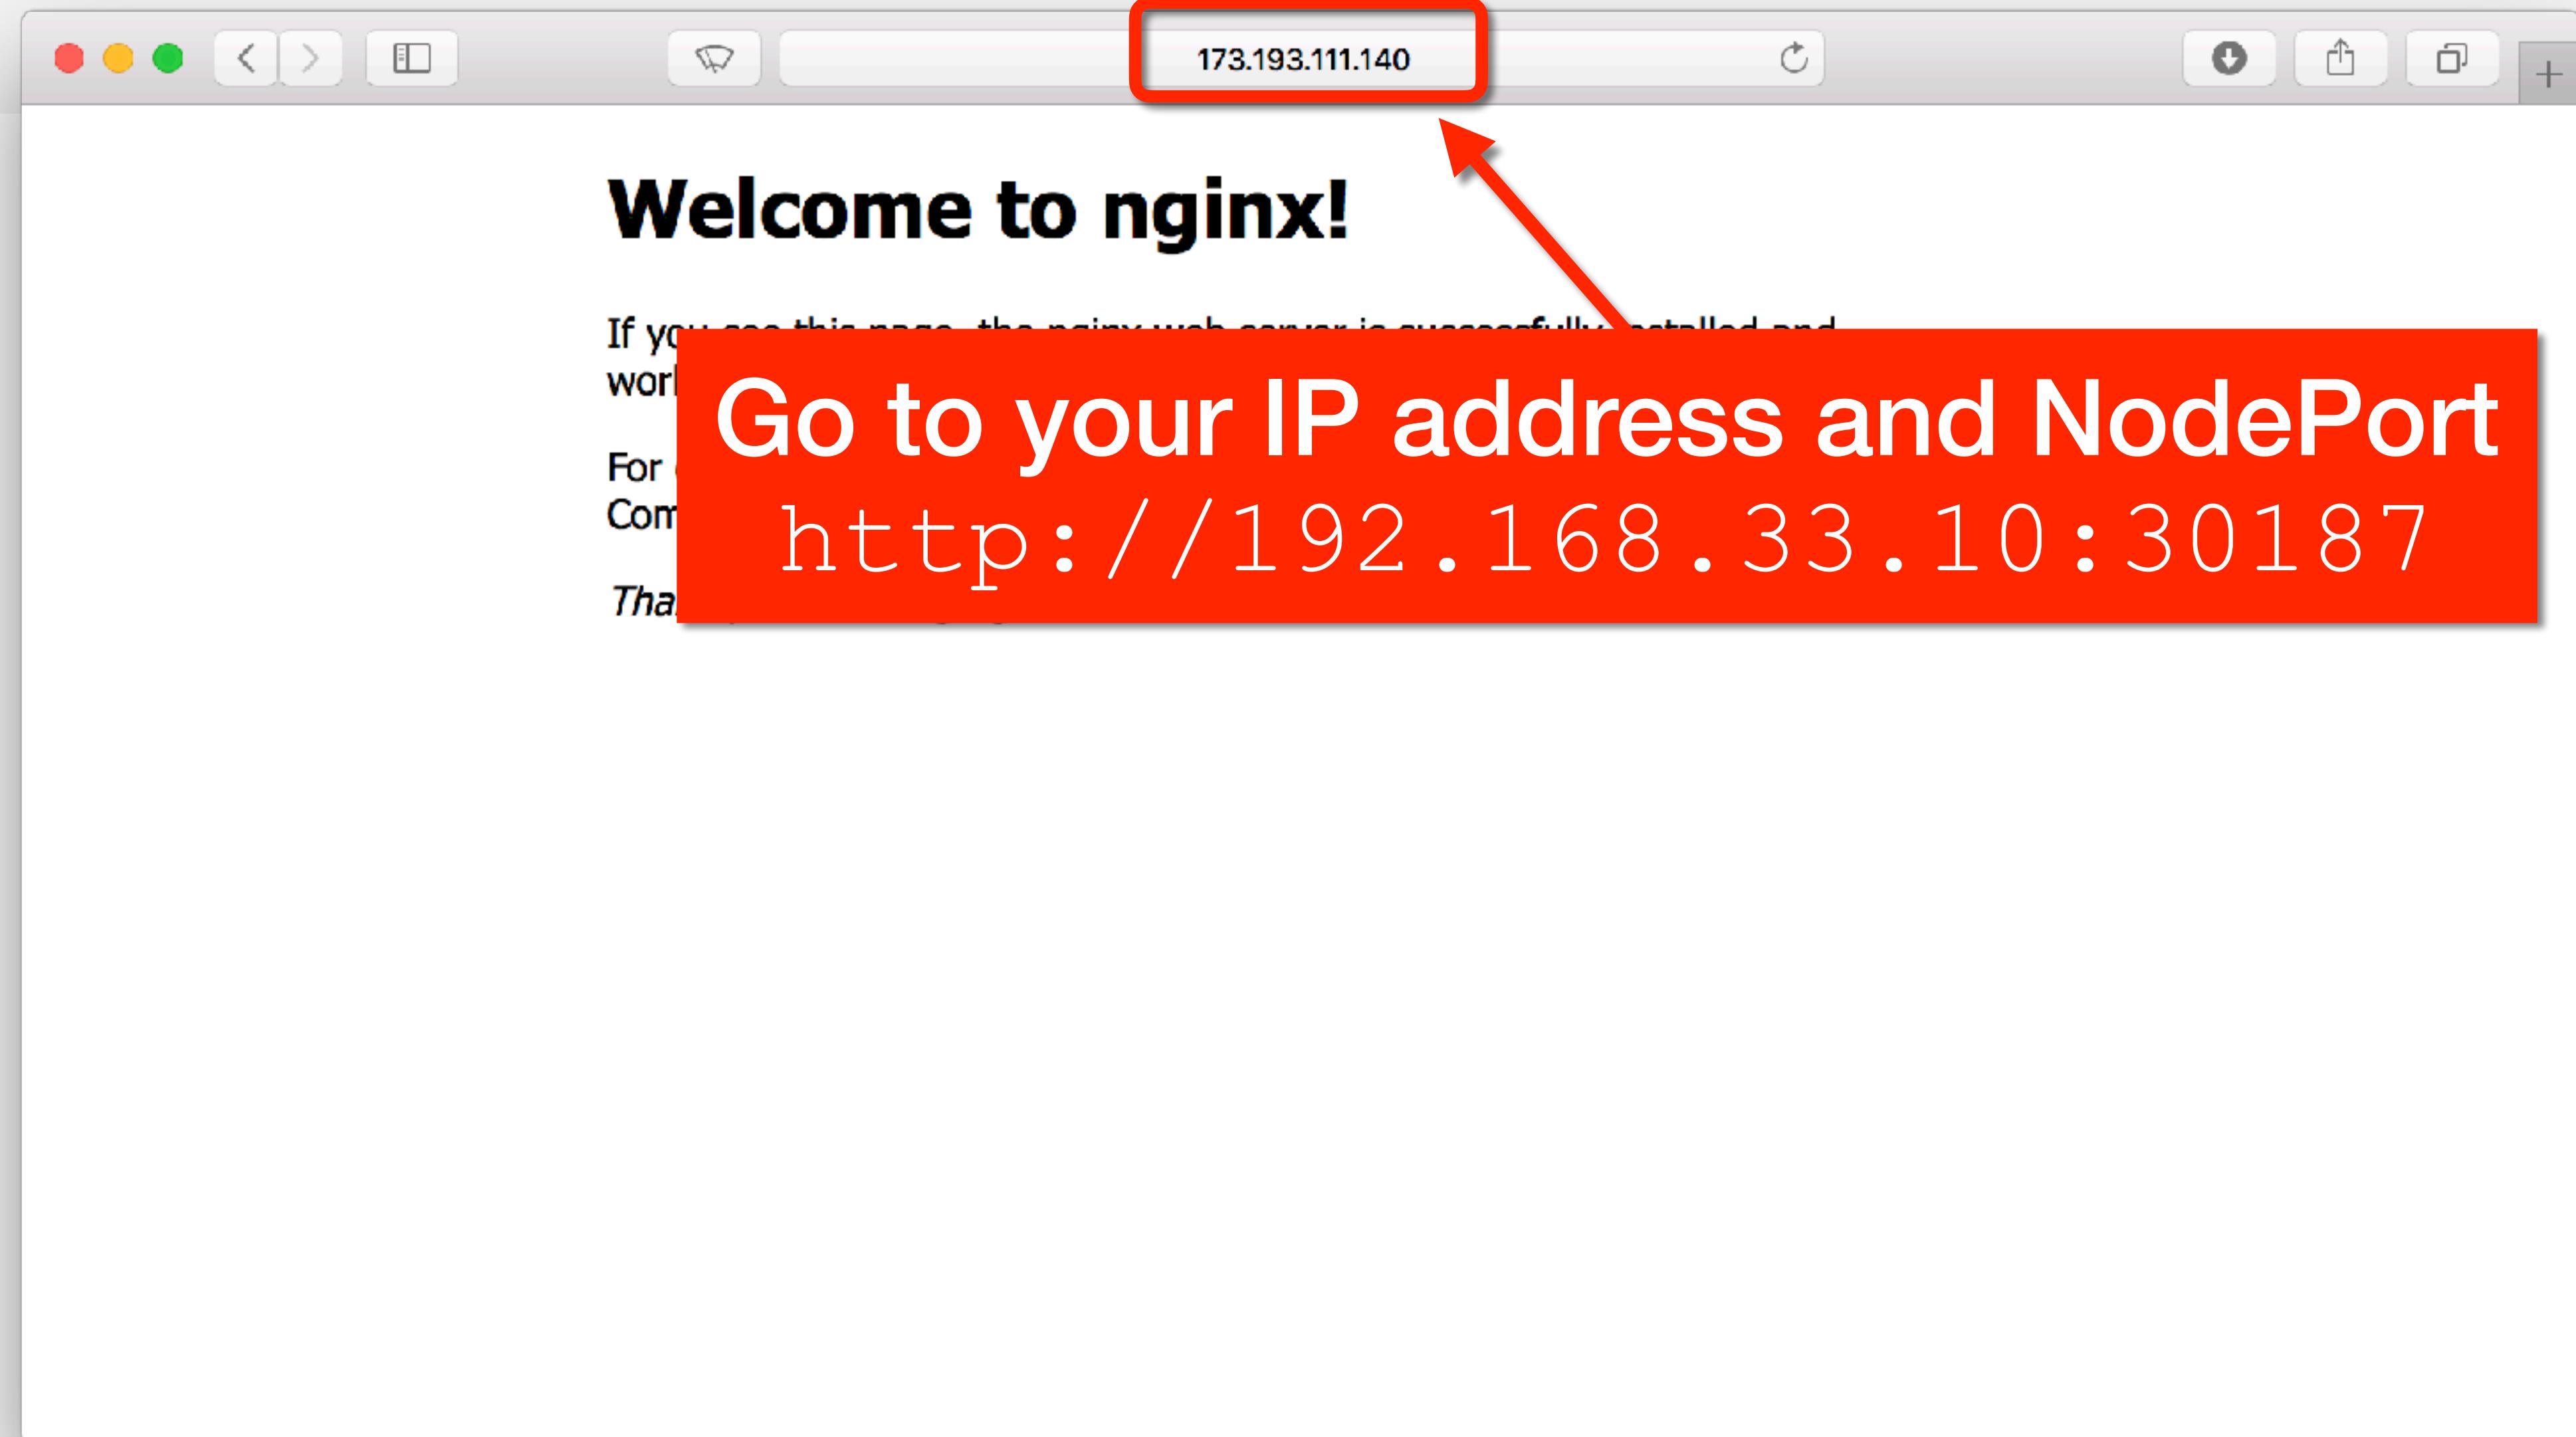
This is the NodePort that was assigned

# Get the Public IP address

The IP address of the Vagrant Virtual Machine is in the Vagrantfile

**192.168.33.10**

# Access NodePort



# Create your Credentials Secret

- Use base64 to encode your secret and add it to the `secret.yaml` file:

```
$ echo "postgres://admin:s3cr3t@postgres:5432/postgres" | base64  
cG9zdGdyZXM6Ly9hZG1pbjpzM2NyM3RAcG9zdGdyZXM6NTQzMj9wb3N0Z3Jlcwo=  
  
$ vi secret.yaml  
  
$ kubectl apply -f secret.yaml  
  
$ kubectl apply -f postgres.yaml
```

# Create A PostgreSQL Service\*

- The pet demo application requires Postgres to store it's state
- We can deploy a Postgres container from the `postgres.yaml` file in the `./deploy` folder of the current repository

```
$ cd /vagrant/deploy
```

```
$ kubectl apply -f postgres.yaml
deployment.apps/postgres created
service/postgres created
```

\* Note: in a "real" deployment we would attach a volume to the Redis container to persist it's data beyond the life of the container

# postgres.yaml

Service

```
apiVersion: v1
kind: Service
metadata:
  name: postgres
spec:
  type: NodePort
  selector:
    app: postgres
  ports:
    - name: primary
      protocol: TCP
      port: 5432
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres
spec:
  replicas: 1
  selector:
    matchLabels:
      app: postgres
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:alpine
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 5432
              protocol: TCP
          env:
            - name: POSTGRES_USER
              value: "admin"
            - name: POSTGRES_PASSWORD
              value: "s3cr3t"
```

Deployment

# postgres.yaml

Service

```
apiVersion: v1
kind: Service
metadata:
  name: postgres
spec:
  type: NodePort
  selector:
    app: postgres
  ports:
    - name: primary
      protocol: TCP
      port: 5432
```

Don't try this at home  
...it's only a demo!

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: postgres
spec:
  replicas: 1
  selector:
    matchLabels:
      app: postgres
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgres
          image: postgres:alpine
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 5432
              protocol: TCP
          env:
            - name: POSTGRES_USER
              value: "admin"
            - name: POSTGRES_PASSWORD
              value: "s3cr3t"
```

Deployment

# Now we can Deploy Our Image

- You can use kubectl apply with the files under /vagrant/deploy

```
$ cd /vagrant/deploy
```

```
$ kubectl apply -f deployment.yaml
```

```
deployment.apps/pet-demo created
```

```
$ kubectl apply -f service.yaml
```

```
service/pet-demo-service created
```

```
$ kubectl apply -f ingress.yaml
```

```
ingress.extensions/pet-demo-ingress created
```

# Check Kubernetes

```
$ kubectl get all
```

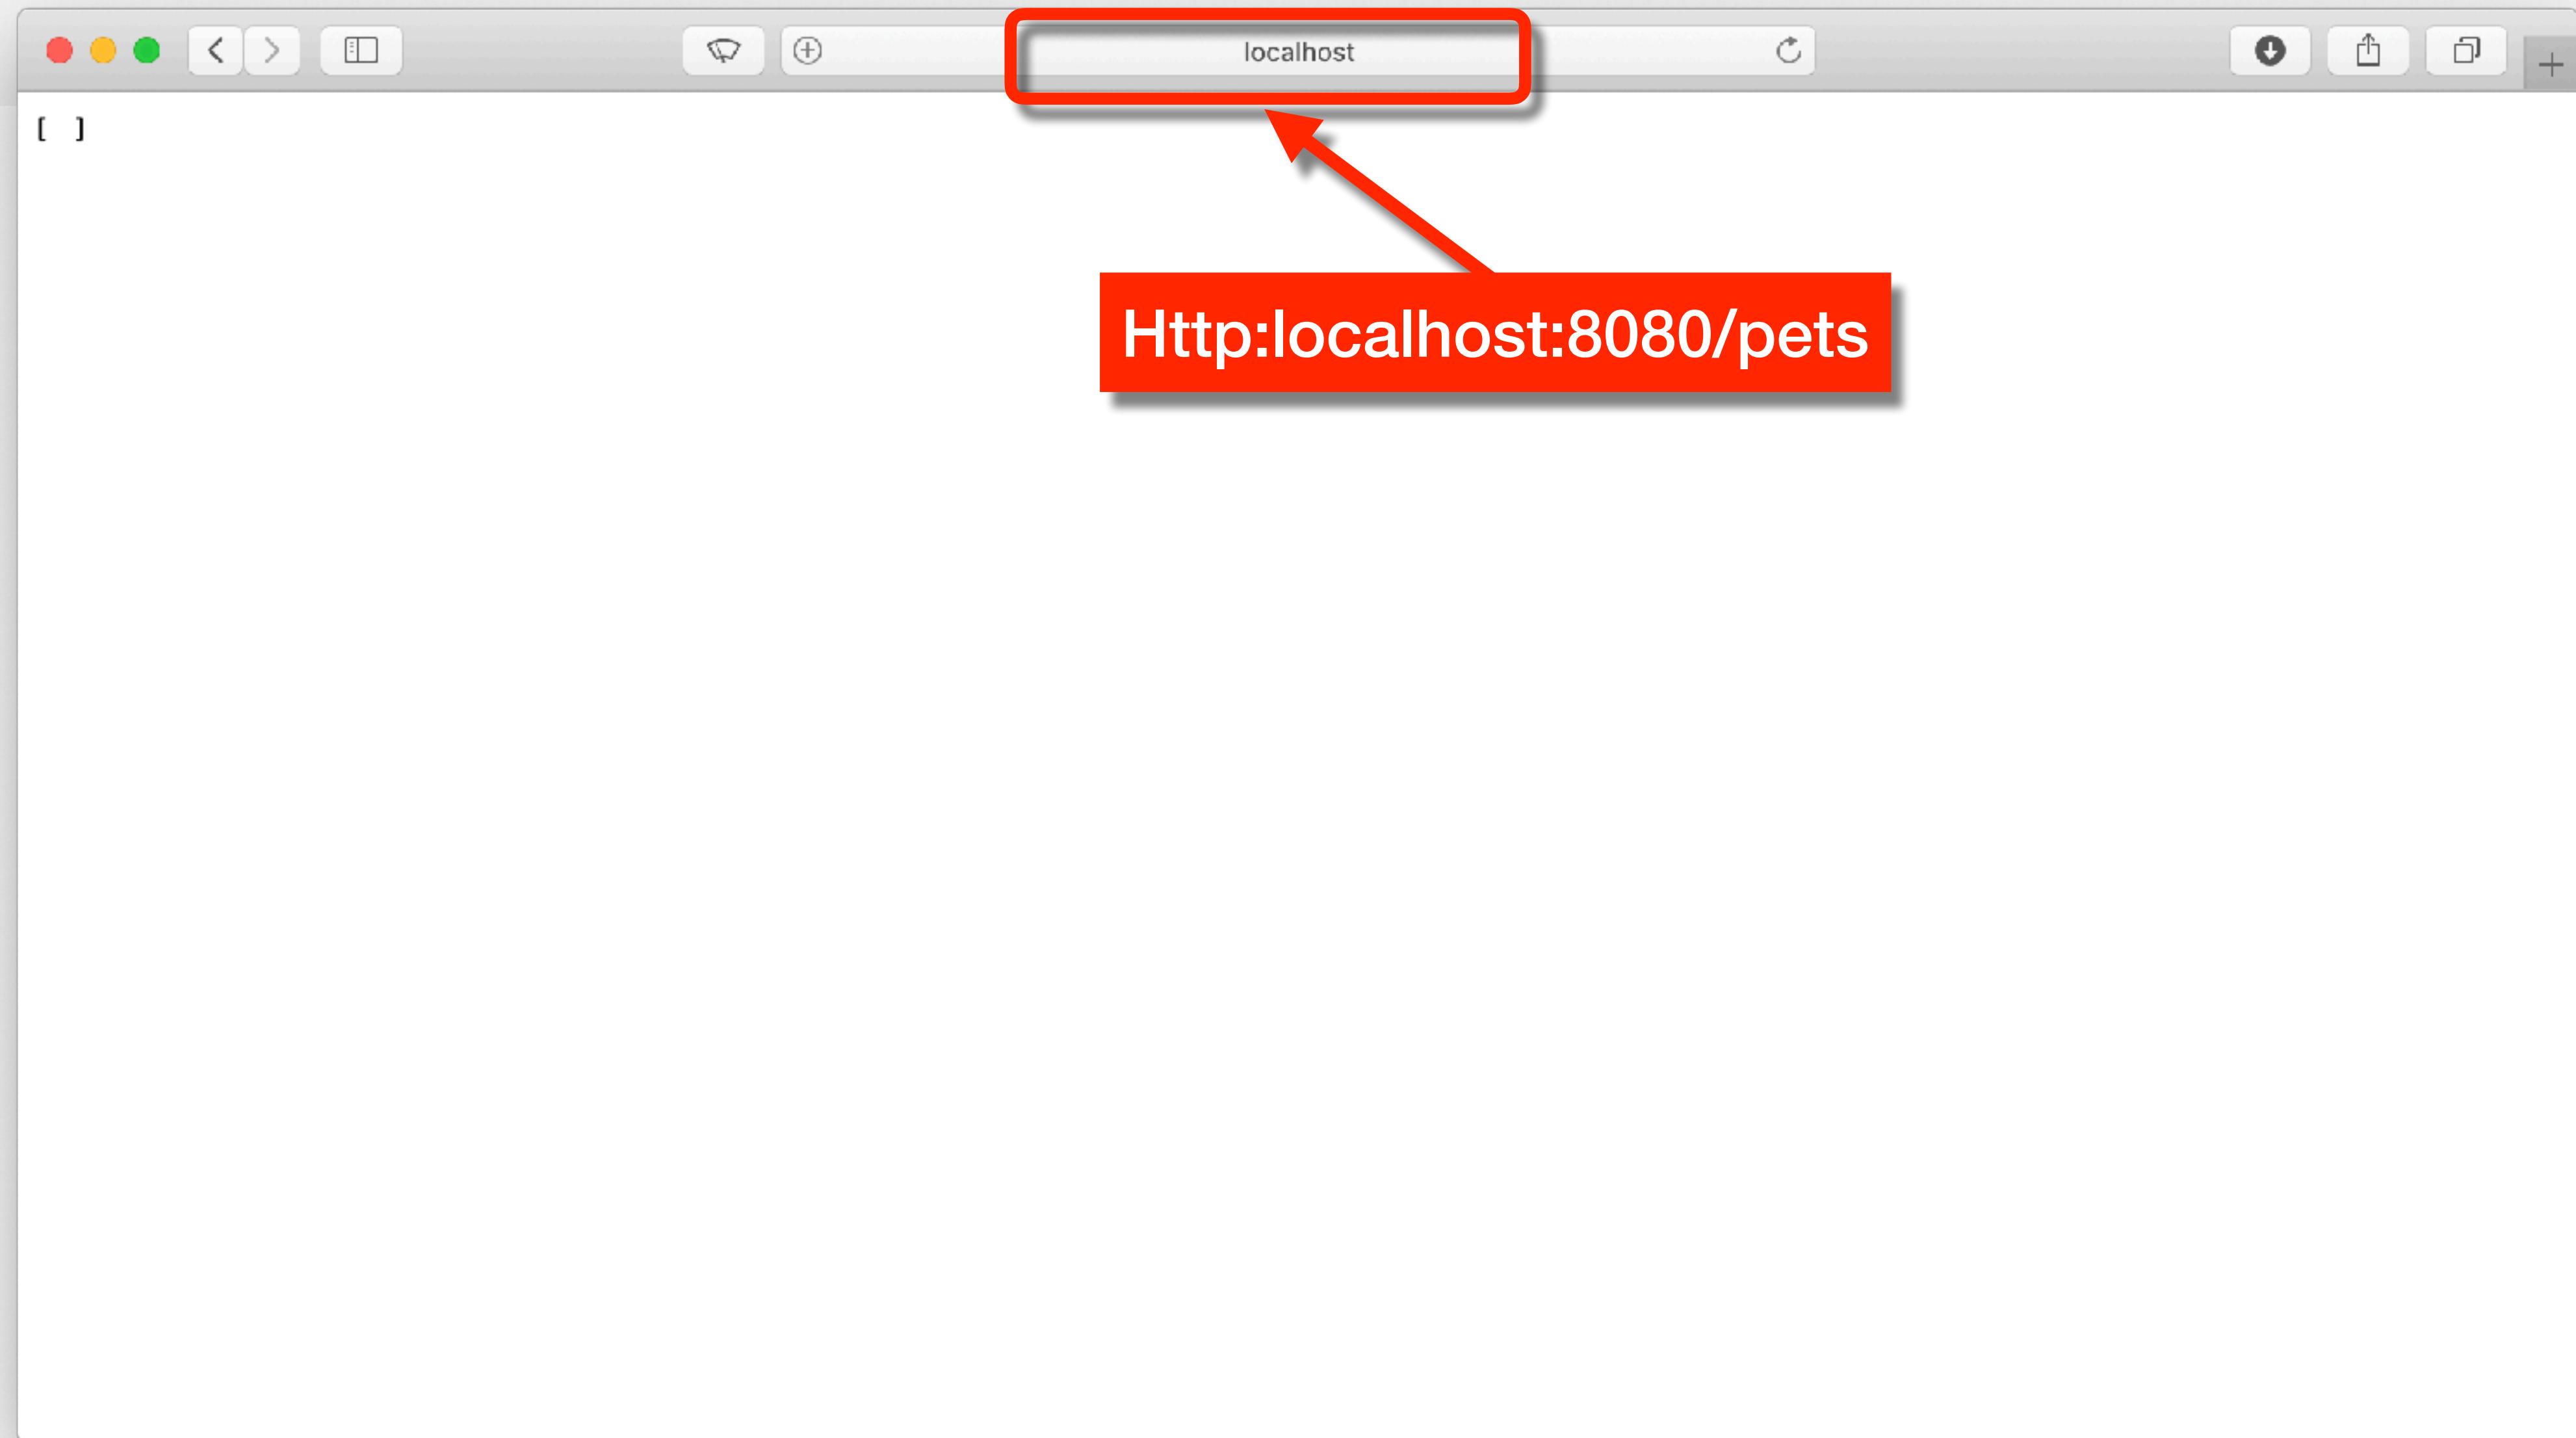
NAME	READY	STATUS	RESTARTS	AGE
pod/pet-demo-b8f9cc44d-7dq5w	1/1	Running	0	2m12s
pod/postgres-6565c9f5b8-bdbm4	1/1	Running	0	65m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	76m
service/pet-demo-service	NodePort	10.98.93.70	<none>	5000:31489/TCP	2m3s
service/postgres	NodePort	10.102.252.68	<none>	5432:31967/TCP	51m

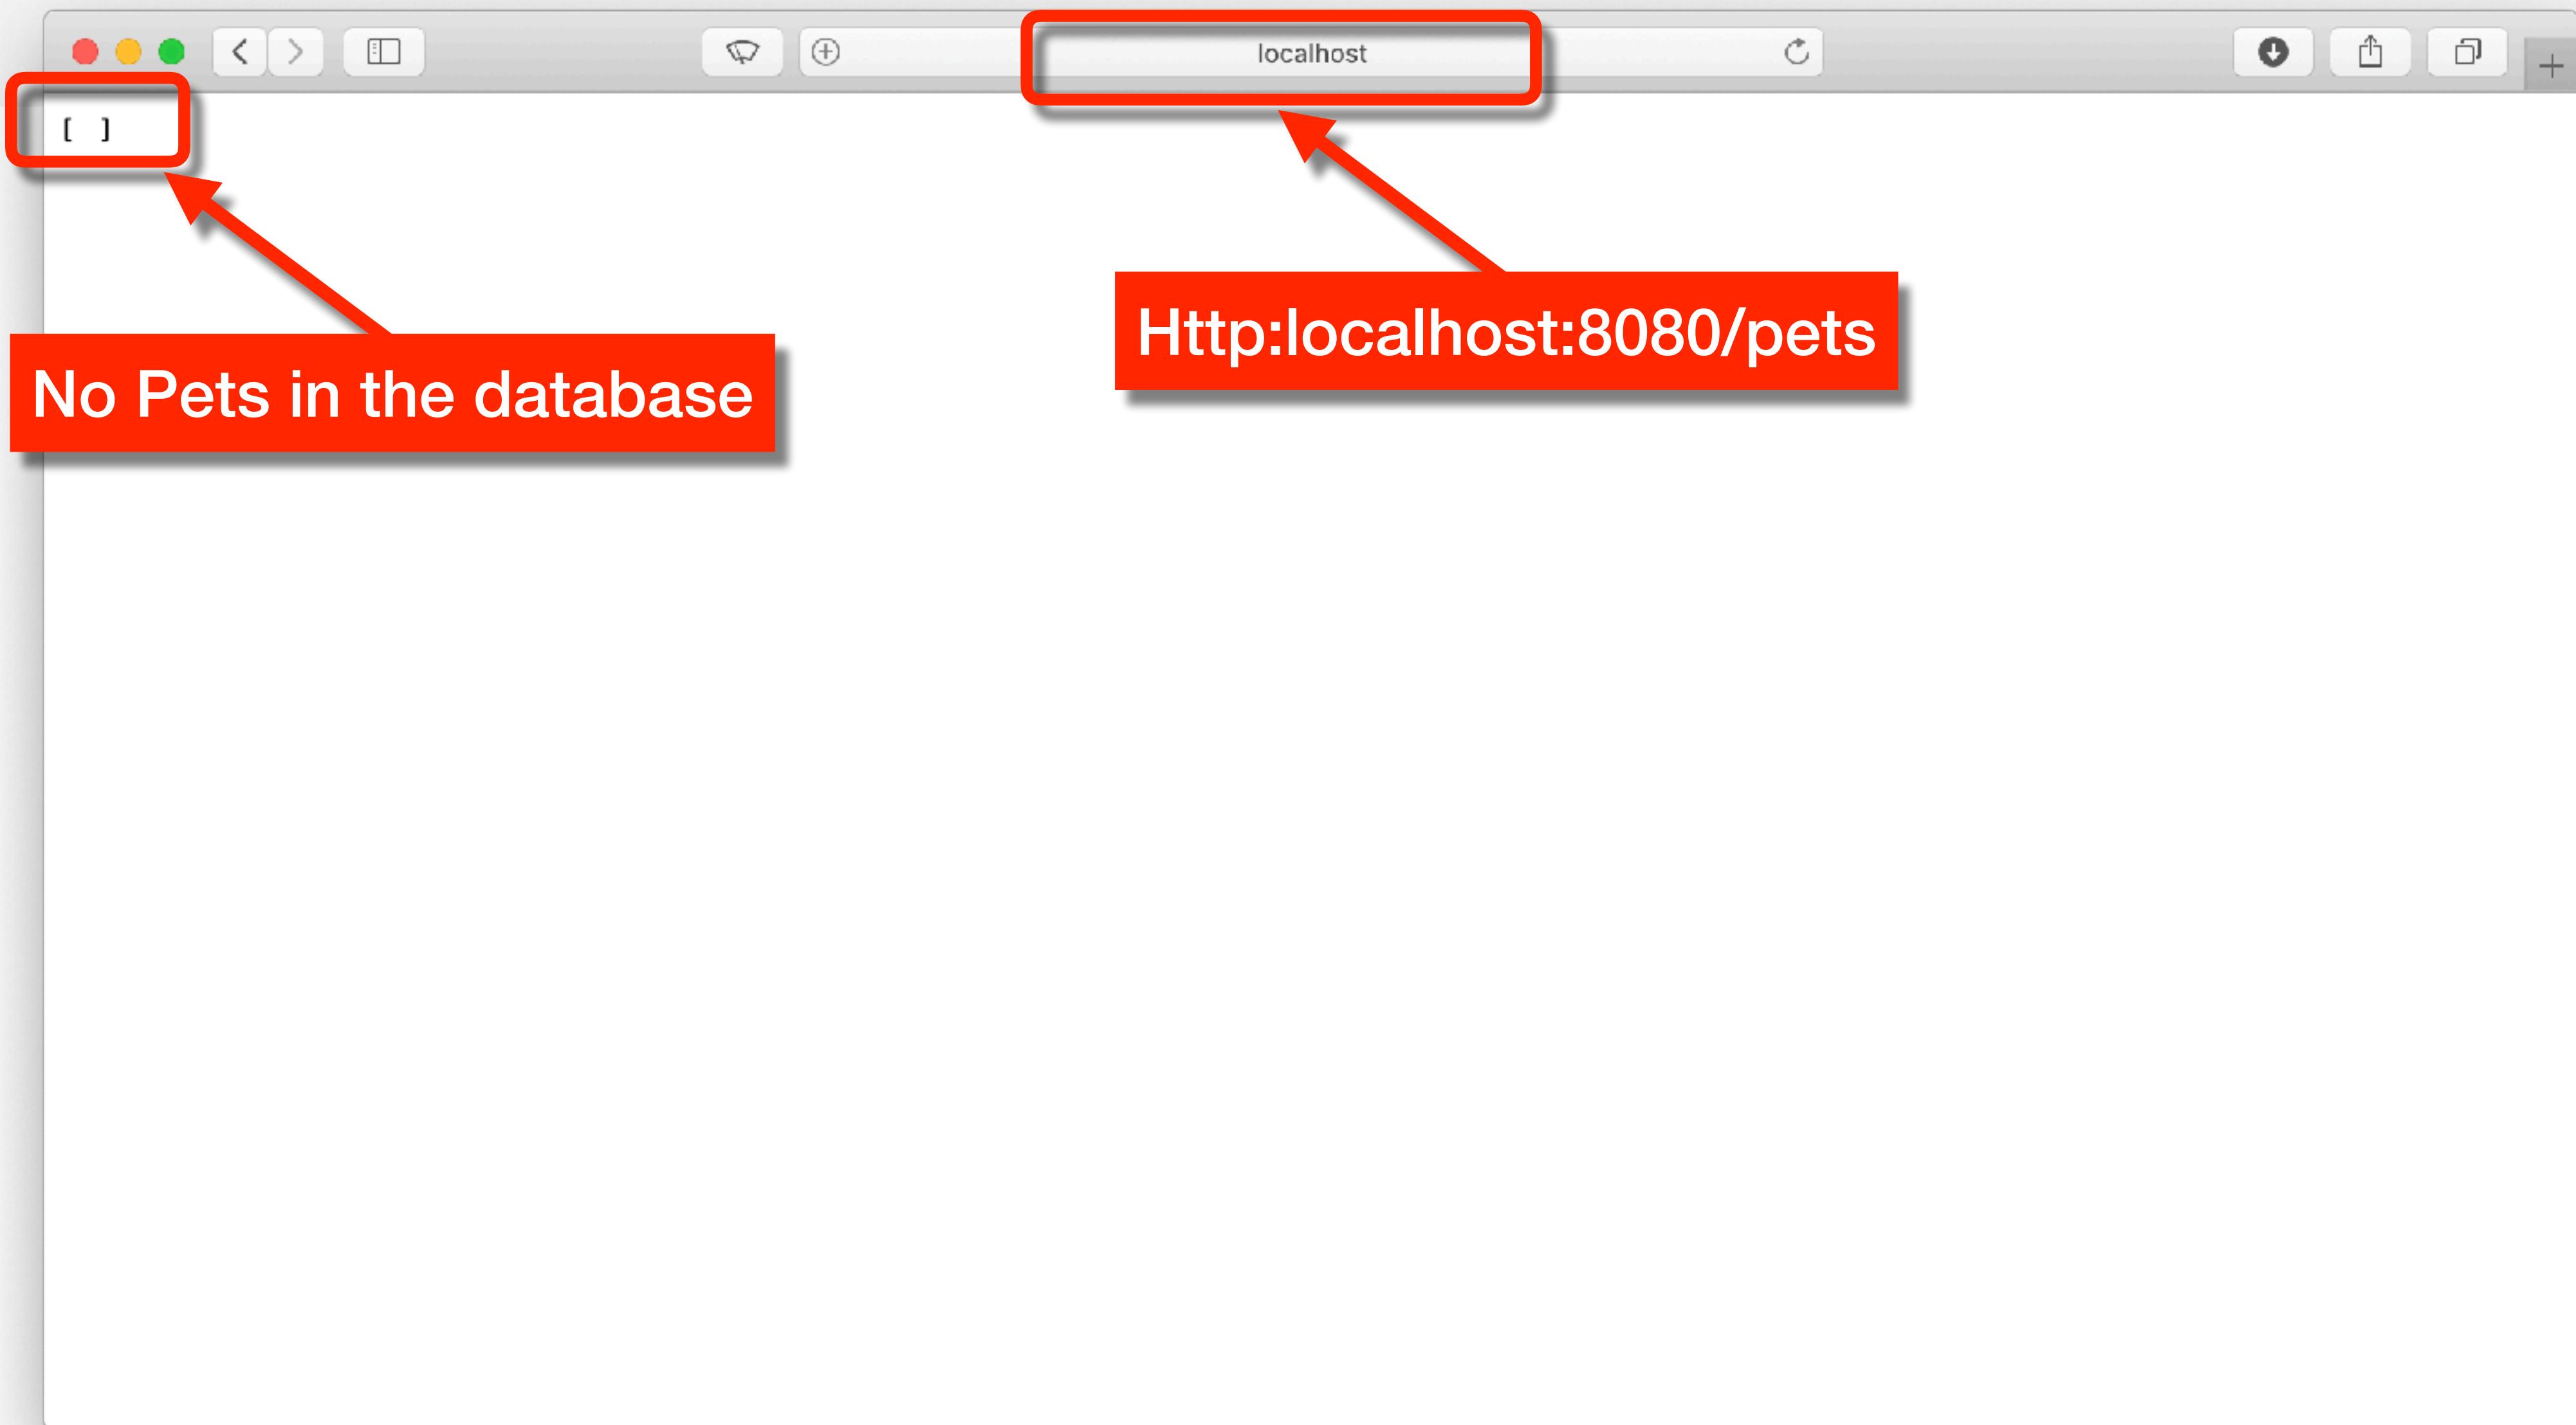
NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/pet-demo	1/1	1	1	2m12s
deployment.apps/postgres	1/1	1	1	65m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/pet-demo-b8f9cc44d	1	1	1	2m12s
replicaset.apps/postgres-6565c9f5b8	1	1	1	65m

# Your App on Kubernetes



# Your App on Kubernetes



# Use HTTP command to create some data

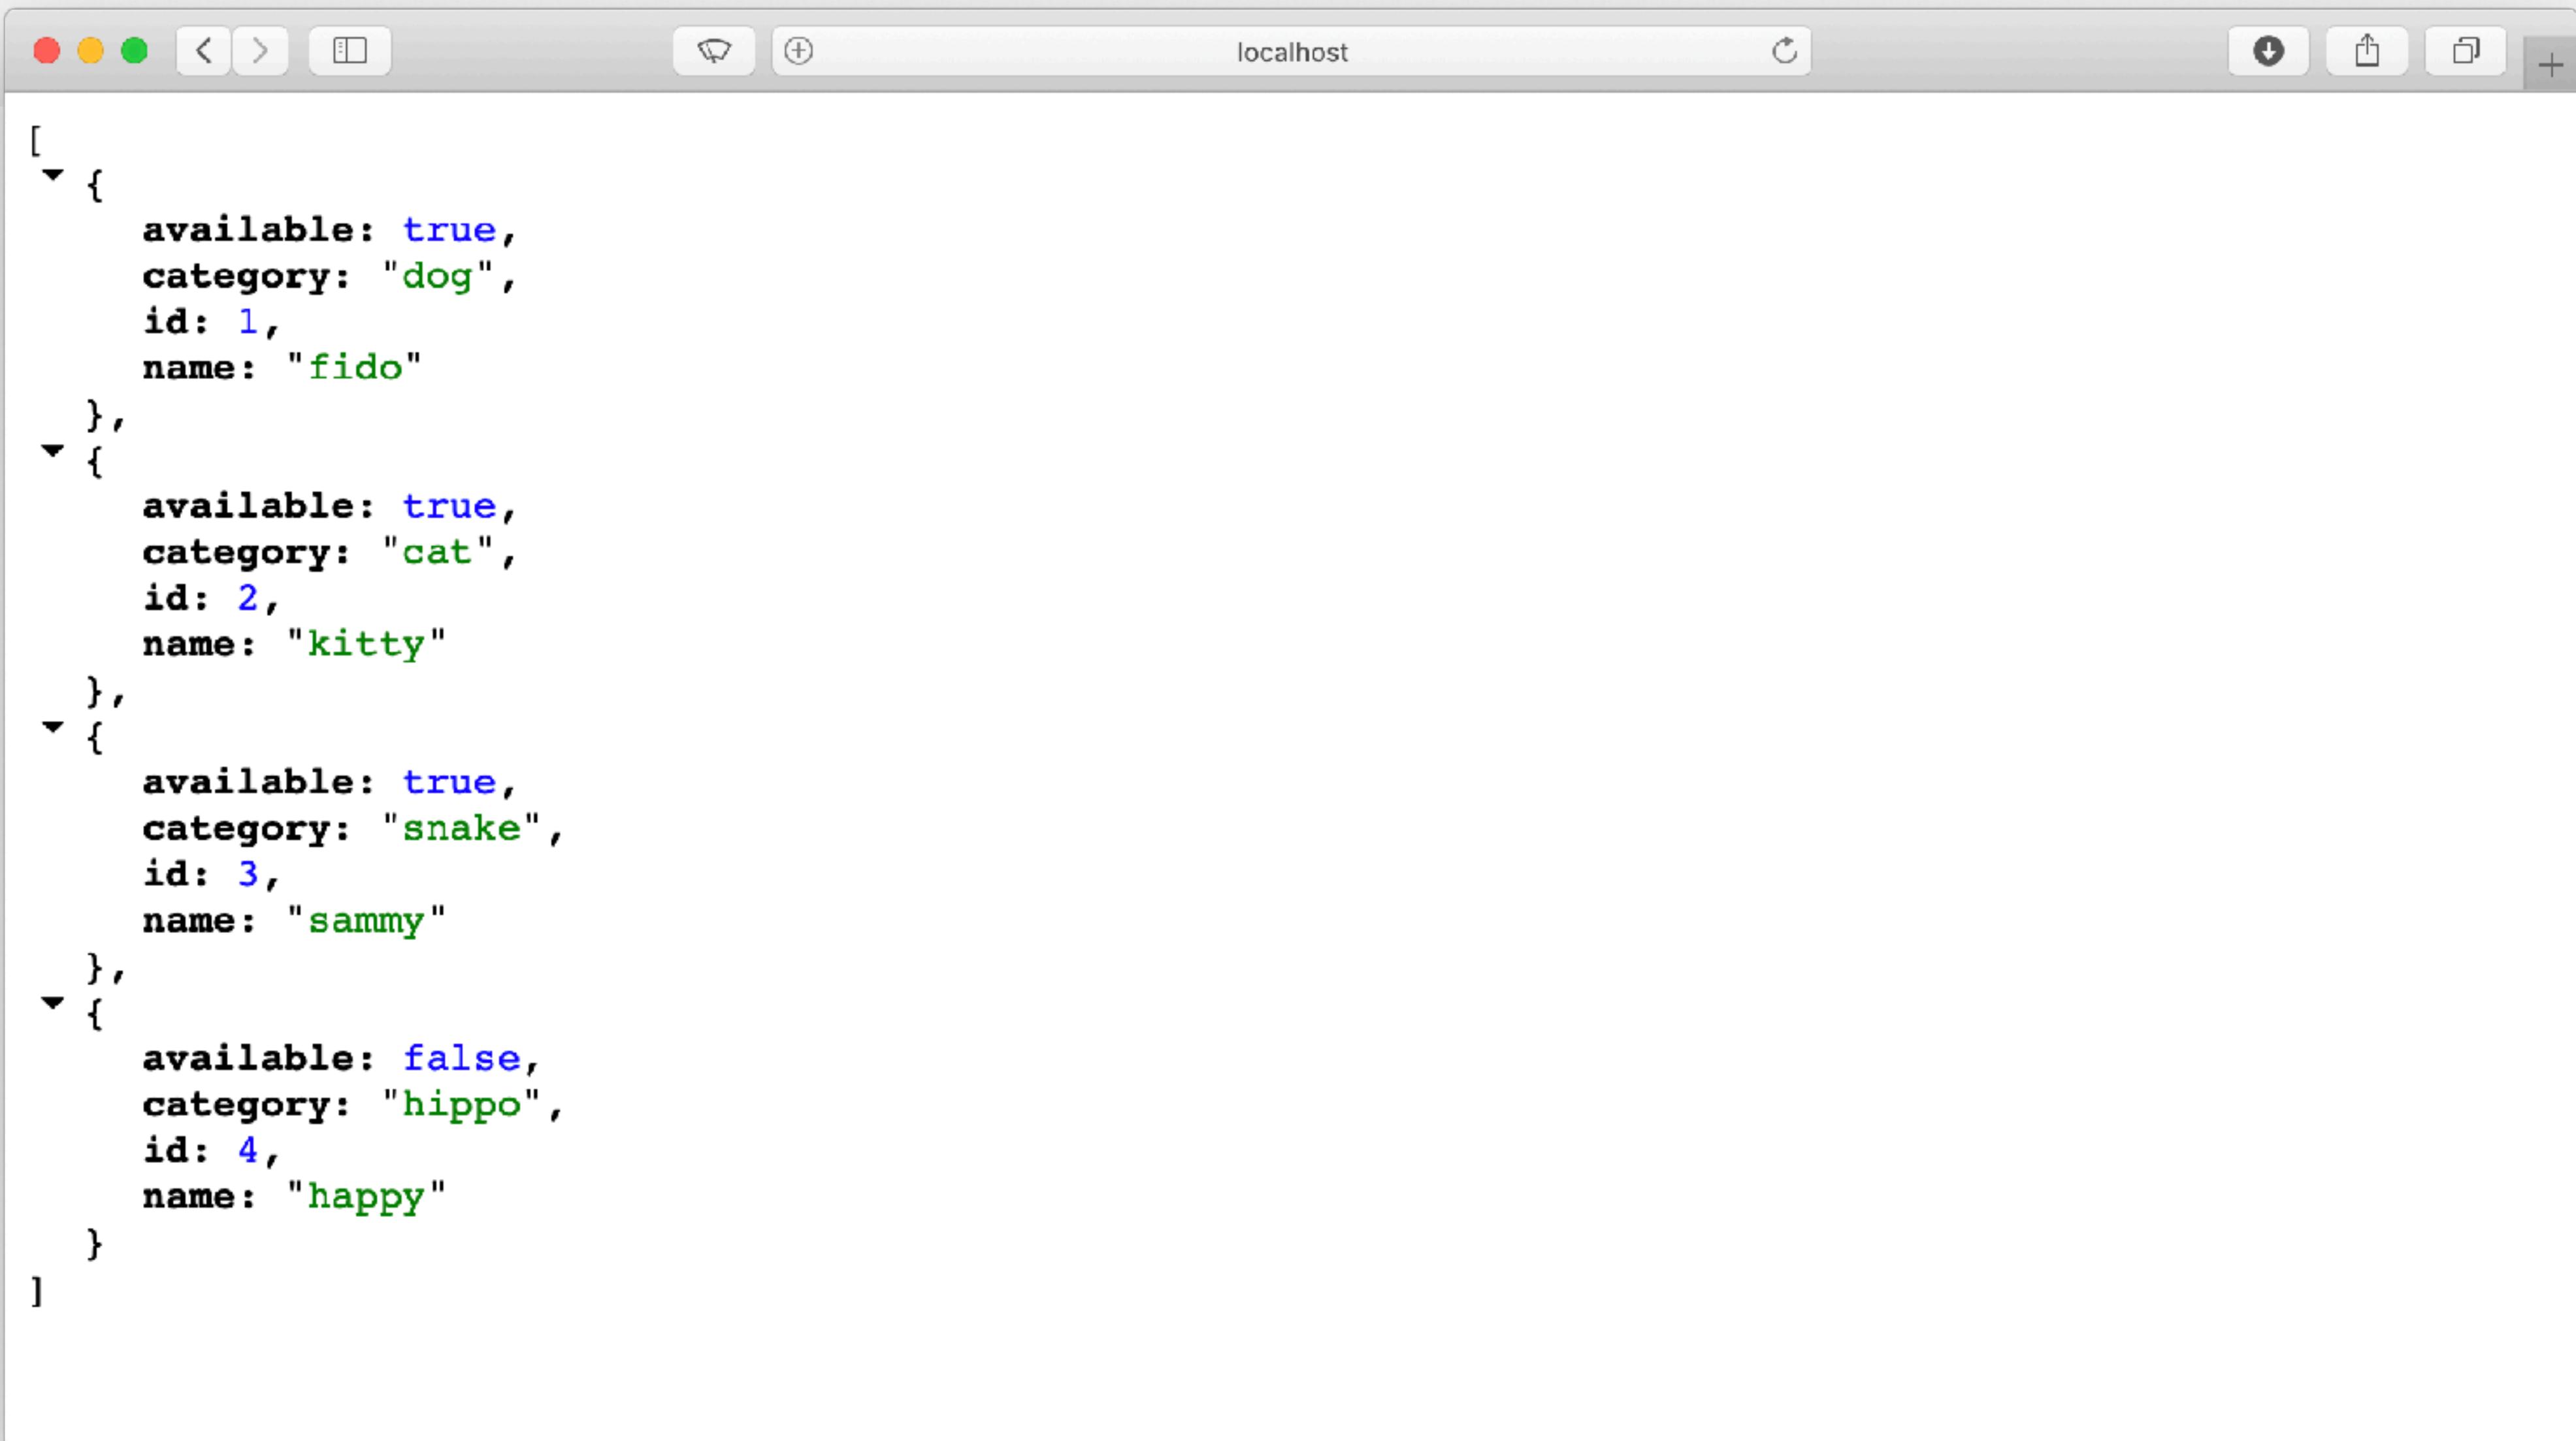
```
$ http post localhost/pets name='fido' category='dog' available:=true  
HTTP/1.1 201 CREATED  
Connection: keep-alive  
Content-Length: 57  
Content-Type: application/json  
Date: Thu, 04 Jul 2019 21:37:46 GMT  
Location: http://localhost/pets/3  
Server: nginx/1.15.9  
  
{  
  "available": true,  
  "category": "dog",  
  "id": 3,  
  "name": "fido"  
}
```

# Add these pets

- Use http to these pets

```
$ http post localhost/pets name='fido' category='dog' available:=true  
$ http post localhost/pets name='kitty' category='cat' available:=true  
$ http post localhost/pets name='sammy' category='snake' available:=true  
$ http post localhost/pets name='harry' category='hippo' available:=false
```

# Refresh for browser



A screenshot of a web browser window titled "localhost". The page displays a JSON array of four objects, each representing an animal with properties: available, category, id, and name. The objects are:

```
[  
  {  
    available: true,  
    category: "dog",  
    id: 1,  
    name: "fido"  
  },  
  {  
    available: true,  
    category: "cat",  
    id: 2,  
    name: "kitty"  
  },  
  {  
    available: true,  
    category: "snake",  
    id: 3,  
    name: "sammy"  
  },  
  {  
    available: false,  
    category: "hippo",  
    id: 4,  
    name: "happy"  
  }  
]
```

# Refresh for browser

localhost

```
[  
  {  
    "available": true,  
    "category": "dog",  
    "id": 1,  
    "name": "fido"  
  },  
  {  
    "available": true,  
    "category": "cat",  
    "id": 2,  
    "name": "kitty"  
  },  
  {  
    "available": true,  
    "category": "snake",  
    "id": 3,  
    "name": "sammy"  
  },  
  {  
    "available": false,  
    "category": "hippo",  
    "id": 4,  
    "name": "happy"  
  }]
```

Here are the Pets we sent  
with the http command  
via the REST API

# Summary

You just created your first Kubernetes deployment

You learned how to create deployments

You learned how to expose deployments as services

You can now deploy your microservices on any Kubernetes cloud

